

Khanh Gia Bui¹, Luong Van Tam¹, Dang Tri Trung², Daud Shahbaz³

Theoretical Machine Learning
Learning Formalism and Double Descent

Theoretical and novel analysis, second draft fixture

Waiting for publication 

¹Department of Physics, Hanoi University of Science, Vietnam National University, Hanoi, Vietnam

²Monash University, Department of Computer Science, Melbourne, Australia

³COMSATS University Islamabad, Islamabad, Pakistan

Contents

Contents	i
List of Figures	v
Introduction	ix
Background and motivation	ix
1 Preliminaries	1
1.1 Concentration inequalities	1
1.1.1 What are concentration inequalities?	1
1.1.2 Relevant toolset	2
1.1.3 Markov’s inequality	2
1.1.4 Chebyshev’s inequality	3
1.1.5 McDiarmid’s inequality	4
1.1.6 The Cramer-Chernoff method	5
1.2 Metric and pseudometric	6
1.3 Graph theory	8
1.3.1 Algebraic representation	9
1.3.2 Permutation invariance and equivariance	11
1.3.3 Graph Fourier Transform	11
2 Classical learning theory	13
2.1 Precursors to learning theoretic	13
2.2 Transparency shift	14
2.2.1 Ladder of functionals	16
2.2.2 Classical statistical learning	18
2.3 Classical setup	20
2.3.1 Instantiation of the learning problem - ordering	25
2.4 Optimal learning - PAC-learning	26
2.4.1 Representation size and PAC-configuration	29
2.4.2 Generalization bounds	31
2.4.3 Agnostic-general PAC-learning	32
2.4.4 Empirical and structural algorithms	34
2.4.5 The PAC-ERM notion	36
2.4.6 Noise, Mohri et al. [2012]	38

2.5	Simplistic learning – Occam’s Razor	38
2.5.1	Occam Learning and Succinctness	39
2.6	Discrete space empiricism – Rademacher complexity	42
2.6.1	Rademacher descriptions	42
2.6.2	Growth function	46
2.7	Equivalent spatial separation – VC-theory	48
2.8	Algorithmic stability	50
2.9	Summary	50
3	Introduction to Double Descent	53
3.1	Introduction	53
3.1.1	Outline	54
3.1.2	Related works	55
3.2	Foundational investigation	55
3.2.1	Preliminary – precursor double descent estimation	56
3.3	The bias-variance – double descent dilemma	64
3.3.1	Bias-variance tradeoff	65
3.3.2	Double descent	69
3.4	What is double descent, then?	72
3.4.1	Underfitting and overfitting	73
3.4.2	Discovery of double descent	73
3.4.3	Theoretical attacks	74
3.4.4	High-dimensional ridgeless interpolation – Hastie et al. [2019]	75
3.4.5	Sample-wise double descent – Nakkiran [2019b]	77
3.4.6	Multiscale feature learning dynamics – Pezeshki et al. [2021]	77
3.4.7	Reinforcement learning and double descent – Brellmann et al. [2024]	78
3.4.8	Variations on analysis	82
4	Experimental investigations	85
4.1	Experimental planning	85
4.1.1	Setting: gradient-based optimization	85
4.1.2	Experiments list	87
4.1.3	Preliminaries on observations – Schaeffer et al. [2023]	89
4.2	Data generation	90
4.3	Linear function class $h(u)$ – Lafon and Thomas [2024]	99
4.3.1	Extension	102
4.4	Polynomial models	106
4.4.1	Structure	106
4.4.2	Scenario schema	108
4.4.3	Procedure	109
4.5	Graph Neural Network (GNN) – Shi et al. [2024]	110
4.5.1	Methodology	110
4.5.2	Double descent and GNN	114
4.6	Support Vector Machine (SVM)	115
5	Interpretation	117
5.1	Structural and operational problems	117

5.1.1	Bias-variance and the definitive form	118
5.1.2	Computing bias and variance	120
5.1.3	The definitions of double descent	123
5.1.4	Gradient descent and metric mismatch	125
5.2	Complexity notion and its pluralism	125
5.2.1	Current complexity measure	128
5.3	The nitpicking set	131
5.3.1	Authoritative statements	131
5.3.2	Model complexity	132
5.3.3	Model structure and the bounding inequality	132
5.3.4	Theoretical insufficiency	134
6	Intersectional insights	137
6.1	Interpreting the learning theory	138
6.1.1	Domain and process	138
6.2	The learning theoretic approach	139
6.2.1	Overfitting and underfitting	143
6.3	System-wide structuralization	146
6.3.1	The idea of neural network formalism	146
6.3.2	Attractor-based insights	150
6.4	Reattaching the neural perspective	151
6.5	Motivation	152
6.5.1	Neurons perspective	152
6.5.2	Why we call it minimal	157
6.5.3	Preliminary Formalisation	157
6.6	Polynomial insights	158
7	Modelling neural network	169
7.1	Multilayer perceptron	169
7.2	Experiments	169
7.2.1	Experimental A - Conditioned MLP	169
8	Statistical physics	173
8.1	Note for team writing member	174
8.1.1	Main structure	174
8.1.2	Notices	175
8.1.3	In detail systematic projection	177
8.2	Approach A - Isomorphism	179
8.2.1	From Analogy to Isomorphism	179
8.2.2	Limitations of Equilibrium Theories	181
8.3	Identification: The Dynamic-Landscape Interaction	182
8.3.1	Experimental Verification: A Toy Model Proposal	184
8.3.2	Physical Observables: Beyond Test Error	185
8.3.3	Experimental Protocol	186
8.3.4	Extension	187
	Index	191

Reevaluation of theorem 4.24	192
Remark	193
Observing gradient descent decomposition	194
Details of supervised learning structure	197
Practical consideration	201
Observational space partitioning	201
Bibliography	203

List of Figures

2.1	An illustration of statistical learning theory on the evaluation of the risks and errors, during learning process. c' is presented in the ‘orbital’ vicinity around c , with its distance of certain metric define how ‘accurate’ the reconstruction from distribution can be. Of the hypothesis set $\mathcal{H}$, there exists the Bayes hypothesis h_B and an arbitrary ‘random’ hypothesis h , and their respective measure.	24
3.1	Illustration of of sample complexity in terms of cardinality $m = \mathcal{S} $, with respect to ϵ or δ measure.	58
3.2	Illustration of of generalizational error/risk in terms of cardinality $m = \mathcal{S} $, with respect to changing $\mathcal{H}$ cardinality and δ -value measure. Blurred line indicates different cases or values, the deviation is rather small, thus the bound is very tight.	59
3.3	Prediction made by certain configurations of Hypothesis 3.2.1. Here, the injected PAC-type decay term provides the global envelope, while the interpolation offset term introduces singular behaviour near capacity-sample equality. The resulting superposition reproduces a family of double-descent-like curves depending on the trajectory in $(m, \mathcal{H})$ -space.	62
3.4	Decomposition of the error term into 3 parts. Respectively, the irreducible error $\mathbb{V}[y(\mathbf{x})]$, the variance (blue) and the bias (dark yellow). All of them are assumed to take up 100% of the error observed in such composition. The proportion is included using the three coefficient $\lambda_\epsilon, \lambda_V, \lambda_B$ where $\lambda_\epsilon + \lambda_V + \lambda_B = 1$	66
3.5	Curves for training risk (dashed line) and test risk (solid line). (a) The classical <i>U-shaped risk curve</i> arising from the bias-variance trade-off. (b) The <i>double descent risk curve</i> , which incorporates the U-shaped risk curve (i.e., the “classical” regime) together with the observed behaviour from using high capacity function classes (i.e., the “modern” interpolating regime), separated by the interpolation threshold. The predictors to the right of the interpolation threshold have zero training risk. Reproduced from Belkin et al. [2019a]	69
3.6	Left: Train and test error as a function of model size, for ResNet18s of varying width on CIFAR-10 with 15% label noise. Right: Test error, shown for varying train epochs. All models trained using Adam for 4K epochs. The largest model (width 64) corresponds to standard ResNet18. Resued from Nakkiran et al. [2019]	71
3.7	Left: Test error as a function of model size and train epochs. The horizontal line corresponds to model-wise double descent-varying model size while training for as long as possible. The vertical line corresponds to epoch-wise double descent, with test error undergoing double-descent as train time increases. Right Train error of the corresponding models. All models are Resnet18s trained on CIFAR-10 with 15% label noise, data-augmentation, and Adam for up to 4K epochs. Reused from Nakkiran et al. [2019]	71

3.8	As the model complexity N/m (for N parameters, m distinct visited states) increases, the MSBE first shows a U-shaped curve, peaking around the interpolation threshold ($N = m$). Double descent refers to the phenomenon for $N/m > 1$ where the MSBE drops once again. Continuous lines (red) indicate the theoretical values, the crosses (blue) are numerical results averaged over 30 instances with their standard deviations after the learning with regularized LSTD on Taxi-v3 for $\gamma = 0.95, \lambda = 10^{-9}, n = 5000, m = 310$. Taken from Brellmann et al. [2024]	79
4.1	The different variations of the noise profile that can be often seen when sampling from a simple regression model. By order, (1) represents mean-scale additive noise, (2) considers heteroskedastic multiplicative noise, (3) considers variance-scaled additive noise, (4) considers absolute additive noise (where scale independent of y), (5) considers heavy-tailed noise (Student- t modelling), and (6) is the standard canonical Gaussian noise (basically the default). All are under reduced dimension of the simplest linear model on feature space as possible for illustration.	91
4.2	The different variations of the noise profile that can be often seen when sampling from a simple regression model, but for different controlling nob for strength of the noise, by σ in range of $[0.1, 0.2, 0.3, 0.4, 0.5, 0.6]$. Again, by order, (1) represents mean-scale additive noise, (2) considers heteroskedastic multiplicative noise, (3) considers variance-scaled additive noise, (4) considers absolute additive noise (where scale independent of y), (5) considers heavy-tailed noise (Student- t modelling), and (6) is the standard canonical Gaussian noise (basically the default). All are under reduced dimension of the simplest linear model on feature space as possible for illustration.	92
4.3	Noise uncertainty creation, aggregated through the metric $\mathcal{U}(\mathcal{S}'_n)$ of the diffused dataset. Here, we can see that for y -only corrupted noise, the uncertainty increased on bounding for $2\sigma^2$, of which diverges several noise profiles from one another. An interesting outlier is in the suppressing effect of mean-scale additive noise, which is not as expected.	93
4.4	The different variations of the noise profile that can be often seen when sampling from a polynomial model, particularly constrained, 5th degree p_5 univariate polynomial on $[-1, 1]$ domain inverse range. By order, (1) represents mean-scale additive noise, (2) considers heteroskedastic multiplicative noise, (3) considers variance-scaled additive noise, (4) considers absolute additive noise (where scale independent of y), (5) considers heavy-tailed noise (Student- t modelling), and (6) is the standard canonical Gaussian noise (basically the default). We can see that each distributive pattern produces various different, often irreducible to Gaussian noise configurations.	94
4.5	Theorem 4.3.1 behaviours on randomized setting. For this model, we have $d = 100, \sigma = 0.5$, and variational n . Full test cases are for $p = [1, 100], n = \{12, 24, 36, 50, 76, 82, 90, 100, 106, 112, 121\}$ accordingly. The test function of the concept itself is the same linear model, but with different input only.	100
4.6	Contrary to Theorem 4.3.1, this is the behaviours on randomized setting for standard parameter-wised mean square error measure (MSE-on-parameter). For this model, we have $d = 100, \sigma = 0.5$, and variational n . Full test cases are for $p = [1, 100], n = \{12, 24, 36, 50, 76, 82, 90, 100, 106, 112, 121\}$ accordingly. The test function of the concept itself is the same linear model, but with different input only.	101

4.7	Contrary to Theorem 4.3.1 and both figure 4.6, this is the behaviours on randomized setting for standard parameter-wised mean square error measure (MSE-on-prediction). Especially, we do not see double descent pattern in this setting, as it is. For this model, we have $d = 100$, $\sigma = 0.5$, and variational n . Full test cases are for $p = [1, 100]$, $n = \{12, 24, 36, 50, 76, 82, 90, 100, 106, 112, 121\}$ accordingly. The test function of the concept itself is the same linear model, but with different input only.	102
4.8	Similar experiment according to setting and result of Theorem 4.24, where p is in the range $[1, 5523]$. Here, we can see small double descent modelled exactly at the interpolation threshold, while the later region exhibit similar phenomenon as bias-variance tradeoff.	102
4.9	Extension general experimentation, using the prediction made of theorem 4.3.1 for verification. Here, $q = 1$ as the general basis case of verification.	105
4.10	Illustration of polynomial model in definition 4.4.2.	107
4.11	Illustrative decomposition of the graph on its simplexes edges and vertices: Each of the decomposed space $\mathcal{V}$ and $\mathcal{E}$ can be encoded separately to their own ordeal, for example, of the edge connection through incident matrix.	111
4.12	A conceptual illustration on the running flow of an n -layer GNN on particular structure of interest. Note that the data section itself has particular embedding structure on its own. . . .	113
4.13	Encoding space and the change of encoding inbetween a static graph (or snapshot-wise) GNN. The encoding space is warped toward arbitrary notions inside a GNN, to the point that after certain layers of processing, the encoding outputs different from expected encoding space (native to the model, not to the designer), though there might be universal notions preserved, like the degree of node represented in a different way.	114
5.1	Double descent risk curves on MNIST for different model classes. Comparison of the phenomenon in (left) a fully connected neural network and (right) a random forest as model complexity increases. We use results from Belkin et al. [2019a] , however, extension is rather natural.	124
6.1	Illustrative dynamic of the learning problem. For an incremental hypothesis space sequence $\mathcal{H}_1, \mathcal{H}_2, \mathcal{H}_3$, we aim to obtain the procedure $\mathcal{T}$ that would either reach the <i>generalization solution</i> h , or the <i>empirical solution</i> h^* for a particular hypothesis class, with respect to the concept c and its observed concept c' . Model selection hence dictates how an algorithm or procedure than choose the best possible hypothesis to approximate the generalization solution from the empirical solution set. Do notice that here we explicitly state that the data would create a proxy concept that overlaps with the concept class and of some arbitrary ‘distance’ from the true concept by $d(c, c')$. In case the two classes overlap, there still exists the arbitrary distance. Also, we can also observe the intuitive notion of increasing hypothesis class - while it indeed can help in getting closer to the concept class, the somewhat intrinsic property of current learning theory being the relative random procedure class make it probabilistically unstable. . .	141
6.2	Minimal neuron structure	154
6.3	Table of monotonicity for the saddle point landscape between $[a, b]$, for both cases of the landscape. . .	161
6.4	Intuition of Hypothesis 6.6.1. Effectively, in the overparameterized range ($p > n$), the problem turns into soft interpolation, where the polynomial is possible to be dense enough such that relative error remains internally low. This can also happen if the dataset itself is intrinsically large.	167

7.1	Diagrammatic illustration of the decomposed neural architecture. In case of non-sparsity, all the graph would have to be fully-connected, and the order of the network is m . The canonical construction on $l = 1$ imposes final compression to (T_F, M_F, P_F) , which is then the admissible form used for evaluation of prediction on based input setting.	171
8.1	NESP Mechanism of Double Descent. At the interpolation threshold (Sharp Minima), high curvature induces strong anisotropic SGD noise (red wavy arrows), driving the system to escape. The system then relaxes to flat minima where low noise ensures stability. This process explains the second descent in test error.	183
1	Risk curves for varying sample size n (dimension $d = 100$, noise $\sigma = 0.5$). Double descent is determined in similar sequence, even though the interpolation theoretical $p = n$ is actually shifted over to the right from the actual interpolation observables, in parameterization. Afterward, correlation fails, and the theoretical p is ineffective, similar to previous theorem-based test run. Sample set size range is $n \in \{12, 24, 36, 50, 76, 82, 90, 100, 106, 112, 121, 150, 231, 250, 256, 280\}$. One additional remark is that the error landscape is relatively thin in either side, making it abnormal in comparison to actual bias-variance curvature.	193
2	(a–e) The different bias-variance decompositions. (f–j) Corresponding theoretical predictions for $\gamma = 0$, $\phi = 1/16$ and $\sigma = \tanh$ with SNR = 100 as the model capacity varies across the interpolation threshold (dashed red). (a,f) The semi-classical decomposition of Hastie et al. [2019] , Mei and Montanari [2019a] has a nonmonotonic and divergent bias term, conflicting with standard definitions of the bias. (b,g) The decomposition of Neal et al. [2018] utilizing the law of total variance interprets the diverging term V_D^c as “variance due to optimization”. (c,h) An alternative application of the law of total variance suggests the opposite, <i>i.e.</i> the diverging term V_P^c comes from “variance due to sampling”. (d,i) A bivariate symmetric decomposition of the variance resolves this ambiguity and shows that the diverging term is actually V_{PD} , <i>i.e.</i> “the variance explained by the parameters and data together beyond what they explain individually.” (e,j) A trivariate symmetric decomposition reveals that the divergence comes from two terms, V_{PX} and $V_{PX\epsilon}$ (outlined in dashed red), and shows that label noise exacerbates but does not cause double descent. Since $V_\epsilon = V_{P\epsilon} = 0$, they are not shown in (j). Taken from Adlam and Pennington [2020]	195
3	An illustration of the (supervised) statistical process. Phase III contains two parts: First is the evaluation $\nabla(h, c)$ according to the data $\mathcal{D}$, and second is the Update process to re-align c to the actual target.	197

Introduction

This thesis-like document, outlines specific contributions, theoretical developments, and experimental investigations on the **theoretical encoding theory**, logic of operations, **statistical learning theory** and **general learning theory**, within the standard outcome also concluded of statistical physics emulations, and explicit specialized movements. Overall, the document posits itself in the *synthesis package*, in which old ideas are reframed, old formalism being look from different angles of perspective, clarification of concepts, restructuralization of different frameworks and their axioms, presumptions, conjectures, hidden dilemma, and structural failure modes, and proposal of newer insights, outlooks, and theories. In specialities, we focus on the *reformation of the theoretical machine-based functions*, of which applies to current practical reduction to computer-dependent choice of construction, *presumption bias* in changing encoding system and the task of encode or decode, hidden assumptions or structural admissions, and ontological framing of different problem sets. On statistical learning, we focus on the proxy of the problem of *double descent* to highlight theoretical-empirical mismatch, theoretical holes and potential failure modes apparent of the old learning theory, reformulation and synthesis of existing frameworks and usages, and novel contribution to a new theory itself, with experimental toy models for investigations. Additionally, we also emphasize the practical notions and selective simulation or emulations encoding on machine systems, plus the general notion of the learning concepts. In conjunction, collaborations and tangent line of thoughts also consider the statistical learning interpretation, of such learning system.

The core authoring contribution comes from **Bui Gia Khanh**, from Department of Physics, Hanoi University of Science, Vietnam National University, with experimental results in majority coming from **Luong Van Tam**, **Dang Tri Trung** and **Daud Shahbaz**. Specifically, Luong Van Tam also participated as the main contributor in the chapter about statistical physical interpretation of the learning dynamics. We would like to thank any particular contributors participating in this document and research progress itself. Until further funding is acquired, we would like to clarify of open-source and self-funded research approach on this document.

Background and motivation

The new advancements of Artificial Intelligence (AI), and thus techniques in conjunction to it, for example, the **learning theory**, or the technique of knowledge encoding, the latent space specifications, modelling constructions, agent-theoretic formalism, **actionable system feedback** techniques, and so on. Within such, resurfaced the utilization of (super)computers and so on, within the baseline of computational theory of construction, and subsequently, computational learning theory / statistical learning theory (CoLT/SLT). Such system relies on the algorithmic approaches and procedural, computable representation of the automata systems to create models with capacities capable of being called as AI, and in a narrower sense, capable of adjusting or self-regulating response upon which it is provided of observa-

tions and situation-aware information. Such is where the word *machine learning* comes, and situated itself as the main anecdotal reason or capability of AI, for an artificial intelligence system to be intelligent, is posited to have the learning process. Of a long history in development, ever since the few first systems and classical notions of AI, McCulloch and Pitts [1943], Rosenblatt [1958a], Minsky and Papert [1969], Hebb [1949], Robbins and Monro [1951], Widrow and Hoff [1960b], Hestenes and Stiefel [1952] of the old paradigm in symbolic and neuronal approaches of discrete formulaic learning, to the looser *statistical, data-based learning*, particularly as of Valiant [1984], Vapnik [1999], Kearns and Vazirani [1994], Hopfield [1982], Fukushima [1980], Cybenko [1989], Bishop [1995], Haykin [1994], Kohonen [1982], Geman et al. [1992], Nemirovski and Yudin [1983], Rattray et al. [1998], Mohri et al. [2012], and up to modern development of such materials, such as, Jacot et al. [2018], Khan and Rue [2024], Buschjäger et al. [2020], Manin and Marcolli [2024], Devlin et al. [2018], Goodfellow et al. [2016], Kaplan et al. [2020], Bronstein et al. [2017, 2021], Luo and Tseng [1992], Tseng [2001], Friedman et al. [2007], Hu et al. [2021b], Wright [2015], Janik and Witaszczyk [2021], and so on, of the topics about the modern consideration of learning theories, optimization techniques and interpretation landscape, formal languages of the learning processes, as well as architectures on which modern systems are based on; more advanced techniques are developed, for example, as *reinforcement learning*, as seen of Sutton and Barto [2018], Bertsekas and Tsitsiklis [1996], Bertsekas [2017], Szepesvári [2010], Lattimore and Szepesvári [2020], Kaelbling et al. [1996], Watkins and Dayan [1992], Bellman [1957], Williams [1992], Konda and Tsitsiklis [2000], Mnih et al. [2015, 2016], Silver et al. [2016], Schulman et al. [2015, 2017], Lillicrap et al. [2015], Haarnoja et al. [2018], Arulkumaran et al. [2017], Ernst et al. [2005], Fu et al. [2020], Kostrikov et al. [2021], or *online learning* of Shalev-Shwartz [2012], Cesa-Bianchi and Lugosi [2006], Zinkevich [2003], Hazan [2016] and Littlestone and Warmuth [1994] as authoritative examples. Even further, is toward Large Language Model, or LLM (Vaswani et al. [2017a], Devlin et al. [2019], Brown et al. [2020a], Zhao et al. [2023, 2024], Radford et al. [2019, 2018], Raffel et al. [2020], Touvron et al. [2023a,b], Chowdhery et al. [2022], Ouyang et al. [2022], Wei et al. [2022c], Kaplan et al. [2020], Hoffmann et al. [2022], Bai et al. [2022]), where cutting edge efforts and so on are to build agentic systems capable of replacing certain labours kind of human in real systems, and so on improving quality of life. The scope is truly large, the field of applications and its limit overwhelmingly extended, and a much richer landscape was formed in those fields.

However, even with such advancements, certain pushback started to emerge. Those are *structural push-backs* against previously held assumptions, foundational problems regarding concepts taken for granted of hand-waived, or simply is not considered, structural behaviours out of scope of the dynamical trajectory that classical theory tracks (i.e. a theoretical sensory problem), and inadequacy of theory against empiricism, practiced in the modern AI landscapes. Such mismatch highlights substantial loopholes and opaque understanding and patchworks workaround that was present within the field, in which making a large portion of important problems simply too hard to solve analytically or conclusively, as for example, grokking, seen of Power et al. [2022a], or *double descent*, as of Belkin et al. [2019a], Nakkiran et al. [2019], in which exposes the loose nature of the definition on model selections and its landscape, model complexity, and evaluation on model theoretic approaches. The push for information theoretic learning, statistical interpretation, discourses among different interpretation of the learning setting and so is also for artificial intelligence in general, plus the concerns on explainability of artificial intelligence, of which the modern paradigm poorly offer such, as seen of such pushes on Doshi-Velez and Kim [2017], Lipton [2018], Ribeiro et al. [2016], Lundberg and Lee [2017], Molnar [2020]. Foundational problems still retain its strength unquestioned or unresolved, such as Dreyfus [1965, 1972], Dreyfus and Dreyfus [1986], Suchman [1987], Brooks [1991], Searle [1980], McCarthy and Hayes [1969], Harnad [1990]. On the more modern side, of contemporary critiques, Pearl [2009], Marcus [2018], Sutton [2019], and the modernized Chinese Room Arguments. The *frame problem*, while problematic and is often ignored by modern prac-

tioners, often resurfaced when discussions about explainability and operability are in range, as well as the epistemological notions of artificial operatives, or proofs that such artificial constructs truly ‘understand’, or ‘learn’, and its particular computational nature. More than all, is the apparent *theoretical entropy* observed from the landscape of development itself. While empiricism pushes forward, such encouraged particular fields, or rather somewhat all of them, to introduce different analogous interpretation, framing, concepts, or treatments in and of their own to the mechanical details of artificial intelligence. While not solving the inherent foundational problems or assumptive framework that it was based on, various different theories are used to tackle, in a heuristically employed manner, to patch different deviations or unexplainable phenomena together.

The goal is then to analyse such stack of complexity, to at the foremost, provide a more structural and insightful analysis into the topic of double descent or related phenomena, and the general issue coming around with it. While it might be overstretching, such problem can be considered the ultraviolet catastrophe of learning theory, especially within high-order construct or contrastive claim classical theory seems to hold against newer form of observation. By Kuhnian cycle of theoretical evolution, this signals a shift in theoretical insufficiency, though not catastrophic enough to trigger a theatrical meltdown and total replacement of theory, but rather serious consideration, and dialectic analysis of every moving components in the working machine. Right now, the field is hinging on operational definitions and heuristics, for example that would appear much in the following up chapters, are *model complexity*, and functional semantic of descriptions. Such is to say, a cross-field recognition of **Rice’s theorem**: semantic of operations can not be extracted from a certain given substrate. On more of the technicality, even the definition of *parameterization* is often fiercely debated, such is to observe there exist many modes of definition, in which heterogeneity by variance of structure is considered a feature, but not a categorical danger. Such is one of the stance we will make explicit, which can be right or wrong with the goal to be falsifiable, underlying the report’s insights.

We hope to provide a provisional view onto applying the theory of statistical physics, as seen of widespread usage in the history of development of artificial intelligence and computational method of such, as seen from Hopfield [1982], Amit et al. [1985], Gardner and Derrida [1988], Seung et al. [1992], Mezard et al. [1987], Mezard and Montanari [2009], Engel and Van den Broeck [2001], Advani et al. [2013], Bahri et al. [2020], Mehta and Schwab [2014], Advani and Saxe [2017a], Li and Sompolinsky [2021], Ferraro et al. [2014], with mentions of such statistical mechanics application in Roberts et al. [2021]. Some of those statistical learning schemes resulted in the works of G. E. Hinton on the variation of neural network called **Boltzmann machine**, as seen of Ackley et al. [1985a] and Hinton [2012]. Nevertheless, such question remains if the statistical scheme overreached of its current applicable form, to indict misleading or only limited improvement and interpretability onto a machine learning and in general an artificial intelligence system dynamics. In addition, assumptions and structural constraints, such as those incurred on Boltzmann machines for it to facilitate an environment interpretable of statistical physics remains a

question of interest¹. Even with such problems, we would like to retain the view that statistical physics can be indeed used to interpret double descent, though with such phenomena, requires extreme cautions of logical assumptions and modelling setting. However, running the risk of the field being theoretically fragmented, with each framework claiming explanatory primacy, there is no easy way to fix, but to for now explicitly consider such possibility a reality enforced by observations.

In all, while not conclusive, the research ultimately aims itself to clarify the epistemological and foundational concepts used of the classical theory. Then, from doing such, we provide a systematic experimental and theoretical clarification of the landscape in which double descent is in, and of the wider theme of classical learning theory friction at large. Doing so, we note that the results are in posture of non-closure, such is to say our works are not *conclusive* or definitive in any restricted form. Hence, sections of the materials are susceptible to reformation, and thus deliberately direct structural and experimental falsification.

¹It is widely concern of the fact that there exists many interpretations of machine learning mechanics under different formal language, or application-biased one. Such is, for example, information theoretic, as seen of [Russo and Zou \[2016\]](#), [Xu and Raginsky \[2017a\]](#), categorical theoretic approaches, of [Manin and Marcolli \[2024\]](#), and accordingly, control theory, topological learning, graph theoretical, and statistical physics. While interesting as to illuminate the same problem under different lens, it is troubling of the messiness in such approach, as different theories claim different hypothesis, principles, theorems and axioms. Furthermore, it is of questionable intent to simply forgive the account of applying one's field on another — there exists the fundamental *domain bias* — of which as can be seen in statistical physics, for certain event to be applicable of using statistical law, one either must adapt the notion and introduce analogous notion (like temperature and relative model diffusing entropy), or restricting the domain, structure, range of expression, and mechanics — similar to how Boltzmann machines are. Additionally, applying without care gives increasingly large amount of domain bias, as concepts typically seen and assumed of a system in one theory might interject wrongfully, or simply conflict, narrow down the applicable range of such application, or in the worst case, invalidate the result.

Chapter 1. Preliminaries

There are several preliminaries to this particular document, many of which are basic or rather facts regarding different notations, different theoretical knowledges, and so on, covering from mathematics of algebra to functional analysis, mathematical simulation and mathematical modelling, optimization theory and convergence bounds, concentration inequalities, coursework about computational theories, and so on. This will be condensed and precise, and does not aim to replace traditional knowledge or textbooks on such matter. For more information, please refer to authoritative sources as ideal.

1.1 Concentration inequalities

A lot of results in statistical learning is proved using bounds. For example, if $X_1, \dots, X_n$ are independent Bernoulli (θ) random variables representing the outcomes of a sequence of n tosses of a coin with bias (probability of head), then for any $\epsilon \in (0, 1)$:

$$\mathbb{P}\left(|\hat{\theta}_n - \theta| \geq \epsilon\right) \leq 2e^{-2\epsilon^2}, \quad \text{for } \hat{\theta}_n = \frac{1}{n} \sum_{i=1}^n X_i \quad (1.1)$$

For the fraction of heads in $X^n = (X_1, \dots, X_n)$. Since $\theta = \mathbb{E}\hat{\theta}_n$, 1.1 says that the sample (or empirical) average of X_i concentrates sharply around the statistical average. Bound like these are fundamental in statistical learning theory (well, mostly for proof, not going to lie). For now, let us learn the technique used to derive such bounds for settings much more complicated than coin tossing.

1.1.1 What are concentration inequalities?

In a probabilistic setting, we are sometimes interested of the random fluctuations of functions of independent random variables. *Concentration inequalities* quantify such statements, typically by bounding the probability that such a function differs from its expected value (or from its median) by more than a certain amount.

The search for concentration inequalities has been a topic of intensive research in the last decades in a variety of areas because of their importance in numerous applications. Typically, this falls into the range of designing non-deterministic (by the definition of such words), dynamic randomized algorithms or procedure of interest, while able to bound them with concentration inequalities to a certain given degree of high probabilistic ‘accuracy’, for example. Among the areas of applications, without trying to be exhaustive, we mention statistics, learning theory, discrete mathematics, statistical mechanics, random matrix theory, information theory, and high-dimensional geometry (while not actually sure why geometry is in here).

Informally, the *concentration phenomenon* asserts that if $X_1, \dots, X_n$ are independent random variables,

then $f(X_1, \dots, X_n)$ does not deviate much from its mean provided that $f(x_1, \dots, x_n)$ is not too sensitive in any of its coordinate x_i . While concentration properties for sums of independent random variables were thoroughly studied and fairly well understood in classical probability theory, powerful tools to handle more general functions of independent random variables were not introduced until the appearance of *martingale methods* in the 1970s, by Yurinskii (1976), Maurey (1979), Milman and Schechtman (1986), Shamir and Spencer (1987) and McDiarmid (1989).

Note 1.1.1. While the topic itself is much more complex than it is guaranteed to be, a *martingale* can be defined as a sequence of random variables (i.e. for a stochastic process) for which, at a particular time, the conditional expectation of the next value in the sequence is equal to the present value, regardless of all prior values. A basic definition can be achieved. For a process $(M_n)_{n \geq 0}$, it is called martingale if:

- For every $n \geq 0$, the expectation $\mathbb{E}[M_n]$ is finite, and hence equivalently, $\mathbb{E}[|M_n|] < \infty$.
- For every $n \geq 0$ and all $m_n, m_{n-1}, \dots, m_0$ we have:

$$\mathbb{E}(M_{n+1} \mid M_n = m_n, \dots, M_0 = m_0) = m_n \quad (1.2)$$

1.1.2 Relevant toolset

Of our book chapter, by a *concentration inequality* we most often mean an upper bound for the probability that a real-valued random variable Z differs from its expected value by more than a certain amount. In other word, we seek upper bounds for tail probabilities of the form

$$\mathbb{P}[Z - \mathbb{E}(Z) \geq t], \quad \mathbb{P}[Z - \mathbb{E}(Z) \leq -t] \quad (1.3)$$

where $t > 0$. Of course, here we assume implicitly that the expected value exists, as for the probability to be included of the measurable space.

To derive a lot of our concentration bounds that will solve such problem, however, we have to develop a few basic tools in the analysis. This will include Markov's inequality, Chebyshev's inequality, the Chernoff bound (often called the Chernoff trick), and Hoeffding's lemma.

1.1.3 Markov's inequality

Markov's
inequality

Let $Y \in \mathbb{R}$ be a nonnegative random variable. Then for any $t > 0$, we have the *Markov's inequality* :

$$\mathbb{P}(Y \geq t) \leq \frac{\mathbb{E}[Y]}{t} \quad (1.4)$$

Before jumping straight into a proof, let us see what are we looking at from the perspective of the inequality.

Proof. The proof is straightforward:

$$\begin{aligned} \mathbb{P}(Y \geq t) &= \mathbb{E}[\mathbb{1}_{(Y \geq t)}] \\ &\leq \frac{\mathbb{E}[Y \mathbb{1}_{(Y \geq t)}]}{t} \\ &\leq \frac{\mathbb{E}[Y]}{t} \end{aligned} \quad (1.5)$$

We use the fact that

$$\mathbb{P}(Y \in A) = \int_A P_Y(dy) = \int_Y \mathbb{1}_{(x \in A)} P_Y(dy) = \mathbb{E}[\mathbb{1}_{Y \in A}] \quad (1.6)$$

Additionally, we also have $Y \geq t > 0$ implies $Y/t \geq 1$, and $Y \geq 0 \implies Y \mathbb{1}_{(Y \geq t)} \leq Y$ which implies their expected value. $\square$

If ϕ denotes a nondecreasing and nonnegative function defined on a possibly infinite interval $I \subset \mathbb{R}$, and if Y denotes a random variable taking values in I , then Markov's inequality implies that for every $t \in I$ with $\phi(t) \neq 0$,

$$\mathbb{P}(Y \geq t) \leq \mathbb{P}[\phi(Y) \geq \phi(t)] \leq \frac{\mathbb{E}[\phi(Y)]}{\phi(t)} \quad (1.7)$$

This can be used to derive the Chebyshev's inequality later.

1.1.4 Chebyshev's inequality

The common application of Markov's probability is Chebyshev's probability. Let X be an arbitrary real random variable. Then for any $t > 0$,

$$\mathbb{P}(|X - \mathbb{E}[X]| \geq t) \leq \frac{\text{Var}[X]}{t^2} \quad (1.8)$$

Where $\text{Var}[X] := \mathbb{E}[|X - \mathbb{E}[X]|^2] = \mathbb{E}[X^2] - (\mathbb{E}[X])^2$ is the variance of X . The prominent role of Chebyshev's inequality is not only explained by historical reasons, but also among all the absolute moments of the form $\mathbb{E}[|Z - \mathbb{E}(Z)|^q]$, the variance is typically the easiest to handle.

Proof. Let X be a random variable, so $(X - \mathbb{E}[X])^2$ is a non-negative random variable. Hence, we can apply Markov's inequality:

$$\mathbb{P}(X - \mathbb{E}[X] \geq t) = \mathbb{P}((X - \mathbb{E}[X])^2 \geq t^2) \quad (1.9)$$

$$\leq \frac{1}{t^2} \mathbb{E}[(X - \mathbb{E}[X])^2] \quad (1.10)$$

$$= \frac{1}{t^2} \text{Var}[X] \quad (1.11)$$

$\square$

The property of Chebyshev's inequality contains several constraints, though arguably better than Markov's. Chebyshev's inequality does not require that the random variable is non-negative. However, it also requires that we know the variance in addition to the mean. Hence, the goal of Chebyshev's inequality is to bound the probability that the RV is far from its mean in either direction. While in principle Chebyshev's inequality asks about distance from the mean in either direction, it can still be used to give a bound on how often a random variable can take large values, and will usually give much better bound than Markov's inequality - since we have more knowledge of the distribution.

We can use Chebyshev's inequality to prove a soft variation of the law of large number, or the *weak law of large number*.

Theorem 1.1.1 (Weak Law of Large Number). *Let $X_1, \dots, X_n$ be a sequence of i.i.d. random variables with mean μ . Let $\bar{X}_n = (1/n) \sum_{i=1}^n X_i$ be the sample mean. Then $\bar{X}_n$ converges in probability to μ . That is, for any $\epsilon > 0$,*

$$\lim_{n \rightarrow \infty} \mathbb{P}(|\bar{X}_n - \mu| > \epsilon) = 0 \quad (1.12)$$

Proof. By the property of the expectation and variance of sample mean consisting of n i.i.d. variables, for $\mathbb{E}[\bar{X}_n] = \mu$ and $\text{Var}(\bar{X}_n) = \sigma^2/n$. By Chebyshev's inequality,

$$\lim_{n \rightarrow \infty} \mathbb{P}(|\bar{X}_n - \mu| > \epsilon) \leq \lim_{n \rightarrow \infty} \frac{\text{Var}(\bar{X}_n)}{\epsilon^2} = \lim_{n \rightarrow \infty} \frac{\sigma}{n\epsilon^2} = 0 \quad (1.13)$$

thus recover our desired equation. $\square$

1.1.5 McDiarmid's inequality

A generalization of Hoeffding's inequality is the McDiarmid's, or Bounded Difference inequality, where the quantity of interest is some function of the data, i.e. $S_n = \phi(X_1, \dots, X_n)$. Some restrictions on ϕ are required to get exponential bounds.

Theorem 1.1.2 (McDiarmid's Inequality). *Let $X_1, \dots, X_n$ be independent random variables, where X_i has range $\mathcal{X}_i$. Let $f : \mathcal{X}_1 \times \dots \times \mathcal{X}_n \rightarrow \mathbb{R}$ be any function with the $(c_1, \dots, c_n)$ -bounded difference property: for every $i = 1, \dots, n$ and every $(x_1, \dots, x_n), (x'_1, \dots, x'_n) \in \mathcal{X}_1 \times \dots \times \mathcal{X}_n$ that differ only in the i -th coordinate, we have:*

$$|f(x_1, \dots, x_n) - f(x'_1, \dots, x'_n)| \leq c_i \quad (1.14)$$

For any $t > 0$:

$$\mathbb{P}\left([f(X_1, \dots, X_n)] - \mathbb{E}[f(X_1, \dots, X_n)] \geq t\right) \geq \exp\left(-\frac{2t^2}{\sum_{i=1}^n c_i^2}\right) \quad (1.15)$$

Proof. Write $\mathbb{E}_i[\cdot]$ to denote expectation conditioned on $X_{1:i} := (X_1, \dots, X_i)$. Therefore, $g_i(X_{1:i}) := \mathbb{E}_i[f(X_{1:n})]$ is, as the notation suggests, a function of $X_{1:i}$. Now define the following random variables:

$$\begin{aligned} Y_i &:= g_i(X_{1:i}) - g_{i-1}(X_{1:i-1}), \\ A_i &:= \inf_{x_i \in \mathcal{X}_i} g_i(X_{1:i-1}, x_i) - g_{i-1}(X_{1:i-1}), \\ B_i &:= \sup_{x_i \in \mathcal{X}_i} g_i(X_{1:i-1}, x_i) - g_{i-1}(X_{1:i-1}). \end{aligned}$$

These random variables satisfy

$$Y_i \in [A_i, B_i], \quad \mathbb{E}_{i-1}[Y_i] = 0,$$

$$\sum_{i=1}^n Y_i = f(X_{1:n}) - \mathbb{E}[f(X_{1:n})].$$

Furthermore, we claim that $[A_i, B_i]$ is always an interval of length at most c_i . We defer the proof of this claim to later.

To bound the probability that the sum of the Y_i 's is at least t , we use the Chernoff bounding method. Let $S_i := Y_1 + \dots + Y_i$. Then

$$\mathbb{E}[\exp(\lambda S_n)] = \mathbb{E}[\exp(\lambda(Y_n + S_{n-1}))] \quad (1.16)$$

$$= \mathbb{E}[\mathbb{E}_{n-1}[\exp(\lambda Y_n)] \exp(\lambda S_{n-1})] \quad (1.17)$$

$$\leq \exp(\lambda^2 c_n^2 / 8) \mathbb{E}[\exp(\lambda S_{n-1})] \quad (1.18)$$

$$= \exp(\lambda^2 c_n^2 / 8) \mathbb{E}[\exp(\lambda(Y_{n-1} + S_{n-2}))] \quad (1.19)$$

$$= \exp(\lambda^2 c_n^2 / 8) \mathbb{E}[\mathbb{E}_{n-2}[\exp(\lambda Y_{n-1})] \exp(\lambda S_{n-2})] \quad (1.20)$$

$$\leq \exp(\lambda^2 c_n^2 / 8) \exp(\lambda^2 c_{n-1}^2 / 8) \mathbb{E}[\exp(\lambda S_{n-2})] \quad (1.21)$$

$$\vdots \quad (1.22)$$

$$\leq \exp(\lambda^2 c_n^2 / 8) \cdots \exp(\lambda^2 c_1^2 / 8) \quad (1.23)$$

$$= \exp\left(\frac{\sum_{i=1}^n c_i^2}{4} \cdot \frac{\lambda^2}{2}\right). \quad (1.24)$$

Above, the inequalities all follow from Hoeffding's inequality. Therefore, S_n is $\sum_{i=1}^n c_i^2 / 4$ sub-Gaussian. The probability bound now follows from the Cramer-Chernoff inequality for subgaussian random variables.

It remains to prove that $[A_i, B_i]$ is an interval of length at most c_i . We show that $B_i - A_i \leq c_i$. By independence of $X_1, \dots, X_n$, we have

$$g_i(X_{1:i-1}, x_i) = \mathbb{E}[f(X_{1:i-1}, x_i, X_{i+1:n}) \mid X_{1:i-1}, X_i = x_i] \quad (1.25)$$

$$= \mathbb{E}[f(X_{1:i-1}, x_i, X'_{i+1:n}) \mid X_{1:i-1}], \quad (1.26)$$

where $X'_{i+1:n}$ is an independent copy of $X_{i+1:n}$. Then

$$B_i - A_i = \sup_{b \in \mathcal{X}_i} g_i(X_{1:i-1}, b) - \inf_{a \in \mathcal{X}_i} g_i(X_{1:i-1}, a) \quad (1.27)$$

$$= \sup_{a, b \in \mathcal{X}_i} g_i(X_{1:i-1}, b) - g_i(X_{1:i-1}, a) \quad (1.28)$$

$$= \sup_{a, b \in \mathcal{X}_i} \mathbb{E}[f(X_{1:i-1}, b, X'_{i+1:n}) - f(X_{1:i-1}, a, X'_{i+1:n}) \mid X_{1:i-1}] \quad (1.29)$$

$$\leq \mathbb{E}\left[\sup_{a, b \in \mathcal{X}_i} |f(X_{1:i-1}, b, X'_{i+1:n}) - f(X_{1:i-1}, a, X'_{i+1:n})| \mid X_{1:i-1}\right] \quad (1.30)$$

$$\leq c_i. \quad (1.31)$$

The last step above uses the bounded differences property of f . $\square$

1.1.6 The Cramer-Chernoff method

One more important bounding method that can be discussed in this section is the Cramer-Chernoff bound. This method determines the best possible bound for a tail probability that one can possibly obtain by using Markov's inequality with an exponential function $\phi(t) = \exp(\lambda t)$. This simple technique leads to surprisingly sharp bounds in many cases. We work out some simple examples, and then will proceed with some of its analytical properties.

Let Z be a real-valued random variable. For $\lambda \geq 0$, Markov's inequality implies

$$\mathbb{P}[Z \geq t] \leq \exp(-\lambda t) \mathbb{E}[\exp(\lambda Z)] \quad (1.32)$$

Since this inequality holds for all values of $\lambda \geq 0$, one may choose λ to minimize the upper bound. Defining the logarithm of the moment generating function as

$$\phi_Z(\lambda) = \log(\mathbb{E}[\exp(\lambda Z)]), \quad \forall \lambda \geq 0 \quad (1.33)$$

and introducing the supremum of such generating function being

$$\phi_Z^*(t) = \sup_{\lambda \geq 0} (\lambda t - \phi_Z(\lambda)) \quad (1.34)$$

Chernoff's
inequality

We obtain the Chernoff's inequality:

$$\mathbb{P}[Z \geq t] \leq \exp \left[- \left(\sup_{\lambda \geq 0} (\lambda t - \phi_Z(\lambda)) \right) \right] = \exp [-\phi_Z^*(t)] \quad (1.35)$$

The function $\phi_Z^*(t)$ is called the *Cramer transform* of Z . Usually, this is also called the Legendre transform (for example, in Ostaszewska and Zajkowski [2014]) or Legendre-Fenchel transform. The Cramer transform is characterized of its shape by satisfying the *Cramer's condition*,

Cramer condition

Definition 1.1.1 (Cramer's condition). Let M_X denote the moment-generating function of a given random variable X , that is, $M_X(t) = \mathbb{E}[\exp(tX)]$. Then the random variable X satisfies the *Cramer condition* if there exists $c > 0$ such that $\mathbb{E}[\exp(|cX|)] < \infty$. If X satisfies the Cramer condition, then we call M_X to be *well-defined* (taking finite values) on a connected neighbourhood, containing the interval $[-c, c]$, of zero and moreover possess the expansion

$$M_X(t) = \sum_{n=0}^{\infty} \frac{\mathbb{E}[X^n]}{n!} t^n \quad (1.36)$$

where $|t| < t_0$ and $t_0 > c$.

The function itself is a nonnegative function, since $\phi_Z(0) = 0$. If $\mathbb{E}[Z]$ exists, then the convexity of the exponential function and Jensen's inequality imply that $\phi_Z(\lambda) \geq \lambda \mathbb{E}[Z]$ and therefore, for all negative values of λ , $\lambda t - \phi_Z(\lambda) \leq 0$ whenever $t \leq \mathbb{E}[Z]$. This means that we may formally extend the supremum over all $\lambda \in \mathbb{R}$ of the Cramer transform,

$$\phi_Z^*(t) = \sup_{\lambda \in \mathbb{R}} (\lambda t - \phi_Z(\lambda)) \quad (1.37)$$

from this, formally, we treat the right-hand expression of the often known, as what we call the *Fenchel-Legendre dual function* of ϕ_Z . Thus, at every $t \geq \mathbb{E}[Z]$, the Cramer transform coincides with the Fenchel-Legendre dual. That is not to say they are the same, however the range of λ is relevant.

1.2 Metric and pseudometric

One of the treatment, or rather insight that we would be looking at, consider the notion of a **metric**. Of those, we think of it as an additional structure upon two sets, to encode specific quantization to the notion of distinguishability, and distance metric. Out of those, there are the space of such admissible constructions, which includes *metric space* and *pseudometric spaces*.

Definition 1.2.1 (Pseudometric space). Suppose X is a set. A function $\rho : X \times X \rightarrow \mathbb{R}$ is said to be a **pseudometric**, if it satisfies:

- i $\rho(x, x) = 0, \forall x \in X,$
- ii $\rho(x, y) = \rho(y, x), \forall x, y \in X,$
- iii $\rho(x, z) \leq \rho(x, y) + \rho(y, z), \forall x, y, z \in X.$

The value of ρ is positive-bounded, such that we can write $\rho : X \times X \rightarrow \mathbb{R}_+$. Suppose ρ is a pseudometric on X , then we say that (X, ρ) is a **pseudometric space**.

We can see that a pseudometric would not impose characteristic on distinguishment upon two points, but only identifiability via $\rho(x, x)$. In fact, if two points are of zero-distance on the metric ρ , such cannot be determined if the two points are the same, but rather only a weaker form of $x \sim y$ as to require $\rho(x, y) = 0$. If, however, we impose that they are the same, then we gain a metric.

Definition 1.2.2 (Metric). Suppose X is a set. A function $\rho^* : X \times X \rightarrow \mathbb{R}$ is said to be a **metric** if it is a pseudometric, and satisfies $\rho^*(x, y) = 0 \Rightarrow x = y$. Then, we say (X, ρ^*) is a **metric space**.

Later on, we would be seeing the difference between the two, and one of the categorical error of choosing metric space condition on loss-functional space. Let us illustrate the construction for now in this section. Let $\mathcal{H}$ be the hypothesis class of arbitrary structure, $\mathcal{Z}$ the admitted observations space, $\ell : \mathcal{H} \times \mathcal{Z} \rightarrow \mathbb{R}_+$ as the loss functional or resultant loss application on $\mathcal{H} \times \mathcal{Z}$, and suppose there exists a quantification of probability distribution $\mathcal{P}$ on $\mathcal{Z}$ of which we suppose to capture generating patterns on the observation pack. Then, the error functional on admitted dataset most often is of the form similar to an empirical error measure,

$$\hat{R}(h) = \mathbb{E}_{z \sim \mathcal{P}}[\ell(h, z)] \quad (1.38)$$

and thus giving us, suppose that there exists the $\mathcal{A}$ -algorithm that produces $h_i, h_{i+1}, \dots$, the following distance metric:

$$d_\ell(h_i, h_{i+1}) := \mathbb{E}_{z \sim \mathcal{P}}[|\ell(h_i, z) - \ell(h_{i+1}, z)|] \quad (1.39)$$

This object satisfies *non-negativity*, *symmetry*, and *triangle inequality* as specified of pseudometric. However, often we do not have the virtue to say that if $d_\ell(h_i, h_{i+1}) = 0$ then $h_i = h_{i+1}$, aside from which we have total exposure of the system itself. Let us consider the setting. Suppose that for the picture $(z, \mathcal{P})$ of the perceived observations space characteristic, it has a true concept $c \in \mathcal{C}$ of arbitrary structure concept class $\mathcal{C}$. For any assumption on oracle access on c , we see that $R(c) \leq 0$ take as arbitrary ϵ , since observation can deviate from the concept resultant observation itself. Hence, we get that

$$\begin{aligned} d(h_i, c) &:= \mathbb{E}_{z \sim \mathcal{P}}[|\ell(h_i, z) - \ell(c, z)|] \leq \mathbb{E}_{z \sim \mathcal{P}}[\ell(h_i, z)] + \mathbb{E}_{z \sim \mathcal{P}}[\ell(c, z)] \\ &= \mathbb{E}_{z \sim \mathcal{P}}[\ell(h_i, z)] + \epsilon \end{aligned} \quad (1.40)$$

which means that even by this, the distance is never guaranteed as to be zero. And that is within oracle access setting itself. Usually, if this is not guaranteed, then we have no access on c , thus we also have no access to complete this distance metric calculation on itself, and thus $d(h, c)$ does not exist. Here is the conflation of the effect. If, we are to assume that we are comparing the notion of $x = y$ for any object of the functional space itself as to be the factorization, or parameterization that governs $z \rightarrow q$ of the input-output observable space on $\mathcal{Z}$, then the loss-functional gives a metric, as for $d(h, c) = 0$ means they phenomenologically approach the same focal point. However, if the target is structural, then $d(h, c) = 0$ does not mean $h = c$, even by assumption that $\mathcal{H} = \mathcal{C}$ i.e. concept space and the hypothesis

space has same scope. Putting that aside, because we never have access to c , the way to interpret, with presumption of c to exist, then $\ell(h, z) = \ell(c, z)$ only guarantees that we settle the model is an equivalence class $h \sim c$ only; this is strengthened if we further consider that $(\mathcal{Z}, \mathcal{P})$ imposes a structural proxy c^* that can deviate from c , and can also admit different structural interpretation outside of class $\mathcal{C}$. Thereby, interpretively, loss-functional space is not metric but pseudometric, when it comes to the target c measure in which the measure comparing to this object simply does not exist. To obtain equality, one typically use quotient space on it, which is a very strong structural addition and assumption, and might, or might not, be guaranteed, though in this direction of admitting quotient decision space is valid. If, however, one approaches the other way, by saying that different hypotheses must be metrically distinguishable, i.e. imposing injectivity on ℓ , full support on distribution $\mathcal{P}$, and identifiability condition, that the assumption set would be very strong and almost never work in practice.

1.3 Graph theory

We retain a section explicitly about graph theory, partially because it would be used as a provisional example in the material, specifically in experimental setting. Thereby, we would like to contain several trivializations of knowledge on such structure.

Definition 1.3.1 (Graph). A **simple graph** is a pair (V, E) , where V is a finite set, and where E is a subset of $\mathcal{P}_2(V)$, with $\mathcal{P}_2(V)$ is the set of all 2-element subsets of V .

A quick way to refers to certain elements of the graph is as followed.

Definition 1.3.2 (Notations). Let $G = (V, E)$ be a simple graph.

1. The set V is the vertex set of G , denoted $V(G)$. The element of V are called the **vertices** (or **nodes**) of G .
2. The set E is called the **edge set** of G , denoted $E(G)$. The element of E are called the **edges** of G . We denoted uv for $\{u, v\}$ as connections between u and v .

Certain properties of such simple graph case follows:

Definition 1.3.3 (Connectedness). If the two graphs are $G_1 = (V(G_1), E(G_1))$ and $G_2 = (V(G_2), E(G_2))$ where $V(G_1)$ and $V(G_2)$ are disjoint, then their **union** $G_1 \cup G_2$ is the graph with vertex set $V(G_1) \cup V(G_2)$ and edge family $E(G_1) \cup E(G_2)$.

Definition 1.3.4 (Adjacency). Two vertices u, v are said to be **adjacent** to each other if $uv \in E$. uv is then said to **join** u with v . The **neighbour** of a vertex $v \in V$ is then the set of all $u \in V$ such that $uv \in E$.

Definition 1.3.5 (Graph Isomorphism). Let G and H be two simple graphs. A **graph isomorphism** (or **isomorphism**) from $G \rightarrow H$ is the bijection $\phi : V(G) \rightarrow V(H)$ that *preserves edges*, i.e., for any two vertices u and v of G , we have

$$(uv \in E(G)) \iff (\phi(u)\phi(v) \in E(H)) \quad (1.41)$$

We say that G and H are isomorphic (denoted $G \cong H$) if there exists a graph isomorphism from G to H .

We also list several types of graphs, specifically, the **subgraph**, null /complete graph, path /cycle graph, Kneser graphs, digraphs and multigraph (or general graph).

Amane

What if, for a general graph, there are multiple non-equal edges and connections, such that each node function more as a small sub-graph, or either an intermediate I/O node?

Definition 1.3.6 (Subgraph). A subgraph of a graph X is a graph Y such that $V(Y) \subseteq V(X)$ and $E(Y) = E(X)$. If $V(Y) = V(X)$ then Y is a *spanning subgraph* of X . A subgraph Y of X is an *induced subgraph* if two vertices of $V(Y)$ are adjacent in Y if and only if they are adjacent in X .

One of the important definition for subgraph, especially for those of recommendation system, and graph construction network, is the notion of **induced subgraph**.

Definition 1.3.7 (Induced subgraph). Let $S \subset V$. The induced subgraph of G on the set S denotes the subgraph

$$(S, E \cap \mathcal{P}_2(S)) \quad (1.42)$$

of G . In other words, it denotes the subgraph of G whose vertices are the elements of S , and whose edges are precisely those edges of G whose both endpoints belong to S .

Definition 1.3.8 (Null graph). A graph whose edge-set $E = \emptyset$ is a **null graph**. we denote the null graph on n vertices by $N_n = (V, \emptyset)$.

Definition 1.3.9 (Complete graph). A simple graph $G = (V, E)$ in which each pair of distinct vertices are adjacent is a **complete graph**. We denote this by K_n . For graph K with $V(G) = n$, we have:

$$V(K_n) = \frac{n(n-1)}{2} \quad (1.43)$$

Definition 1.3.10 (Path graph). For each $n \in \mathbb{N}$, we define the n -th **path graph** P_n to be the simple graph

$$P_n = (\{1, 2, \dots, n\}, \{i(i+1) \mid 1 \leq i < n\}) \quad (1.44)$$

This graph has n vertices and $n - 1$ edges (unless $n = 0$)

Definition 1.3.11 (Cycle graphs). For each $n > 1$, we define the n -th **cycle graph** C_n to be the simple graph

$$C_n = (\{1, \dots, n\}, \{i(i+1) \mid 1 \leq i < n\} \cup \{1n\}) \quad (1.45)$$

This graph has n vertices and n edges, unless $n = 1$, then it has only 1 edge. For multigraph, however, the C_2 graph can have two edges.

1.3.1 Algebraic representation

The tool of computation is still linear algebra, for any large computational system under the neural network system. Hence, there are several useful way to represent such graph on linear algebraical notation and system.

Definition 1.3.12 (Adjacency matrix). Let G be a simple graph with $V(G) = \{1, \dots, n\}$ and $E(G) = \{e_1, \dots, e_m\}$. The adjacency matrix of G , denoted $A(G)$, is the $n \times n$ matrix defined as

$$A_{dj} = A \in \mathbb{R}^{n \times n} : \quad A_{ij} = \begin{cases} 1 & v_i v_j \in E, i \neq j \\ 0 & \text{otherwise} \end{cases} \quad (1.46)$$

The modification (somewhat familiar) of the adjacency matrix is by the notion of **incidence matrix**, which takes into account the case of directed graph.

Definition 1.3.13 (Incident matrix). Let G be a graph with $V(G) = \{1, \dots, n\}$ and $E(G) = \{e_1, \dots, e_m\}$. Suppose each edge $e \in E(G)$ is an orientation, which is arbitrary but fixed. The (vertex-edge) **incidence matrix** of G , denoted $Q(G)$, is the matrix of size $n \times m$, defined as followed:

$$Q(G) \in \mathbb{R}^{n \times m} : \quad Q_{ij} = \begin{cases} 1 & (i, e_j) \in G, e_j \text{ originates at } i \\ -1 & (i, e_j) \in G, e_j \text{ terminates at } i \\ 0 & \text{otherwise (both)} \end{cases} \quad (1.47)$$

For undirected graph, the matrix is defined as

$$Q(G) \in \mathbb{R}^{n \times m} : \quad Q_{ij} = \begin{cases} 1 & \exists k, e_j = \{v_i, v_k\} \\ 0 & \text{otherwise} \end{cases} \quad (1.48)$$

For multidirected graph (or multigraph), we use the **Laplacian matrix** to represent such graph in matrix form. Notice that of the coming definition, the Laplacian matrix does not encode incidences in its matrix, as it is not the main focus of the definition.

Definition 1.3.14 (Laplacian matrix). Let G be a graph with $V(G) = \{1, \dots, n\}$ and $E(G) = \{e_1, \dots, e_m\}$. The **Laplacian matrix** of G , denoted $L(G)$, is the $n \times n$ matrix defined as followed:

$$L(G) \in \mathbb{R}^{n \times n} : L_{ij} = \begin{cases} 0 & i \neq j \\ \deg(i) & i = j \end{cases} \quad (1.49)$$

The Laplacian matrix and the incidences matrix can be related, with the fixed orientation assumption, by

$$L(G) = Q(G)Q'(G) \quad (1.50)$$

where $Q'(G)$ is the transpose of Q . This identity suggests that the Laplacian might depends on the orientation, although it is evident from the definition that the Laplacian is independent of orientation.

We present more subsidiary theory on graph theory.

Theorem 1.3.1 (Handshaking Lemma). *In any graph G , the sum of all the vertex-degree is an even number. Or, the vertex degree add up to *twice* the number of edges:*

$$\sum_{v \in V(G)} \deg_G(v) = 2|E(G)|$$

Proof by induction. The base case here is $m = 0$. The graph with no edge then $\deg(v) = 0$.

Assume that the degree sum formula holds for all $(m - 1)$ -edge graphs. Let G be a graph with $m \geq 1$ edges, and let xy be any edge of G . We can apply the inductive hypothesis to $G - xy$ (the graph we get from deleting xy from G), a graph with $m - 1$ edges.

Both x and y have one extra incident edge in G they don't have in $G - xy$ - the edge xy itself. Hence

$$\deg_G(x) = 1 + \deg_{G-xy}(x)$$

and

$$\deg_G(y) = 1 + \deg_{G-xy}(y)$$

For any other vertex v , $G - xy$ and G have the same number of edges, so we have

$$\deg_G(v) = \deg_{G-xy}(v)$$

Furthermore, $G - xy$ and G have the same set of vertices. So if we add up the vertex degrees in $G - xy$ and G , the result is that

$$\sum_{v \in V(G)} \deg_G(v) = 2 + \sum_{v \in V(G-xy)} \deg_{G-xy}(v)$$

Applying the inductive hypothesis, we get that the degree sum in $G - xy$ is $2(m - 1)$, so the degree sum in G is $2(m - 1) + 2 = 2m$. By induction, the degree sum formula holds all graphs. $\square$

1.3.2 Permutation invariance and equivariance

In mathematical terms, any function f that takes an adjacency matrix $\mathbf{A}$ as input should ideally satisfy one of the two following properties:

$$f(\mathbf{PAP}^\top) = f(\mathbf{A}) \quad (\text{Permutation Invariances}) \quad (1.51)$$

$$f(\mathbf{PAP}^\top) = \mathbf{P}f(\mathbf{A}) \quad (\text{Permutation Equivariance}) \quad (1.52)$$

where $\mathbf{P}$ is a permutation matrix. Permutation invariance means that the function does not depend on arbitrary ordering of the rows/columns in the adjacency matrix, while permutation equivariance means that the output of f is permuted in a consistent way when we permute the adjacency matrix.

1.3.3 Graph Fourier Transform

In graph theory, especially for the **spectral graph theory**, we present an encoding standard, called the **graph Laplacian**. Then we will focus on the Fourier analysis of such graph representation, such as the *graph Fourier transform*. For the purpose, however, we believe that the spectral graph theory alongside its analysis can get us to the layered encoding appropriated of analysing the aggregation contour that the graph is taken into.

Graph Laplacian

Given the adjacency matrix $A \in \{0, 1\}^{n \times n}$ and degree matrix $D \in \mathbb{Z}^{n \times n}$, for a graph $G = (V, E)$, we define the graph Laplacian as the matrix L such that:

$$L := D - A \quad (1.53)$$

Intuitively, it is the analog for the Laplacian operator $\Delta f(x) = \nabla \cdot \nabla f(x)$. In this form, it is quite different from the above Laplacian matrix, but the same idea applies. In this form, it is perhaps better for analysis, though.

We list several conclusive properties of the Laplacian matrix:

1. L is a real symmetric matrix, there by $\lambda_k > 0, \forall \lambda \in \{\lambda_k\}_{k=0}^{N-1}$ with associated orthonormal eigenvectors $\{\varphi_k\}_{k=0}^{N-1}$.
2. If G is finite and connected, then

$$0 = \lambda_0 < \lambda_1 < \dots < \lambda_{N-1} \quad (1.54)$$

3. The spectrum of the Laplacian $\sigma(L)$ is fixed but one's choice of eigenvectors $\{\varphi_k\}_{k=0}^{N-1}$ can vary. Throughout the paper, we assume that the choice are fixed.
4. Let Ψ denote the $N \times N$ matrix where the k -th column is precisely the vector φ_k , then we can easily show that $\varphi_0 = 1/\sqrt{N}$.

Fourier transformation

Let $G(V, E)$ be a weighted graph, $\mathbf{L}$ be its corresponding graph Laplacian, and $f : V \rightarrow \mathbb{R}$ a function defined on the vertices of G . The **Graph Fourier Transform** (GFT) is defined as

$$\mathcal{GF}[f](\lambda_l) = \hat{f}(\lambda_l) = \langle f, \varphi_l \rangle = \sum_{i=1}^n f(i) \varphi_l(i) \quad (1.55)$$

The inverse Fourier transform is then given by

$$f(n) = \sum_{i=0}^{N-1} \hat{f}(\lambda_i) \varphi_i(n) \quad (1.56)$$

This also leads it to the graph convolution (used in graph convolutional network) as

$$f * g(n) = \sum_{l=0}^{N-1} \hat{f}(\lambda_l) \hat{g}(\lambda_l) \varphi_l(n) \quad (1.57)$$

for $f, g : V \rightarrow \mathbb{R}$ as functions on the continuous domain.

Chapter 2. Classical learning theory

Classical learning theory, as most popular of Valiant [1984], Vapnik [1999], Vapnik and Chervonenkis [1971], is the cornerstone of modern learning landscape. It was largely adopted, along the side of computational empiricism, as the main interpretation framework of machine learning since the 2000s. To understand the limitations in application, and different out-of-bound failure modes (as already mentioned of unresolvable phenomena, and structural requirement failures), it is suggested that a thorough coverage of the topic is required.

What we have done, what we have been investigated, including mathematical modelling, can be said to be the *static construction*. What this means is that, albeit intrinsic of collecting and transforming resources, information, or any given kind of input – given the black box, in-out system treatment. This comes as a cost – the model requires human touches to even conform to certain structures, knowledge, tasks, and utilities. Mitigate this issue requires we to come up with a notion that automates the acquisition and relative understanding (butchering the word understanding, since our model currently does not understand *anything*, strictly speaking), and it is called (and found) as the *learning theory* of machines.

2.1 Precursors to learning theoretic

Before the introduction of learning theory as a practical and theoretical process, functional models, usually mathematical models, mostly relies on acute, specific solution, correct formulation of given form, or rather being done solely in a static sense. What this means must be followed through by historic development to understand why it is called static, and subsequently what is then considered dynamic.

Mainly mathematics (but not restricted to such, sometimes usual construction is alright) has been used to formalize, simplify, express, and encapsulate the fraction of reality through the language of quantification, to make **models** that befit the knowledge and the respective fraction observed. Simply said, considering mathematics as the foremost language of abstraction (removing physical realization), quantification (give us the notion of quantity), and the build up of structures within said language, it is fundamentally useful in decoding and expressing reality in the way that makes it understandable, interpretable, logical (the word logical brings around the roundabout), easy to quantify components, actions, factors, objects, and perhaps many further benefits of using such language. Thus, one of the most fundamental example of using the language of mathematics to formalize and quantize reality to be a functional construct, is physics. And its development throughout history also says plenty bit on what is static.

A theory is often considered a static construct, by virtue of how it is formed. It begins by setting up axioms, assumptions, simplification, the target of the theory, and development on which it is built on upon. Further development will follow through the main foundation upward, without many changes to the theory. If one is to modify the foundation, it would mean usually to create new theory on said basis. This process of transition is typically thought in philosophy as the *negation of the negation*, or in

propositional logic, double negation. Fortunately, the transition is perhaps, if not totally very natural of natural science and overall interpretation of the world that we often use. We can attribute this to the fact that we do not understand the world in its entirety, Furthermore, mathematical modelling on itself is not perfect, and cannot capture intrinsically every given thing available, for specific methods – basically, our hands are tied, and is not perfect either way. That is why there must be, and would be this type of changes throughout the theory itself. Building a structure is the same, and so much for a system.

Most of the current, widely perceived general model construction and the way that artificial intelligence was created (under the umbrella term of automated thinking machine, artificial is one, and machine learner later on was subjected to be the equivalent, more *advanced* development) belongs to either the specific type of inference parameterized model, or symbolic model (AI) – which utilizes the propositional logic ϕ of the truth functionality $\mathcal{V}^n \rightarrow \mathcal{V}$. Under said notion, precursor AI, for even specialized task, was conducted in the way that it is designed by experts, for particular problem of knowledge implementation, and then use logic to deduce from said logic pool reasonable course of action, or rationalization; historically, [Minsky and Papert \[1969\]](#) and the symbolic AI practitioners encode hard logics and baseline logical learning to itself. This is partially because of the computational limit of the time – at the time of 1950s to 1970s, and the lack of networks of information and data. Unfortunately, such design choice led to a limitation of the *frame problem* widely present in predicate-based artificial intelligence.

One major theme of the disadvantage of the phenomenological approach on statistically-assumed space, is the fact that it is sensitive to changes. That is, given the holding parameters k for the Boole's coefficient of a spring under stress, for certain deviating interference, by for example, $2k$, then the effective phenomenological model usefulness is reduced by the same factor. Trivially, one can then simply expand the system to accommodate the various setting of k , but such task requires human-inferred choices and modifications, and loss of generalization, in comparison to the mechanistic modelling approach. It is well-understood that without certain dynamic process, and objectively the primitive notion of learning, a phenomenological sometimes might be significantly difficult to handle, based on the immense dimensionality requirement of the parameters used in relation to the system.

Afterward, it is perceived that expert inputs and traditional knowledge acquisition and their form as logical proposition is not adequate – in fact, it is unable to be scaled up. Yes, it is good – even by today's standard – but given a larger, more comprehensive and more complex problem, the system break down to be unsustainable and unattainable. Thereby, we need something else. Something of the quality that makes the model *learn for itself*, rather than waiting for developer's implementation and expert input. That comes of as, eventually not so surprising, an ability that human possesses: *learning*.

In another sense, learning can also be conjured to be the action of *computational reasoning with uncertainty*, as opposed to the general notion of artificial agent, a type of model which makes use of, as previously mentioned, logics and relational discrete spaces. We would look into this aspect later, however, it is also suggested to note that, learning can also still happen, within the logical region of a model. By that, we think of the internal space and external space – the classical definition for learning theory and its distinction in such case is typically the classification between *external-dependent reasoning* and *internal-dependent reasoning* (which depends only on the propositional space).

2.2 Transparency shift

The learning theory started from the early 1970s. By the sense of classical, we do not mean by, as currently called, of non-deep learning theory, but rather the old treatment of theoretical boundaries and rigorous formalization of the concept itself. The nature of the learning theory can also be brought back from certain proposition or argument, most probably, the fact that it is mostly from the black-box inter-

pretation – hence making almost all model being phenomenological by default.

In historical term, such learning problem is derived out from the problem of **adaptation**. Supposedly, given a system M , in an environment E , can it derive itself toward a goal p ? Or rather, can an *agentic machine* capable of existing own their own, in the environment produced by the perceived world? Such requires implicitly the objective of survival p for the machine in such space, and the desired pattern for the respective survival. Typically, this is reduced to finding the right answer from information, or deduction logics toward accurate functionals, as seen in symbolic-based chain logical rules. Without constructing and considering the internal mechanism of the model, we ask the following basic questions:

1. What is the environment? What mode of existence the model is enabled in it? For example, if the world itself is considered symbolically, then it means that the environment it sees happens in the form of causality, as every observation is connected to each other of a symbolic decision graph, thus make use of *first-order predicate calculus* on causal chain, as seen in mostly Clocksin and Mellish [2003], Sterling and Shapiro [1994], Norvig [1992], Brachman and Levesque [2004], Philip C. Jackson [1985], Lucas [1989], Giarratano and Riley [1998], with looser version scaled to *expert systems*, where objects are encoded into the environment via human expert's consideration. The downside on this approach is that observational spaces are limited to what the environment truly encodes – for example, different discrete type of features and decision domains. Thus, the environment itself is *not universal*, in the sense that the connection from reality or larger domain is simplified to the decision network of the environment itself. On the other hand, statistical models encapsulate *observational encoding*, potentially on infinite feature's value range, and use the observed statistics as truth instead. Such removes the rigid structural consideration, and enforce phenomenological interpretation of the data-centric view as production of actual operations and situations in real life. While still limited to sensors or measurement scope, it provides more freedom, patterns, and so on, in exchange for the lack of mechanistic extraction.
2. What is the constraint in which the environment forces on the existence of the model? What is the interaction between the model and its world? What are its allowed dynamics and actions? What is the specified state of the world, static, dynamic, mutable, or objectively considered as facts? Those are questions from the internal construct of the environment, as well as actionable domain of the model of interest, which lives in it. For example, a model in symbolic sense can only be a *predicate processor* with weighted path, such that certain predicative decisions are enforced instead of others, for morality (wrong or right on ethics), and so on.
3. What are the failure modes of the environment that *discourage* the model from doing certain parts? In typical setting, we assume that in the operation of such model inside the environment, it will encounter *problems*; however, the question is to encode it, and toward such, the question elaborates to what kind, how, and the role of such inhibitions.

Those questions are then answered by giving adaptation in a form of **trial and errors**. Effectively, this means learning a scenario or target concept by treating them in a form that can be given of *causal relationship*. The simplest and most straightforward of such is the construct to be function, i.e. $f : x \rightarrow y$. This permits partial realization of the system at worst of such framework, and total predictability of such system at best. Suppose a system contains $(x, y, z, q_1, \dots, q_k)$ for relative set of $q_i \leq k$ elements of quantification. Then, the system can have the class of all *directed function* mapping from $x \rightarrow y$, thus in a directed graph can be represented using directed edge, with the arrow pointing to dependent value in respective frame. In general, however, such relationships are often two-way, thus guarantee the existence of the inverse function. Then, a system can have binary multivariate relationship to be encoded of such

causal form, for example, $v_1 : x \rightarrow z$, $v_2 : (x, y) \rightarrow q_3$, and so on. The set of all of such function in the system can be considered as the edge class (in a sense, a graph structure without fixed numerical representation is really effective for this type of assumption). The conjecture is this: the model can, estimate certain thread of this space, and in ideal case, find a causal chain good enough to infer different information about it. We call this specific node the **universal node**. Whether this node exists or not in the structure is unknown, and much like other form of assumptions and typical structural view, it is not totally guaranteed but of the bias in constructing such structure by itself. Nevertheless, much of learning theory exists, or rather, use this for approximation of systems.

Symbolic systems, in such sense, can be said to trade flexibility and dynamics, for standard rigid logical procedures. Such can be used in compiler designs, rigidly defined from optimized knowledge system where deviation cannot be touched upon, and many procedure-adaptive only systems, of which we expect the class of operables to be restricted. On the other side of the spectrum, statistical approaches trades such structural richness and clarity for flexibility, dynamics, and inferencing ease, with more scalability. Such defines the operational transparency window of which we will call as the **transparency ladder**, and the action of juggling between the two extreme is called the *transparency shift*.

transparency

Definition 2.2.1 (Transparency ladder). Categorization of system can be defined on the **transparency measure** of the system, by their structural commitments on such. We set the **ontological maxima** of such measure representative of classical finite logical-predicate systems, and the **ontological minima** of the transparency scale as the full-distribution, statistical based system as the default. Any system in operation would be defined within this ladder.

Almost all systems contain certain form of structural information, bias, assumptions, or simply implementation choices. Such are grouped into a term **bias**, of which infers situational structural changes from the absolute neutral baseline. Such bias term can be many, including in dynamics or the structural formation itself. To get it further, we might as well investigate the bias-stack of the functionals ladder.

2.2.1 Ladder of functionals

Due to the above observation itself, the minimal baseline on assumption of any given exposed system, is that potentially, such systems do change but would rather be unattainable of all sufficient details of the hidden variable, or the graph theoretical mapping of all given quantities or properties expressible and assigned, from the perspective of the designer and the hypothesis model class. Thereby, the first assumption in worst-case scenario would be to remove this. However, before that, we need to clarify what are the ladder of functional in effective operation.

There are two nontrivial system that would give rise to the observation typically seen in symbolic AI, and classical learning theory. One of which is the separation of the **mechanistic layer** and the **phenomenological layer**. Specifically, symbolic or any other structure admitted of the more sophisticated modelling, takes on the form of mechanistic layer – i.e. structural commitment on what the system operates, either is the substrate by itself, or a rather mechanistic components of it. Phenomenological layer can be given as the upper layer, of which is largely tied into the **observation** of the specific system in operationalization, or the action of either giving a system existing on paper or writing, or examine a natural one, in its own implementation and existence, and of its operation throughout (i.e. *dynamics*). Per such, the following setup would preside both the classical systems, and the newer statistical one.

mechanistic layer

phenomenological
layer
operationalisation

1. **Layer N** (Nulla): The ontological oracle itself. None can be said about such.

2. **Layer I:** Here, we have the **object in consideration** itself and of its details. Supposedly, let us denote such object being $c \in \mathcal{C}$ of a specific class of objects, or arbitrary similar-type classification. Such object is not object in definitive form, but can be a system, ‘world’ of dynamics in which there exists *observables* that can be used to infer properties of the object space itself. Here, we assume that such object can be contained of the descriptions $\mathcal{W}\{(\Theta, \Gamma)\}$ for structural and actionable components, and the dynamical operations, thus giving datasets or observations on actions from $\mathcal{W}$. This is the system baseline. This is considered the *partial oracle*, or *frame-defined oracle*, a degradation of the oracle, of which we suppose the **consideration** above into a frame Q containing all specification of such system.
3. **Layer II:** The layer reserved for *observables*. Such is to say, during $\mathcal{W}$ observations, there exists a set of conditions, $Q = (q_1, \dots, q_n)$ that is finite on conditioned experiment, and there exists a set of *measurable observations*, or observable, such that the theorized structural ‘production’ of such measurable is of the form $c' : Q \rightarrow Y$ for Y the observable sets. We distinguish between the theorized c' and c . Per usual modelling system, the assumption typically considers $c' \sim c$, and in the best case possible of mechanistic outcomes, $c' \equiv c$. Using such, we say of the statistical and symbolic approach as to presuppose layer II directly, with symbolic AI suppose that symbolic systems are created from logical derivation and truth logics on supposed actionable objects, while statistical methods using structural inference from data patterns instead.
4. **Layer III:** From here, we suppose the beginning of the starting implementation of the subject itself. The concept c' exists, but we do not yet clarify the assumption on which $c' \equiv c$. However, let us take such as true. We have the observable objects instantiated into the **dataset** D , of which is most often in the form $(\mathbf{X}, \mathbf{Y})$. Sequentially, even text information can be considered such special ordering on dataset form, but implicit. Such dataset is supposed by the statistical observation under limitation, from a structural-based distribution $\mathcal{D}$. We note that this means we assume $(\mathcal{X}, \mathcal{Y}) \sim_{\mathcal{D}}$ $\mathcal{D}$ of such distribution, however, $\mathcal{D} \approx c'$ and would not have guarantees on c . Several assumptions can also be made, for example the collapse of inter-observables interaction, or hidden conditioning transparency, into i.i.d. (independently and identically sampled distribution), such is to say the setup being single-object on itself¹. Starting from here, Symbolic AI no longer works within such stack. Or rather, Layer III of symbolic AI is the presupposition of the predicate system, where that there exists a universal process that can be inferred from layer II construct, that is consistent with any sample. Hence, they presume the stability that is identical to **infinite observations** (implicitly on logical conditions) of the statistical scheme with full oracle exposure, which are both realistically unexpected in statistical assumption set, and too strong oracle presumption of system dynamics. This makes Symbolic AI very structurally correct, but non-trivial and unable to scale exponentially, with brittle formulation starting from its own predicate expectation modes.

We might want to pause here and clarify a few consequences of the modelling systems as this point. Defining this point means admitting that **probability** is presupposed from the transition between II and III. More specifically, since layer II has only $c' : Q \rightarrow Y$ as the *total partial oracle*, we can observe that up to

¹Under scrutiny, any subject of which i.i.d. does not hold, means that the system itself is not perfect, which is a natural conclusion. However, several reasons remain, for example the existence of other samples in the distribution does not affect of others, and so on. That is, the system is collapsed into **single-object** ensemble sampling, which restricts itself to instantiation of the subject class under static formation. This has implications for what type of stable system that can expose effectively under such assumption; including exclusion of memory effects, system feedbacks, adaptive behaviours, and correlated hidden noise. However, it is still an effective simplification.

layer II, there is no probability, but only when we move to operational level of layer III, that we encounter specific hard points. We would say that the *gap* between II–III is called the *operational gap*. Under such gap, probability enters in circumstances with limited or no structural access, with existing hidden variable, and uncertain frame Q . Such is where probability can be used for structural assumption on the distributor $\mathcal{D}$ itself. However, under partial dataset, as another problem is that with finite many datasets, there is no guarantee that D itself is the fully-capable or possible dataset for the problem, thus there exists the $\mathcal{P}$ -distribution of the dataset itself. Thus, the generator $\mathcal{D} \rightarrow \mathcal{P}$ pipeline produces further probabilistic interpretation, by solving the partial-knowledge problem. Formally, we say that there exists a *structural commitment* where we suppose of the existence for the $\text{Sm} : \mathcal{C}' \rightarrow \mathcal{D}$ sampling operator, of such can also be injected with i.i.d. assumptions. The consequence is thus randomness collapse of structures, replacing unknown mechanistic variations with a measure of inferred distributive shape, and treat observables as stochastic samples than structured outputs. Those are commitments made by statistical learning.

2.2.2 Classical statistical learning

Classically, we can see that the problem of double descent that would be later introduced in later sections is not inherently a problem that can be explained only by considering model complexity and the error rate. As such, it requires a framework that consider multiple factors together, and interpreting different hypothesis. Nevertheless, classical theory does indeed point to certain factors or aspect that should be of concern. For example, we have seen plenty of analysis that indicates that *sample complexity* might be the biggest factor, considering the behaviour is intrinsically operable and perhaps observable in almost any model, except for certain model that is specialized of the data structure itself. Even by then, the error still can be obtained given large enough model disparity to the concept class. There are attempts to, indeed, analyse double descent from a VC perspective (Lee and Cherkassky [2022]), however, the paper falls into the same pitfall as others in their interpretation and analysis. Nevertheless, within the spirit of this section on the introduction to classical learning theory, it is still worth it to analyze the internal structures encapsulated in such theory itself.

More specifically, even though we have said plenty of the disadvantage of the classical theory in terms of statistical learning, their insight is valuable in determining the possibility of dynamic in question. There is a reason why they are versatile and structurally valuable enough that our majority of understanding are formed from such framework itself. In simplicity, within the principle of working, our model is called the **hypothesis** $h \in \mathcal{H}$ of certain **hypothesis class** $\mathcal{H}$. The target objective is typically concerned of, denoted by $c \in \mathcal{C}$, as to be of a structural class $\mathcal{C}$, in which observations creates information in the form of the observable set $\mathcal{S}$. The set of observation $\mathcal{S}$ is presentable in the space of 2-tuple $(\mathcal{X}, \mathcal{Y}) \subset (\mathbb{R}^d, \mathbb{R}^p)$. Here, $\mathcal{X}$ is called the *ground* or input space, and $\mathcal{Y}$ is called the *target* or output space. By default, in general, the input space $\mathcal{X}$ is supposed to follow some special distribution $\mathcal{D}$. This treatment branches from the ambiguity of the origin of $x \in \mathcal{X}$ for elements of the input space, that is, $x \sim \mathcal{D}$ by a random sampling process. This assumption also ensures relative population representability². By this, we mean that the observation captured fully describe what was happening with the entire concept – even including the pathological and outlier. We also assume that this sample of observation is **identically and independently sampled** of a given sampling process.

The task for the hypothesis, and the designer (us) is to figure out how to best approximate c , assuming arbitrary notion of knowledge, prior information, such that h correctly *mimic* the behaviour and phenomenon observed by c . We call this loosely as the **learning problem**. Naturally, learning problem is

²By population representability, we mean that the density is well-illustrated, such that the data fully captures the entire operational space that is relevant and meaningful of the underlying concept.

a statistical problem, since their setting is fully observational.

The **learning problem** divides the act, with respect to configuration the chain of component into two parts. The first one is for any measurable space of the model, after defining the system, the dependency, and the question, we have to gain the *accuracy measure*, or *phenomenological aberration* of a given model, and for the respective problem. Depends on the specific requirement of the problem, the measure can change, but, they have the same landscape. The second one requires the development of a *scheme* to dynamically move the given hypothesis, under fixed description, into a supposed optimal point in which under specific threshold, our model is 'perfect' for the concept. That is, *mimic c by h in a perfect way*. It is perceived, though, that the objective of learning the concept perfectly is impossible. Hence, we will often opt for the latter half of approximating it to a given threshold θ

Setting in the second part, most of our problem is concerned with the issue of predicting the unseen states of being, or rather, the *generalization problem*. This then perhaps, classify the problem of statistical learning theory to *mostly* predictive, or rather, we are concerned of prediction model. This is subsidized by the inner *estimation problem* for an estimation model. What is the difference between the two? First, the distinction comes from the trivial on-sight fact that estimation problems and their models, comes of as a closed-form working space. That is, they work on a closed space of what is observed, and extract properties, determine statistically the behaviours, characteristic of that system alone and only. In other word, an estimator *uses data* to guess and learn about the true state, and properties of the setting (usually also said to be the *true state of nature*). However, we notice a discrepancy - not always will we get the entirety of the problem that we are interested at. For some reason, for example, our data is only quite a portion, lacking a bit valley, a bit of patterns here and there by the blank spaces where observations left behind. Then, what if we can generalize pass that? What if we also want to somewhat estimate correctly for unseen portion of our concept class? This is called the predicting problem we have said earlier, in which we use the estimation already established, and add some flavours into what can possibly constitute a larger view - a more entire view of the learning concept. From a statistical perspective, the model itself in statistical learning theory, would have to both estimate the intrinsic empirical distribution and concept $(c_S, \mathcal{D}_S)$, and then somehow, leaves room for generalization to the actual distribution and concept $(c, \mathcal{D})$, which is particularly of more interest than not in practical settings. Here, the notation simply means two thing: the facade of which c is the true concept, and $\mathcal{D}$ is the uncertain, probabilistic presumption of the structure in which c seems to exhibit, thus there then also exists $c_C \equiv c'$ of the proxy concept c' that 'seems right' in comparison to the true concept, and also its observation-based distribution.

With just that much narrow scope of the claims, we can then, at least formalising ourselves of particularly this layer of understanding, toward which we will inevitably have to procure from mathematical formulation, of the learning process, as followed.

Definition 2.2.2 (Learning process, provisional). A procedure $\mathcal{P}$ is called a **learning process**, given two **constructs** or **model** $h \in \mathcal{H}$, $c \in \mathcal{C}$ if it follows an **error measure** or operator $\ell(\cdot, \cdot)$ such that, for $\mathcal{P}_i \in \mathcal{P}$ indicating the i th iteration of the procedure $\mathcal{P}$, h_i the i th hypothesis iteration of $\mathcal{P}$, minimize the operator that it chooses, that is,

$$\exists n, i \in \mathbb{N}^*, n \geq i, \quad \ell(h_n, c) = \min_{h \in \mathcal{H}} \ell(h_i, c), \quad (2.1)$$

In most results of the statistical learning theory, the majority would, and will use the setting of *supervised learning framework*. The reason why such can be justified can be expressed in simple terms. Consider unsupervised learning. Under the scope of unsupervised learning, the model essentially goes into a free-mode, where there exists no metric that is non-context enough for evaluating the theoretical form of the model. For a given density estimation model ψ , there can be many ways to 'think of evaluating it' on

how well it groups objects together. This means there are not so much meaningful information can be extracted from examining the behaviour of the model, including its performance. Furthermore, since it is random as much as a random process, no conclusive bounds or theoretical theorem can be made on the maximal-minimal problem of a given model. Hence, most of the learning theory is concerned of supervised learning setting, because it is a controlled, closed, well-defined system of interest. For what it is good at of uniform convergence bounding, however, the theory itself established a controlled region of analysis in which models are often guaranteed of behaviours and patterns, fitting what was assumed.

2.3 Classical setup

hypothesis class

In simplicity, in classical learning theory in the modern paradigm (data-centric), our model is called the **hypothesis** $h \in \mathcal{H}$ of certain **hypothesis class** $\mathcal{H}$. This basically means that h is configured to hypothesize of the one-one relationship of the observational space. The other ingredient of the learning theory includes:

1. The concept class versus function class can be fully incorporated into one as $\mathcal{C}$, if one is to assume the form of any given concept as $c : X \rightarrow Y$ for another set Y which can lie in the sigma-algebra $\mathcal{S}$, or the measurable mapping as f .
2. The task for the hypothesis, and the designer (us) is to figure out how to best approximate c , assuming arbitrary notion of knowledge, prior information, such that h correctly *mimic* the behaviour and phenomenon observed by c . We call this loosely as the **learning problem**. Naturally, learning problem is a statistical problem, since their setting is fully observational.

We are given our information of the system by the form of a dataset $\mathcal{S}$, called *observation*. We first define the **observations**. The set of observation $\mathcal{S}$ is presentable in the space of 2-tuple $(\mathcal{X}, \mathcal{Y}) \subset (\mathbb{R}^d, \mathbb{R}^p)$. Usually, hence, our data will be realized in the Borel σ -algebra $\mathcal{B}$, and so are most of our numerical spaces. The full dataset is *discrete*, contain of n instances of observations, and is denoted by

$$\mathcal{S} = \{(x_1, y_1), \dots, (x_n, y_n)\} \subset (\mathcal{X} \times \mathcal{Y})^n \quad (2.2)$$

Here, $\mathcal{X}$ is called the *ground* or input space, and $\mathcal{Y}$ is called the *target* or output space. By default, in general, the input space $\mathcal{X}$ is supposed to follow some special distribution $\mathcal{D}$. This treatment branches from the ambiguity of the origin of $x \in \mathcal{X}$ for elements of the input space, that is, $x \sim \mathcal{D}$ by a random sampling process. This assumption also ensures relative population representability. By this, we mean that the observation captured fully describe what was happening with the entire concept – even including the pathological and outlier. We also assume that this sample of observation is **identically and independently sampled** of a given sampling process.

The **learning problem** divides the act, with respect to configuration the chain of component into two parts. The first one is for any measurable space of the model, after defining the system, the dependency, and the question, we have to gain the *accuracy measure*, or *phenomenological aberration* of a given model, and for the respective problem. Depends on the specific requirement of the problem, the measure can change, but, they have the same landscape. The second one requires the development of a *scheme* to dynamically move the given hypothesis, under fixed description, into a supposed optimal point in which under specific threshold, our model is 'perfect' for the concept. That is, *mimic c by h in a perfect way*. It is perceived, though, that the objective of learning the concept perfectly is impossible. Hence, we will often opt for the latter half of approximating it to a given threshold θ .

The learning problem – which is most importantly described by the phase of *phenomenological aberration* (or data-based replication) – is then formulated as followed. Given the machine learning model expressed

a hypothesis h of the hypothesis class $\mathcal{H}$, the learning theory aims for creating a procedure to learn either elements of the concept class $\mathcal{C}$ of all concepts $c : \mathcal{X} \rightarrow \mathcal{Y}$, or the function class $\mathcal{F}$ of all functions $f : \mathcal{X} \rightarrow \mathcal{Y}$, which is usually set to $\{0, 1\}$ or $[0, 1]$. This distinction is trivial, hence, if it is clear, we will talk about the concept class $\mathcal{C}$ only as representative form. The learner $\mathcal{L}(h)$ consider the set of possible hypothesis $\mathcal{H}$, in which might not coincide with $\mathcal{C}$. It receives a partial image of sample $S = (x_1, \dots, x_2, \dots, x_n)$ drawn i.i.d. according to $\mathcal{D}$ as well as the label $(c(x_1), \dots, c(x_n))$. This constitutes the dataset $\mathcal{S}$, which are based on specific concept $c \in \mathcal{C}$ for the hypothesis to learn. The task is then to use (or *extract*) meaningful information to select a hypothesis $h_S \in \mathcal{H}$ that accurately mimic c , with marginal error Θ . The notation h_S stands for all hypothesis that can be inferred from the range of the dataset (the first argument, $\mathcal{X}$).

The marginal error Ξ is considered of two parameters, the empirical error $\hat{R}(h)$ and the generalization error $R(h)$. This can be justified as following.

To compare between h and c , we use the metric provided by $\mathcal{Y}$. Then, define certain type of function called *loss function*, also called objective, error function, denoted by $\ell\{h, c\}$ such that it measures the difference between h and c . This measure is arbitrary, and we assume no form of it, aside from the fact that it also acts and return values on certain real subset of values. The following definitions provide us with the notion of *empirical error* and *generalization error*. Notice that one is empirical – indeed, since we only can work on certain fixture of c , i.e. the dataset provided, and the other one is generalized – meaning the true error toward the actual target concept $c \in \mathcal{C}$.

Definition 2.3.1 (Empirical risk). Given a hypothesis $h \in \mathcal{H}$, a target concept $c \in \mathcal{C}$, and a sample $S = (x_1, \dots, x_m)$. For some particular $\epsilon > 0$, the **empirical guarantee** or *empirical bound* of h is defined by

empirical risk

$$\hat{R}_S(h) \cong \mathbb{P}_{x \in S \sim \mathcal{D}}[\ell\{h(x), c(x)\} \leq \epsilon] \quad (2.3)$$

of which it estimates the probability statement on the guarantee of errors for normalized probabilistic-measure of empirical in-dataset correspondence from h to c under the operative setting. If computed heuristically, the **cummulative empirical measure**, or the *empirical risk*, as

$$\hat{R}_S(h) = \frac{1}{m} \sum_{i=1}^m \ell\{h(x_i), c(x_i)\} \quad (2.4)$$

Similarly, the generalization risk defined in the same fashion, but extending beyond the dataset S that the empirical risk is usually conditioned on, and thus presume the posture on *infinite access* of the oracle itself. That is, for the concept c , the finite observation space $[a, b]$ is dense for any arbitrary $a, b \in \mathbb{R}$.

Definition 2.3.2 (Generalization risk). Given a hypothesis $h \in \mathcal{H}$, a target concept $c \in \mathcal{C}$, and an underlying distribution $\mathcal{D}$ on $\mathcal{X}$. For some particular $\epsilon > 0$, the generalization error or *risk* of h is defined by

generalization risk

$$R(h) = \mathbb{P}_{x \sim \mathcal{D}}[\ell\{h(x), c(x)\} \leq \epsilon] = \mathbb{E}_{x \sim \mathcal{D}}[\ell\{h(x), c(x)\}] = \int_{x \in \mathcal{X}} \ell\{h(x), c(x)\} dP(x) \quad (2.5)$$

Thus, we say the generalization risk is **dense** on observation space.

The notion of empirical and generalization error clearly indicates the issue of two concepts – one observable ‘component’ concept c' presented by the dataset, and the actual concept c itself. Hence, we assume, of the universal $\mathcal{X}$ set belongs to c , there exists $\mathcal{X}_c$ and $\mathcal{X}_{c'}$, hence we have two distributions, for example, denoted by $\mathcal{D}_c$ and $\mathcal{D}_{c'}$. The process of learning is then rather simple. To effectively find the

model h that fits c , it is imperative that both empirical error $\hat{R}$, and the generalization error R remains low. This essentially posits the configuration space where there exists two absolute convex hulls - one with its critical point near $\hat{R}(h) = 0$, and one near $R(h) = 0$. Usually, we note that $R(h) \neq \hat{R}(h)$ unless under certain configuration, often by the statistical analysis that $m \rightarrow \infty$ of the sample size will make the relationship an equality. Such means, further, that we are to optimize between $R(h)$ and $\hat{R}(h)$ on the hypothesis-loss slice, and ideally get both small, though in reality this is not always the case. Suppose we use a procedure or algorithm $\mathcal{A}$ to achieve this goal. Let $\mathcal{A}_t = a_{i,t}$ be the i th run at t times of this algorithm such that each time, such algorithm produces $h_{i,t}$ of the different hypothesis in this space. Then, the setting is *suitable* for an *effective procedure* $\mathcal{A}$ if there exists $t = n$ for $n > N$ of the *reduction threshold*, such that

$$\begin{aligned} c \approx_\epsilon h_{\mathcal{A}} &\equiv \arg \min_{\ell \in \{L\}} \mathbb{E}_{i < M} \mathbb{E}_{t < T} \mathcal{L}_\ell[\mathcal{A}_{i,t}([\mathcal{H}], \nabla, [q_1, \dots, q_p])] \\ &\equiv \arg \min_{t \leq N, i \leq M} \mathbb{E}_{\ell \in \{L\}} \{\mathcal{L}_\ell[h_{i,t}]\}_{\mathcal{A}}, \quad \mathcal{L}_\ell[h_{M,N}] \leq \epsilon \end{aligned} \quad (2.6)$$

can be satisfied, for ϵ is the threshold that exceeding N iteration gives, generally speaking, the suitable hypothesis model that can approximately correct up to specified accuracy as required. Here, $[\mathcal{H}]$ is used to denote the *space characterization* of the variations in the hypothesis class, thus allowing for the algorithm to change the structure of the hypothesis in working correctly. However, this is a characterization that is often reduced in classical setting. For example, typical literature handles $|\mathcal{H}|$ as a fixed object - where the set of functions the algorithm can output is not determined of its dynamic on how it changes as the algorithm runs, in many sufficient dynamical setup. Characterization then unfortunately should not be overcommitted to a specific formalism, as such would be particularly contradictory toward what can be said of the arbitrary choice given in many statistical notions.

Equivalently, in addition, it means that the setting itself converges such that the series on ℓ -metric reduces to deviation ϵ only from c . This is extended on $\mathbb{E}$ accordingly toward the family of relevant loss-injection functional space, and it is used in the loss evaluator $\mathcal{L}$. ∇ here is the *proxy sight* of $c \in \mathcal{C}$; if one switch $\nabla \rightarrow \mathcal{C}$, it means the algorithm has full access to comparison to the concept itself. Finally, $[q_1, \dots, q_p]$ is the set of initial conditions or **specific parameter**, all of which initialize the learning algorithm and the setting in which the algorithm would be reserved for. We note that this is not the later established system hyperparameters, since they are two different conditioning, but also can overlap. Reducing $t = 1$ means constraining the algorithm to a specific set of analytical algorithm, for example, least square methods in regression analysis. Take from the series, we call $\mathcal{L}$ in its formal name the **general risk**, of which is taken as

$$\mathcal{L}_\ell(h_{M,N}) = \underbrace{\frac{1}{m} \sum_{j=1}^m \ell\{h_{M,N}(x_j), c(x_j)\}}_{\hat{R}_S(h_{M,N})} + \underbrace{\int_{x \sim \mathcal{D}} \ell\{h_{M,N}(x), c(x)\} dP(x)}_{R(h_{M,N})} \quad (2.7)$$

of which we can replace M, N with i, t under expressions of the iterative form. From there, under the same decomposition, we recover the case

$$\begin{aligned} \arg \min_{t \leq N, i \leq M} \mathbb{E}_{\ell \in \{L\}} \{\mathcal{L}_\ell[h_{i,t}]\}_{\mathcal{A}} &\Rightarrow \mathcal{L}_\ell[h_{M,N}] \leq \epsilon \\ &\equiv \hat{R}_S(h_{M,N}) + R(h_{M,N}) \leq \epsilon \end{aligned} \quad (2.8)$$

which recovers the $R(h)$ and $\hat{R}(h)$ problem, i.e. optimization on Ξ -marginal error landscape. This now appears as rather a **budget allocation problem**, because often in settings, the condition $\mathcal{R} = \hat{\mathcal{R}}$ is rarely true, and only under dense dataset, i.e. for $m \rightarrow \infty$, $\mathcal{S}$ dataset representational on population. Thereby, one must either set up a proxy where this is satisfied in practice, or required to control this budget allocation logic such that the algorithm $\mathcal{A}$ being picked can navigate correctly that the problem can be solved. Alternatively, the consequence also follows that there exists specific **algorithmic space**, where the cost spent per algorithm initialization would be taken into account. Robust algorithms can complete the sufficiency-guaranteed setup in lesser cost, or more, depending on the natural alignment between problem and algorithmic space. However, it also implies that if sufficiency setting is not guaranteed, the conclusion does not follow that the lesser cost models under sufficiency can be better than the more expensive one, as it exposes hidden **structural bias assumption** on constrained condition space.

There exists a problem in the underlying theory of data-based statistical learning. The question is as simple. In the statistical learning setting, one gain the dataset $(\mathcal{X}, \mathcal{Y})$. Because we are using the *black-box interpretation*, the concept that is supposed to be in $c : \mathcal{X} \rightarrow \mathcal{Y}$ is said to be governed by a probability distribution $\mathcal{P}$ on said data. Now, depends on the type that you want, it can either be **dense**, that is, there are infinitely many $x \in \mathcal{X}$ as the phase space, and the probability is on y , that is, $\mathcal{P}(Y)$. Or, you can remove the strictly dense assumption and get the probability $P(\mathcal{X}, \mathcal{Y})$ of joint probability between the two instead. We know from the start that the dataset is not always representative, neither captures the entirety of the concept c . Hence, there must exist a particular $\mathcal{D}$ of true probability distribution that describes c . In a perfect world, we would not need to see c as probability distribution – since in said perfect world, everything might as well be deterministic, but the best one can do is to be able to get close to the presumptions probability distribution of the event. Then, the question is, what is the *inherent correlation* between $\mathcal{P}$ and $\mathcal{D}$? To answer this question, we might as well need some divergence metric to see their diverging property. This will eventually be the KL-divergence measure, which we will discuss later. A disadvantage of KL-divergence, however, is its inability to take into account the structural problem of the system, as well as its model. We might as well address this in our later section. Nevertheless, as stated, the KL-divergence can be used to capture

$$D_{KL}(\mathcal{P}||\mathcal{D}) = \int \mathcal{P}(x) \log \frac{\mathcal{P}(x)}{\mathcal{D}(x)} dx \quad (2.9)$$

of the problem set. However, in nominal scenario, such is often being done using empirical/proxy KL divergence measure. Using the above notion of generalization, we can separate the learning problem into two goals – one to minimize $\hat{R}(h)$, and one to minimize $R(h)$. As we have been saying, and for what would be of Theorem 6.2.3, one of our main assumption, as well as a guarantee under uniform convergence, is the fact that we must have $\hat{R}(h) \approx R(h)$ for larger and larger dataset size, or the empirical space. Generally, this branches off to be two problems. First is the problem on condition of optimizing for the empirical risk, i.e. minimizing on dataset $\mathcal{S}$. The second problem considers the next step, where one test the optimized model on empirical to generalization risk, and thus choosing models that also minimize this condition as well. In all, the problem of learning can be defined on the phase:

1. Begin optimization of h on $\hat{R}$.
2. Test on $R(h)$ to recover generalization error.
3. Calibrate until obtainig minimal functional h -approximation.

This corresponds to the **Empirical Risk Minimization** method, or ERM. Under our same argument in the general risk and learning procedure, ERM effectively fixes $R(h)$ uncontrolled, while optimizing for

$\hat{R}(h)$ on conditioned spatial set $\mathcal{S}$. The consequence, of which we would explore later, of ERM on this approach, is that even with cross-validation, which is equal to proximal-Monte Carlo statistical schema on training procedure, artefact of post-processing generalization proxy remains, and double descent can be hypothesized as the artefact of this uncontrolled growth. We note, however, that often time, under the condition that ERM is often applied to, it is rather perfect. That is, optimizing the game plan

$$\mathcal{A}_{ERM} = \arg \min_{h \in \mathcal{H}} \hat{R}_{\mathcal{S}}(h) \quad (2.10)$$

is often enough for a variety set of problems, that the assumption or hope that $R(h) - R_{\mathcal{S}}(h) \leq \gamma$ of small γ -deviation is often enough. However, under volatile setting, unclear hypothesis space and unconditioned models, for example, complex many-dimension neural network structure, the setting is unfounded, and often time, might give rise to pathological patterns.




Figure 2.1: An illustration of statistical learning theory on the evaluation of the risks and errors, during learning process. c' is presented in the 'orbital' vicinity around c , with its distance of certain metric define how 'accurate' the reconstruction from distribution can be. Of the hypothesis set $\mathcal{H}$, there exists the Bayes hypothesis h_B and an arbitrary 'random' hypothesis h , and their respective measure.

Bayes model

For the generalization problem's best model in a fixed $\mathcal{H}$, we have the notion of a **Bayes model**.

Definition 2.3.3 (Bayes risk). Given a distribution $\mathcal{D}$, the Bayes error R^* is defined as the infimum of the errors achieved by measurable functions $h : \mathcal{X} \rightarrow \mathcal{Y}$:

$$R^* = \inf_{h \in \mathcal{H}} R(h) \quad (2.11)$$

Definition 2.3.4 (Bayes model). A hypothesis h with $R(h) = R^*$ is called a *Bayes hypothesis* or Bayes classifier, and is taken by

$$\forall x \in \mathcal{X}, \quad h_B = \arg \max_{y \in \mathcal{Y}} \mathbb{E}[\ell(y, \mathcal{Y}) | x] \equiv \arg \max_{y \in \mathcal{Y}} \mathbb{P}[y | x] \quad (2.12)$$

for the second equivalence on h_B happens only in 0-1 regime loss.

The Bayes model is understood to be the minimal of $R(h)$. And, by our observations, $\hat{R}(h) \neq R(h)$ by a set of assumption on generalization access. In all, the dynamics depend on the commitment on Ξ as the marginal error measure containing both itself. For now, we denote the risk as such, the goal of any system can be interpreted to minimize Ξ in the learning setting. Thus, reframing the learning problem, we can give the following structural observation.

Theorem 2.3.1 (The general learning problem). *We suppose the statistical, classical learning problem setting. Then, the learning problem of optimizing (h, c) correlation, for $\mathcal{H}, \mathcal{C}$ the hypothesis spaces and concepts space, on comparative (pseudo)measure $L\{\ell\}$ of which denotes a comparative objective functional, possibly composite, whose components include empirical and generalization terms, is the optimization for procedure $\mathcal{A}$, of which it produces the following model,*

$$\arg \min_{h \in \mathcal{H}}^{\epsilon \geq 0} L_\ell(h, c) \equiv \mathcal{A}(\mathcal{H}, \Xi) = \mathcal{A}(\mathcal{H}, \hat{R}_S(\cdot) + R(\cdot)) \quad (2.13)$$

where $\arg \min_{h \in \mathcal{H}}^\epsilon$ is the ϵ -minimum natural irreducible error without origin admission. If such model is efficient in $\text{poly}(\cdot)$ time, we say it is an *efficient algorithm* of the learning problem up to ϵ -floor. Under such, $\mathcal{P} \equiv \mathcal{D}$.

The notion would be generalized in the ongoing chapters, and the notation are being conveyed not particularly clear. However, as said, the main insight is this – the problem is between different ϵ -tolerance class, of which affects the model and the algorithm process itself. Furthermore, as said, this learning problem setting presupposes **oracle enumerative access**, which is infinite sampling and distribution abstraction $\mathcal{D}$. Those are rather strong assumptions, and should be made explicit within the real, practical learning problem itself. For such, we would like to present another definition, with only a slight change of hand.

2.3.1 Instantiation of the learning problem - ordering

When instantiate the setting of the learning problem, we are required to specify at least a sufficient domain validity of which the finite experiment can work on. Hence, while we can non-trivially but rather effortlessly define certain set of admissible working mechanisms, for example floating point precision abstraction, there are deliberately setting controlling methods on variations of the experimental setup that should be clarified. This can be said as for determining the **order of hyperparameter change order**, of which affects randomization patterns, reproducibility, and comparative class differences. We note this from now before moving on to later sections, as it is standard of the experimental preparation.

Under the learning setting, there exists two types of **global hyperparameter**. We recall hyperparameter as a name assigned for non-learning mass of the model, which means they are conditioned on the setting of the learning setting itself. The two types are then accordingly, the **system hyperparameter**, and the **dynamic hyperparameter**, where the first one controls the space, for example, the space $\{\mathcal{S}, \mathcal{H}, \mathcal{C}, [H_1, \dots, H_k]\}$ and so on for hyperparameters of the hypothesis, sample space, concept space, and another system parameters and specifications, which can includes the algorithm class $\mathbb{A}$. The second one, deliberately works on the instantiation of the hypothesis h , the concept space c and its exposure, proxy c' , the parameters considering hypothetical $\mathcal{D}$, aggregated measures $\mathcal{P}$ on dataset-distribution, the controlling parameters of the design on $\mathcal{A}$ -algorithm, and so on.

The **specific hyperparameter**, contains hyperparameters that work upon the operation of specific algorithm or specific in operation construction. This includes step-size schedule, data feed, iterative constraints, where the global hyperparameter only specifies N iteration, but we choose to only run $t = 1$ time – which reduces it to analytical capable algorithm if one wants to justify the learning criteria. It also includes local outlier handling, procedural constraints, individual regularization patterns or the absence

thereof, weighted decision template, kernel bandwidths, and so on, including dataset splitting in runtime, for example the 80–20 ratio between train and testing set. Those are often the region where it overlaps with the initial conditions specified by $[q_1, \dots, q_p]$, since initial condition can be considered the specific hyperparameter operable of the algorithmic input, however in certain case, such can or cannot be the case. We note this tension of categorization in our attention. In all, the purpose of this separation is to showcase that there exists deliberate **clustered and nested layer** of construction, typically produced or controlled via parameterization into hyperparameters. Hence, specifically, to control experimental variations, the order is from global hyperparameter to downward to specific hyperparameters. This includes changing the m sample size. This lies totally on the global hyperparameter, hence if one deliberately changes m and also simultaneously change iterative patterns, but then change the high-level order of the global algorithm hyperparameters, it produces artefact of misalignment in hyperparameter control that cannot be taken into account. Hence, using such, in later section, we can also categorize the wide canonical family of double descent result into such order, where one then observe that epoch-wise double descent is of lower-order than the traditional double descent simply by hyperparameter ordering itself.

2.4 Optimal learning - PAC-learning

The PAC-learning process (Shalev-Shwartz and Ben-David [2014], Mohri et al. [2012], Hajek and Raginsky [2021]), or rather, *Probably Approximately Correct* learning, connects three complexity measures – the sample space complexity, representation complexity, and optimization complexity, to the evaluation of the statistical learning process. It refers to the categorization of various learning procedure that satisfies the Probably Approximately Correct criteria in its name, as such being followed. The theory was proposed first by Valiant [1984], of whom then won the Turing Award, but his contribution stayed in the shadow in much of the end of the century, until the formal foundation of machine learning theory, of which then it becomes the stay ground on the standard discussion on the story of machine learner and algorithmic guarantees. As such, it is the de-facto foundation of statistical-based learning. The materials here based heavily on Mohri et al. [2012] and Kearns and Vazirani [1994], aside from interpretations.

For now, we assume no structure of h and c . They can be functions, partial functions, relations, complex algorithms, or others. In a typically learning setting, we also have the argument of *preliminary knowledge*, presented in literatures of the term *inductive bias*. For example, if the concept class $\mathcal{C}$ is assumed, then we say the setting is **model-specific**. If there exists no hard assumption on $\mathcal{C}$, then the learning setting is said to be **model-free**. For clarification, we also use in the examples the 0–1 loss function, $\ell = \delta(h, c)$ for the Kronecker delta, such that the expected empirical risk is calculated as

$$\hat{R}(h) = \frac{1}{n} \sum_{i=1}^n \delta_{h(\mathbf{x}_i) \neq y_i}, \quad \delta_{h(\mathbf{x}_i) \neq y_i} = \begin{cases} 1, & h(\mathbf{x}_i) \neq y_i \\ 0, & \text{o.w.} \end{cases} \quad (2.14)$$

We can later on replace δ and the binary setting with others, albeit the complexity will increase on the formal derivation for others loss functions. One of the main reason why we often choose 0 – 1 loss error in our analysis, comes from a variety of reasons. But, from what can be seen, it reduces the complexity of the problem, as well as raises the absoluteness of the loss measure. Indeed, naturally, from a soft point of view, the mean square loss measure is more natural of $\mathcal{L}_2$ norm on a representation space. As we have been saying, the loss function is chosen can influence the problem setting that one might expect the model to act on. 1 – 0 measure effectively reduces the problem to the absolute *binary classification problem*, where it is either correct or not.

The *learning algorithm* $\mathcal{A}$ can be interpreted to belong to a specific set of algorithm $\mathcal{A}$ such that it gives a sequence $\{A_i\}$. If there exists such that $|\{A_i\}| = n$ for a fixed n , and can be explicitly defined and

expressed, then the algorithm is *deterministic*. Otherwise for arbitrary $m > 1$ varying in magnitude, and with added uncertainty – that is, no sequence $\{A_i\}$ is the same, the algorithm is said to be *probabilistic*, and by extension, a *stochastic process*. We will give the definition of the stochastic process later on. Then, the algorithm is denoted by

$$\mathcal{A} = [h]\{A_1, \dots, A_n\} = \{A_i\}_{i=1}^n$$

Under this line, we separate the classification of c into two types: model-specific means that c is deterministic – it is supposed to belong to certain class $\mathcal{C}$, or rather, its description can be contained into certain concept class (which ultimately is some priori). Model-free means that c is not apparent to be able to separated to a concept class, which adds uncertainty, though all of them are still supposed to be at least a transformation $c : V \rightarrow U$ of arbitrary space V, U , by the setting of mathematical functional relationship.

Definition 2.4.1 (Model-specific learning). A learning scenario is called *model-specific* learning if, for $R(h)$ being a measure-theoretic metric on the hypothesis h , we are given the dataset $D = \{(\mathcal{X}, \mathcal{Y})\}$ with distribution $x \sim \mathcal{D}$, problem is to find a hypothesis $h \in \mathcal{H}$ that satisfies:

model-specific
learning

$$R(h) = \mathbb{P}_{x \sim \mathcal{D}}[h(x) \neq c(x)] = \mathbb{E}_{x \sim \mathcal{D}}[1_{h(x) \neq c(x)}] \quad (2.15)$$

In the same way, we define the model-free learning, where the distribution is extended to $\mathcal{Y}$. This is to facilitate the uncertainty of the concept class. If the concept class $\mathcal{C}$ is deterministic, then there would be no ambiguity of the concept $c \in \mathcal{C}$ – we would only have to learn $\mathbb{P}(y | x)$, to predict or reconstruct accordingly. In such case, the only factor in consideration is the ground space $\mathcal{X}$, or rather, the relationship between x and y itself.

model-free
learning

Definition 2.4.2 (Model-free learning). A learning scenario is called *model-free* learning if, for $R(h)$ being a measure-theoretic metric on the hypothesis h , we are given the dataset $D = \{(\mathcal{X}, \mathcal{Y})\}$ with distribution $(x, y) \sim \mathcal{D}$. Problem is to find a hypothesis $h \in \mathcal{H}$ that satisfies

$$R(h) = \mathbb{P}_{(x,y) \sim \mathcal{D}}[h(x) \neq y] = \mathbb{E}_{(x,y) \sim \mathcal{D}}[1_{h(x) \neq y}] \quad (2.16)$$

In both setting, we take the stance of absoluteness, using measure $\neq$, similar to a type of discrete decision process called *classification*, for a finite class counts of the target class. This notion can be aggravated to be generalized to other classes, simply by replacing $h(x) \neq y$ by a valued-function, or by considering specific measure (for example, *Bregman measure* for strictly convex functions). In the context of n -dimensional Euclidean space, and with the isomorphism $\mathbf{E} \simeq \mathbb{R}^n$ of suitable basis, the measure is simply the finite-accuracy Euclidean distance $d(h(x), c(x))$.

Before analysing this problem in a PAC-like fashion, let us recall the setting in which we are operating upon.

- The goal of the learning game is to learn an unknown target set, but the target set is not arbitrary. Instead, there is a known and rather strong constraint on the target set – it is a rectangle in the plane whose sides are parallel to the axes. Under certain analysis, we can call this as having indeed a strong *priori* to the concept $c \in \mathcal{C}$ possible.
- Learning occurs in a probabilistic setting. Examples of the target rectangle are drawn randomly in the plane according to a fixed probability distribution which is unknown and unconstrained.

- The hypothesis of the learner is evaluated relative to the *same* probabilistic setting in which the training takes place, and we allow hypotheses that are only approximations to the target. The tightest-fit strategy might not find the target rectangle exactly, but will find one with only a small probability of disagreement with the target.
- We are interested in a solution that is efficient: not many examples are required to obtain small error with high confidence, and we can process those examples rapidly – that is, for time efficiency t . For it to be controllable, it is in our interest to know if we can run the learning process in polynomial time, that is, with P-complexity class.

Ideally, we would want our model to be able to do the following:

- The number of calls to the concept c , or by extension, the oracle process $EX(c, \mathcal{D})$, is small, in the sense that it is bounded polynomially by some parameters.
- The amount of computation performed is small.
- The algorithm outputs a hypothesis h such that Θ is small. Recall that $\Theta_h = (\hat{R}(h), R(h))$, but generally, the ideal goal would be to reduce to the infimum of $R(h)$.

We are now ready to specify the definition of Probably Approximately Correct learning. For now, we designate this as a preliminary definition, because we will add more features into it.

Definition 2.4.3 (PAC Learning, Preliminary Definition). Let $\mathcal{C}$ be a concept class over X . We say that $\mathcal{C}$ is **PAC-learnable** if there exists an algorithm $\mathcal{L}$ that for every concept $c \in \mathcal{C}$, for every distribution $\mathcal{D}$ on X , and for all $0 < \epsilon < 1/2$ and $0 < \delta < 1/2$, if $\mathcal{L}$ is given access to $EX(c, \mathcal{D})$ for its output, and inputs ϵ, δ , then with probability at least $1 - \delta$, outputs a hypothesis $h \in \mathcal{H}$ satisfying $R(h) \leq \epsilon$. This probability is taken over the random examples drawn by calls to $EX(c, \mathcal{D})$, and any internal randomization of $\mathcal{L}$.

If $\mathcal{L}$ further run in time polynomial in $1/\epsilon$ and $1/\delta$, we say that $\mathcal{C}$ is *efficiently PAC learnable*. We will sometimes refer to the input ϵ as the **error parameter**, and the input δ as the **confidence parameter**.

The hypothesis $h \in \mathcal{H}$ of the PAC-learning algorithm is thus approximately correct, with high probability, hence we got the name. Usually, $\mathcal{H}$ can not coincide with $\mathcal{C}$. However, a stronger form of the preliminary PAC-learning can set $h \in \mathcal{C}$, which guarantee the existence of a near-perfect hypothesis class.

For the preliminary PAC learning model, there are two important problems with it, or rather, remarks on how it is constructed.

First, the parameter pair ϵ, δ control two types of failure to which a learning algorithm in the PAC model is inevitably susceptible. The error parameter ϵ is necessary since there may be only a negligible probability that a small random sample will distinguish between two competing hypotheses that differ on only one improbable point in the instance space. The confidence on only one improbable point in the instance (ground) space. The confidence parameter δ is necessary since the learning algorithm may occasionally be extremely unlucky, and draw a terribly "unrepresentative" sample of the target concept – for example, a sample consisting only of repeated draws of the same instances. The best we can hope for is that the probability of both types of failure can be made arbitrarily small at a modest cost.

Second, we notice that we demand a PAC learning algorithm perform well with respect any distribution $\mathcal{D}$. This is a very strong requirement, which is moderated by the fact that we only evaluate the hypothesis of the learning algorithm with respect to the same distribution $\mathcal{D}$. Whether this is reduced or not, depends on the shape of which the distribution is induced via determination of distribution structure.

2.4.1 Representation size and PAC-configuration

An important issue that has not been mentioned in our previous formation of PAC learning, is the fundamental distinction between a concept and its representation.

Consider a class of concepts defined by the satisfying assignments of certain formulae. A concept from this class that satisfies such formulae, can be represented by a formula f , a truth table, or any given formulae that is tautologically equivalent formulae f' to f . An example can be shown in high-dimensional Euclidean space $\mathbb{R}^n$, we may choose to represent a convex polytope (the hell is this?), we can specify the problem by either specifying its vertices, or specifying linear equations for its faces, and these two representation scheme, while being of the same problem, can differ in size.

In each of these examples, we are fixing some representation scheme – that is, a precise method for encoding concepts – and then examining the size of the encoding for various concepts. Other natural representation schemes that we are familiar it could be the vanilla neural network, or decision trees, though we cannot say exactly that the representation of neural network is indeed of the ‘vanilla’ representation. Taken example of the boolean formulae, for example, the boolean parity function $f(x_1, \dots, x_n) = x_1 \oplus \dots \oplus x_n$ which can be computed by a circuit of $\wedge, \vee$ and $\neg$ gates whose size is bounded by a fixed polynomial in n , but to represent this same function as a disjunctive normal form (DNF), requires size exponential in n . In these representation schemes, there is an obvious mapping from the representation (can be decision tree, or neural network) to the set or boolean function that is being represented. There is also a natural measure of the size of given representation in the scheme (for example, the number of weights, or neuron in a neural network).

representation
scheme

Since our PAC-learning algorithm for its model, follows the phenomenological interpretation, hence it is experimental data-only, it has absolutely no information about which, if any, of teh many possible representations is actually being used to represent the target concept in reality. However, it matters greatly which representation the algorithm chooses for its hypothesis, since the time to write this representation down is obviously a lower bound on the running time of the algorithm. At the end though, note that this is purely a computational concern at a glance, yet, further analysis into the problem might stem out certain consideration that previously unseen, for example, that fact that the representation scheme actually matters more in sense of a system analysis and its dynamic.

Formally, a *representation scheme* for a concept class $\mathcal{C}$ is defined as followed to capture this notion of representation size, and can be defined as below.

Definition 2.4.4 (Representation scheme). For a given concept class $\mathcal{C}$, a **representation scheme** for such concept class is a function $\mathcal{R} : \Sigma^* \rightarrow \mathcal{C}$, where Σ is the finite alphabet of symbols. We call any configuration $\sigma \in \Sigma^*$ such that $\mathcal{R}(\sigma) = c$ a **representation** of c , under $\mathcal{R}$. For any given c , the representation scheme might not be unique.³

To capture the notion of representation size, we assume that associated with $\mathcal{R}$, there is a mapping size: $\Sigma^* \rightarrow \mathbb{N}$ that assigns a natural number $size(h)$ to each representation $h \in \Sigma^*$. Note that we allow $size(\cdot)$ to be any such mapping results obtained under a particular definition for $size(\cdot)$ will be meaningful only if this definition is natural. In a more simple case, and arguably realistic setting, we can take $\Sigma = \{0, 1\}$, thus, we have a binary encoding of concepts, and define $size(h)$ to be the length of h

³Perhaps representation scheme can be further realized by applying a Borel σ -algebra on top, and put on it a perspective measure to handle the notion of representation size. For example, in the traditional example of axis-aligned rectangle, one such representation scheme could be the real number schema $\mathcal{R} : (\Sigma \cup \mathbb{R})^* \rightarrow \mathcal{C}$, which then can utilize the extended notion of a Lebesgue measure over $\mathbb{R}$ and its spaces.

in bits. Although we will use other definitions of size when binary representations are inconvenient, our definition of $size(\cdot)$ will always be within a polynomial factor of the binary string length definition.

So far, this definition is applicable only to representations, that is, to strings $h \in \Sigma^*$. We would like to extend this to $c \in \mathcal{C}$. Since the learning algorithm has access only to the input-output behaviour of c , in the worst case, it must assume that the simplest possible mechanism is generating this behaviour. Thus, we define $size(c)$ to be the infimum, that is:

representation
size

$$size(c) = \min_{\mathcal{R}(\sigma)=c} \{size(\sigma)\}$$

In other words, $size(c)$ is the size of the smallest representation of the concept c in the underlying representation scheme $\mathcal{R}$. While proceeding to the next section, a caution has to be made. Our notion of representative complexity lies fully within compressible string-expression system. That is, a model is reduced to its *constructive description*, of which are records and protocols written on the computable tape. However, dynamics of the computation system lies within the changing state, thus there exists a blind gap in between this description of script length, and the actual instantiation. Rather, what would be present in a modern treatment, would be the initial-value-similar, *static representational complexity* of a hypothesis class or concept class instead, with them being extension for the governed dynamical motions in dynamics, often ruled by algorithms or any specific active routines.

Definition 2.4.5 (Representation complexity). Given a concept $c \in \mathcal{C}$, the **representation complexity** of c is defined to be the size of the smallest representation of the concept c in the underlying representation scheme $\mathcal{R}$, that is,

$$size(c) = \min_{\mathcal{R}(\sigma)=c} \{size(\sigma)\}$$

for any string σ of the representation scheme $\mathcal{R}$ that still represents c .

Under such definition, the more "complex" the concept c is with the respective chosen representation scheme, the larger $size(c)$ is. This is also where the concept of the concept class $\mathcal{C}$ comes from – it comes from the *representation class induction* that we have in mind some fixed concept classes we study by their representation scheme.

From this consideration of the representation size for the language in which the problem setting is expressed in, we gain the complete definition of PAC-learning for a fixed representation priori of the concept class.

Definition 2.4.6 (PAC-learning). A concept class $\mathcal{C}$ is said to be **PAC-learnable** if there exists an algorithm $\mathcal{A}$ and a polynomial function $poly(\cdot, \cdot, \cdot, \cdot)$ of 4-argument such that for any $\epsilon > 0$, $\delta > 0$, for all distribution $\mathcal{D}$ on $\mathcal{X}$ and for any target concept $c \in \mathcal{C}$, the following holds for any sample size $m \geq poly(1/\epsilon, 1/\delta, n, size(c))$:

$$\Pr_{S \sim \mathcal{D}^m} [R(h(S)) \leq \epsilon] \geq 1 - \delta$$

for a given error measure $R(h_S)$. If $\mathcal{A}$ further runs in $poly(1/\epsilon, 1/\delta, n, size(c))$, then $\mathcal{C}$ is said to be *efficiently PAC-learnable*. When such algorithm exists, we call $\mathcal{A}$ a PAC-learning algorithm.

Finally, to compress the observations above. Current system operates within expression on binary systems, or rather computable substrate on $(0, 1)^n$ string representation. Each string as such, encodes specific details, though not perfect, of the concept class within frame Q . Hence, any concept representation c' in practice, now that we separate between c and c' , can be considered exhibiting its structural details in

different aspects from the dataset within such encoding. However, most of our construction in enumerative and numeral domains are *of depth*, in the sense that the frame is rather complex and large to begin with. In most circumstances, the operating frame Q_f for the financial frame, as such, encodes macroeconomic and certain simplification techniques, making it being sophisticated internally yet miniscule global of the uncharted descriptions. Natural extension to the question of general computing substrate than this, and whether or not dataset can be meaningfully and when constructed with full accuracy to the concept class, remains the open question, but generally efforts are fragmented and small (see inductive inference and extension for automata-based system and further on - we note that Turing machine is a rather poor choice in this problem, as there exists other representation of Turing machine that would not omit relevant details). Secondly, the algorithm constrain on polynomial factor $size(\cdot)$ is a polynomial equivalence condition. Different reasonable size measures as such are computationally equivalent, in the sense that *generalizing in-between algorithm class* is possible, such that a polynomial-time algorithm under one size measure is a polynomial-time algorithm under any other, and the reverse on the algorithm class holds. This means that PAC learnability results, are invariant to the choice of size measure up to polynomial factors. This transfer across reasonable representation schemes, and thus condition a rather strong assumption about the description space, especially pre-dynamics. Most of the issues here, from identification of Layer II to Layer III, can be seen almost as artefact from Layer I to Layer II transition and choice.

2.4.2 Generalization bounds

The theory based its goal around the target of **generalization learning problem**, hence, most of the bounds can be focus on achieving generalization target. Superficially, we assume the following is true of specific h when trying to resolve the generalization problem.

Assumption 2.4.1 (Structural assumptions). *We assume that $h \in \mathcal{H}$ is expressive enough on the hypothesis class, and algorithm $\mathcal{A}(h)$ is capable enough to extrapolate h^* of the empirical best on the empirical space $\mathcal{S}$.*

Under such assumption, to introduce the explicit objective of generalization bounds, we first need to introduce the categorization of the hypothesis based of one criterion: consistency. We say that an algorithm is a consistent learner for a concept class C using hypothesis class $\mathcal{H}$, if for all n , for all $c \in C_n$ and for all m , given $\mathcal{S}$, as input, $x_i \in X_n$ outputs, then the algorithm L outputs a hypothesis $h \in \mathcal{H}_n$ such that $h(x_i) = c(x_i), i = 1, \dots, m$. Then h is called the **consistent learner**. In other word, a consistent hypothesis is the hypothesis that is consistent with all the training data provided to it from the concept c . The following section contains the debate between consistent hypothesis (i.e., for example, admitting no error on training data, almost perfect - maybe even so), and inconsistent hypothesis, of which have general defects from the actual concept, or even different by a margin.

Famous results contain two types of bound for classical PAC: the finite $\mathcal{H}$, and consistent hypothesis, and finite $\mathcal{H}$, but inconsistent. Infinite $\mathcal{H}$ assumes different stochasticity, and in consistent bound, generalization depends on the variances between empirical and generalization space, which is not covered in the theory, though in general setting the metric-induced distance can be small and controllable, which makes PAC applicable on said domain. For consistent case, we will consider the general sample complexity bound, or equivalently, a generalization bound for consistent hypothesis in the case where the cardinality $|\mathcal{H}|$ is finite. We will assume, as such, that the target concept c is in $\mathcal{H}$.

Theorem 2.4.2 (Learning bound - finite $\mathcal{H}$, consistent case). *Let H be a finite set of functions mapping from $\mathcal{X} \rightarrow \mathcal{Y}$. Let $\mathcal{A}$ be an algorithm that for any target concept $c \in H$ and i.i.d. samples S returns a consistent*

hypothesis H_S , such that $\hat{R}(h_S) = 0$. Then for any $\epsilon, \delta > 0$, the inequality $\Pr_{S \sim D^m}[R(h_S) \leq \epsilon] \geq 1 - \delta$ holds if

$$m \geq \frac{1}{\epsilon} \left(\log |H| + \log \frac{1}{\delta} \right)$$

This sample complexity result admits the following equivalent statement as a generation bound: for any $\epsilon, \delta > 0$, with probability at least $1 - \delta$,

$$R(h_S) \leq \frac{1}{m} \left(\log |\mathcal{H}| + \log \frac{1}{\delta} \right) \quad (2.17)$$

Theorem 2.4.3 (Learning bound – finite $\mathcal{H}$, inconsistent case). *Let $\mathcal{H}$ be a finite hypothesis set. Then, for any $\delta > 0$, with probability at least $1 - \delta$, the following inequality holds:*

$$\forall h \in \mathcal{H}, \quad R(h) \leq \hat{R}_S(h) + \sqrt{\frac{\log |\mathcal{H}| + \log 2/\delta}{2m}} \quad (2.18)$$

The proofs are left in the Appendix section. In general, theorem 2.4.2 shows that when the hypothesis set $\mathcal{H}$ is finite, a consistent algorithm $\mathcal{A}$ is a PAC-learning algorithm, since the sample complexity is dominated by a polynomial in $1/\epsilon$ and $1/\delta$. The generalization error of consistent hypothesis is upper bounded by a term that decrease w.r.t m . The decreases rate of $O(1/m)$ is guaranteed by this theorem. For the inconsistent theorem 2.4.3, it indicates that for a finite hypothesis set $\mathcal{H}$, then

$$R(h) \leq \hat{R}_S(h) + O \left(\sqrt{\frac{\log_2 |\mathcal{H}|}{m}} \right) \quad (2.19)$$

The term $\log |\mathcal{H}|$ can be interpreted as the number of bit needed to represent $\mathcal{H}$. A larger sample size m guarantees better generalization, and the bound increases with $|\mathcal{H}|$, but only logarithmically. Note that the bound suggests seeking a trade-off between reducing the empirical error versus controlling the size of the hypothesis set: a larger hypothesis set is penalized by the second term but could help reduce the empirical error, that is the first term. But for a similar empirical error, it suggests using a smaller hypothesis set. This can be viewed as an instance of the so-called *Occam's Razor*, of which can be said as: *Plurality should not be posited without necessity*, or, the simplest explanation is best.

2.4.3 Agnostic-general PAC-learning

For the definition of PAC-learning in the above setting, for model-specific notion, there are some issues with it, more specifically, with the type of *strong assumptions* that it holds.

1. The *representation size* and boundary only works for priori given of the concept class $\mathcal{C}$. In case that we cannot determine this, and there are a large potential for that to fail in capturing the minimal representation scope of the concept, the bound simply fails to a certain extent – given that our definition includes the polynomial bound of the concept class itself.
2. The assumption that the target concept c belongs to $\mathcal{C}$ means that we are trying to fit a hypothesis to data, which are a priori known to have been generated by some member of the model class defined by $\mathcal{C}$. However, in general, we may not want to assume much about the data generation process, and instead would like to find the best fit to the data at hand using an element of some model class of our choice.

3. The assumption that the training features are labelled noiselessly rules out the possibility of noisy measurements or observations.
4. Even if the above assumptions were somehow true, we would not necessarily have a priori knowledge of the concept class $\mathcal{C}$, containing the prior knowledge of the concept class $\mathcal{C}$ containing the target concept or function. In that case, the best we could hope for is to pick our own model class and seek the best approximation to the unknown target concept among the elements of that class.

Agnostic case is considered the most general scenario of supervised learning, where the distribution $\mathcal{D}$ is defined over $\mathcal{X} \times \mathcal{Y}$, and the training data is a labelled sample S drawn i.i.d. according to $\mathcal{D}$:

$$S = \{(x_1, y_1), \dots, (x_m, y_m)\}$$

The learning problem is to find a hypothesis $h \in \mathcal{H}$ with small generalization error

$$R(h) = \mathbb{P}_{(x,y) \sim \mathcal{D}}[h(x) \neq y] = \mathbb{E}_{(x,y) \sim \mathcal{D}}[1_{h(x) \neq y}] \quad (2.20)$$

This more general scenario is referred to as the stochastic scenario. Within this setting, the output target is a probabilistic function of the input. The stochastic scenario captures many real-world problems where the label of an input point is not unique⁴. That said, the case of model-free is can be constructed in a simple manner, where we forgo any assumptions on the behaviours or structures of $\mathcal{C}$ and its associated c -model, and thus the observable pairs is simply $(x, y) \sim \mathcal{D}_{mf}$. Here, we denote $\mathcal{D}_{mf}$ to separate it from $\mathcal{D}$ of which is a singular x -distribution; however, we can also write $\mathcal{D}_{ms}$ and $\mathcal{D}_{mf}$ in certain specific cases. Per [Hajek and Raginsky \[2021\]](#), the treatment of model-agnostic (no model c assumption), targets the ingredient sets of,

1. Sets $\mathcal{X}, \mathcal{Y}$ and $\mathcal{U}$. Here, $\mathcal{U}$ is the correlated h -functional result, i.e. the hypothesis model, covering the case of **probabilistic models**.
2. A class $\mathcal{P}$ of probability distribution on $Z := \mathcal{X} \times \mathcal{Y}$. Here, the distribution covers both.
3. A class $\mathcal{H}$ of functions $h : \mathcal{X} \rightarrow \mathcal{U}$.
4. A loss function as objective, $\ell : \mathcal{Y} \times \mathcal{U} \rightarrow [a, b]$, preferably $[0, 1]$. Additional overhead measures are also required.

The *agnostic PAC-learning* format is then modified by considering $\mathcal{Y}$ as also a random variability, and thus can be stated as followed.

Definition 2.4.7 (Agnostic PAC-learning). A concept class $\mathcal{C}$ is said to be **PAC-learnable** if there exists an algorithm $\mathcal{A}$ and a polynomial function $poly(\cdot, \cdot, \cdot, \cdot)$ of 4-argument such that for any $\epsilon > 0$, $\delta > 0$, for all distribution $\mathcal{D}$ on $\mathcal{X} \times \mathcal{Y}$ and for any target concept $c \in \mathcal{C}$, the following holds for any sample size

⁴Considering two set X and Y . A probabilistic function from X to Y assign to each $x \in X$ a relevant sub-distribution of elements of Y , rather than a single value $y \in Y$. Such a sub-distribution is a function $\delta : Y \rightarrow [0, 1]$ such that

$$\sum_{y \in Y} \delta(y) \leq 1$$

In other words, a probabilistic function $X \rightarrow Y$ is a function $f : X \rightarrow (Y \rightarrow [0, 1])$, interpreted as $f : X \times Y \rightarrow [0, 1]$ such that for all $x \in X$,

$$\sum_{y \in Y} f(x, y) \leq 1$$

$m \geq \text{poly}(1/\epsilon, 1/\delta, n)$:

$$\Pr_{S \sim \mathcal{D}^m} \left[R(h(S)) - \min_{h \in \mathcal{H}} R(h) \leq \epsilon \right] \geq 1 - \delta \quad (2.21)$$

If $\mathcal{A}$ further runs in $\text{poly}(1/\epsilon, 1/\delta, n)$, then $\mathcal{C}$ is said to be *efficiently agnostic PAC-learnable*. When such algorithm exists, we call $\mathcal{A}$ an agnostic PAC-learning algorithm.

Under such lens, PAC-on-agnostic can be realized also using a more enumerative bound-on-definition instead:

Definition 2.4.8 (agnostic-PAC learning, Hajek and Raginsky [2021]). We say that a learning algorithm for a problem $(\mathcal{X}, \mathcal{Y}, \mathcal{U})$ and interpretations $(\mathcal{P}, \mathcal{H}, \ell)$ is PAC-learnable to accuracy $\epsilon > 0$, if, we have

$$\lim_{n \rightarrow \infty} \sup_{P \in \mathcal{P}} P^n \left(Z^n : \hat{R}_P(h_n) > \inf_{h \in \mathcal{H}} R_P(h) + \epsilon \right) = 0 \quad (2.22)$$

An algorithm that is PAC-learnable to accuracy ϵ for every $\epsilon > 0$ is said to be a PAC-learning algorithm. A learning problem is said to be model-free learnable if there exists such algorithm, for every $A \in \mathcal{A}$ of such class to be PAC-learning algorithm class.

PAC-learning bound argues in the sense of *computational complexities*, and *space complexities*. Specifically, the term $\text{poly}(1/\epsilon, 1/\delta, n, \text{size}(c))$ hopes to bound the required learning process to polynomial time, specified by 4 parameters. Here, n is the *input dimension*, $\text{size}(c)$ is the encoding complexity of the target concept, $\epsilon > 0$ is the *accuracy bound*, and δ is the confidence bound – the probability that h fails.

2.4.4 Empirical and structural algorithms

Suppose that for a concept class $\mathcal{H}$, we can then partition it to $\mathcal{H} = \bigcup_{k \in \mathbb{N}} \mathcal{H}_k$ for finite k hypotheses. Assume that they satisfy the uniform convergence (Shalev-Shwartz et al. [2010]), which means in the arbitrary path landscape, one can guarantee ℓ -reduction, of m samples on $\mathcal{S}$ to:

$$\lim_{m \rightarrow \infty} \sup_{\mathcal{D}} \mathbb{E}_{S \sim \mathcal{D}^m} \left[\sup_{h \in \mathcal{H}} [R(h) - R_S(h)] \right] = 0 \quad (2.23)$$

or, can also be expressed of equal to the following expression, for functions ϵ_k satisfies that for all k and $\delta \in (0, 1)$, $\lim_{m \rightarrow \infty} \epsilon_k = 0$.

$$\forall k \in \mathbb{N}, \quad \Pr_{S \sim \mathcal{D}^m} \left(\sup_{h \in \mathcal{H}_k} R(h) - \hat{R}_S(h) \leq \epsilon_k(m, \delta) \right) \geq 1 - \delta \quad (2.24)$$

Then such is said to be possible as a ruling dynamic when for a learning problem, the empirical risks of hypotheses in the hypothesis class converges to their population risk uniformly, with a distribution-independent rate. The result of this is that a problem then can be considered learnable with the ERM rule, for any given characterization. We state the following definition.⁵

⁵The reason why $r(\cdot)$ is not characterized in this preliminary section, is because of their uncertainty of effect. Hypothetically, $r(\cdot)$ is an additional structure (bias) of which apply a new gradient field in conjunction to the loss gradient field itself. This addition is uncertain of its content and analysis, since it depends on various stochastic dynamics of the model. We reserve from using regularization for now because of such reason.

Definition 2.4.9. A learning problem is learnable if there exist a learning rule $\mathcal{A}$ and a monotonically decreasing sequence $\epsilon_{cons}(m)$ such that $\epsilon_{cons}(m) \xrightarrow{m \rightarrow \infty} 0$ and

$$\forall \mathcal{D}, \quad \mathbb{E}_{S \sim \mathcal{D}^m} [F(\mathcal{A}(S)) - F^*] \leq \epsilon_{cons}(m), \quad (2.25)$$

with $F^* = \inf_{h \in \mathcal{H}} R(h)$. A learning rule $\mathcal{A}$ for which this hold is denoted as a universally consistent learning rule.

There are certainly a fundamental class of rule or algorithm $\mathcal{A}$ to resolve this particular uniform convergence of classes. One of the major assumption or observation from the point of the hypothesis class, is that the hypothesis cannot know the generalization error. Thereby, one strategy is the *empirical risk minimization* (ERM) method, minimizing only the empirical risk to the empirical best; then, certain measures to ‘generalize’ this heuristically is apprehended on the hypothesis.

Definition 2.4.10. A rule $\mathcal{A}$ is an ERM (Empirical Risk Minimizer), denoted $\mathcal{A}_{ERM}$ if it minimizes the empirical risk $\hat{R}(h) = (1/m) \sum_{c \in \mathcal{C}} \ell(h, c)$, over the observational space S ,

$$\hat{R}_S(\mathcal{A}_{ERM}(S)) = \hat{R}_S(h_S) = \inf_{h \in \mathcal{H}} \hat{R}_S(h) \quad (2.26)$$

Hence, the ERM tries to obtain $\text{ERM}(h') = \arg \min_{h \in \mathcal{H}} \hat{R}(h)$.

The problem of overfitting still somewhat resides in the setting of the method. Hence, usually, we have an implicit term in addition, called the *regularizer*, of which is as followed.

Definition 2.4.11. A rule $\mathcal{A}$ is an SRM (Structural Risk Minimizer), denoted $\mathcal{A}_{SRM}$ if it minimizes the empirical risk $\hat{R}(h) = (1/m) \sum_{c \in \mathcal{C}} \ell(h, c)$, over the observational space S ,

$$\hat{R}_S(\mathcal{A}_{SRM}(S)) = \hat{R}_S(h_S) + \lambda r(h) = \inf_{h \in \mathcal{H}} \hat{R}_S(h) + \lambda r(h) \quad (2.27)$$

of the regularizing term $r(\cdot)$. Hence, the SRM tries to obtain $\text{SRM}(h') = \arg \min_{h \in \mathcal{H}} \hat{R}(h) + \lambda r(h)$, while taking into consideration a forcing term r , usually defined on the model complexity.

Hence, the SRM is inherently stronger than ERM. The regularizer $\lambda r(h)$, in specific case of convex loss landscape and reasonably well-behaved irregularities between domains, and of the structural regularization term $r_s(h) = \sum_{i,j} w_{ij}^2$ for weights of ij index, then we have the following theorem.

Conjecture 2.4.1 (ERM-SRM growth). Consider the regularizer $\lambda r(h)$ for $h \in \mathcal{H}$. Assuming structural regularization $r_S(h) = \sum_{i,j} w_{ij}^2$ for linear-layering of ij index, for ERM and SRM optimizer classes,

$$\begin{aligned} & \text{ERM} \left(\inf_{h \in \mathcal{H}} \hat{R}_S(h), \ell \right) \\ & \approx o \left(\text{SRM} \left(\inf_{h \in \mathcal{H}} \hat{R}_S(h) + \lambda r(h), [\ell, \lambda r(h)] \right) \right) \end{aligned} \quad (2.28)$$

Where the little- o notation is used for saying of the optimization strength per arbitrary time-similar notion of algorithmic evolution. The reason for this change is because of the restriction domain that SRM indicted on the model optimization landscape. Effectively, we introduce the structural bias $\lambda r(h)$

that forces the morph of the total loss landscape. In well-behaved case, this is usually effective to a point, however in realistic system and pathological cases, such can lead to the c' approximation in c -space, just as we introduced in the reformulation of learning system. In this paper, we would like to presume the ERM method in question as the famous *gradient descent*. We refer to Achlioptas, Ruder [2017] for general view of the algorithm, and Zhang [2019] for a source on particularly gradient descent algorithm on deep learning models.

One particularly troublesome and should take serious consideration into interpreting it, is the fact of the insufficiency on ERM. The following proposition is rather questionable.

Proposition 2.4.4. *For any sample S , the following inequality holds for the hypothesis returned by ERM:*

$$\mathbb{P}\left[R(h'_S) - \inf_{h \in \mathcal{H}} R(h) > \epsilon\right] \leq \mathbb{P}\left[\sup_{h \in \mathcal{H}} |R(h) - \hat{R}_S(h)| > \frac{\epsilon}{2}\right] \quad (2.29)$$

We will prove this proposition in later section. For now, this proposition bounds the error of ERM, with respect to the probability of the generalization distance $|R(h) - \hat{R}_S(h)|$ between the two errors. As we might have suspected, the performance of ERM is typically very poor. This is because the algorithm disregards the complexity of the hypothesis set $\mathcal{H}$: in practice, either it is not complex enough, in which case the error can be large, or $\mathcal{H}$ is so rich, that the complexity makes the model shoot out of the complexity range. Additionally, in many cases, determining the ERM solution is computationally intractable.

2.4.5 The PAC-ERM notion

Without restricting ourselves to specific ERM construction, we can still apply PAC on the set of ERM algorithms in its generality of itself. Within such understanding, the empirical risk minimization strategy is often motivated by the hope that $R(h)$ and $\hat{R}(h)$ are not too far-distant apart. Particularly, it also implicitly agree with statements conjectured with law of large number, for $n \rightarrow \infty$, $|R(h) - \hat{R}(h)| \approx \epsilon$ of irreducible noise. To what extent is this true, we have to take certain boundary on which might have been previously shown. We refer to authoritative introductory texts like Alquier [2024], for both PAC-ERM and later sections on PAC-Bayes variation.

Proposition 2.4.5 (Alquier [2024]). *For any $h \in \mathcal{H}$, for any $\delta \in (0, 1)$,*

$$\mathbb{P}_S \left(R(h) > \hat{R}(h) + C \sqrt{\frac{\log(1/\delta)}{2n}} \right) \leq \delta \quad (2.30)$$

Proof. Applying Hoeffding's inequality onto $U = \mathbb{E}[\ell_i(h)] - \ell_i(h)$,

$$\mathbb{E}_S \left[\exp tn[R(h) - \hat{R}(h)] \right] \leq \exp [(nt^2 C^2)/8] \quad (2.31)$$

For any $s > 0$, we have by Markov's inequality,

$$\begin{aligned} \mathbb{P}_S(R(h) - \hat{R}(h) > s) &= \mathbb{P}_S(\exp nt[R(h) - \hat{R}(h)] > \exp(nts)) \\ &\leq \frac{1}{\exp(nts)} \mathbb{E}_S \left[e^{nt[R(h) - \hat{R}(h)]} \right] \\ &\leq \exp \left(\frac{nt^2 C^2}{8} - nts \right) \end{aligned} \quad (2.32)$$

We can make this bound as tight as possible, by optimizing the choice of t . For $t = 4s/C^2$, we have

$$\mathbb{P}(R(h) - \hat{R}(h) > s) \leq \exp\left(\frac{-2ns^2}{C^2}\right) \quad (2.33)$$

This means that, for a given h , the empirical risk $\hat{R}(h)$ cannot be much smaller than the risk $R(h)$. The order of this "much smaller" can be better understood by introducing $\delta = \exp(-2ns^2/C^2)$, of which substituting δ to s yields the required proposition. $\square$

Proposition 2.4.5 bounds the asymptotic term such that $R(h)$ will exceed $\hat{R}(h)$ by certain margin, at worst $1/\sqrt{n}$. However, this is not enough to justify the use of ERM (which indeed use the near-space approximation from empirical to general risk), since the proof only applies for h that was fixed above in terms of complexity. For those that is implicit of the data itself, we cannot apply them⁶. The usual approach though, is to use the inequality

$$R(h_{ERM}) - \hat{R}(h_{ERM}) \leq \sup_{h \in \mathcal{H}} [R(h) - \hat{R}(h)] \quad (2.34)$$

and to prove a version of Proposition 2.4.5 that would hold uniformly on $\mathcal{H}$.

Theorem 2.4.6. Assume that $\text{card}(\mathcal{H}) = M < +\infty$, for any $\delta \in (0, 1)$,

$$\mathbb{P}_S \left(R(h_{ERM}) \leq \inf_{h \in \mathcal{H}} \hat{R}(h) + C \sqrt{\frac{\log(M/\delta)}{2n}} \right) \geq 1 - \delta \quad (2.35)$$

Proof. We upper bound the supremum in the inequality by

$$\begin{aligned} \mathbb{P}_S \left(\sup_{h \in \mathcal{H}} [R(h) - \hat{R}(h)] > s \right) &= \mathbb{P}_S \left(\bigcup_{h \in \mathcal{H}} \{ [R(h) - \hat{R}(h)] > s \} \right) \\ &\leq \sum_{h \in \mathcal{H}} \mathbb{P}_S(R(h) - \hat{R}(h) > s) \\ &\leq M \exp\left(\frac{-2ns^2}{C^2}\right) \end{aligned} \quad (2.36)$$

This time, put $\delta = M \exp(-2ns^2/C^2)$ we gain

$$\mathbb{P}_S \left(\sup_{h \in \mathcal{H}} [R(h) - \hat{R}(h)] \leq C \sqrt{\frac{\log(M/\delta)}{2n}} \right) \leq \delta \quad (2.37)$$

Thus it satisfies

$$\mathbb{P}_S \left(\sup_{h \in \mathcal{H}} [R(h) - \hat{R}(h)] \leq C \sqrt{\frac{\log(M/\delta)}{2n}} \right) \geq 1 - \delta \quad (2.38)$$

⁶This means models such as *support vector machine* (SVM) or *Gaussian process regression* of which the empirical optimization function and the hypothesis itself, relies on the number of data points available.

Using the inequality

$$\mathbb{P}_S \left(R(h_{ERM}) \leq \hat{R}(h_{ERM}) C \sqrt{\frac{\log(M/\delta)}{2n}} \right) \geq 1 - \delta \quad (2.39)$$

As $\mathcal{H}$ is finite, $\hat{R}(h_{ERM}) = \inf_{h \in \mathcal{H}} \hat{R}(h)$. □

This bound coincides for the *Probably Approximately Correct* PAC-bounds, specifically for algorithm of ERM principle, since it approximates $R(h_{ERM})$ within $C \sqrt{\log(M/\delta)/2n}$ with probability $1 - \delta$, and once again, under very strong arguments. Certain version of this theorem's result can be derived to 'predict' double descent, but not giving its the diagnostic view.

2.4.6 Noise, Mohri et al. [2012]

noise Using the notion of the Bayes classifier and Bayes error, one can define the *noise* of a given learning setting under PAC-learning consideration, as followed.

Definition 2.4.12 (Noise). Given a distribution $\mathcal{D}$ over $\mathcal{X} \times \mathcal{Y}$, the *noise* at point $x \in \mathcal{X}$ is defined by

$$noise(x) = \min_{y \in \mathcal{Y}} \{\mathbb{P}[1 | x], \mathbb{P}[0 | x]\}$$

The average noise or the noise associated to $\mathcal{D}$ is then the minimum possible error aside from intrinsic model error, and is calculated by $\mathbb{E}[noise(x)]$

Thus, the average noise is precisely the Bayes error, that is, $noise = \mathbb{E}[noise(x)] = R^*$. The noise is a characteristic of the learning task indicative of its level of difficulty. A point $x \in \mathcal{X}$, for which noise is close to $1/2$, is sometimes referred to as *noisy* and is of course a challenge for accurate prediction. Such is understandable, and is often enough, however, we note that structurally, it is missing particular components or decomposition components that would eventually make up the bulks of its claims.

2.5 Simplistic learning - Occam's Razor

The PAC model that we introduced defined learning directly in terms of the predictive power of the hypothesis output by the learning algorithm. That is, given a hypothesis $h \in \mathcal{H}$ and the learning algorithm $\mathcal{L}$, the learning involves switching this dynamic of the predictive power of the hypothesis, over certain setting and problem. We can, as it is possible to apply this measure, to a learning algorithm because we made the assumption that the instances are drawn, from a probabilistic perspective, from a fixed distribution $\mathcal{D}$ then measured the predictive power with respect to the same distribution.

Our new notion, dubbed as *Occam learning* (our source is mainly taken from the exposition of Kearns and Vazirani [1994]), would let us consider a different definition of learning that makes no assumptions about how the instances in a labelled sample are chosen, though we will assume that the labels are generated by a target concept chosen from a known class. Instead of measuring the predictive power of a hypothesis, the new definition judges the hypothesis by how succinctly it explains the observed data (or a labelled example). The crucial difference between PAC learning and Occam learning, is then that in PAC learning, the random sample drawn by the learning algorithm is intended only as an aid for reaching an accurate model of some external process, while in the new definition, we are concerned only with the fixed sample before us, and not any external process. By that, we simply reduce it to somewhat similar of the statistical estimation problem setting, though not exactly the same.

This also contains the substance of which we often call the *Occam's razor*, telling that overly complex scientific theories should be subjected to a simplifying knife. If we equate "simplicity" with representational succinctness, then another way to interpret Occam's razor would be to say learning is the act of finding a pattern in the observed data that facilitates a compact representation or compression of this data. In our simple concept learning setting, succinctness is measured by the size of the representation of the hypothesis concept. Equivalently, we can measure succinctness by the cardinality of the hypothesis class used by the algorithm, for if this class is small then a typical hypothesis from the class can be represented by a short binary string, and if this class is large then a typical hypothesis must be represented by a long string. Thus, an algorithm is an *Occam algorithm* if it finds a short hypothesis consistent with the observed data.

Additionally so, it would be plenty interesting to see that Occam's razor somehow falls into the range of what would be expected of the *bias-variance tradeoff*. This perhaps, will prompt us to push further into the inquiry of double descent, which empirically, breaks the principle of the Lord William Occam.

2.5.1 Occam Learning and Succinctness

Let $X = \cup_{n \geq 1} X_n$ be the instance space, let $\mathcal{C} = \cup_{n \geq 1} \mathcal{C}_n$ be the target concept class, and let $\mathcal{H} = \cup_{n \geq 1} \mathcal{H}_n$ be the class of hypothesis representation. The notation is preferably clarified: X_n is under the binary representation typically concerned $\{0, 1\}^n$, or a boolean string of length n . Hence, $\mathcal{H}_n$ and $\mathcal{C}_n$ are subsequently concept and hypothesis class specified for such string space. In this part, we will assume, unless explicitly stated otherwise, that the hypothesis representation scheme of $\mathcal{H}$ uses a binary alphabet, and we define $size(h)$ to be the length of the bit string h . Also, recall that for a concept $c \in \mathcal{C}$, $size(c)$ denotes the size of the smallest representation of c in $\mathcal{H}$. Let $c \in \mathcal{C}_n$ denote the target concept. A labelled sample S of cardinality m is a set of pairs:⁷

$$S = \{(x_1, c(x_1)), \dots, (x_m, c(x_m))\}$$

An **Occam algorithm** L takes as input a labelled sample S , and outputs a short hypothesis h , "relatively so", that is consistent with S . By consistent, we mean as one definition above, $h(x_i) = c(x_i)$ for each i , and for short we mean that $size(h)$ is a sufficiently slowly growing function of n , $size(c)$ and m . This is formalised in the following definition.

Definition 2.5.1 (Occam algorithm). Let $\alpha \geq 0$ and $0 \leq \beta < 1$ be constant. L is an (α, β) -**Occam algorithm** for $\mathcal{C}$ using $\mathcal{H}$ if on input a sample S of cardinality m labelled according to $c \in \mathcal{C}_n$, L outputs a hypothesis $h \in \mathcal{H}$ such that:

- h is consistent with S , or $h(x_i) = c(x_i)$ for all $x_i \in S_x$.
- $size(h) \leq (n \cdot size(c))^\alpha m^\beta$

We say that L is an **efficient** (α, β) -Occam algorithm if its running time is bounded by a polynomial in n, m and $size(c)$.

The growth complexity $\mathcal{O}(n \cdot size(c))^\alpha m^\beta$ is somewhat difficult to scale down. However, we can somewhat try to prove if it is convergence or not.

Theorem 2.5.1. Fix n . Then $\mathcal{O}(n \cdot size(c))^\alpha m^\beta$ is convergence or divergence depends on the ratio of α and β .

⁷This section is taken, and studied, from a pretty old book. It would be then observed that a lot of the concept, including the representation, is taken in a rather computer science fashion such as for the hypothesis $\mathcal{H}$ to be binary representation.

Proof. Fixed n , then we have the class $\mathcal{C}_n$ to be finitely restricted of n representation class. Then, $\text{size}(c)$ for $c \in \mathcal{C}$ is also constant. We then reduce the first complexity measure to be $\mathcal{O}(\lambda^\alpha m^\beta)$. Since $0 \geq \beta < 1$, it follows that as β increase, $\text{size}(h)$ monotonically decrease. This is also the case for α in range $[0, 1)$. Hence, it converges to at least n , and diverges accordingly for arbitrary $\alpha > 1$, though the trend will not as strong if $m > n$, or $m \gg n$. $\square$

Now, in what sense can we say that the output h of an Occam algorithm succinct? First, let us assume that $m \gg n$, so that the above bound can be effectively simplified to $\text{size}(h) < m^\beta$, for some $\beta < 1$. Since the hypothesis h is consistent with the sample S , H allows us to reconstruct the m labels $c(x_1) = h(x_1), \dots, c(x_m) = h(x_m)$, and is given only the unlabelled sequence of instance $x_1, \dots, x_m$. Thus the m bits $c(x_1), \dots, c(x_m)$ have been effectively *compressed* into a much shorter string h of length at most m^β . Note that the requirement $\beta < 1$ is quite weak, since a consistent hypothesis of length $\mathcal{O}(mn)$ can always be achieved by simply storing the sample S in a table (at a cost of $n + 1$ bits per labelled example) and giving an arbitrary answer for instances that are not in the table. We would be certainly not expecting such a hypothesis to have any predictive power.

By sheer observation of the definition we have been constructing, the $(1, 0)$ -Occam algorithm is the largest deviation, with the bound $\text{size}(h) \leq n \cdot \text{size}(c)$. This means it can at most, have the representation size equal to *all* representation partition of $c \in \mathcal{C}$. It is also good to note the misconception, or rather the misinterpretation: $\text{size}(c)$ is different from n , even though they are particularly the same from first glance. It is because n is actually the representation size of the instance string, and not the concept representation string, which utilizes another alphabet entirely. Though, suppose they are of the same alphabet, then $X = \{\Sigma\}_n$, then again, there exists a number q such that $q \geq n$. In the case where $q = n$, it is most likely that the representation is constrained of the linear transformation order, of which does not change the string length.

Let us observe that even in the case $m \ll n$, the shortest consistent hypothesis in $\mathcal{H}$ may in fact be the target concept, and so we must allow $\text{size}(h)$ to depend at least linearly on $\text{size}(c)$. We will see cases where this makes it easier to efficiently find a consistent hypothesis — by contrast, computing the shortest hypothesis consistent with the data is often a computationally hard problem.

Theorem 2.5.2 (Occam's Razor). *Let L be an efficient (α, β) -Occam algorithm for $\mathcal{C}$ using $\mathcal{H}$. Let $\mathcal{D}$ be the target distribution over the instance space X , let $c \in \mathcal{C}_n$ be the target concept, and $0 < \epsilon, \delta \leq 1$. Then there is a constant $a > 0$ such that if L is given as input a random sample S of m examples drawn from $EX(c, \mathcal{D})$, where m satisfies:*

$$m \geq a \left(\frac{1}{\epsilon} \log \frac{1}{\delta} + \left(\frac{(n \cdot \text{size}(c)^\alpha)}{\epsilon} \right)^{1/1-\beta} \right) \quad (2.40)$$

then with probability at least $1 - \delta$ the output h of L satisfies $\text{error}(h) \leq \epsilon$. Moreover, L runs in time polynomial in n , $\text{size}(c)$, $1/\epsilon$ and $1/\delta$.

Notice that as β tends to 1, the exponent in the bound for m tends to infinity. This corresponds with the assumption of intuition that the length of the hypothesis approaches that of the data itself, then the predictive power of the hypothesis is diminishing.

To apply this theorem into variational setting, we turn to the arguably more general definition, measuring representational succinctness by the cardinality of the hypothesis class rather than the representationally supposed bit length $\text{size}(h)$. Then, theorem 2.5.2 will be the special case of such theorem. To make this precise, let $\mathcal{H}_n = \cup_{m \geq 1} \mathcal{H}_{n,m}$.

Theorem 2.5.3 (Occam's Razor, Cardinality Version). *Let $\mathcal{C}$ be a concept class and $\mathcal{H}$ a representation space. Let L be an algorithm such that for any n and any $c \in \mathcal{C}_n$, if L is given as input a sample S of m labelled examples of c , then L runs in time polynomial in n , m , and $\text{size}(c)$, and output an $h \in \mathcal{H}_{n,m}$ that is consistent with S . Then there is a constant $b > 0$ such that for any n , any distribution $\mathcal{D}$ over X_n , and any target concept $c \in \mathcal{C}_n$, if L is given as input a random sample from $EX(c, \mathcal{D})$ of m examples, where $|\mathcal{H}_{n,m}|$ satisfies:*

$$\log |\mathcal{H}_{n,m}| \leq b\epsilon m - \log \frac{1}{\delta} \quad (2.41)$$

or equivalently, where m satisfies $m \geq (1/b\epsilon)(\log |\mathcal{H}_{n,m}| + \log(1/\delta))$ then L is guaranteed to find a hypothesis $h \in \mathcal{H}_n$ that with probability at least $1 - \delta$ obeys $\text{error}(h) \leq \epsilon$.

We do not mention, or necessarily claim that L is an efficient PAC learning algorithm. In order for the theorem to apply, we must pick m large enough so that $b\epsilon m$ dominates $\log |\mathcal{H}_{n,m}|$.

Proof of Theorem 2.5.3. We first give a proof of theorem 2.5.2, by proving the more general form in theorem 2.5.3, of which then is considered a special case to prove. Firstly, we say that a hypothesis $h \in \mathcal{H}_{n,m}$ is *bad* if $\text{error}(h) > \epsilon$, where the error is of course measured with respect to the target concept c and the target distribution $\mathcal{D}$. Then, by the independence of the random examples, the probability that a fixed bad hypothesis h is consistent with a randomly drawn sample of m examples from $EX(c, \mathcal{D})$ is at most $(1 - \epsilon)^m$. Using the union bound, this implies that if $\mathcal{H}' \subseteq \mathcal{H}_{n,m}$ is the set of all bad hypotheses in $\mathcal{H}_{n,m}$, then the probability that some hypothesis in $\mathcal{H}'$ is consistent with a random sample of size m is at most $|\mathcal{H}'|(1 - \epsilon)^m$. We want this to be at most δ , since $|\mathcal{H}'| \leq |\mathcal{H}_{n,m}|$ we get a stronger condition if we solve for $|\mathcal{H}_{n,m}|(1 - \epsilon)^m \leq \delta$. From this, we obtain

$$\log |\mathcal{H}_{n,m}| \leq m \log(1/(1 - \epsilon)) - \log(1/\delta) \quad (2.42)$$

Using the fact that $\log(1/(1 - \epsilon)) = \Theta(\epsilon)$, we get the statement of the theorem. $\square$

Proof of Theorem 2.5.2. Let $\mathcal{H}_{n,m}$ be the set of all possible hypothesis representation that the (α, β) -Occam algorithm L might output when given as input a labelled sample S of cardinality m .

Since L is an (α, β) -Occam algorithm, every such hypothesis has bit length at most $(n \cdot \text{size}(c))^\alpha m^\beta$, thus implying that $|\mathcal{H}_{n,m}| \leq 2^{(n \cdot \text{size}(c))^\alpha m^\beta}$. By Theorem 2.5.3, the output of L has error at most ϵ with confidence at least $1 - \delta$, provided

$$\log |\mathcal{H}_{n,m}| \leq b\epsilon m - \log \frac{1}{\delta} \quad (2.43)$$

Transposing, we want m such that

$$m \geq \frac{1}{b\epsilon} \log |\mathcal{H}_{n,m}| + \frac{1}{b\epsilon} \log \frac{1}{\delta} \quad (2.44)$$

The above condition can be satisfied by picking m such that both

$$m \geq \left(\frac{2}{b\epsilon}\right) \log |\mathcal{H}_{n,m}|, \quad m \geq \left(\frac{2}{b\epsilon}\right) \log \frac{1}{\delta} \quad (2.45)$$

will effectively hold. Choosing $a = 2/b$ yields the statement in the theorem. ⁸ $\square$

⁸Why is the pairing of (a, b) value is as such, we are not so sure.

2.6 Discrete space empiricism - Rademacher complexity

The jump toward discrete space is natural, as information are often discrete on infinite and continuous phase space (i.e. the domain, for example, $\mathbb{R}$ -values). The hypothesis sets typically used in machine learning are infinite, even in toy model, as the rectangular learning problem itself can be expressed by infinitely many rectangles. But the sample complexity bounds are uninformative when dealing with infinite hypothesis sets - for arbitrary construction of $\mathcal{H}$. One could ask whether efficient learning from a finite sample is even possible when the hypothesis set $\mathcal{H}$ is infinite. The basic goal with the Rademacher complexity is this exact question, for the first complexity treatment. The general idea for doing so consists of reducing the infinite case to the analysis of finite sets of hypotheses, and then proceed as previous sections. Authoritative sources provided mainly the insight and theories in this section, as of [Mohri et al. \[2012\]](#) and [Bartlett and Mendelson \[2002\]](#).

2.6.1 Rademacher descriptions

We will continue to use $\mathcal{H}$ to denote a hypothesis set as in the previous chapters. Many of the result of this section are general and hold for arbitrary loss function, or $L : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$. In what follows, $\mathcal{G}$ will generally be interpreted as the *family of loss functions associated to $\mathcal{H}$* , mapping from $\mathcal{Z} = \mathcal{X} \times \mathcal{Y}$ to $\mathbb{R}$. That is,

$$\mathcal{G} = \{g : (x, y) \mapsto L(h(x), y) : h \in \mathcal{H}\}$$

However, the definitions are given in the general case of a family of arbitrary functions $\mathcal{G}$, mapping from an arbitrary input space $\mathcal{Z}$ to $\mathbb{R}$. Thus, we can say that we endowed upon the landscape a measure set $\mathcal{G}$ of all singular mapping.

ERC The **Rademacher complexity** captures the richness of a family of functions by measuring the degree to which a hypothesis set can fit random noise. This is perhaps one of the very *empirical complexity notion* that would be seen, as for the random noise idea, that is.

Let G be a family of functions mapping from $Z \rightarrow [a, b]$, $S = (z_1, \dots, z_m)$ a fixed sample of size m with elements in Z . Then, the **empirical Rademacher complexity** of G with respect to the sample S is defined as:

$$\hat{\mathfrak{R}}_S(\mathcal{G}) = \mathbb{E} \left[\sup_{g \in G} \frac{1}{m} \sum_{i=1}^m \sigma_i g(z_i) \right]$$

where $\sigma = \{\sigma_1, \dots, \sigma_m\}^\top$ with σ_i s independent uniform random variables taking values in $\{-1, +1\}$. The random variable σ_i are called **Rademacher variables**.

The empirical Rademacher complexity can be rewritten as:

$$\hat{\mathfrak{R}}_S(\mathcal{G}) = \mathbb{E} \left[\sup_{g \in G} \frac{\sigma \cdot \mathbf{g}_S}{m} \right]$$

for $\mathbf{g}_S$ denoting the vector of values taken by g over the sample S : $\mathbf{g}_S = (g(z_1), \dots, g(z_m))^\top$. The inner product measures the correlation of $\mathbf{g}_S$ with the vector of random noise σ . The supremum $\sup_{g \in G} (\sigma \cdot \mathbf{g}_S / m)$ is a measure of how well the function class $\mathcal{G}$ correlates with σ over the sample S . Thus, the empirical Rademacher complexity measures on average how well the function class correlates with random noise on the dataset. What will the random noise be like? Well, not much. For a fixed sample size, it captures how **random** the loss function distribution is. A random noise system is inherently chaotic by

default, even if we restrict its value in $\{-1, +1\}$ ⁹. Hence, if the model is good enough on certain dataset, that the noisy, useless data configuration performs "somewhat pretty well", then from that, we can know the complexity of the model by its flexibility in such noisy fitting. That is, in the topic of fitting, at least. So the richer or more complex families $\mathcal{G}$ can generate more vectors $\mathbf{g}_S$ and thus better correlate with random noise on average.

We come to the definition of *the* Rademacher complexity.

Rademacher
complexity

Definition 2.6.1 (Rademacher complexity). Let $\mathcal{D}$ denote the distribution according to which samples are drawn. For any integer $m \geq 1$, the **Rademacher complexity** of $\mathcal{G}$ is the expectation of the empirical Rademacher complexity over all samples of size m drawn according to $\mathcal{D}$.

$$\mathfrak{R}_m(\mathcal{G}) = \mathbb{E}_{S \sim \mathcal{D}^m} [\hat{\mathfrak{R}}_S(\mathcal{G})] \quad (2.46)$$

We are now ready to present our first generalization bound based on Rademacher complexity.

Theorem 2.6.1. Let $\mathcal{G}$ be a family of functions mapping from $\mathcal{Z}$ to $[0, 1]$. Then, for any $\delta > 0$, with probability at least $1 - \delta$ over the draw of an i.i.d. sample S of size m , each of the following holds for all $g \in \mathcal{G}$:

$$\mathbb{E}[g(z)] \leq \frac{1}{m} \sum_{i=1}^m g(z_i) + 2\mathfrak{R}_m(\mathcal{G}) + \sqrt{\frac{\log(1/\delta)}{2m}} \quad (2.47)$$

and

$$\mathbb{E}[g(z)] \leq \frac{1}{m} \sum_{i=1}^m g(z_i) + 2\mathfrak{R}_m(\mathcal{G}) + 3\sqrt{\frac{\log(2/\delta)}{2m}} \quad (2.48)$$

The final term in both equations is typically much smaller than the Rademacher complexity. Note that they are one-sided uniform deviation bounds, and that they are also *data-dependent* bound, just like how Gaussian process regression is also functionally data-dependent.

Proof. For any sample $S = (z_1, \dots, z_m)$, and any $g \in \mathcal{G}$, we denote $\hat{\mathbb{E}}_S[g]$ by the empirical average of g over S . The proof then consists of applying McDiarmid's inequality to function Φ defined for any sample S by:

$$\Phi(S) = \sup_{g \in \mathcal{G}} (\mathbb{E}[g] - \hat{\mathbb{E}}_S[g]) \quad (2.49)$$

This is the supremum of the different between the average and the empirical average for all $g \in \mathcal{G}$. If there is anything, this is similar to the following rational choice: just simply change the inequality to:

$$\begin{aligned} \mathbb{E}[g(z)] - \frac{1}{m} \sum_{i=1}^m g(z_i) &\leq 2\mathfrak{R}_m(\mathcal{G}) + 3\sqrt{\frac{\log(2/\delta)}{2m}} \\ \mathbb{E}[g(z)] - \hat{\mathbb{E}}_{S'}[g(z)] &\leq 3\mathfrak{R}_m(\mathcal{G}) + a\sqrt{\frac{\log(2/\delta)}{2m}} \end{aligned} \quad (2.50)$$

⁹Actually, this is indeed a problem. What will happen if the value is totally synchronous for a specific configuration that is, though statistically impossible, yet is the solution to the finite space?

in which we then take the supremum. The supremum is for the pretty trivial choice: for this equality to hold true then the largest element of the LHS $\mathbb{E}[g(z)] - \hat{\mathbb{E}}_{S'}[g(z)]$ to be less than the RHS, for any δ (for equation 2.47) and $\delta/2$ (for equation 2.48). Note that the same thing is applicable.

Let S and S' be two samples differing by exactly one point, say z_m in S and z'_m in S' . Then, since the difference of suprema does not exceed the supremum of the difference, we have:

$$\Phi(S') - \Phi(S) \leq \sup_{g \in \mathcal{G}} \left(\hat{\mathbb{E}}_S[g] - \hat{\mathbb{E}}_{S'}[g] \right) = \sup_{g \in \mathcal{G}} \frac{g(z_m) - g(z'_m)}{m} \leq \frac{1}{m} \quad (2.51)$$

Similarly, we can obtain

$$\Phi(S') - \Phi(S) \leq \frac{1}{m} \implies |\Phi(S) - \Phi(S')| \leq \frac{1}{m} \quad (2.52)$$

By McDiarmid's inequality, for any $\delta > 0$, with probability at least $1 - \delta/2$, the following holds.

$$\Phi(S) \leq \mathbb{E}[\Phi(S)] + \sqrt{\frac{\log 2/\delta}{2m}} \quad (2.53)$$

We next bound the expectation of the right-hand side as follows:

$$\mathbb{E}_S[\Phi(S)] = \mathbb{E}_S \left[\sup_{g \in \mathcal{G}} \left(\mathbb{E}[g] - \hat{\mathbb{E}}_S(g) \right) \right] \quad (2.54)$$

$$= \mathbb{E}_S \left[\sup_{g \in \mathcal{G}} \left(\mathbb{E}_{S'}[\hat{\mathbb{E}}_{S'}(g)] - \hat{\mathbb{E}}_S(g) \right) \right] \text{ from 6.2.3} \quad (2.55)$$

$$= \mathbb{E}_S \left[\sup_{g \in \mathcal{G}} \mathbb{E}_{S'} \left[\hat{\mathbb{E}}_{S'}(g) - \hat{\mathbb{E}}_S(g) \right] \right] \quad (2.56)$$

$$\leq \mathbb{E}_{S,S'} \left[\sup_{g \in \mathcal{G}} \left(\hat{\mathbb{E}}_{S'}(g) - \hat{\mathbb{E}}_S(g) \right) \right] \quad (2.57)$$

$$= \mathbb{E}_{S,S'} \left[\sup_{g \in \mathcal{G}} \frac{1}{m} \sum_{i=1}^m (g(z'_i) - g(z_i)) \right] \quad (2.58)$$

$$= \mathbb{E}_{\sigma,S,S'} \left[\sup_{g \in \mathcal{G}} \frac{1}{m} \sum_{i=1}^m \sigma_i (g(z'_i) - g(z_i)) \right] \quad (2.59)$$

$$\leq \mathbb{E}_{\sigma,S,S'} \left[\sup_{g \in \mathcal{G}} \frac{1}{m} \sum_{i=1}^m \sigma_i g(z'_i) \right] + \mathbb{E}_{\sigma,S} \left[\sup_{g \in \mathcal{G}} \frac{1}{m} \sum_{i=1}^m -\sigma_i g(z_i) \right] \quad (2.60)$$

$$= 2 \mathbb{E}_{\sigma,S} \left[\sup_{g \in \mathcal{G}} \frac{1}{m} \sum_{i=1}^m \sigma_i g(z_i) \right] = 2\mathfrak{R}_m(S). \quad (2.61)$$

We note that there are several implicit choices, especially for the properties of the subadditivity of the supremum, $\sup A + B = \sup A + \sup B$.

The reduction to $\mathfrak{R}_m(\mathcal{G})$ in equation 2.61 yields the bound in the previous equation (the first one), using δ instead of $\delta/2$. To derive a bound in terms of $\hat{\mathfrak{R}}_S(\mathcal{G})$, we observe that, by definition, changing one

point in S changes $\hat{\mathfrak{R}}_S(\mathcal{G})$ by at most $1/m$. Then, using again McDiarmid inequality, with probability $1 - \delta/2$, the following holds:

$$\mathfrak{R}_m(\mathcal{G}) \leq \hat{\mathfrak{R}}_S(\mathcal{G}) + \sqrt{\frac{\log 2/\delta}{2m}} \quad (2.62)$$

Finally, we use the union bound to combine inequalities together, which yields

$$\Phi(S) \leq 2\hat{\mathfrak{R}}_S(\mathcal{G}) + 3\sqrt{\frac{\log 2/\delta}{2m}} \quad (2.63)$$

which matches the inequality. $\square$

The above inequality can be somewhat extended to be formulated, using the empirical Rademacher complexity rather than the expected Rademacher complexity.

Corollary 2.6.1. *Suppose that a sample S of size m is drawn according to distribution $\mathcal{D}$. Then for any $\delta > 0$, with probability at least $1 - \delta$, the following holds for all $g \in \mathcal{G}$:*

$$\mathbb{E}[g(z)] = L(g) \leq \underbrace{L_S(g)}_{\hat{\mathbb{E}}_S[g]} + 2\mathfrak{R}_S(\mathcal{G}) + O\left(\sqrt{\frac{\log 1/\delta}{m}}\right) \quad (2.64)$$

Proof. We may consider the empirical Rademacher complexity $R_S(G)$ as a function of the points $z_1, \dots, z_m$ that comprise the sample S . Changing one of the z_i to a new value z'_i changes $R_S(G)$ by at most $1/m$. Applying McDiarmid's inequality with $c = 1/m$, and $\epsilon = \sqrt{\log 2/\delta/2m}$ we have that with probability at least $1 - \delta/2$,

$$R_S(G) \leq R_m(\mathcal{G}) + \sqrt{\frac{\log 2/\delta}{2m}} \quad (2.65)$$

By union bound, with probability at least $1 - e\delta$, the inequality both holds. This implies that our corollary holds for all $g \in \mathcal{G}$ with probability at least $1 - 2\delta$. Replace δ by $\delta/2$ yields the required result. $\square$

The following result relates the empirical Rademacher complexities of a hypothesis set $\mathcal{H}$ and to the family of loss functions $\mathcal{G}$ associated to $\mathcal{H}$ in the case of binary loss (zero-one loss).

Lemma 2.6.2. *Let $\mathcal{H}$ be a family of functions taking values in $\{-1, +1\}$, and let $\mathcal{G}$ be the family of loss functions associated to $\mathcal{H}$ for the zero-one loss:*

$$\mathcal{G} = \{(x, y) \mapsto 1_{h(x) \leq y} : h \in \mathcal{H}\}$$

For any sample $S = \{(x_i, y_i)\}_m$ of elements in $\mathcal{X} \times \{-1, 1\}$, let $S_{\mathcal{X}}$ be the projection over $\mathcal{X}$: $S_{\mathcal{X}} = (x_1, \dots, x_m)$. Then, the following relation holds:

$$\hat{\mathfrak{R}}_S(\mathcal{G}) = \frac{1}{2}\hat{\mathfrak{R}}_{S_{\mathcal{X}}}(\mathcal{H}) \quad (2.66)$$

Proof. For any sample $S = ((x_1, y_1), \dots, (x_m, y_m))$ of elements in $\mathcal{X} \times \{-1, +1\}$, by definition, the empirical Rademacher complexity of $\mathcal{G}$ can be written as:

$$\hat{\mathfrak{R}}_S(\mathcal{G}) = \mathbb{E}_{\sigma} \left[\sup_{h \in \mathcal{H}} \frac{1}{m} \sum_{i=1}^m \sigma_i \mathbb{1}_{h(x_i) \neq y_i} \right]$$

$$\begin{aligned}
&= \mathbb{E}_\sigma \left[\sup_{h \in \mathcal{H}} \frac{1}{m} \sum_{i=1}^m \sigma_i \frac{1 - y_i h(x_i)}{2} \right] \\
&= \frac{1}{2} \mathbb{E}_\sigma \left[\sup_{h \in \mathcal{H}} \frac{1}{m} \sum_{i=1}^m -\sigma_i y_i h(x_i) \right] \\
&= \frac{1}{2} \mathbb{E}_\sigma \left[\sup_{h \in \mathcal{H}} \frac{1}{m} \sum_{i=1}^m \sigma_i h(x_i) \right] \\
&= \frac{1}{2} \hat{\mathfrak{R}}_S(\mathcal{H}),
\end{aligned}$$

where we used the fact that $\mathbb{1}_{h(x_i) \neq y_i} = (1 - y_i h(x_i))/2$ and the fact that for a fixed $y_i \in \{-1, +1\}$, σ_i and $-\sigma_i y_i$ are distributed in the same way. $\square$

The lemma then implies by taking expectations, that for any $m \geq 1$,

$$\mathfrak{R}_m(\mathcal{G}) = \frac{1}{2} \mathfrak{R}_m(\mathcal{H}) \quad (2.67)$$

These connections between the empirical and average Rademacher complexities can be used to derive generalization bounds for binary classification in terms of the Rademacher complexity of the hypothesis set $\mathcal{H}$. This leads to the following result on the Rademacher complexity bounds on binary classification.

Theorem 2.6.3 (Rademacher complexity bounds - binary classification). *Let $\mathcal{H}$ be a family of functions taking values in $\{-1, +1\}$ and let $\mathcal{D}$ be the distribution over the input space $\mathcal{X}$. Then, for any $\delta > 0$, with probability at least $1 - \delta$ over a sample S of size m drawn according to $\mathcal{D}$, each of the following holds for any $h \in \mathcal{H}$.*

$$R(h) \leq \hat{R}_S(h) + \mathfrak{R}_m(\mathcal{H}) + \sqrt{\frac{\log 1/\delta}{2m}} \quad (2.68)$$

and

$$R(h) \leq \hat{R}_S(h) + \mathfrak{R}_m(\mathcal{H}) + 3\sqrt{\frac{\log 2/\delta}{2m}} \quad (2.69)$$

Proof. This follows immediately from the above theorem and lemma. $\square$

2.6.2 Growth function

growth function

Here, we will show how the Rademacher complexity can be bounded in terms of the *growth function*. The idea encoded of growth function simply asks for *how many* truly different hypotheses are there when we only observe m input from infinitely discrete possible inputs of $\mathcal{X}$.

Definition 2.6.2 (Growth function). Given a binary class of functions $\mathcal{H}$, we define the *growth function* $\Pi_{\mathcal{H}} : \mathbb{N} \rightarrow \mathbb{N}$ as:

$$\Pi_{\mathcal{H}}(m) = \max_{\{x_1, \dots, x_m\} \subseteq \mathcal{X}} |\{h(x_1), \dots, h(x_m) : h \in \mathcal{H}\}| \quad (2.70)$$

In other words, $\Pi_{\mathcal{H}}(m)$ is the maximum number of distinct ways in which m points can be classified using hypotheses in $\mathcal{H}$. This growth function is then also a measurement of the richness of a class of function, and each one of these distinct classification is called a *dichotomy*. The growth function then counts the number of dichotomies that are realized by the hypothesis. This differs from the Rademacher complexity, in the way that it is independent of the distribution of points sampled, hence is purely combinatorial. Also, the growth function is conducted for all $h \in \mathcal{H}$ of the finite hypothesis space. If it is the binary case, then there is at most 2^m categorization configuration. To relate these two together, theoretical consideration uses the Massart's lemma.

Lemma 2.6.4 (Massart's lemma). *Let A be a finite subset of $\mathbb{R}^m$, and let*

$$\max_{\vec{y} \in A} \|\vec{y}\|_2 \leq r$$

Then

$$\mathbb{E}_{\sigma_1, \dots, \sigma_m} \left[\sup_{\vec{y} \in A} \sum_{i=1}^m \sigma_i y_i \right] \leq r \sqrt{2 \log |A|} \quad (2.71)$$

This lemma offers not much aside from an interesting bound. Using this result, we can now bound the Rademacher complexity in terms of the growth function.

Corollary 2.6.2. *Let $\mathcal{G}$ be a family of functions taking values in $\{-1, +1\}$. Then the following holds:*

$$\mathfrak{R}_m(\mathcal{H}) \leq \sqrt{\frac{2 \log \Pi_{\mathcal{H}}(m)}{m}} \quad (2.72)$$

Proof. For a fixed sample S , we denote by $\mathcal{G}_{|S}$ the set of vectors of function values $(g(x_1), \dots, g(x_m))^{\top}$ for $g \in \mathcal{G}$. Since $g \in \mathcal{G}$, $|g'| \in \mathcal{G}_{|S}| \leq \sqrt{m}$. We can then apply Massart's lemma as:

$$\begin{aligned} \mathfrak{R}_m(\mathcal{G}) &= \mathbb{E}_S \left[\mathbb{E}_{\sigma} \left[\sup_{u \in \mathcal{G}_{|S}} \frac{1}{m} \sum_{i=1}^m \sigma_i u_i \right] \right] \\ &\leq \mathbb{E}_S \left[\frac{\sqrt{m} \sqrt{2 \log |\mathcal{G}_{|S}|}}{m} \right] \end{aligned} \quad (2.73)$$

By definition, $|\mathcal{G}_{|S}|$ is bounded by the growth function. Thus,

$$\begin{aligned} \mathfrak{R}_m(\mathcal{G}) &\leq \mathbb{E}_S \left[\frac{\sqrt{m} \sqrt{2 \log \Pi_{\mathcal{G}}(m)}}{m} \right] \\ &= \sqrt{\frac{2 \log \Pi_{\mathcal{G}}(m)}{m}} \end{aligned} \quad (2.74)$$

which concludes the proof. $\square$

From this and the generalization bound, we will also yield the following generalization bound in terms of the growth function.

Corollary 2.6.3 (Growth function generalization bound). *Let $\mathcal{H}$ be a family of function taking values in $\{-1, +1\}$. Then, for any $\delta > 0$, with probability at least $1 - \delta$, for any $h \in \mathcal{H}$:*

$$R(h) \leq \hat{R}_S(h) + \sqrt{\frac{2 \log \Pi_{\mathcal{H}}(m)}{m}} + \sqrt{\frac{\log(1/\delta)}{2m}} \quad (2.75)$$

This growth function bound can also be derived directly. The computation of the growth function is also, pretty weird, and is often not always convenient, especially because of its combinatorial nature.

2.7 Equivalent spatial separation - VC-theory

VC, or Vapnik–Chervonenkis theory (Vapnik and Chervonenkis [1971], Vapnik [1999], Sauer [1972], Shelah [1972]), base a lot of its results and theories on both the growth function, and the combinatorial nature of which it takes form. The problem of the formation of VC theory is simple. We have seen that a consistent learner can be used to design a PAC-learning algorithm, provided the output hypothesis comes from a class that is not too large, in particular, of any polynomial class $\text{poly}(\cdot, \cdot, \cdot, \cdot)$. However, one certain assumption and condition on the PAC-learning algorithm measure is that it only works on (the result, that is) finite hypothesis class or concept class. Concept classes that are uncountably infinite are often used, for example, linear halfspaces, or more familiar, the hypothesis on such space is the **linear classifier** or **perceptron**. In this section, we would particularly see how we can use a new capacity measure called VC dimension of a concept class to make a consistent learner to be PAC-learnable. All theories and practices using the VC-dimension as one of its main ingredient is then called the VC theory, from authoritative sources. We also note that from this point on, there exists more fundamental extensions of the VC-theory, or alternatives. This includes algorithm-dependent bounds, of Sachs et al. [2023], treatments of learning bounds based on stochastic algorithmic stability as in Bousquet and Elisseeff [2002], similar discrete and combinatorial view as of Natarajan’s dimension (Natarajan [1989]), pseudo-dimension (Pollard [1984]), information theoretic approach (Xu and Raginsky [2017b]), or representational complexity form into *approximate description length*, as seen in Daniely and Katzhendler [2022].

The main concept of VC-theory lies in the apparatus of capacity as the **VC-dimension**. The VC-dimension is often regarded as a purely combinatorial notion, but it is often easier to compute than the growth function, or the Rademacher Complexity, nevertheless holding equivalence to both and statements of PAC-learning. As we shall see, the VC-dimension is a key quantity in learning, and is directly related to the growth function.

shattering

To define the VC-dimension, we will have to define the notion of *shattering*. As we recalled of the previous sections on the growth function, on the hypothesis set $\mathcal{H}$, a dichotomy on a set S is one of many possible ways of labelling the points in S using a hypothesis $\mathcal{H}$. A set S of $m \geq 1$ points is said to be shattered by a hypothesis set $\mathcal{H}$ when $\mathcal{H}$ realizes all possible dichotomies of S , that is when $\Pi_{\mathcal{H}}(m) = 2^m$.

Definition 2.7.1 (Shattering). We say that a finite set $S \subset \mathcal{X}$ is **shattered** by $\mathcal{C}$, if $|\Pi_{\mathcal{C}}(S)| = 2^{|S|}$. S is shattered by $\mathcal{C}$ if all possible dichotomies over S can be realized by $\mathcal{C}$.

VC-dimension

We now can define the notion of the *dimension* for a concept class $\mathcal{C}$.

Definition 2.7.2 (VC-dimension). The **VC-dimension** of a hypothesis set $\mathcal{H}$ is the size of the largest set that can be shattered by $\mathcal{H}$,

$$\text{VCdim}(\mathcal{H}) = \max_{m \in \mathbb{N}} \{m : \Pi_{\mathcal{H}}(m) = 2^m\} \quad (2.76)$$

Note that, by definition, if $\text{VCdim}(\mathcal{H}) = d$, then there exists a set of size d that can be shattered. However, it does not mean that all sets of size d or less are shattered, because this only get you the existential criteria. Similar to the growth notion, if we are to look at the definition of VC-dimension, we see the similar problem: it is very computationally heavy for experimental purposes. However, its treatment of the learning theory through discretized notion, and separation conjunctive space is powerful, making it effective to express many hypothesis classes.

A famous example illustrating this concept is the problem of halfspace partition. The **halfspace** hypothesis space is the set of hypotheses that consist of a hyperplane in a d -dimensional coordinate space that classifies a feature vector $\phi(x) \in \mathbb{R}^{d+1}$ as either $-1, 1$ based on which side the of the hyperplane it lies. Here, d represents the number of features on item x .

Mathematically, to define the halfspace hypothesis space, consider the domain $\mathcal{X} = \mathbb{R}$ and concept set $\mathcal{Y} = \{-1, 1\}$. The class $\mathcal{H}_{gl}$ of general linear classifiers, or linear halfspaces, is defined as

$$\mathcal{H}_{gl} = \{h_{\mathbf{w},b} \mid \mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R}\} \quad (2.77)$$

where

$$h_{\mathbf{w},b}(x) = \text{sgn}(\langle x \cdot \mathbf{w} \rangle + b) = \text{sgn}\left(\left(\sum_{i=1}^d x_i w_i\right) + b\right) \quad (2.78)$$

General linear halfspaces in $\mathbb{R}^d$ can be viewed as *homogenous linear halfspaces* in $\mathbb{R}^{d+1}$ via a simple transformation. We define the class of homogenous linear halfspaces $\mathcal{H}_l$ as

$$\mathcal{H}_l = \{h_{\mathbf{w}} \mid \mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R}\}$$

where

$$h_{\mathbf{w}}(x) = \text{sgn}(\langle x \cdot \mathbf{w} \rangle) = \text{sgn}\left(\sum_{i=1}^d x_i w_i\right) \quad (2.79)$$

Let $x = h_{\mathbf{w},b}(x) \in \mathbb{R}^d$, for $\mathbf{w} = (w_1, w_2, \dots, w_d) \in \mathbb{R}^d$ and $b \in \mathbb{R}$ be given, we then define

$$\mathbf{x}' = (1, x_1, \dots, x_d) \in \mathbb{R}^{d+1} \quad (2.80)$$

$$\mathbf{w}' = (1, w_1, \dots, w_d) \in \mathbb{R}^{d+1} \quad (2.81)$$

Then we get $\langle x \cdot \mathbf{w} \rangle + b = \langle \mathbf{x}' \cdot \mathbf{w}' \rangle$ for all $x \in \mathbb{R}$, and thus

$$h_{\mathbf{w},b}(x) = h_{\mathbf{w}'}(\mathbf{x}')$$

Let $\mathcal{X}$ be a set of $d + 2$ points. By Radon's theorem, it can be partitioned into two sets $\mathcal{X}_1$ and $\mathcal{X}_2$ such that their convex hulls intersect. Observe that when two sets of points $\mathcal{X}_1, \mathcal{X}_2$ are separated by a hyperplane, their convex hulls are also separated by that hyperplane. Thus, $\mathcal{X}_1$ and $\mathcal{X}_2$ cannot be separated by a hyperplane and $\mathcal{X}$ is not shattered. Combining our lower and upper bounds, we have proven that $\text{VCdim}(\mathcal{H} \in \mathbb{R}^d)$ is $d + 1$.

The VC-dimension of many other hypothesis sets can be determined or upper-bounded in a similar way. In particular, the VC-dimension of any vector space of dimension $r < \infty$ can be shown to be at most r . The next result, known as *Sauer's lemma*, clarifies the connection between the notions of growth function and VC-dimension.

Sauer's lemma

Theorem 2.7.1 (Sauer's lemma). *Let $\mathcal{H}$ be a hypothesis set with $\text{VCdim}(\mathcal{H}) = d$. Then, for all $m \in \mathbb{N}$, the following inequality holds:*

$$\Pi_{\mathcal{H}}(m) \leq \sum_{i=0}^d \binom{m}{i} \quad (2.82)$$

2.8 Algorithmic stability

2.9 Summary

In the above section, we have introduced the notion of the learning problem, in its simplistic form and mathematical formalism. With such, we developed a setting of which is reasonable to guarantee learning and the existence of a hypothesis, PAC-learning, complexity analysis (Rademacher complexity, growth function, VC-dimension) and application of them on gaining learning bounds, and some of the examples on typical learning problems. Further than such, we have some remarks and problems to discuss.

The learning theory result itself is limited to particular class of result. Most of the results from classical learning setting either are optimization bound, "learning bound", or convergence criteria bounds. Some of them are concerned of the complexity of the model and vice versa, and a lot of them then are concerned with connecting the complexity notion to the learning bound itself - more like a dual relationship. Those results are good, for example, PAC-learning is *asymptotic learning bound* of time complexity, which makes sense of the time and the bounding method is restricted to (ϵ, δ) pair, or by some versions of VC-learning which is the *combinatorial discrete complexity bounding* method for those model itself. However, aside from the setting is very much limited, the notion of generalization is more or less generally just being empirical and generalization error (or risk) without sufficient practical heuristic special case, the results from which does not give us much idea about what to do and how to achieve the capable result in such bound. They *guarantee* results as probable, or measurable, but give us no effective ruler to measure it. Some of the measure is improbable, for example, the growth function complexity in various case, and generality of such. Furthermore, complexity is very not so well-defined, especially in the numerical space encoding. Some of the topic, while very interesting, is not discussed, or rather cannot be incorporated into the existing consideration.

The learning setting itself, from this setting, misses certain requirement to define it in a concrete sense. For example, its inclusion set contains the set X of instances in observational space, σ -algebra $\mathcal{S}$ on X , the family $\mathcal{P}$ of probability measure on $(X, \mathcal{S})$ (which is quite optional, give it or take), and subset $\mathcal{C} \subseteq \mathcal{S}$ called the *concept class*, or else a family $\mathcal{F}$ of measurable mapping $X \rightarrow [a, b]$ of closed set on the Borel σ -algebra, called the *function class*. These formalisms are often implicitly defined, yet uncertain of their actual property and interpretation. For example, in some testing cases, we would want to know the deterministic degree of the probabilistic measure space $\mathcal{P}$, for whether p_X will collapse into singular behaviour - which indeed probabilistically can be done using elementary probability, however, for different system and different concept, different setting, this proves to be rather difficult. Usually, $[a, b] = [0, 1]$. Furthermore, there are also the definition for the *encoding space* $\mathcal{E}(\cdot, \dots)$ of all fundamental components, of the ambient space $\mathbb{A}$ and the representation space $\mathcal{R}$ together, with a distinct encoding rule that is not included. In utmost importance, encoding space needs to be one of the main focus of the theory, however, reality often proves to be unpleasant. Some of the definitions and further clarification is not also included in the rigours, leaving questions that is unattainable of the setting provided. Granted, the theory is minimized, but for example, such is to introduce the distribution $\mathcal{D}$ into the preliminary setting, one must clarify and indeed, concretely provide the fundamental setting of the probabilistic nature of the system - in which sense and when will the system becomes probabilistic, and under what circumstances

the probabilistic nature of it varies. Couples that with another theoretical idea of control/process theory, in which the learning theory is often settled to be discrete process, leaves much clarification to be found. Worst, philosophically speaking, and developmental speaking, the notion of *learning* is still very premature of interest – which lead directly to the case where everyone has frameworks and all, but not a lot actually works, and many overlaps up to almost identical.

Aside from all of those, we have an inconsistent framework to deal with. Usually, as we see, there exists no concrete structure to the model that we are asking of. Most of the time, the landscape of practical application changes so fast that we also do not have the ability to structuralize it. As of an interdisciplinary field, bring ideas of a maturing subject with uncertainty also come with difficulty of raising a good hallmark to the problem. While it is true that we have two domains – the *classical structure* of model with design in mind closer to inference, and the *modern structure* which almost identical now to *deep neural structure*, focus on something akin to prediction rather than not, most of the time, this classification is a poor classification in the grand scheme of all, especially if we take into account others model and structures that lie outside of it (Markov chain, generative evolutionary structure, etc., programming synthesis).

Thereby, a new construction is recommended to exist, to formalize the theory itself, and to further down increase the development speed since we are pretty much lagging behind, both in formalization and practical implementation. Should the new theory arise in this book, it must fix almost all the classical theory's pitfall, and further on of the *concurrent theory* being discussed.

Chapter 3. Introduction to Double Descent

The analysis of learning action, machine learning and related practices in the field theoretically has been made by utilizing Computational Learning Theory (CoLT) [Valiant \[1984\]](#) and Statistical Learning Theory (SLT) [Vapnik \[1999\]](#). These two theories, while overlapped, provides a general framework in analysing and justifying learning actions and learning model constructions, aiding in the formation of modern practices. One of the famous insight using such framework is the *bias-variance tradeoff* [Geman et al. \[1992\]](#), which states that model complexity and generality inversely affect each other, thus guarantee the need for a safe bracket between them. However, recent literatures, [Belkin et al. \[2019a\]](#) has indicated the fallout of such dilemma by the new phenomenon called *double descent*, where there exists an interpolation threshold that renders the current statistical justification for bias-variance inconsequential to a given region. Various ‘anomaly’ has also been detected similarly, with varying degree of sophistication and potential. In this paper, we would like to review through progress made in this particular field of research, and the phenomena potential explanations up until now.

3.1 Introduction

The modern machine learning landscape is dominated and guaranteed by the foundation of theoretical machine learning, specifically computational learning theory (CoLT) and statistical learning theory (SLT). The theory of machine learning and its modern practice has been developed and researched, of a substantial portion by empirical and heuristic approach, either by advancing new practices, architectures or by try-and-test modification. From its early onset of a regression estimator $\theta(x, y)$ on the linear regression problem, machine learning has developed substantially. Certain model concept with great successes includes regression-classification model, Bayesian modelling, generative learning model, support vector machine, Gaussian processes, and more. Of all such, the more formal and complex model architecture created, is the concept of a *neural network*. With increasingly sophisticated architecture, heuristic approach become popular, the fast-paced advancement of the field comes with new method, new results, new observations, and its far-reaching application which led to even bigger and larger scale deployment, there has been questions about the formation and status of a theoretical ground, a rigorous matter on the side of **theoretical machine learning**.

While rigorous and well-formulated in a sense, classical and theoretical machine learning was dwarfed by the modern advancement of machine learning as a whole, leading to several anecdotal problems regarding the interpretation of phenomena, the re-evaluation of the theory to fit the more updated analysis, and explanation to more sophisticatedly designed system. This and many more, plus as present, many of such advancements and improvements are heuristic, and the general theory and conceptual understanding remain limited, led to the choice to often opt for analogies and empirical workaround. Ultimately, this resulted in multiple ‘failures’ in explaining, interpreting, and predicting behaviours of new phenom-

ena appeared in the modern landscape of machine learning. While there exists novelty of analysis, and different theoretical schemes to analyse such, particular problems remain in the form of irregularity and conflict of interpretations between frameworks.

Our main focus is to review on the topic, analyse existing analysis and observations, and understanding of the umbrella term classical learning theory, for the recent phenomena observed named ‘double descent’. Statistical learning theory (SLT) and computational learning theory (CLT) Vapnik [1999], Mohri et al. [2012], Shalev-Shwartz and Ben-David [2014], Hajek and Raginsky [2021], Bousquet et al. [2020] has been prominent in constructing a well-rounded formal theory surrounding learning problems, models, and machine learners analysis. Valuable insights have been dissected from treatment of statistical theory and mathematical modelling on models, including the *bias-variance tradeoff* Geman et al. [1992], Domingos [2000c], which serves as a bound for efficient learning and model configuration (or complexity). However, recently there has been observations of *double descent* Belkin et al. [2019a], Schaeffer et al. [2023], Nakkiran et al. [2019], Lafon and Thomas [2024] which refute the famous tradeoff assumption, and hence brings question to the establishment of the theory, as well as several assumptions and insight in the framework. Further events and phenomena observed also includes grokking and triple, to n -descent Davies et al. [2023], d’Ascoli et al. [2020]. Furthermore, many problems of defining and formalizing notions used in designing and implementing machine learning models are inconclusive, such as, for example, *model complexity* and others.

3.1.1 Outline

The main contributions of the paper lie in the heart of the problem of *machine learning theoretical*, and specifically *double descent*. As such, we focus on the following:

1. Theoretical analysis of the machine learning framework, and reinterpretation of the empirical-generalization framework. We explicitly redetermine perspectives and formulation, about the problem of generalization learning, and model selection principles. Additionally, we also formalize the algorithmic objective of the objective system.
2. Observations and insight reviews and critical analysis. We analyse current knowledge, different analysis of the phenomena, and provide critiques, critical analysis on limitation, constraints, different assumptions and fallacy observed in those results, and verification or extension of experiments. This means analysing existing papers specific of details, and intrinsic of certain functional class.
3. Providing our own insight to the phenomena of double descent, and hence the need for new analysis or different theory scheme to be utilized in solving the problem.

Even though we already have the **theoretical attacks and analysis** of the learning theory in much greater details above, such details are simply too large to be totally relevant to our formal investigation on double descent. Largely speaking, it is much more relevant in our further sections targeting the reformulation of the descent theory rather than not. Thus, we will *reintroduce* some of those notions in a much denser form, as followed in the following critique.

While this is non-conclusive of the problem of double descent, we want to emphasize the importance of such issue to be further than simply the dilemma between model complexity and risk error. It hints at the mismatch of knowledge, fuzzy nature of certain concepts, failure modes being apparent that indicates the intricacy of the sensitivity of the theory when encountering general structures, and the immaturity in which the theoretical landscape contains many noises.

3.1.2 Related works

On itself, the phenomena *double descent* has been investigated somewhat comprehensively, firstly introduced by Belkin et al. [2019a], which gives it distinctive saddle escape illustration. Further analysis was made by several literatures, particularly attributed the existence of double descent to the concept of model complexity and inductive bias, and potential explanations of such. Nakkiran et al. [2019] expanded the phenomena into deep neural network models. Their conclusion is reached by considering the perturbation of a learning procedure $\mathcal{T}$ on the effective model complexity $\text{EMC}_{\mathcal{D},\epsilon}(\mathcal{T})$ defined in the paper, separating the eventual phenomena into regions of observations, thereby in one way or another, predicting the tendency of double descent. This is also the first, arguably, "concise" definition of double descent, even though its nature is an empirical definition, including the notion of model complexity. Recently, there is also the Neural Tangent Kernel (Jacot et al. [2018]) which can be used to explain double descent, however further researches are still required. Preliminary works are done by Lafon and Thomas [2024], Schaeffer et al. [2023], and Liu and Flanigan [2023] on the role of optimization in double descent. Davies et al. [2023] attempted to unify grokking with double descent, theorized a possibility of similarity during the generalization phase transition of the inference period, and Olmin and Lindsten [2024] attempted to explain epoch-wise double descent within the model of two-layer linear network. However, it is mentioned that it can also be expanded into deep nonlinear networks.

Further literatures on double descent varies, but fairly new, ever since their discovery articles in Belkin et al. [2019a] and Nakkiran et al. [2019]. Since then, we have Shi et al. [2024], Schaeffer et al. [2023], Quéto and Tartaglione [2023a], Lafon and Thomas [2024], Neal [2019], Davies et al. [2023], Quéto and Tartaglione [2023b], Liu and Flanigan [2023], Mei and Montanari [2020, 2019a], Adlam and Pennington [2020], Transtrum et al. [2025], Liu et al. [2021], Allerbo [2025], Pezeshki et al. [2021], Brellmann et al. [2024], Nakkiran [2019b], Mei and Montanari [2019b], Heckel and Yilmaz [2020], Nakkiran et al. [2020], Belkin et al. [2020], Yang and Suzuki [2024], Spiess et al. [2023], Nakkiran et al. [2021], Zhang [2024], Cherkassky and Lee [2024], Nakkiran [2019a]. An even more exotic version of double descent, dubbed *triple descent* (or more), is discovered in d'Ascoli et al. [2020].

3.2 Foundational investigation

Several foundational notions must be clarified before it is able to take in the major works about double descent. Specifically, we need to clarify the structure of classical learning, relevant structures, notions of underfitting, overfitting, overparameterization and underparameterization, double descent and what it is counteracting, and so on. Such will also be followed by a section strictly the particular solutions and empirical corrections of theories.

For the need of analysing the phenomenon of double descent, we begin by listing and reviewing different foundational notions that describe the typical theory on the learning dynamics, and the learnability problem. Such ranges from statistical learning theory, as seen of Vapnik [1999], Hajek and Raginsky [2021], Mohri et al. [2012], and the more practical learning theory, as seen in James et al. [2009].

The setting of typical machine learning system comes in two-fold, either from the outer objective of the *observation space* $\{x_i, y_i\}_{i \leq n}^n$, denoted $\mathcal{S}$ for any 2-tuple of that form, or the general, oracle-induced view in which there exists the main form, or the oracle $c \in \mathcal{C}$, of which form the functional spaces. We would like to review these two outlooks.

Because we are to introduce certain formalism, or at least interpretation of the learning process, as well as the different dynamical components associated, the preliminary would be first reserved to clarify such notions, alongside fixing usual notations that would be used accordingly. Historically speaking, the

learning theory started from the early 1970s. By the sense of classical, we do not mean by, as currently called, of non-deep learning theory, but rather the old treatment of theoretical boundaries and rigorous formalization of the concept itself. In such development, lies the formal setup of learning, and then, the task of which developing it, until we can even reach the *Probably Approximately Correct* learning (PAC) and Vapnik-Chervonenkis theory of maximal dimension, Neural Tangent Kernels (Jacot et al. [2018]), and so on.

The nature of the learning theory can also be brought back from certain proposition or argument, most probably, the fact that it is mostly from the black-box interpretation – hence making almost all model being **phenomenological** by default. An interesting aspect of such, though, is the learning process itself, or we call it so.

3.2.1 Preliminary - precursor double descent estimation

In a glance, we can fairly setup a condition on a type of specific learning criterion specification, to predict the existence of the pattern for double descent. The PAC-learning process (Shalev-Shwartz and Ben-David [2014], Mohri et al. [2012], Hajek and Raginsky [2021]), or rather, *Probably Approximately Correct* learning, connects three complexity measures – the sample space complexity, representation complexity, and optimization complexity, to the evaluation of the statistical learning process. It refers to the categorization of various learning procedure that satisfies the Probably Approximately Correct criteria in its name, as such being followed. The theory was proposed first by Valiant [1984], of whom then won the Turing Award, but his contribution stayed in the shadow in much of the end of the century, until the formal foundation of machine learning theory, of which then it becomes the stay ground on the standard discussion on the story of machine learner and algorithmic guarantees. As such, the de-facto foundation is such.

PAC-learning procedure gives, by assumption of a **consistent learner**, two different generalization bounds for two different cases. We first define the notion of a consistent learning algorithm, or consistent learner, for a concept class C and hypothesis h .

Definition 3.2.1 (Consistent Learner). We say that an algorithm is a consistent learner for a concept class C using hypothesis class $\mathcal{H}$, if for all n , for all $c \in C_n$ and for all m , given

$$S = \{(x_1, c(x_1)), (x_2, c(x_2)), \dots, (x_m, c(x_m))\}$$

as input, $x_i \in X_n$ outputs, then the algorithm $\mathcal{A}$ outputs a hypothesis $h \in \mathcal{H}_n$ such that

$$h(x_i) = c(x_i), i = 1, \dots, m$$

A consistent hypothesis is the hypothesis that is consistent with all the training data provided to it from the concept c . The two bounds then contain the debate between consistent hypothesis (i.e., for example, admitting no error on training data, almost perfect – maybe even so), and inconsistent hypothesis, of which have general defects from the actual concept, or even different by a margin. This equates to observational noise σ , typically used in simulation of noisy feature and spread, and also for randomization. Inconsistent then means the further this particular noise is increase, the higher the lower bound of the error can be for any particular measure, in comparison to the true concept c in consistent case. We need to note that this inherently pushes the consistency as for both the algorithm L , and the hypothesis $\mathcal{H}$. Decomposing such definition, we can clearly see two assumptions:

Assumption 3.2.1 (Consistent learner, condition 1). *The algorithm is capable of extrapolate the entire coverage structure of class $\mathcal{H}$. Thus, $\mathcal{A}(R, \ell)$ of the objective condition is guaranteed upon all class $\mathcal{H}$.*

Assumption 3.2.2 (Consistent learner, condition 2). *There exists $h \in \mathcal{H}$ that absolutely equals to c . That is, for the concept c of interest, on the arbitrary expression space of which contains $\mathcal{H}, \mathcal{C}$, then for a subset $\mathcal{C} \supset \mathcal{C}' \supset c$ as the size of $|\mathcal{C}'| \rightarrow 0$, $\mathcal{C}' \subset \mathcal{H}$.*

On such guarantee of the formalization and assumption of coverage, PAC-learning famously leaves two results on the learning bound of the learning setup. The first one concerns of the form as to be consistent, and the second case is for where the second assumption fails. However, we have no correct formalization of this failure mode. Specifically, the learning bound intricacy states that of the **global correctness**, that is, absolutely fitting. Usually this is not true, and the inconsistent case is much more possible, but yet again loose as to only examine the case where consistency fails, not *how much* or how many.

Theorem 3.2.3 (Learning bound – finite $\mathcal{H}$, consistent case). *Let H be a finite set of functions mapping from $\mathcal{X} \rightarrow \mathcal{Y}$. Let $\mathcal{A}$ be an algorithm that for any target concept $c \in H$ and i.i.d. samples S returns a consistent hypothesis h_S , such that $\hat{R}(h_S) = 0$. Then for any $\epsilon, \delta > 0$, the inequality $\Pr_{S \sim D^m}[R(h_S) \leq \epsilon] \geq 1 - \delta$ holds if*

$$m \geq \frac{1}{\epsilon} \left(\log |H| + \log \frac{1}{\delta} \right)$$

This sample complexity result admits the following equivalent statement as a generation bound: for any $\epsilon, \delta > 0$, with probability at least $1 - \delta$,

$$R(h_S) \leq \frac{1}{m} \left(\log |\mathcal{H}| + \log \frac{1}{\delta} \right) \quad (3.1)$$

Proof. See [Mohri et al. \[2012\]](#) for the original proof. □

The theorem is illustrated as followed in Figure 3.1 and Figure 3.2:

As for inconsistent case, the margin of error exists, thus $\mathcal{H}$ is bounded like the following.

Theorem 3.2.4 (Learning bound – finite $\mathcal{H}$, inconsistent case). *Let $\mathcal{H}$ be a finite hypothesis set. Then, for any $\delta > 0$, with probability at least $1 - \delta$, the following inequality holds:*

$$\forall h \in \mathcal{H}, \quad R(h) \leq \hat{R}_S(h) + \sqrt{\frac{\log |\mathcal{H}| + \log 2/\delta}{2m}} \quad (3.2)$$

Proof. This proof is based on parts of [Mohri et al. \[2012\]](#) and the section about PAC theorems. In the most general case, there may be no hypothesis in $\mathcal{H}$ consistent with the labelled training sample. This, in fact is the typical case in practice, where the learning problems may be somewhat difficult or the concept classes more complex than the hypothesis set used by the learning algorithm.

To derive the learning guarantees in the more general setting, we would use the following corollary, of which relates the generalization error and empirical error of a single hypothesis.

Corollary 3.2.1. *Fix $\epsilon > 0$. Then, for any hypothesis $h : X \rightarrow \{0, 1\}$, the following inequalities hold:*

$$\mathbb{P}_{S \sim \mathcal{D}^m} \left[\hat{R}_S(h) - R(h) \geq \epsilon \right] \leq \exp(-2m\epsilon^2) \quad (3.3)$$




Figure 3.1: Illustration of sample complexity in terms of cardinality $m = |\mathcal{S}|$, with respect to ϵ or δ measure.

$$\mathbb{P}_{S \sim \mathcal{D}^m} [\hat{R}_S(h) - R(h) \leq -\epsilon] \leq -\exp(-2m\epsilon^2) \quad (3.4)$$

By the union bound, this implies the following two-sided inequality:

$$\mathbb{P}_{S \sim \mathcal{D}^m} [|\hat{R}_S(h) - R(h)| \leq \epsilon] \leq 2\exp(-2m\epsilon^2) \quad (3.5)$$

Proof. This can be proved using Hoeffding inequality. Recall that the inequality is stated for $Z_1, \dots, Z_n$ independent random variable, for range $z_i \in [a_i, b_i]$ with probability one, $S_n = \sum_{i=1}^n Z_i$ then for all $t > 0$:

$$\mathbb{P}(S_n - \mathbb{E}(S_n) \leq -t) \leq \exp\left(\frac{-2t^2}{\sum (b_i - a_i)^2}\right)$$

and

$$\mathbb{P}(S_n - \mathbb{E}(S_n) \geq t) \leq \exp\left(\frac{-2t^2}{\sum (b_i - a_i)^2}\right)$$

We can derive this for our case of the two empirical and generalization risk. Remember from 6.2.3 that we have $R(h) = \mathbb{E}[\hat{R}_S(h)]$, we can transform the phrase to be:

$$\mathbb{P}_{S \sim \mathcal{D}^m} [\hat{R}_S(h) - R(h) \geq \epsilon] = \mathbb{P}_{S \sim \mathcal{D}^m} [\hat{R}_S(h) - \mathbb{E}[\hat{R}_S(h)] \geq \epsilon]$$




Figure 3.2: Illustration of of generalizational error/risk in terms of cardinality $m = |\mathcal{S}|$, with respect to changing $\mathcal{H}$ cardinality and δ -value measure. Blurred line indicates different cases or values, the deviation is rather small, thus the bound is very tight.

with respect to $\mathcal{D}^m$. The form of $\hat{R}_S(h)$ has an additional $1/m$, hence, for simplification, we will take the form

$$\mathbb{P}_{S \sim \mathcal{D}^m} \left[\hat{R}_S(h)/m - \mathbb{E}[\hat{R}_S(h)]/m \geq \epsilon m \right]$$

By Hoeffding inequality, we have:

$$\mathbb{P}_{S \sim \mathcal{D}^m} \left[\hat{R}_S(h)/m - \mathbb{E}[\hat{R}_S(h)]/m \geq \epsilon m \right] \leq \exp \left(\frac{-2m^2\epsilon^2}{\sum (b_i - a_i)^2} \right) \quad (3.6)$$

which is

$$\mathbb{P}_{S \sim \mathcal{D}^m} \left[\hat{R}_S(h)/m - \mathbb{E}[\hat{R}_S(h)]/m \geq \epsilon m \right] \leq \exp \left(\frac{-2m\epsilon^2}{(b-a)^2} \right)$$

Here, our $\sum (b_i - a_i)^2$ was replaced by a more general bound which contains only the sufficient amount of m samples. This is because the sum follows from m samples, which are taken as random variable of choice, then, for $a \leq Z_i \leq b$

$$\frac{m^2}{\sum (b_i - a_i)^2} = \frac{m^2}{m(b-a)^2} = \frac{m}{(b-a)^2}$$

Since our hypothesis is mapping to $\{0, 1\}$, hence our result space is discretely bounded in the normal range, the inequality reduces to

$$\mathbb{P}_{S \sim \mathcal{D}^m} \left[\hat{R}_S(h)/m - \mathbb{E}[\hat{R}_S(h)]/m \geq \epsilon m \right] \leq \exp(-2m\epsilon^2) \quad (3.7)$$

Rearranging the left-hand-side to the original yields the inequality. Doing the same for the second case, and using the union bound, we get the two-sided inequality for $|\hat{R}_S(h) - R(h)|$. $\square$

Such bound can be naturally extended toward any arbitrary range $[a, b]$, as for Hoeffding inequality being defined on the arbitrary closed interval instead. Then, we have the general bound

$$\mathbb{P}_{S \sim \mathcal{D}^m} \left[|\hat{R}_S(h) - R(h)| \leq \epsilon \right] \leq 2 \exp - \frac{2m\epsilon^2}{(b-a)^2} \quad (3.8)$$

Setting the right-hand side of the final expression to be equal to δ and solving this for ϵ yields immediately the following bound for a single hypothesis.

Corollary 3.2.2 (Generalization bound – single hypothesis). *Fix a hypothesis $h : \mathcal{X} \rightarrow \{0, 1\}$ or more generally onto $h : \mathcal{X} \rightarrow [a, b] \subset \mathbb{F}$. Then, for any $\delta > 0$, the following inequality holds with probability at least $1 - \delta$:*

$$R(h) \leq \hat{R}_S(h) + \sqrt{\frac{\log 2/\delta}{2m}} \quad (3.9)$$

The desired bound can then be realized as extending the proved corollary bound, by noticing that,

$$\hat{R}_S(h) + \sqrt{\frac{\log 2/\delta}{2m}} < \hat{R}_S(h) + \sqrt{\frac{\log |\mathcal{H}| + \log 2/\delta}{2m}} \quad (3.10)$$

for all $|\mathcal{H}|$. $\square$

For a looser bound, we can ultimately use an ϵ -consistent learner criteria in such case. PAC-learning setting raises some peculiar insight toward the problem itself. First, is the learning bound itself of which raises the importance of several factors — the dataset size, m , of which is used as the **total resources required** for approximating such function. Furthermore, the set complexity $\mathcal{H}$ is also used, albeit questionably inconsistent in the term $\log |\mathcal{H}|$, which implies at least consideration on the cardinality of the hypothesis class. In a finite domain, we notice that PAC-learning of consistent case is exactly the case of interpolation, assuming non-smooth and noisy data, and the inconsistent case describe the wider range of approximation-based case, but does not consider an approximation-aware bound. The design of $|\mathcal{H}|$ also does not allow for certain exact hypothesis class size measurement, except for combinatorial case where $\mathcal{H}$ is indeed finite. Nevertheless, the bound indeed indicates that for m increase, the hypothesis class often do much better in general. One particularly troublesome notion of which we will meet, in the inconsistent case, is the reversal of guarantee compared to bias-variance form. Specifically, the inconsistent, model-agnostic version of Theorem 3.2.4 is as followed:

Theorem 3.2.5 (Inconsistent case, model-agnostic, finite $\mathcal{H}$). *For finite $\mathcal{H}$, then for any $\delta > 0$, up to ϵ -accuracy, with probability at least $1 - \delta$, the following inequality holds:*

$$R(h_S) \leq \min_{h_S \in \mathcal{H}} \hat{R}_S(h_S) + \sqrt{\frac{\log |\mathcal{H}| + \log 2/\delta}{2m}} \quad (3.11)$$

such that

$$\sqrt{\frac{\log |\mathcal{H}| + \log 2/\delta}{2m}} \leq \epsilon \quad (3.12)$$

for $m > 0$ samples. We also say that the error varies by $\mathcal{O}(\sqrt{1/m})$, such that

$$R(h_S) \leq \min_{h_S \in \mathcal{H}} \hat{R}_S(h_S) + \mathcal{O}\left(\sqrt{\frac{\log |\mathcal{H}|}{m}}\right) \quad (3.13)$$

We denote, again, $\hat{R}_S$ the empirical risk given on set S of observations, and h_S the hypothesis obtained from dynamics regarding such observational set.

This theorem directly indicates that, for bigger $|\mathcal{H}|$, the rate of which m increases can also mitigate bias-variance, thus bigger dataset can mitigate such issues of overfitting. Indeed, we notice that $\mathcal{O}(\log |\mathcal{H}|) \leq \mathcal{O}(m)$ linear terms, thus if the algorithm $\mathcal{A}$ exhibits such dynamics in which dataset varies, bigger hypothesis space should not increase significantly of the error measure. However, this contradicts, in a sense, with the usual term of notion to bias-variance, of which suggest otherwise, hence overfitting. Or, rather insidiously speaking, the interpretation is wrong, such that bias-variance concern fixed m situation, while such bound indicates, or at least suggest, of variational m .

While this does not cover the case of which $|\mathcal{H}|$ is argued to be infinite, such is then also be bounded, as for any measure, temporary called $\mu : \mathcal{H} \rightarrow \mathcal{F}$ for field $\mathcal{F}$ of which can be evaluated to $\mathbb{R}$ or else, then such is bounded in the contribution terms, and thus of the framework in which we use to express the learning dynamics. It is also then to be concerned of different classes of $\{\mu\}$ itself, as a limiting factor. This, perhaps, will lead to a potential solution of double descent.

Indeed, a loose analogy, or rather speculation can follow. While criticisms of classical uniform-convergence PAC learning correctly observe that one-dimensional capacity sweep indeed, fail to predict or explain the modern double descent curve, such cannot be said in certain cases of the inconsistent setting. Suppose we introduce the parameter ρ of the growth regime on $\mathcal{M}(\mathcal{H}) : \mathcal{H} \rightarrow \mathbb{F}^k$ of the complexity measure or the size of the hypothesis space, in which usually $k = 1$ and the value is on the real field, for $\gamma = |S|$ as the sample size, such that we have τ correlates to a particular trajectory in which hence the escape trajectory is then controlled subsequently. We can then interpret the τ parameter as encoding the moving **interpolation threshold**. The parameter τ then can be computed, accordingly for the optimization term as

$$\mathcal{O} \left(\sqrt{\frac{\log |\mathcal{H}|}{m}} \right) \equiv \sqrt{\frac{\log \tau + \log 2/\delta}{2\tau}} \quad (3.14)$$

This particular trajectory, can indeed, create certain cases where the right-half of the curve is similar to the trajectory of double descent. In general, investigation would see for

$$\mathcal{O} \left(\sqrt{\frac{\log |\mathcal{H}|}{m}} \right) \equiv \sqrt{\frac{\log q(\tau) + \log 2/\delta}{2q(\tau)}} \quad (3.15)$$

for arbitrary function $q(\tau), p(\tau)$ where it should exhibit double descent terminal terms. However, this task is really hard aside from simply guessing. A simple guess on structural reformulation would be that

$$\min_{h_S \in \mathcal{H}} \hat{R}_S(h_S) \approx \frac{1}{m} \left(\log |\mathcal{H}| + \log \frac{1}{\delta} \right), \quad (3.16)$$

$$\mathcal{O} \approx I \sqrt{\frac{\log(\tau - p) + \log(2/\delta)}{2(\tau - p)^\epsilon + 1}} \quad (3.17)$$

Thus the sum would be defined in a more orthodox way, as

$$\frac{1}{m} \left(\log |\mathcal{H}| + \log \frac{1}{\delta} \right) + I \sqrt{\frac{\log(\tau - p) + \log(2/\delta)}{2(\tau - p)^\epsilon + 1}} \quad (3.18)$$

for arbitrary constant I, p , of which now the addition operator would be slightly adjust, such that the sum will contain only the LHS of the operator if the RHS is in its undefined domain, and otherwise they add on together. The factor $(\tau - p)$ appears in both previously the sample size m and the hypothesis space size $|\mathcal{H}|$. This result is highly speculative, as it is an educated guess on the error landscape, however, still produces the approximated term on double descent. Such is then of the following hypothesis.

Hypothesis 3.2.1 (The flawed PAC-descent). *Let $\mathcal{H}$ be the hypothesis class, for moving parameter τ called the interpolation threshold. We define the interpolation offset region, $\Delta(\tau) \equiv \Delta := \tau - p$, and the regularized deviation term $\eta, \kappa > 0$. Then, for specific $q(\tau)$, where we set $\tau = |\mathcal{H}|$ of the hypothesis class size, and $\tau = m$ case of the number of samples on $\mathcal{S}$, then $R(h_{\mathcal{S}})$ is bounded by, for at least probability of $1 - \delta$:*

$$R(h_{\mathcal{S}}) \leq \begin{cases} \frac{1}{m} (\log |\mathcal{H}| + \log \frac{1}{\delta}) & \Delta \leq -\gamma \\ \frac{1}{m} (\log |\mathcal{H}| + \log \frac{1}{\delta}) + \lambda(\Delta) I \sqrt{\frac{\log(\tau-p) + \log(2/\delta)}{2(\tau-p)^{\epsilon} + 1}} & |\Delta| < \gamma \\ \frac{1}{m} (\log |\mathcal{H}| + \log \frac{1}{\delta}) + I \sqrt{\frac{\log(\tau-p) + \log(2/\delta)}{2(\tau-p)^{\epsilon} + 1}} & \Delta > \gamma \end{cases} \quad (3.19)$$

for $I > 0$, $\lambda(\Delta) \in [0, 1]$ the smooth terms on transition region, where taking $\gamma = 0$, $\lambda(0) = 1/2$ returns the discrete separation case.




Figure 3.3: Prediction made by certain configurations of Hypothesis 3.2.1. Here, the injected PAC-type decay term provides the global envelope, while the interpolation offset term introduces singular behaviour near capacity-sample equality. The resulting superposition reproduces a family of double-descent-like curves depending on the trajectory in $(m, |\mathcal{H}|)$ -space.

While imperfect, this encodes in the simplest form of PAC-system, why such bound can somewhat say of the double descent phenomena, in terms of the singular trajectory error instead. But it also reveals the problematic trouble within doing this, essentially move the theory into the untested and unprepared zone in which orthodox choices has to be made, but no interpretation is available. The ineffective side of such derivation is that:

1. Ill-defined terms, $\log(\tau - p)$ is extremely brittle and thus makes for an interpolation threshold, but theoretically, this is inexplicable, also for $(\tau - p)^\epsilon$ for the error ϵ bounded. Even with the regularized term, such is forced on because of necessity, not structural reason.
2. Circularity of the concept is that we are treating on the assumption of the escape trajectory, i.e. $(\tau - p)$, will fit both the hypothesis space size and the sample size. Indeed, such procured the double descent curve, but does not explain how and why.
3. Continuity and unit mismatch is indeed troublesome, as there exists no way to tell specific $q(\tau)$ that preserves the parameterization to τ . Continuity means having to form piece-wise terms, hence we result in an abrupt bounding instead.

Inherently, τ is also empirical, and thus $q(\tau)$. Such expression in general appears in KL-convergence of PAC-Bayes, and linear models isotropic design gives $\sigma^2 p / (m - p - 1)$, which is fairly similar to what we have. Nevertheless, one should realize that current theories are not equipped to express this in particularly detailed insight. In certain cases, it only serves as heuristics or empirical presumptions, or rather theoretical speculation in a non-constructive way. Imitating double descent aside, this prompts us to seriously tackle the double descent problem, or otherwise the messiness of analysis will continue. For a brief explanation on why this is true, we can agree on an exercise. Despite appearing different and using different control nobs, the major classical theories all rely on the same structural primitives, where the dataset is assumed to be sampled from independent $\mathcal{D}$ -distributing sampling operator for m samples of dataset $\mathcal{S}$, hypothesis class $\mathcal{H}$ arbitrarily, risk $R(h)$, and the empirical risk $\hat{R}(h)$, or $\hat{R}_{\mathcal{S}}(h)$ on strictly $\mathcal{S}$ -conditioned evaluation space. The theory only attempts to control $R(h) - \hat{R}_{\mathcal{S}}(h)$ or its variants in the general risk expression, using a capacity measure. This can be for (1) VC theory, as VC dimension; (2) Rademacher theory as the Rademacher complexity; (3) Occam learning and compression scheme using description length, and PAC-Bayes for KL-divergence between posterior and prior. Those are different regularizers on the same skeleton, but they also inherit the same structural blind spot of PAC. More specifically, there exists no variables in-theory that describe interpolation rank or the existence thereof, feature geometry, optimisation trajectory, and formal condition on parameterization structure and organization. Usually, the bounds are always

$$R(h) \leq \hat{R}_{\mathcal{S}}(h) + \mathcal{O}(|\mathcal{H}|), \quad \mathcal{O}(|\mathcal{H}|) \in \left\{ \sqrt{\frac{\text{comp}(h)}{m}}, \frac{\text{KL}(Q||P)}{m}, \mathfrak{R}_m(\mathcal{H}) \right\} \quad (3.20)$$

of which presuppose complexity condition. Even NTK as a modern analogue suffers the same problem, but with different reason than the PAC-family of derivatives. For a quick review, we see that consider a neural network $f(x; \theta)$ on parameterized complexity scheme $\theta \in \mathbb{R}^P$ of parameter space, Neural Tangent Kernel (NTK) is defined as

$$\Theta(x, x', \theta) = \nabla_{\theta} f(x; \theta) \cdot \nabla_{\theta} f(x'; \theta) \equiv \sum_{p=1}^P \frac{\partial f(x; \theta)}{\partial \theta_p} \frac{\partial f(x'; \theta)}{\partial \theta_p} \quad (3.21)$$

for x' as another input point in the input space, thus NTK is a local comparison scheme of the kernel to show how the network output changes with respect to its parameter. NTK can be said to determine the training dynamics in function space, and its main result, where in [Jacot et al. \[2018\]](#) is stated as

$$\Theta(x, x' \theta(t)) \rightarrow \Theta_{\infty}(x, x'), \quad n_{\ell} \rightarrow \infty \quad (3.22)$$

where ℓ denotes the layer widths, and the result takes that it remains constant during training. This collapses the network into linearized initialization condition, for training becomes kernel version of gradient descent, and furthermore, prediction is essentially collapsed to kernel regression with the kernel Θ_∞ . Such is to say, in the NTK regime, parameters move infinitesimally, for the kernel to remain **fixed**, and static feature space, the model becomes kernel regression in fixed feature space. Structural consequence is that it freezes the feature representation, the learning problem is reduced to the condition on

$$\text{ERM}_{\text{NTK}}(h) \equiv \min_{\alpha \in \mathbb{R}^m} \hat{R}(\alpha) \equiv \min_{\alpha \in \mathbb{R}^m} \sum_{i=1}^m \left(\sum_{j=1}^m \alpha_j \Theta(x_i, x_j) - y_i \right)^2 \quad (3.23)$$

for empirical risk minimization, with closed form solution as for α learnable coefficients being $(K + \lambda I)^{-1}y$, which is essentially kernel ridge regression. The behaviours are now flat into eigenvalue of K , the sample size m , and the noise level. Because of such reduction, NTK struggles with phenomena, either by structural incompatibility or bloats, in explaining representation learning, hierarchical feature formation, deep network scaling laws in loose form, neural collapse phenomenon, large model generalization behaviours, and specifically relevant as **double descent empirical regimes**. Thus, the reason is the same absence of variables, but the reason differs within simplifications.

That said, what do we gain from this modification? Specifically, we have shown that the coupling on the term, loosely said as **interpolation threshold** for the lack of better word, though much of the opposite, can be used to inject structures into the conditioned space that PAC bound still works, but cannot or rather wasn't equipped of its analytical toolkit to analyse, but attains double descent. Even in PAC-setting, the effective counteraction would only to recognize of which traversing paths are there of the $(m, |\mathcal{H}|)$ -space we are operating on, and the smooth convergence curve of PAC-setting is also consistent with the explanation, as it restricts into a landscape, but not remove the probability of extension on itself. By then, we can identify three fixes required of a deeper theoretical analysis: the hypothesis growth trajectory $q(\tau)$, the sample growth $p(\tau)$, and the interpolation offset, all of which can be dependent on various properties of extended system in hindsight, and of which are dependent on various transitional assumptions that cannot be structurally derived on PAC, even if it allows for such. This is extended also to VC-theory and a class of assumption-based, Rademacher-similar heuristics on model complexities. Lastly, we emphasize from this observation that, (1) geometry is not inherently only in $|\mathcal{H}|$, especially if we proceed to more dynamically complex and often pathological work cases; (2) worse-case bounds are loose enough for us to inject geometry into it, though it means conditions are still possibly consistent with the later analysis; (3) deep networks are inherently different, however, an action of streamlining model interpretability and structural functions of such, different from simply mathematical reduction is necessary. An example of the last point is explained in the later chapters, where we investigate the polynomial system, where of which it is actually that we are using a very poorly conditioned Vandermonde polynomial basis, much less of the simple mathematics on model descriptions.

3.3 The bias-variance - double descent dilemma

Before continuing further, we must review the problem statement that led to the problem of double descent. We will have a look at the theory, the original conflict, and the illustrative phenomenon in action. In general, the learning setting in consideration of statistical learning theory is often supervised learning, for guarantee of a plausible upper bound result.

3.3.1 Bias-variance tradeoff

In classical sources, [James et al. \[2009\]](#), [Goodfellow et al. \[2016\]](#), [Hajek and Raginsky \[2021\]](#), [Mohri et al. \[2012\]](#), [Shalev-Shwartz and Ben-David \[2014\]](#), the principle of bias-variance tradeoff is a classical principle used in the categorization of technique called *model selection*. Specifically, this refers to the progress of choosing the best predictor h^* that minimize the generalization error that is the main goal of a particular machine learning model, during the process of training and inference. Originally, the theory that bias and variance often come in to conjunction is **Estimation theory**, of which there exists an estimator $\hat{\theta}$ that use a set of observations, to estimate the concept, or the underlying mechanics of certain system that outputted the observation set, by the **probabilistic perspective** - that is, it is governed by appearance by an independently and identically distributed process of parameterized probability $\theta \in \Theta$. Here, parameterized means being expressed by a set, often finite, of parameters, by [E. L. Lehmann \[1998\]](#), [Paninski \[2005\]](#). It is only when [Geman et al. \[1992\]](#) introduced it to machine learning that we have an equivalent structure that is similar to bias-variance in machine learning.

Precursor ([Geman et al. \[1992\]](#))

The original definition of bias-variance tradeoff by [Geman et al. \[1992\]](#) is first constructed using the means-square error, which is regarded as a normal measure in the real encoding space. Their approach is to justify bias-variance via decomposition of the loss function ℓ , for such to find an alternative reasonable form of such loss landscape. Suppose of a regression problem to construct a hypothesis function $f(x)$ from $(x_1, y_1, \dots, x_N, y_N)$ for the purpose of generalization - that is, predicting unseen variational values for different pair $(x_j, ?)$ such that $? = y_j + \epsilon$ for a conceivable implicit error. To be explicit about the relation of this problem, or f on the given data $\mathcal{D} = \{(x_i, y_i) \mid i \leq N\}$, denote $f(x; \mathcal{D})$ instead of f , the natural mean-square measure as a predictor is:

$$\mathcal{M}(f, y) = \mathbb{E} [((y - f(x; \mathcal{D})))^2 \mid x, \mathcal{D}] \quad (3.24)$$

for $\mathbb{E}[\cdot]$ the expectation wrt to a distribution P . Decomposing the right-hand side, we have:

$$\begin{aligned} \mathcal{M}(f, y) &= \mathbb{E} [((y - f(x; \mathcal{D})))^2 \mid x, \mathcal{D}] \\ &= \mathbb{E} [(y - \mathbb{E}[y \mid x])^2 \mid x, \mathcal{D}] + (f(x; \mathcal{D}) - \mathbb{E}[y \mid x])^2 \end{aligned} \quad (3.25)$$

Here, $\mathbb{E} [(y - \mathbb{E}[y \mid x])^2 \mid x, \mathcal{D}]$ does not depend on $\mathcal{D}$, but simply the statistical variance of y given x . The term $(f(x; \mathcal{D}) - \mathbb{E}[y \mid x])^2$ is considered a natural measure of effectiveness on $\mathbb{R}^n$ as a singular predictor of y . Now, for $\mathbb{E}_{\mathcal{D}} [(f(x; \mathcal{D}) - \mathbb{E}[y \mid x])^2]$ which depends on the training set $\mathcal{D}$ in its computation, is decomposed into the form of *bias-variance decomposition* terms, by derivation:

$$\begin{aligned} \mathbb{E}_{\mathcal{D}} [(f(x; \mathcal{D}) - \mathbb{E}[y \mid x])^2] &= \underbrace{\{\mathbb{E}_{\mathcal{D}}[f(x; \mathcal{D})] - \mathbb{E}[y \mid x]\}^2}_{\text{bias term}} \\ &\quad + \underbrace{\mathbb{E}_{\mathcal{D}} \{(f(x; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[f(x; \mathcal{D})])^2\}}_{\text{variance term}} \end{aligned} \quad (3.26)$$

We summarize this in the following statement.

Theorem 3.3.1 (Bias-variance decomposition). *Suppose the model $f(x; \mathcal{D})$ for the data $\mathcal{D} = (x_i, y_i)$ and its parameter x is defined. For y_i of the target concept's responses y , and consider a regression problem with the loss*

measure $\mathcal{M}(f, y)$ of mean squared risk, the following statement is true:


$$\begin{aligned} \mathbb{E}[\mathcal{M}(f, y)] &= \mathcal{B}(f, y) + \mathcal{V}(f, y) \\ &\quad + \mathbb{E} \left[\mathbb{E} \left[((y - f(x; \mathcal{D}))^2 \mid x, \mathcal{D}) \right] \right] \end{aligned} \quad (3.27)$$

for $\mathbb{E}[\cdot \mid x, \mathcal{D}]$ any expression with dependencies on x and $\mathcal{D}$. The bias and variance term is subsequently expressed by

$$\mathcal{B}(f, y) = \underbrace{\{\mathbb{E}_{\mathcal{D}}[f(x; \mathcal{D})] - \mathbb{E}[y \mid x]\}^2}_{\text{bias}}, \quad (3.28)$$

$$\mathcal{V}(f, y) = \underbrace{\mathbb{E}_{\mathcal{D}} \{ (f(x; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[f(x; \mathcal{D})])^2 \}}_{\text{variance}} \quad (3.29)$$

Figure 3.4: **Decomposition of the error term into 3 parts.** Respectively, the irreducible error $\mathbb{V}[y(\mathbf{x})]$, the variance (blue) and the bias (dark yellow). All of them are assumed to take up 100% of the error observed in such composition. The proportion is included using the three coefficient $\lambda_{\epsilon}, \lambda_V, \lambda_B$ where $\lambda_{\epsilon} + \lambda_V + \lambda_B = 1$.



The above decomposition principle is often expressed into a form where there exists the intrinsic noise [Brown and Ali \[2024\]](#):

$$\begin{aligned} \mathbb{E}_{\mathcal{D}} \left[\mathbb{E}_{xy} (y - \hat{f}(x))^2 \right] &= \mathbb{E}_x \left[\left(y^* - \mathbb{E}_{\mathcal{D}}[\hat{f}(x)] \right)^2 \right] \\ &\quad + \mathbb{E}_x \left[\mathbb{E}_{\mathcal{D}} \left(\hat{f}(x) - \mathbb{E}_{\mathcal{D}}[\hat{f}(x)] \right)^2 \right] \\ &\quad + \mathbb{E}_{xy} \left[(y - y^*)^2 \right] \end{aligned} \quad (3.30)$$

In all, almost all representation of bias-variance decomposition are the same. For simplicity of notation, we adopt the similar form in the standard case of the paper [Adlam and Pennington \[2020\]](#):

$$\mathbb{E} [\hat{y}(\mathbf{x}) - y(\mathbf{x})]^2 = (\mathbb{E}\hat{y}(\mathbf{x}) - \mathbb{E}y(\mathbf{x}))^2 + \mathbb{V}[\hat{y}(\mathbf{x})] + \mathbb{V}[y(\mathbf{x})] \quad (3.31)$$

A main common theme of criticism toward bias-variance tradeoff is the fact that the decomposition is much more general, and intrinsic for the class of *mean squared loss*. However, when considering the naturalness of an error measure and then its direct applicant, mean square error on real space of sufficient support and measure comes of as a very natural choice of an error measure, at least in consideration of the regression setting; for that, we then can also define classification as another top-layer above a regression's similar real continuous space, i.e. a decision layer. Furthermore, it can also be shown [Brown and Ali \[2024\]](#), [Pfau \[2013\]](#) that it also holds for the class of Bregman divergence measure.

What does the decomposition say about model selection principle in general? In general, bias-variance is typically presented of its decomposition only, and the asymptotic prediction of its behaviours. In fact,

one of the reason that it became the rule-of-thumb for ML practitioner, as well as generally statistical learning (Lafon and Thomas [2024] provides a quite rigorous treatment of bias-variance tradeoff in the section on statistical learning theory) solidify the trade-off as a particular model selection principle. Generally, this tradeoff can be summarized as followed:

Theorem 3.3.2 (Bias-variance tradeoff). *For the expected loss of any given hypothesis h , the bias $\mathcal{B}(f, y)$ and variance $\mathcal{V}(f, y)$ is inversely proportional, that is, $\mathcal{B}(f, y) \propto \lambda^{-1} \mathcal{V}(f, y)$ for some proportionality λ that may or may not be constant. In the most general case possible, $\lambda = -1$ on the entire error range.*

The theorem can be typically seen as a conjecture instead. Regardless, the tradeoff is then stated to be of inverse proportionality. Indeed, statistically, we have such tradeoff on a statistical framework in a more concrete sense. For the bias to increase, variance will increase, of which the criterion is inverse - we would like to have more bias but lower variance, and vice versa, according to such theory. Such drives the theory to choose on its model basing on the two interpretations - of increasing capacity but lead to accumulating too much noise in typical statistically noisy systems, or reducing them too much that leads to incapacity to realize higher-order patterns. Thus, variance aims to reduce the *volatility* of the model performance across admitted region of operation, and bias aims to reduce the inherent errors for under-configured models. That said, if we do not want to simply let them be related of strict equality, we can establish an inequality instead, that is, for

$$\mathbb{E}_D \left[\mathbb{E}_{xy} \left(y - \hat{f}(x) \right)^2 \right] \geq \mathbb{E}_x \left[\left(y^* - \mathbb{E}_D[\hat{f}(x)] \right)^2 \right] + \mathbb{E}_x \left[\mathbb{E}_D \left(\hat{f}(x) - \mathbb{E}_D[\hat{f}(x)] \right)^2 \right] \quad (3.32)$$

We do not inherently claim that noise will add to the fix of the equality as a case, intentionally to capture this gap between them. We do note that the equality often only holds in cases of mean-square error on Euclidean metric space. Per usual, we would like to use the abstract notion between $\mathcal{B}$ and $\mathcal{V}$ instead of explicit metric, and thus the following will outline several variations of sort, for the better part of the terms themselves.

Expansion and variations

The bias-variance decomposition and tradeoff is not general. Indeed, works, by Geman et al. [1992], Sharma and Aiken [2014], Domingos [2000d], Adlam and Pennington [2020], Yang et al. [2020] and more all attempted to reconstruct and reformulate bias-variance, and it did leave out a picture of uncertainty regarding the concept. Specifically, there are certainly different way to do the bias-variance decomposition, as classified in Adlam and Pennington [2020] paper. Some of this are particularly specialized in *neural network* training, basing on the main structure of modern machine learning in neural network formalism.

1. **Semiclassical approach:** The semiclassical approach is fairly simple. We can rewrite to another variation of the classical approach to be

$$\begin{aligned} \mathbb{E}_{h \in \mathcal{H}} [\ell(h)] &= \mathbb{E}_x \mathbb{E}_\epsilon [\hat{y}(x) - y(x)]^2 \\ &= \mathbb{E}_x \left(\mathbb{E}_\epsilon [\hat{y}(x)] - y(x) \right)^2 + \mathbb{E}_x \mathbb{V}_\epsilon [\hat{y}(x)] \end{aligned} \quad (3.33)$$

and instead, consider the *structural conformation* of the decomposition, or rather, taking the expectation within $\theta \in \Theta$ of the structural description of $h \in \mathcal{H}$. Then, we shall have

$$\begin{aligned} \mathbb{E}_{SC} [\ell(h)] &= \mathbb{E}_x \mathbb{E}_\theta \mathbb{E}_\epsilon [\hat{y}(x) - y(x)]^2 \\ &= \mathbb{E}_x \mathbb{E}_\theta \left(\mathbb{E}_\epsilon [\hat{y}(x)] - y(x) \right)^2 + \mathbb{E}_x \mathbb{E}_\theta \mathbb{V}_\epsilon [\hat{y}(x)] \end{aligned} \quad (3.34)$$

Do note that the order matters, since the expectation can be constant of a given ordering. And expected value over constant is just the constant.

2. **Symmetric decomposition:** This approach is fairly difficult, we refer to the paper itself of its own justification on the decomposition. The symmetric decomposition relies on the shape of the input, $\mathcal{X} = \{P, D, \epsilon\}$, of which you can remove the ϵ term and retain some decomposition on P and D per distribution.

Generalized decomposition

According to Domingos [2000b,a], the bias-variance decomposition can usually be treated as being controlled by three factor error, with their respective coefficients $\lambda = (\lambda_1, \lambda_2, \lambda_3)$ for such:

$$\begin{aligned} \mathbb{E}_{\mathcal{D}} [d((f, \mathbb{E}[y | x]))] &= \lambda_1 \text{Bias}(f, y) + \lambda_2 \text{Var}(f, y) + \lambda_3 \epsilon(\mathcal{D}) \\ &= \lambda_1 \underbrace{\{\mathbb{E}_{\mathcal{D}}[f(x; \mathcal{D})], \mathbb{E}[y | x]\}}_{\text{bias term}} + \lambda_2 \underbrace{\mathbb{E}_{\mathcal{D}} \{(f(x; \mathcal{D}), \mathbb{E}_{\mathcal{D}}[f(x; \mathcal{D})])\}}_{\text{variance term}} + \underbrace{\lambda_3 \epsilon}_{\text{irreducible error}} \end{aligned}$$

where $\epsilon(\mathcal{D})$ is the irreducible error of the system, depends (theoretically) on the intrinsic imperfection of the dataset $\mathcal{D}$. This method of which the bias-variance-irreducible error are ‘fit’ into the total error by three coefficients $\lambda_1, \lambda_2, \lambda_3$ instead. This does not only scale up the B-V-IE triplet, but also leaves out the unexplainable gap between the scaling, and the actual normal measure.

Approximation-estimation tradeoff

Another closely related notion to bias-variance is the concept of approximation-estimation tradeoff. We refer to Mohri et al. [2012], Lafon and Thomas [2024] for some analysis and mentions.

Let $\mathcal{H}$ be a family of functions mapping $\mathcal{X} \rightarrow \{1, -1\}$. This is the particular case of **binary classification**, in which can be straightforwardly extended to different tasks and loss functions. The *excess error* of a hypothesis $h \in \mathcal{H}$, is the difference between its error $R(h)$ and the Bayes error R^* . This can be decomposed to be the following:

$$R(h) - R^* = \left(R(h) - \inf_{h \in \mathcal{H}} R(h) \right) + \left(\inf_{h \in \mathcal{H}} R(h) - R^* \right) \quad (3.35)$$

The first bracket contains the **estimation error**, and the second bracket contains what is called the **approximation error**. The estimation error depends on the hypothesis h selected. It measures the error of h with respect to the infimum of the error achieved by hypotheses in $\mathcal{H}$, or that of the best-in-class hypothesis h^* when that infimum is reached. The approximation error measures how well the Bayes error can be approximated using $\mathcal{H}$. It is a property of the hypothesis set $\mathcal{H}$, a measure of its richness.

Model selection consists of choosing $\mathcal{H}$ with a favourable trade-off between the approximation and the estimation error. However, in the most general case, this will be done, but not in practice, as it requires the underlying distribution $\mathcal{D}$ to be known to determine R^* , which is not possible. In contrast, the estimation error can be bounded, or can be analysed, using particular metric and analysis. It is however, worth to note that bias-variance and approximation-estimation have a very complicated relationship, for example, in Brown and Ali [2024] shown that they are indeed not the same, and is in fact two different decomposition, where one is the others’ component. To see this, let us compare the two of them in details.




Figure 3.5: Curves for training risk (dashed line) and test risk (solid line). (a) The classical U-shaped risk curve arising from the bias-variance trade-off. (b) The double descent risk curve, which incorporates the U-shaped risk curve (i.e., the “classical” regime) together with the observed behaviour from using high capacity function classes (i.e., the “modern” interpolating regime), separated by the interpolation threshold. The predictors to the right of the interpolation threshold have zero training risk. Reproduced from [Belkin et al. \[2019a\]](#).

3.3.2 Double descent

Nevertheless, of bias-variance and the analogous statistical learning theory concept, the target is the same. It is the dilemma of which is presented in **Occam’s razor**, for choosing the sufficient model of good complexity, or bias, for tradeoff of its generalization ability, or variance. Then there must exist a sweet spot between the axis of bias and variance, since they are as exhibited above inversely proportional to each other. However, double descent seemingly broke the status quo, and insists on an interesting phenomenon – under the same setting, if we ‘crank’ the complexity high enough, we will then reach a point then called the **interpolation threshold**, such that the trend reverse and the error rate, instead of being theorized to go up, goes down to a certain line of lower bound.

The first identification of the double descent phenomena dated back to the paper of Belkin – [Belkin et al. \[2019a\]](#), in which the title is literally “reconciling” modern machine learning practice and the bias-variance tradeoff. In modern machine learning practice, or state-of-the-art developments, models are now bigger than ever. If to notice, we will see that currently models are inherently large, for example, a normal large language model will have from 900 millions (900M) to a few billions, for example 10 billions (10B) parameters. That is not taking into account the overall dynamics and structure of the model, which dictates the operating range and efficiency of the model itself. These model, based on the neural network architecture are somewhat trained to exactly fit (or interpolate) the data, almost certainly so that it turn from a prediction setting to an estimation setting. By statistical learning theory, this would be considered overfitting, and yet, they often obtain very high accuracy on test data.

The main finding that Belkin found is a pattern for how the apparent performance on unseen data depends on model capacity and the mechanism underlying the emergence of double descent. When function class capacity is below the “interpolation threshold”, learned predictors exhibit the classical U-shaped curve from Figure 3.5. The ‘modern’ interpolating regime marks the opposite trend to the right, where the risk starts to decrease up to a lower bound, which then can be called the *optimal descent bound*.

The bottom of the U is achieved at the sweet spot which balances the fit to the training data and the susceptibility to over-fitting: to the left of the sweet spot, predictors are under-fit, and immediately to the right, predictors are over-fit. When we increase the function class capacity high enough (e.g., by increasing the number of features or the size of the neural network architecture), the learned predictors achieve (near) perfect fits to the training data—i.e., interpolation. Although the learned predictors obtained at the interpolation threshold typically have high risk, we show that increasing the function class capacity beyond this point leads to decreasing risk, typically going below the risk achieved at the

sweet spot in the “classical” regime. (Belkin et al. [2019a])

At this point, we would like to also clarify the term interpolation threshold in this case. Usually, under the pretext of overfitting and underfitting, interpolation threshold means the following:

Definition 3.3.1 (Interpolation threshold). Assuming a typical supervised learning setting, for a hypothesis h within class structure $\mathcal{H}$, and a supposed concept c of the class $\mathcal{C}$. Suppose $[X, Y] \subset c \in \mathcal{C}$ being the dataset, or *observable space* of c on which h is tested on. Under the **empirical generalization learning** procedure, partitioning $[X, Y]$ into $[X, Y]_1$ and $[X, Y]_2$, for error evaluated on the first set, $\hat{R}_1(h)$ called training error, and the second set, $\hat{R}_2(h)$ called test error, with optimization applied on the first set, then evaluated statically on the second, the **interpolation threshold** with respect to a particular complexity measure of h , $\mathcal{M}(h)$, is the point $\mathcal{M}_0(h)$ where it is possible for $\hat{R}_2(h)/\hat{R}_1(h) \approx 1$, for $\hat{R}_1(h) \ll \epsilon$, $\hat{R}_2(h) \ll \epsilon$, with $\epsilon \geq 0$.

This basically means that what is learnt, with respect to the model complexity that gauge *how much* and *what* the model can learn, perfectly fits the concept c itself, from the observation space on its own. We can say that, that specific sense, we gain the assumption that the observation space is *fully representative* of the concept class, and the concept class is then fully realized by the model. Such point where that is possible, or there then such behaviour is possible, is called the interpolation point. The definition relies on the statement ‘where it is possible’, hence it still, does not give us a particularly sound notion on to how and what constitute such possibility. The term empirical generalization learning comes from the fact that the procedure in practice of estimating and optimizing train-test-validation set in itself, is an imperfect empirical approximation of generalization testing. While the fact that we cannot gauge generalization error, in a purely blind black-box setting is settled (Mohri et al. [2012], Sugiyama [2015], Shalev-Shwartz and Ben-David [2014]), the method is then to try to mimic such by using generalization-via-hidden observations. To facilitate this, deliberately, the empirical learning procedure is only dynamic during training session, and static throughout the whole testing and validation testing section, and further on.

By default, in general theory, such interpolating point is dangerous. Not only because it is inherently difficult to know exactly what the dataset constitute, whether it is actually representative of the concept that gives rise to it, or the structural information, noises, and interferences in between. It is hence impossible, to a certain point, to truly know the effectiveness of the model beside interpolating the current dataset which forms the observable details.

Another prominent discovery to look at is Nakkiran et al. [2019], on the double descent of deep learning models. The paper serves as the first observation upon which double descent is observed on modern, deep learning theoretic models, especially complex one like Transformer or ResNet networks.

They define *effective model complexity* of $\mathcal{T}$ (w.r.t. distribution $\mathcal{D}$) to be the maximum number of samples n on which $\mathcal{T}$ achieves on average ≈ 0 training error. This is an entirely empirical definition, similarly per definition as VC-dimension.

Definition 3.3.2 (Effective Model Complexity). The *Effective Model Complexity* (EMC) of a training procedure $\mathcal{T}$, with respect to distribution $\mathcal{D}$ and parameter $\epsilon > 0$, is defined as:

$$\text{EMC}_{\mathcal{D}, \epsilon}(\mathcal{T}) := \max \{n \mid \mathbb{E}_{S \sim \mathcal{D}^n} [\text{Error}_S(\mathcal{T}(S))] \leq \epsilon\}$$

where $\text{Error}_S(M)$ is the mean error of model M on train samples S .

Using this definition, their main hypothesis can then be stated as the following three-fold regions:




Figure 3.6: **Left:** Train and test error as a function of model size, for ResNet18s of varying width on CIFAR-10 with 15% label noise. **Right:** Test error, shown for varying train epochs. All models trained using Adam for 4K epochs. The largest model (width 64) corresponds to standard ResNet18. Resued from [Nakkiran et al. \[2019\]](#).




Figure 3.7: **Left:** Test error as a function of model size and train epochs. The horizontal line corresponds to model-wise double descent—varying model size while training for as long as possible. The vertical line corresponds to epoch-wise double descent, with test error undergoing double-descent as train time increases. **Right** Train error of the corresponding models. All models are Resnet18s trained on CIFAR-10 with 15% label noise, data-augmentation, and Adam for up to 4K epochs. Reused from [Nakkiran et al. \[2019\]](#)

Hypothesis 3.3.1 (Generalized Double Descent hypothesis, informal). *For any natural data distribution $\mathcal{D}$, neural-network-based training procedure $\mathcal{T}$, and small $\epsilon > 0$, if we consider the task of predicting labels based on n samples from $\mathcal{D}$ then:*

Under-parameterized regime. *If $\text{EMC}_{\mathcal{D},\epsilon}(\mathcal{T})$ is sufficiently smaller than n , any perturbation of $\mathcal{T}$ that increases its effective complexity will decrease the test error.*

Over-parameterized regime. *If $\text{EMC}_{\mathcal{D},\epsilon}(\mathcal{T})$ is sufficiently larger than n , any perturbation of $\mathcal{T}$ that increases its effective complexity will decrease the test error.*

Critically parameterized regime. *If $\text{EMC}_{\mathcal{D},\epsilon}(\mathcal{T}) \approx n$, then a perturbation of $\mathcal{T}$ that increases its effective complexity might decrease or increase the test error.*

It is to notice that this definition is also only observational. By that, it only outlines the specific region of interest where the paradigm shift is identified through experimental results. It has no effective predictive power, or rather descriptive power more than setting up hypothesis with respect to the effective model complexity, and some arbitrary perturbation. [Nakkiran et al. \[2019\]](#) also noticed this difficulty in

providing a theoretical definition and theorem regarding such hypothesis, as said in their manuscript. The existence of the critical regime is also not defined.

The behaviour itself has been particularly investigated, in certain settings. For example, before [Belkin et al. \[2019a\]](#), [Advani and Saxe \[2017b\]](#) investigated the generalization error in neural networks of high-dimensional measure. Of such, various behaviours that bear similarity to double descent can be observed. [Belkin et al. \[2018\]](#) expanded on his previous research on the particular subset of kernel learning models, providing a theoretical analysis of specifically the Laplacian kernel on standard neural network. [Mei and Montanari \[2020\]](#) investigated random feature regression in such regard, and also found similar results.

The fact that double descent is not well-defined itself is a problem on its own. Up to the author's knowledge and of myself, there has been no effective definition or given description that outlines specifically how double descent can be structured. Instability in reproducing experiments and the effective range of the phenomena remains a question on its own.

3.4 What is double descent, then?

The question against double descent is still rather elusive. Ever since the discovery in [Belkin et al. \[2019a\]](#), despite the general chaotic view of double descent that seems to hold up pretty well, many researches have been conducted throughout, trying to observe or explain it, as we briefly mentioned [Nakkiran et al. \[2019\]](#) as one of such main paper. Before attempting to do anything, and to further contribute with hypotheses forward, we will have to resolve to review previous coursework on such matter. Most of those works include some on theoretical analysis, some on dissection of optimization method then speculation of possible hindsight, some on experiments and observable effects, and so on. Furthermore, sections specifically designed to host experiments and observations will also offer a clearer insight into how this can be done. The task is then for us to review those papers, to grip a sense on what is actually operating inside the domain of analysis. Theoretical insights are often of less emphasis, as double descent occurs at the mainly underdeveloped area of deep and machine learning, and as also for theoretical attacks to simplify their system for convergences bound, leaving mechanisms behind in determining such structures.

The main goal of attacking from theoretical side, of double descent stemmed from the observation in which there exists irregularity in the model expression behaviours. As such, the first thing to do is to contrive the consensus on 'what' constitutes double descent, at least per observable results, by delivering a conjecture. We deliver the phenomenological conjecture as below.

Conjecture 3.4.1 (Probing double descent). *On the landscape of double descent, model complexity versus generalization error, there exists a fundamental point $\mathcal{M}_0$ of model complexity axis, where $dR(h)/d\mathcal{M} < 0$ for $\mathcal{M} > \mathcal{M}_0$, which is the interpolation point.*

While this is not perfect, it provides us with a general idea on where to look for double descent, as the *resultant action* of the model complexity observation on learning error. Usually, double descent introduces the critical point where trends are reduced, unlike asymptotic growth in bias-variance, certain factor makes such pattern to collapse and invert the divergence to the convergence. The assumption here, with respect to previous bias-variance disposition, is also that such error correlates strictly to the general risk. Typically, question remains as to see:

1. What is exactly double descent? As far as we can say, there exists the interpolation threshold in which the dynamic changes, but don't know why. Such exists and occur within the complexity of deep neural networks, and so is also analogous form in classical models. It is suspected, reasonably

so, that it is related to the terms of bias–variance, however, it is not fully developed and no further inquiry is made beside the above conjecture as empirical observations.

2. Where does it lead? What is the setup? In truth, the setup problem in machine learning has been quite troubling. As mentioned, statistical theory lives in the idealized system, and thus many of its bounds hold in such domain. However, there exists the fundamental mismatch between how models are in practice, or rather in a semi-theoretical regime, compared to the rigidity of classical systems. The setup itself is then impractical, for example of the assumption in which the distribution drew out of x is fully Gaussian, or isotropically Gaussian, which is not helpful in practice.
3. What counts as the measure, and what we even measure? This is more of the concern toward what counts as model complexity, and the structural constraints that comes with it. Usually this is directed toward neural network, but even then, parameter-based complexity is often a poor choice of representation, thus begs the question what is even it, and why suddenly it appears there, within the dynamic of such?
4. Why would such an optimization space induce double descent? Surpassing what we have, the bounds that we are having indicates the strong trajectory in which downward trend is the only way, as for the generalization capacity is increased by increased transparency, i.e. $|\mathcal{S}|$ of the observation set, and by increased capacity $|\mathcal{H}|$ on certain induced or ‘natural system’ (convex system topology). However, this seems to not be the case.

With much questions going on, we are ready to follow up the literatures in provision of said questions. Most of them are in consideration of all the questions above, but under different strategies and either phenomenological, or deliberately simplified theoretical studies. Let us also recite what happens within the setting of many double descent results.

3.4.1 Underfitting and overfitting

3.4.2 Discovery of double descent

While double descent is discovered by [Belkin et al. \[2019a\]](#) and further on modern deep learning landscape of [Nakkiran et al. \[2019\]](#), it remains that the conclusion is not reached if there exists some fundamental models that would not have double descent, the consistency of double descent through different models entirely and so on. This is mainly because of a particular ‘shock effect’, since before double descent, bias–variance holds the candle as the main principle of choosing and optimizing models to its optimal saddle point. Suddenly by then we observe double descent, which introduces the interpolating threshold $\mathcal{M}_0$, that breaks this principle. There is then the open question on how far-reaching is double descent in consideration of different architectures, and this is reflected in several papers above. As of now, the main structures that have been verified includes

- **Classical models** [Yang and Suzuki \[2024\]](#), [Nakkiran \[2019b\]](#), [Schaeffer et al. \[2023\]](#), [Lafon and Thomas \[2024\]](#) (linear regression and classical polynomial), [Li and Sonthalia \[2024\]](#) (linear regression upon least square), [Liu et al. \[2021\]](#), [Zhang \[2024\]](#), [Allerbo \[2025\]](#) (kernel methods/regressions), [Brellmann et al. \[2024\]](#) (Classical MSBE-LSTD¹ optimization), [Harzli et al. \[2023\]](#) (Gaussian process). Some papers also have experiments on Random Fourier Feature Networks, and so on, but not substantial (including [Belkin et al. \[2019a\]](#) paper).

¹Mean-Squared Bellman Error and Least-Squared Temporal Difference

- **Deep neural networks:** Luo and Dong [2023], Yang and Suzuki [2024], Pezeshki et al. [2021] (CNN, VGG, ResNet, and DenseNet), Nakkiran et al. [2019] (Transformer), Shi et al. [2024] (GNN).

Most of the experiments are concerned of seeing and analysing double descent, basic analysis and experimental observations, test case scenarios for different models and setting and so on. All of the above models are then considered to be, concretely, reported of observing double descent, through various testing scenario on them. Review papers and literatures can be found as Schaeffer et al. [2023], Hu [2021], and Shi et al. [2024] on specifically, graph neural network. However, GNN would require its separated discussion, so we will only briefly mention it as such.

3.4.3 Theoretical attacks

Because of its intrinsic nature, where double descent breaks the long-standing principle of learning selection, theoretical resolutions has been employed to tackle such problem. Indeed, for example, Cherkassky and Lee [2024] and Lee and Cherkassky [2022] aim toward such direction by understanding the behaviours of double descent using the classical VC-theory approach. However, such papers only introduce, or relate double descent to VC-theory via phenomenological approach – for example, reducing the VC-dimension while fixing or reducing training error, or by examining a potential proxy between the VC-bounds and potential factors encapsulated in the quantities involved. They however, do raise voices of certain inherent problems lies within the theory of learning and in general, learning models.

Beyond such, certain discovery and inquiries help in having a clearer picture on the behaviour of double descent through various structures. For example, Hastie et al. [2019], under pretext of studying high-dimensional ridgeless (ℓ_2 norm) interpolation, analyse the case of *isotropic features*, *correlated features*, misspecified model (feature observation inhibited), on ridgeless interpolation of linear model, and soon as nonlinear model and latent space model. While strong, their results are restricted to a specific class of assumptions and structure, and is doubtful to be extended further. Allerbo [2025] investigated and find a particular double descent trace when changing kernel during training, specified in the term of the minimum allowed bandwidth. The kernel bandwidth $\sigma \geq 0$ of a specific kernel filter determines the information smoothing, hence implies the complexity abides by the kernel structure, as well as determine the smoothness of information indicted by the term.

The sentiments that are shared in between researchers on the topic of double descent vary, but have some common roots, most of them coming from the difficulty to fully assess double descent and its factor components. While double descent concerns the evaluation of *generalization (test) error*, and the *model complexity* notions, the error in question is often implicitly defined or loosely considered in one way or another, and model complexity has no consensus on its definition. In fact, this points out a very potential, if not already mismatch in the theoretical development. While optimizations and phenomenological aberration of models, as well as building and constructing models at will are well-developed, the definitions on which they are defined on is not fully understood, if not undefined. As of now, literatures on model complexities like et al., Hu et al. [2021b], Janik and Witaszczyk [2021], Truong [2025], Luo et al. [2024] introduce their variations or a copula of different definitions altogether, where some are concerned with its effective complexity, some are concerned with its expressive or structural complexity with different ways to define the model underneath, from parameters to units, to random standard variables (Neal et al. [2018]). There is also the rift in consideration between classical theory, of which mainly is for classical models, and modern deep neural network structures of which do not conform to such analysis.

3.4.4 High-dimensional ridgeless interpolation - [Hastie et al. \[2019\]](#)

Aligning with our previous experiments, [Hastie et al. \[2019\]](#) conducted investigation into the class of ridgeless approximation criterion on linear-nonlinear classes disparity. The investigation itself is thorough in parameter-term. Specifically, they investigated the case of which we have **isotropic features**, where asymptotic risk depends on the weight w only through its norm $\|w\|_2$; **anisotropic features**, where the model risk depends on the geometry of the pair (Σ, w) for $\Sigma = \text{Cov}(x_i)$ on the dataset $(y_i, x_i)_{i \leq n} \in \mathcal{S}$; and the case of nonlinear model, where the model itself is tested with added structure - x_i is obtained by passing $z_i \sim \mathcal{N}(0, I_d)$ through a one-layer neural network with random first layer weights.

Overall, the research itself is rather insightful, however, one problem with this research is that, it is fully theoretical with only deterministic prediction and numerical computation on heavily specified, synthetically generated data. Such is to say that the potential for such assumptions to hold in large system is uncertain, and whether if it will apply or transfer to higher structure, for example, similar construction but of neural network formalism. More specifically, the setting of this ridgeless case is constrained. Assuming the dataset as $(\mathbf{X}, \mathbf{Y})$ in form of $(\mathbb{R}^{n \times p}, \mathbb{R}^n)$ for the design/input space matrix $\mathbf{X}$ and response compression vector $\mathbf{Y}$, and parameter vector $\beta \in \mathbb{R}^p$. The fundamental model is of the form

$$y_i = x_i^\top \beta^* + \epsilon_i, \quad y'_i = x_i^\top \beta \quad (3.36)$$

From such, a **ridgeless regression estimator** is defined by the following optimization scheme, where

$$\hat{\beta} = \arg \min_{\beta \in \mathbb{R}^p : \mathbf{X}\beta = \mathbf{Y}} \left\{ \|\beta\|_2^2 : \exists \{\lambda_k\}_{k=1}^\infty, \lambda_k \downarrow 0 \text{ such that } \beta = \arg \min_{\tilde{\beta} \in \mathbb{R}^p} \left(\|\mathbf{X}\tilde{\beta} - \mathbf{Y}\|_2^2 + \lambda_k \|\tilde{\beta}\|_2^2 \right) \right\}. \quad (3.37)$$

Of the penalized complexity on L_2 norm. Note that we should not confuse the ordering of the interpolant to treat λ ridge term as equal to fit; that is, the requirement of the minimization problem is to say, among all interpolants, choose the one that would be preferred by an infinitesimal amount of ridge regularization. Thus, removing λ -ridge bears no side effect to the minimization structure, but only removes a particular selection criterion; within the ridge-based analytical framework, this criterion serves to single out a canonical representative among interpolants. In other word, we say that ridgeless impose the minimal Occam's Razor case of ridge-specification singular model canonical optimizer, and only then hidden structure by ridge is realized. By construction equivalently, this suggests that a ridge regression model requires such condition on any differential-based optimization algorithm. Differentiate with respect to β of the trainable parameters (i.e. actual moving average),

$$\nabla_\beta (\|\mathbf{X}\beta - \mathbf{Y}\|_2^2 + \lambda \|\beta\|_2^2) = 2\mathbf{X}^\top (\mathbf{X}\beta - \mathbf{Y}) + 2\lambda\beta \quad (3.38)$$

On this differential landscape, the stability points that would be present means that any points admissible on such will result in

$$\mathbf{X}^\top \mathbf{X}\beta + \lambda\beta = \mathbf{X}^\top \mathbf{Y} \quad (3.39)$$

$$(\mathbf{X}^\top \mathbf{X} + \lambda)\beta = \mathbf{X}^\top \mathbf{Y} \quad (3.40)$$

and thus leading to the invertability observation for $\lambda > 0$ of the matrix encoding, thus

$$\hat{\beta}_\lambda = (\mathbf{X}^\top \mathbf{X} + \lambda)^{-1} \mathbf{X}^\top \mathbf{Y} \quad (3.41)$$

is an admissible form that should be satisfied at such landscape on itself. Ridgeless appears when you enforce the following limit of the solution, where we get to the minimum norm least squares estimator

of

$$\hat{\beta} = \lim_{\lambda \rightarrow 0^+} (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{Y}, \quad \hat{\beta} = \lim_{\lambda \rightarrow 0^+} \hat{\beta}_\lambda$$

on the matrix design space. Ridgeless here means that *regularization is taken zero from the positive side*, or rather that we take the regularization effect to be minimal yet present in the optimization scheme. A problem can be identified. Without ridge or small negligible ridge, this assumes the most hardwired-free form of the optimization problem, where optimization can be considered without structural penalization, but **only in this sense**. The reason why direction 0^+ matters can also be easily explained as there exists no meaningful way of encoding, in a strict simple sense, negative structural complexity measure. However, we note that the optimization here assumes that the loss-landscape is admissible of the L_2 form on Euclidean metric. In reality, what we have constructed of the learning template suggests a different kind of metric, a pseudometric. Such is to say, we are assuming these two metrics to be equivalent in the same form or shape as to be admissible of both. This is a very strong assumption, and overall, the system itself contains mismatch and extremely high structural assumption. If in ideal environment, such algorithm and formation process might turn out to be sufficient and correct. However, if this is not always the case, and the pseudometric equivalent condition implicitly assumed here can break, then pathological behaviours occur naturally of the unseen structure outside the framework consideration. Those pathologies are not emergent; moreover, they are precluded by definition. If we are to observe more, then the problem of multiple critical points might also be only an artefact of the Euclidean mismatch, outside of linear-enforced, convex setting. Thus, we can further observe another layer of the strong assumption, this time on the properties of the space. Extending this ideal setting, especially to more complex models, where the system are free and chaotic enough even in constrained operation (layer-frame and equivalent neuron units) is ill-posed, and can be reasonably considered a potential pressure point toward why there exists the mismatch between realistic models and theoretical presumptions.

Furthermore, from the model itself, and further on of the linear regression setting, is the fact that we have lossless information availability (true fit data) with sufficiently small $\epsilon_i > 0$. Assuming uniform ϵ further, we can see that under such view, any model can be reconstructed sufficiently, as for there exists only $\mathcal{O}(1)$ growth on parameter space. With small mismatch between β^* and β of the learner, such error can be high during optimization interval due to compression of parameters into a smaller set of parameters, while on high-dimensional, over-parameter case, dilution of parameters contains admissible dynamic of diluting functional form of a smaller set of parameters to a larger one, thus also explains in spirit the formation of double descent. However, this is not rather reflected in the framework. If ϵ_i are high, we still only consider a portion of it into variance control of the model, and nowhere there exists generalization. In fact, linearized model is very poor in gauging generalization behaviours, as the model itself, with the dataset synthetically generated in a supposed adversarial case (two linear models against each other), assumes quietly the structural relational space of f and f^* to be monotonic, and often with a narrow deviation corridor. Meaningful noise also is restricted, and thus can be considered into variance control. Adding this with the full-transparent optimization problem where there exists no statement about meaningful divergent or unseen data, generalization is hardly discovered here, but only ad-hoc parameter stabilization, non-permutating in order measure. That is, model can reconstruct parameters, but they are almost never the same, both in order and magnitude if they are of same expression class.

The effectiveness of ridge observations and so on are basically phenomenal, and contains no sufficient theoretical insight on partition dynamics. In spirit, the theoretical treatment is not structural-free, nor the definition of generalization or unseen structure is sound. This is also troublesome as to rely on this to take on the problem of double descent. If we are to notice, the problem of double descent bears no resemblance to this setting in any conceivable way, at least when it was discovered. This pulls back to the

dilemma of the definition of double descent, where on one research it is about descent on sample space, and on one hand it is on fully-admissible, population representative, full-access data as the paper we are discussing.

3.4.5 Sample-wise double descent - [Nakkiran \[2019b\]](#)

Regarding double descent, [Nakkiran \[2019b\]](#) explored the problem, in which the increment of sample actually hurts linear regression process. As always, the setting considers isotropic $x \sim \mathcal{N}(0, I_d)$. Here, the observation and claim is that the number of samples can indeed, hurt the end result for test MSE on the model. This results into two two claims, mainly of the decomposing bias-variance analysis. First is the overparameterized risk claim, in which is defined of:

Claim 3.4.0.1 (Overparameterized Risk). *Let $\gamma := \frac{n}{d} < 1$ be the underparameterization ratio. The bias and variance are:*

$$B_n = (1 - \gamma)^2 \|\beta\|^2 \quad (3.42)$$

$$V_n \approx \gamma(1 - \gamma) \|\beta\|^2 + \sigma^2 \frac{\gamma}{1 - \gamma} \quad (3.43)$$

And thus the expected excess risk for $\gamma < 1$ is:

$$\mathbb{E}[\bar{R}(\hat{\beta})] \approx (1 - \gamma) \|\beta\|^2 + \sigma^2 \frac{\gamma}{1 - \gamma} \quad (3.44)$$

$$= (1 - \frac{n}{d}) \|\beta\|^2 + \sigma^2 \frac{n}{d - n} \quad (3.45)$$

and the second one is about the underparamterized case, which is as the second claim.

Claim 3.4.0.2 (Underparameterized Risk, [Hastie et al. \[2019\]](#)). *Let $\gamma := \frac{n}{d} > 1$ be the underparameterization ratio. The bias and variance are:*

$$B_n = 0 \quad , \quad V_n \approx \frac{\sigma^2}{\gamma - 1}$$

Those claims are troublesome, in the sense that they indict heavily on the assumptions of which, for example, suddenly the bias collapse to zero without any indication or further justification. The variance is specifically designed for a replica-style curve fitting instead, of which the approximating relation suggest the not-so-useful notion, in which such deviation from the claim can be wide and far yet still right in terms of descriptive behaviours (converges to 0). Vanishing terms transition from the first claim to the second claim inherently accept bias of the system itself, of which is then trouble some to proceed against. And the list of assumptions on fine structures are too many in spirit, especially the linear algebra mapping in which derivations of computation are used in the paper itself. To truly justify this, we need something beyond what is described here.

3.4.6 Multiscale feature learning dynamics - [Pezeshki et al. \[2021\]](#)

In the setting of [Pezeshki et al. \[2021\]](#), they are concerned of the multiscale model, but still rather simply configured, as stated. The object of study is a supervised linear regression problem in which the training labels y , are generated by a noisy linear model,

$$y := y^* + \epsilon, \quad y^* := \mathbf{z}^\top \mathbf{W}, \quad z_i \sim \mathcal{N}(0, 1/\sqrt{d}), \quad (3.46)$$

where $z \in \mathbb{R}^d$ is the teacher's input and $y^*, y \in \mathbb{R}$ are the teacher's noiseless and noisy outputs, respectively. Thus, the setting itself is rather familiar, within which the structure of the oracle is assumed and embedded of the observation space clarified noise ϵ . In a sense, this is a case of **external fixture** of missing information, such that to posit no information package is not present, but most are noise intrinsic of the setting or retrieval purpose. $\mathbf{W} \in \mathbb{R}^d$ represents the weights, assumed as **fixed** in the setting of the teacher and $\epsilon \in \mathbb{R}$ is the label noise. Here, both W_i and ϵ are assumed to be drawn i.i.d. from Gaussian distributions with zero means and variances of 1 and σ^2 , respectively. Within this assumption, though, it means practically nothing as the noise is treated simply as intrinsic and not structural. Admission of such simply gives i.i.d. and Gaussian form a kind of statistical expectation, nothing more, of which can be reasonably extrapolated from the hypothesis class statistics itself. Despite this structural shallowness, the result still has to be analysed to be further criticized, especially considering it is of about multiscale feature learning.

3.4.7 Reinforcement learning and double descent - Brellmann et al. [2024]

temporal
difference

In this paper, the author themselves, contributed to the analysis of double descent inside reinforcement learning, specifically LSTD - Least-squares (LS) for the class of **temporal difference** (TD) algorithms. Interestingly, one of the main contribution of the paper, is the focus on the *double asymptotic regime* for a new way of defining **model complexity**, in specificity. The analysis is strictly heuristic and simple, as seen of linear model and nonlinear random features for approximation of overparameterized single-hidden-layer network in the lazy training regime, for that we do not actually quite sure how it is overparameterized in what way. However, one of the fairly suspicious component, lies in the asymptotic regime itself. Fortunately, the authors working on the paper provided a rather meticulous set of assumptions, in which we will make use for the clear picture of what is in the framework itself, though under full dissection, we have a rather large system for criticism underlying such. First, let us set up the scenario. LSTD is part of a class of solution, for Markov reward processes, essentially a rather complex under mathematical representation, but simple technically speaking transition reward learning optimization scheme, basing itself on the global observed state space restriction. Of such, LDA, or **Linear Function Approximation**, is the family of techniques LSTD is based on, because of the practical infeasibility to solve

$$\mathbf{V} = \bar{\mathbf{r}} + \gamma \mathbf{P} \mathbf{V}, \quad \text{MSBE}(\theta) = \|\bar{\mathbf{r}} + \gamma \mathbf{P} \Sigma_{\mathcal{S}}^{\top} \theta - \Sigma_{\mathcal{S}}^{\top} \theta\|_{D_{\pi}}^2, \quad (3.47)$$

for it being the unique fixed-point Bellman's equation, where $\mathbf{V} \in \mathbb{R}^{|\mathcal{S}|}$ is a vector whose i -th element is the value function of the i -th state $\mathcal{S}_i$; and $\bar{\mathbf{r}} \in \mathbb{R}^{|\mathcal{S}|}$ is the vector containing the expected rewards, for which $\bar{r}_i = \mathbb{E}_{\mathcal{P}}[r | \mathcal{S}_i]$ for all $i \in [|\mathcal{S}|]$. The setting itself is deliberate, when fixing the LSTD estimator in the paper, the analysis reduces to reinterpretation on **lazy training and random features** analysis, thus the TD operator becomes

$$\mathbf{A} = \Sigma_{\mathcal{S}} \mathbf{D}_{\pi} (\mathbf{I}_{|\mathcal{S}|} - \gamma \mathbf{P}) \Sigma_{\mathcal{S}}^{\top} \quad (3.48)$$

which equals to random matrix with structured covariance case. The fixture of the LSTD in the paper also reduces the observations made in the paper into $N/m \approx 1$ on the matrix singularity transition behaviours, not much different than linear regression double-descent interpretation we analysed earlier.

Assumption 3.4.1 (Double Asymptotic Regime, Brellmann et al. [2024]). *As $N, m, d \rightarrow \infty$, we have:*

1. $0 < \liminf \left\{ \frac{N}{m}, \frac{d}{m}, \frac{m}{|\mathcal{S}|} \right\} < \limsup \left\{ \frac{N}{m}, \frac{d}{m}, \frac{m}{|\mathcal{S}|} \right\} < \infty$.
2. *There exists $K_{\mathcal{S}}, K_r > 0$ such that $\limsup_{|\mathcal{S}|} \|\mathbf{S}\| < K_{\mathcal{S}}$ and $r(\cdot, \cdot)$ is bounded by K_r .*

Assumption 3.4.2 (Bounded Eigenspectrum, Brellmann et al. [2024]). *There exist $0 < \xi_{\min} < \xi_{\max}$ such that for every m , all the eigenvalues of $H(\hat{\mathbf{A}}_m)$ are in $[\xi_{\min}, \xi_{\max}]$.*




Figure 3.8: As the model complexity N/m (for N parameters, m distinct visited states) increases, the MSBE first shows a U-shaped curve, peaking around the interpolation threshold ($N = m$). Double descent refers to the phenomenon for $N/m > 1$ where the MSBE drops once again. Continuous lines (red) indicate the theoretical values, the crosses (blue) are numerical results averaged over 30 instances with their standard deviations after the learning with regularized LSTD on Taxi-v3 for $\gamma = 0.95$, $\lambda = 10^{-9}$, $n = 5000$, $m = 310$. Taken from Brellmann et al. [2024]

Assumption 3.4.3 (Run-off growth, Brellmann et al. [2024]). As $n, m \rightarrow \infty$, we have

$$\|\mathbf{A}_n - \mathbf{D}_\pi [\mathbf{I}_{|S|} - \gamma \mathbf{P}]\| = \mathcal{O}\left(\frac{1}{\sqrt{n}}\right) \quad (3.49)$$

Those assumptions, while looking rather innocent, are rather not trivial from what is assumed. For Assumption 3.4.1, the first sub-assumption is structurally the family of assumptions used in random matrix theory, and extensively utilized under context of modern high-dimensional statistics. For the baseline assumption, it decouples any form of *affendants* in between the quantities, or any form of behaviours of such large enough to get one term to dominate. Listed quantities are normalized on N/m , d/m , $m/|S|$, stays $O(1)$ on asymptotic behaviours. This is well-understood and widely-accepted, and we have no issues on such, but only the domain of extension and **interpretative restriction** – toward where and how the constraints of $N, m, d, |S|$ are true enough and will they be extensible. While it keeps the dynamic possible and thus reducible into a small- δ near static version of the bound in fixed regime, those parameters are rather sensitive, and would have structural predicates more than what was assumed in certain cases. The action of having $d \gg |S|$ or $N \gg m$ incurs certain regimes that might elucidate the problem than not. However, those are structurally valid restriction. The second sub-assumption of such, however, is in active effect of covering for a structural issue. The assumption on $\limsup_{|S|} \|\mathbf{S}\| < K_S$ on bounding the spectral norm of the feature matrix means assuming a normalized, non-vacuous only in specific region of the normalization convention. This assumption is used to bound the term $\|\Sigma_S\|/\sqrt{m}$ to be concentrated, but in turns, it directly imports it both in those structural conditions, and also all other results, including the bound on

$$\Pr\left(\frac{1}{m}\|\Sigma_X\| \geq t\right) \leq CNe^{-\frac{cm^2t^2}{2n\|\mathbf{X}\|^2}}, \quad \forall t \geq 4t_0, t_0 = |\sigma(0)| + \|\mathbf{X}\|\sqrt{\frac{d}{m}} \quad (3.50)$$

which is from another paper, which compounds the hidden chain of assumptions in which double descent is very sensitive in. Moving back, the issues contain also the condition that the state matrix has a

specific condition in which it must be **controlled coherent**, meaning the columns are configured in such a way that singular value does not accumulate during operations. The experiments used in the paper, however, also contradicts this by setting $d = 50$ fixed for condition one requires $d/m \rightarrow \text{const}$, which means the experiment itself at this stage, while tractible, does not amount to much but crude strict-case manufactured double descent.

Assumption 2 is rather deliberate on conditioning $H(\hat{\mathbf{A}}_m)$ on specificity of the empirical transition model matrix $\hat{\mathbf{A}}_m$. While not much can be said about this design matrix, nonsymmetric, and is so far not in operator theory application or thereof. Symmetrization happens when the author needs $Q_m(\lambda)$, which means in such $\frac{1}{m}(\hat{\mathbf{U}}_n - \gamma \hat{\mathbf{V}}_n)^T \Sigma_{\mathcal{S}}^T \Sigma_{\mathcal{S}} \hat{\mathbf{U}}_n + \lambda \mathbf{I}_n$ is invertible. The symmetrization as such is derived into the symmetric component $H(\hat{\mathbf{A}}_m)$, which is positive definite to get that the eigenvalues of this matrix has nonnegative real parts. Note that positive definiteness of $H(\hat{\mathbf{A}}_m)$ is not similar to on strictly $\hat{\mathbf{A}}_m$, but only operable via the inequality

$$\text{Re}(\nu_j(\hat{\mathbf{A}}_m)) \geq \nu_{\min}(H(\hat{\mathbf{A}}_m)) \geq \xi_{\min} \quad (3.51)$$

for uniform positive definiteness. For a nonsymmetric matrix, it can be highly non-normal, such is why when extracting eigenvalues and parts of the symmetrization, the carving itself potentially removes the actual components and only take on residue of the transient amplification of the mechanics. For example, as the skew component $S(\hat{\mathbf{A}}_m) = \frac{1}{2}(\hat{\mathbf{A}}_m - \hat{\mathbf{A}}_m^T)$, then $H(\hat{\mathbf{A}}_m) \succ 0$ does not guarantee same behaviours for $S(\hat{\mathbf{A}}_m)$. The inequality itself bounds only the real part on this projection onto spectral complex, which means the bound is rather extremely weak on this complex space projection of operator Hermitian. Furthermore, it removes also the domain of off-policy learning, function approximation fall-off, by constraining $I - \gamma P$ instability, which perhaps is a simplification choice, but left no way to extend results cleanly toward a larger regime than to completely rederive everything. The assumption is then not self-contained as evidentially provided of, but unstated of structural constraints.

The third assumption is rather also radical, as it is very hard to justify the jump into reduction of convergence rate, to $\mathcal{O}(1/\sqrt{m})$ from $\mathcal{O}(1/\sqrt{n})$. This implicitly requires n to grow at least as fast as m in such way. The quantity on $\hat{\mathbf{A}}_n$ itself involves the full state space decomposition, not just the operational visited states, and of which ask the norm bound on the matrix that includes structurally zero rows, to approximate nonzero rows, which only is plausible when $m/|\mathcal{S}| \rightarrow 1$, where coverage is near-complete. This is because the decomposition to terms $U_n, V_n \in \mathbb{R}^{|\mathcal{S}| \times n}$ requires controlling $|\mathcal{S}|$ for the operator norm bound in the stated condition, of which for unvisited states, the corresponding rows of U_n are zero, so their empirical estimate of transition behaviour is exactly zero, while the true proportion on $\mathbf{D}_{\pi} [\mathbf{I}_{|\mathcal{S}|} - \gamma \mathbf{P}]$ being nonzero weight assignment.

There are other assumptions hidden in the setup. Random features here are assigned with Gaussian weights, and the condition also gives system on K_{σ} -Lipschitz condition and activation. The $\sqrt{n}$ normalization term of which appears everywhere, however not stated of its broader structural roles aside from convenience sake. On-policy sampling is deliverately restricted, however not fully realized of its effect, where such condition TD learning actually informs the third assumption itself. In a sense, the paper itself is not of any particular aspect TD learning with double descent identification of structures. Rather, it is a theory of on-policy LSTD with Gaussian random features in a proportional-growth regime with near-complete coverage, which is as similar of methods to the extremely reduced case of linear regression on standard double descent singular matrix collapse. The intrusion of mapping onto the representative structure of the operator theory domain, and more specifically on complex representation is also non-trivial. This jump from a real system to the encoding on Ostrowski lemma for $\mathbf{A}$ Hermitian, and the proof of several

lemmas in the paper, of which invoked $\nu_{\min}(H(\hat{\mathbf{A}}^{-1}))$ are flowed as if spectral theory for Hermitian operator applies on such system, which on itself constraints a rather very deliverate interpretative spectral domain, but also $\hat{\mathbf{A}}$ was previously a real asymmetric matrix, and no explanation was provided on such part for the logical transformation of such. In a sense, the paper invokes Hermitian/PSE spectral theory whenever convenient by symmetrization and truncation toward such regime, and there exists the absence of such clarification on why and what was lost under the transition itself, as also seen by the skew-aspect criticized earlier. In Appendix C, the reformulation requires Mercer's theorem. For such theorem to apply, you require several assumptions: (i), that the compactness is assumed to be consistent before such transition of problem space, $\mathcal{C}$ compact such that $\mathcal{S} \subset \mathcal{C}$ is rather very non-trivial. The import of L^2 inner products, integral operators, and so on. The Hilbert space $L^2(\mathcal{C}, \pi)$ is unwarranted and exists only because of the forcing constraints to operate Mercer's tools, and is not forced by the problem. (ii), is that the asymptotic feature space of which has M potentially infinite eigenfunctions $\{\omega_i\}$ in the paper, is accompanied by N random features on algorithm. The approximation error from this truncation is not justified or quantified in the derivations. (iii), the operator norm and spectral theory section are Hermitian by convenience, of which none is forced. In a sense, such mapping correlates to kernel machine learning actions where one maps this domain of linearized system into self-adjoint kernel space, which on itself is the property admitted of the kernel, not other components. And (iv), the stationary distribution π is also the measure defining the Hilbert space, but also the definitive factor for MRP and MSBE norm. This transition and whether the coupling makes much sense is not specifically treated during role compounding, which risks that the analysis is conflating the two simultaneously.

If we are to push further, we can deliberately say that the paper itself both oversold, and broke its own thesis. For example, the paper's stated purpose is to explain double descent in RL with TD learning scheme, specifically derived toward the complexity N/m term and where it produces a phase transition. However, aside from the term N/m itself being again, another heuristic created in the spatial definition axis without backward compatibility toward other notions of model complexity that operates in conjunction within the model space that the metric is, without justification and without knowing why (which, under specific instance we can say that the phase transition is engineered by the predicted asymptotic ordering, not forced under experimental details), the MSBE being analysed is not consistent of what was used. The true MSBE is defined by

$$\widehat{\text{MSBE}}(\boldsymbol{\theta}) = \frac{1}{n} \|\mathbf{r} + \gamma \Sigma_{\mathbf{X}_n}^\top \boldsymbol{\theta} - \Sigma_{\mathbf{X}_n}^\top \boldsymbol{\theta}\|^2, \quad (3.52)$$

while parts of the analysis uses the second asymptotic MSBE as

$$\text{MSBE}(\boldsymbol{\theta}) = \|\bar{\mathbf{r}} + \gamma \mathbf{P} \Omega_S^\top \mathbf{v}_\theta - \Omega_S^\top \mathbf{v}_\theta\|_{D_\pi}^2 + f_\Delta \left(D_\pi^{1/2} (\mathbf{I} - \gamma \mathbf{P}) \Omega_S^\top \right) \quad (3.53)$$

where Ω_S is the infinite-dimensional eigenfunction feature matrix, $\mathbf{v}_\theta$ is the weight vector in asymptotic feature space with rescaled regularization parameter stated in the paper being $\lambda(m(1+\delta))/N$. These are inherently not the same object, and the double asymptotic equivalence between the two are very fragile, of which is only true for $N, m, d \rightarrow \infty$. Under normal finite circumstance, this fails immediately on consequence. The term f_Δ in such derivation is attributed to double descent, but it is a kernel property on the operationalized space, not the finite random feature model. Whether they actually coincide is a problem that is very much open for exploration, but within such length of the paper one would expect this to be seen or at least partially provided, both for this gap and the assumption thereof. The regularization change from the original algorithm of $\lambda = \lambda_{m,n}/mn$ was silently changed also changed the status of the derived results into informal extrapolation from a formally different structure system, and is not

conditioned on anything, which correlates to the same flawed PAC-descent method that we have seen. And lastly, the double descent explanation itself is circular in its own nature. The paper states for once that the correction terms δ and Δ are responsible for double descent, however, δ is derived from deterministic equivalent of the resolvent $Q_m(\lambda)$, which are inherited of two separated papers on the structure precondition and phase transition behaviours. Δ itself being the second-order term involving δ also makes it so that the structural preconditioning presupposes double descent by directly engineer a space where such is susceptible for the expression used, yet also claimed to elucidate double descent using those, which is a circular problem in its own right. Whether if this is the correct or wrong reading of the paper and its problem, such is remained to be seen from revisions, but under the current status of the paper, such cannot be contained, as it requires tracking structural overcommitment and hidden assumption, proliferation of multiple matrices, Gram/kernel representation, resolvents, regularization parameters, two LSTD solutions, dual- π usage and interpretation, and more than thirty different structural predicates in general of the setting, makes the analysis non-trivial on such grounding for little results.

3.4.8 Variations on analysis

Many sub-variations and analysis alongside major substantiated discoveries listed above can help in analysis. Some of them are toward linking different existing dynamics to large neural network in non-linear fashion, or generalizing approximation of linear network toward nonlinear systems. There are many analyses, typically of interesting being weak feature double descent - [Belkin et al. \[2020\]](#), insights on particularly out-of-Distribution Detection - [Ammar et al. \[2025\]](#), Gaussian processes and neural network - [Harzli et al. \[2023\]](#), and the proliferation of Grokking, Emergence, and Double Descent - [Huang et al. \[2025\]](#). While interesting, we do not specifically cover them in analysis or the variations, as they exhibited the same structural deficiency or experimental validation patterns not so much contributed to more understanding on such matter beyond what has been said.

Dropout and double descent - [Yang and Suzuki \[2024\]](#)

Dropout is a relatively well-established regularization technique for training deep neural networks, usually by exclusion of different neurons during processes, as seen in [Srivastava et al. \[2014a\]](#) to control overfitting. Accordingly, [Yang and Suzuki \[2024\]](#) predicted or stated that regularization techniques can nullify double descent, and thus also dropout. The model considered is linear regression model, for p covariates $x \in \mathbb{R}^p$ and response or output $y \in \mathbb{R}$ are related by

$$y = x^\top \beta_0 + \epsilon, \quad \epsilon \sim \mathcal{N}(0, \sigma^2) \quad (3.54)$$

The situated dropout is as followed. Given training data $z^n = \{(x_i, y_i)\}_{i=1}^n$, for $X = [x_1, \dots, x_n]^\top$ and $y = [y_1, \dots, y_n]^\top$, the estimation of weight β by $\hat{\beta}(z^n)$ (model's weight), to minimizing the training error

$$\mathbb{E}_{r \sim \text{Ber}} \|y - (r \odot X)\beta\|_2^2 = \|y - \gamma X\beta\|_2^2 + (1 - \gamma)\gamma \|\Gamma\beta\|_2^2 \quad (3.55)$$

for $\odot$ denoting the element-wise product on L^2 norm space, each element of $R \in \mathbb{R}^{n \times p}$ takes one and zero with probabilities γ and $1 - \gamma$, for the distribution $\text{Ber}(\gamma)$ corresponding to such, and $\Gamma = \text{diag}(X^\top X)^{1/2}$. The final description gives the form of the effective dropout objective. Originally from the paper, they consider it exactly of a Tikhonov regularization method. Set $\beta' = \gamma\beta$, then

$$\begin{aligned} & \|y - X\beta'\|_2^2 + \frac{1 - \gamma}{\gamma} \|\Gamma\beta'\|_2^2, \\ \hat{\beta}_{n\gamma} &= \left(X^\top X + \frac{1 - \gamma}{\gamma} \Gamma^\top \Gamma \right)^{-1} X^\top y \end{aligned} \quad (3.56)$$

minimization. While experiments regarding such regularization works in such domain and was illustrated as to be reasonably effective, fallacies in claims and insight remains, especially in the nullification of double descent. The problem here is that, any given amount of dropout is, accordingly to the author, similar to Tikhonov or ridge regularization on structural complexity. As explained in the c' optimization of c -space in preliminary sections, such induces a different, much smaller system in which the additional structural attractor points toward the region of low regularizing complexity throughout the domain. Such superposition between the $\{\lambda r(h)\}(\mathbb{R})$ landscape and the classical loss $\hat{R}$ only morph the dynamic setting to a different region, and thereby bypass double descent in such case for absolute non-overfitting convergence to $\min(\hat{R}_S(h))$. This works well on well-behaved and loosely stable case, but in a lot of scenarios, such cannot be guaranteed, especially for pathological systems. Furthermore, Speaking in terms of the *inductive bias*, mismatch between model and effective regularization becomes a problem, as there exists no effective complexity measure for the regularization term $r(h)$. Some are ill-conditioned, for example, $r(h) = \sum_j w_{ij}^2$ for square terms ij -indexed weights only take into account the reasonable weight system, without the semantic itself, or meaningful weights. A forced dropout might potentially in certain less ideal cases, be more destructive, and introduces double descent back again, until the stochastic process shifts meaningful weight to a more resource constrained system. This means we have to induce, or at least formalize if stochastic process can recover meaningful semantic after dropout randomly. We can test this somewhat, by introducing a small toy test on the proposed dropout estimator:

Theorem 3.4.4 (Expected test error, [Yang and Suzuki \[2024\]](#)). *Let $\alpha = C < \frac{1}{(1+\sqrt{p/n})^2}$, the expected test error is*

$$f(\alpha) = \left\{1 - a\alpha + 2\alpha^2 \frac{p}{n}\right\} \|\beta_0\|^2 + \sigma^2 \alpha^2 (\alpha + 1) \frac{p}{n} + O\left(\frac{1}{n^2}\right) \quad (3.57)$$

with $\alpha = \gamma/(1 - \gamma)$.

While the theorem is proved, we then must know if such theorem correctly correlates to even small dynamic systems. However, as it stands, we should not read the paper as to say that dropout resolves double descent, but rather that *double descent* region of solution space is effectively ignored by reduced solution space on dropout condition, especially within ideal setting, and linear functional class, not even multilinear models. Alignment outside of this ideal setting introduces the problem of brittle prediction, and might as well nullify this result.

Chapter 4. Experimental investigations

For understanding and analysing double descent and bias-variance trade-off, and furthermore in later section on identifying double descent, we would like to use test models specifically for exhibiting bias-variance, as well as testing hypothesis and forming theoretical conjectures. This is because of the troublesome nature of the phenomena itself, which proves difficult to analyse without sufficient experimental observations. Furthermore, specific analysis on these types of model might prove useful for generalization toward more difficult and complex structure.

Our major focus is on three types of model. First two are classical models and system, linear models on the real space $\mathbb{R}^n \rightarrow \mathbb{R}$. Second, is the type of **support vector machine** (SVM) models to highlight the analysis toward data-based fundamental analysis. The third one is the class and collection of different *neural network structures* (MLP, or variations). Our goal is then to expansively analyse, formalize, making conjectures and hypothesis and test them, to find, force double descent out, or identify it on the learning landscape. We note that structures of those models are not well-generalized, hence analysis from one might not apply or apply only in small conjunction. This approach is rather quite helpful in its own merit, as its choice and preference will gauge our parameters than we gauge the phenomena of the parameter it induces.

4.1 Experimental planning

This section of the manuscript will explicitly contain instructions, setup, methodology, and so on, with regard to the experimental details of the experiments itself. In all, we would be covering the scheme of optimization algorithm responsible for dynamical modelling, and the other one for one-shot analytical optimization of the least-square form. Furthermore, justification and experimental list, with detailed analysis of each case of the system in consideration of testing would be included so much so to condition the test on the ground of generality as far as possible, and proliferate the experiment coverages. We also cover interpretation, and dataset generation, if such would be elucidated, and the concept of **stress-testing** under formal experimental package, if such is included in later analysis.

4.1.1 Setting: gradient-based optimization

In our usual experiment, we would most specifically, use a closely-regarded family of gradient-based algorithms, such as gradient descent (Cauchy [1847], Robbins and Monro [1951], Widrow and Hoff [1960a], Hestenes and Stiefel [1952], Fletcher and Reeves [1964]), mirror descent, natural descent, and so on (Nemirovski and Yudin [1983], Amari [1996, 1998], Yang and Amari [1998], Rattray et al. [1998], Gunasekar et al. [2021], Luong et al. [2016], Xiao and Zhang [2014], Sohl-Dickstein [2012], Raskutti and Mukherjee [2013]). While implementation and details of the specific algorithm can be different, they all shares some of the same concepts of which making them possible for optimization in the specified

iterative domain. The most general formulation possible is the following algorithm on convexity of the gradient descent family.

Algorithm 1: Gradient descent on general (convex) function

Data: $(x_i, y_i) \in \mathcal{S}$, η , $\nabla(\cdot)$ operator, hypothesis model $h \in \mathcal{H}$, arbitrary construction on complexity d .

Result: Expected optimization to the minimal error on $\ell(\cdot, \cdot)$.

begin

$\mathbf{w} \leftarrow$ randomly initialized values on operation;

$\mathbf{w}_1 \leftarrow$ first state of the state space $\mathbb{W} \subseteq \mathbb{F}^q$ for q specified by d complexity.;

for $t \leftarrow 1$ to T **do**

$\mathbf{y} \leftarrow \{h(\mathbf{x})\}^{|S|}$;

$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t \odot \eta \nabla_{\mathbf{w}} M_\ell(h[\mathbf{w}](x))$, arbitrary $\odot \equiv (-)$ under normal operator, loss evaluation configuration M_ℓ .

end

end

return Model $h[\mathbf{w}_T](x)$ such that $\arg \min_{\mathbf{w} \in \mathbb{W}, h \in \mathcal{H}} M_\ell(h[\mathbf{w}](x)) = h[\mathbf{w}_T](x)$.

Specifically, the scaffolding of the gradient descent algorithm is on the encoding of all numerical entities – we presuppose as the framework a form $h(x)$ fixed functional composition, such that the space $\mathcal{H}$ is a function space of all admissible descriptions on $[qh_1, \dots, qh_n]$ of the hypothesis element construction, such space $\mathcal{H}$ incurs $h(x)$ with its parameters form θ or $\mathbf{w}$ (depends on notation) of sort in compressed notation, hence the model lives in a flat parameter initialization, typically can be compressed to 1-D slice value representation. Often time, we have the parameter or any form of structural descriptions of a model without the second pair on composition and operations Γ_h to live on the state space $\mathbb{W}$, but on presentation status, we can reduce it to repeatable many-point domain. Adding the second dimension on the space to be the conjunctive (y, y_h) -domain, we have the now applied $[\mathbf{w}, \mathbf{y}]$ -space. For grounding, the last representation contains $\mathbf{x}$ for the numerical setting that are inputs to both. In all, we have a 3D space on $[\mathbf{x}, (\mathbf{y}, \mathbf{y}'), \mathbf{w}]$, of which can be used to experiment and observe different tracing of paths.

Nevertheless, one should be cautioned of the intrinsic assumption on which gradient descent of this type assumes. It assumes *convexity*, which can only be possible, if the landscape itself is supposedly, assumed to contain such global minima. Said global minima in a richer language can be described as the epistemic collapse of the learning structure, under the arbitrary optimization space. This is a *very serious assumption*, and should it fail, as seen in a lot of non-populated or general observation space, the entire procedure is invalid. It is also troublesome as to rely on the singular objective (numerically encoded) space in which it is trivial to realise that there are indeed, local minima. This scaling, with respect to both locality and global consideration, must be considered carefully, since *every single contribution* might distort the experiment in a very unexpected way. In a more sophisticated explanation, we would be more aware of the structural conformation that gradient descent make of the encoding space. If one wish, the formulation can be changed to

$$\arg \min_{\epsilon \geq 0; \mathbf{w} \in \mathbb{W}, h \in \mathcal{H}} M_\ell(h[\mathbf{w}](x)) = h[\mathbf{w}_T](x) \quad (4.1)$$

for the arbitrary condition of ϵ -fitness, which relaxes the argument. However, it introduces the determination value of ϵ , whether such is achievable or not which means the algorithm requires **termination**

mode, which lies in the hand of computability design, and how ϵ behaves and specific instance of forcing on such direction.

4.1.2 Experiments list

The following describes and outline specific requirements updated of experiments, aside from done experiments. Currently, there remains the *neural network* in plain form functional approximator experiment, *polynomial neural network (customized)* experiments, validation of *support vector machine (SVM)* on generalized task space experiments, different large-scale neural networks (VGG, ResNET, LLM) experiments set, and toy model for investigations. Specifically:

1. **Linear models on abstract $\mathbb{R}$ -conjunction space**¹: Different kind of components-based on linear behaviours, in different setting and so on, of different testing scenario (regression, classification). Linear model are often the canonical example of double descent, as they offer a rather simplified setting that most often would be tractable for analysis. Thereby, we will start with experiments and surveillance of different settings under such spaces, and actual implementation and testing rigorously thereof for linear regression.
2. **Polynomial-type models on abstract conjunction space**: Similar to the above linear case, this one however allow for polynomial-style components, i.e. degree k terms up to n th degree.
3. **Support Vector Machine (SVM) on abstract spaces and sample space**: The SVM model is rather interesting (within which it has regression and classification as main functionality), because it goes reverse from perspective – instead of creating models to fit data, it uses data to create fitting model.
4. **Neural network on-layer, abstract functional spaces**: This is the set of experiments for layer-based neural networks, on abstract functional spaces of analysis. The list includes vanilla (l, n) -mass neural network with mixed activation patterns, CNN, VGG, ResNet, spatial neural networks, RNN, LSTM, and if able, certain precursor of LLM systems.

All of those models base their purposes in *component-based interpretation*, such that the scope in which their components are allowed, change the dynamics and so on of the model. The theoretical problem of which those experiments are targeted on is to mark the singular hindsight – *simply how* model behaves under controlled structural changes and setting changes. Because of the lacking in such deference of approaches, it also depends heavily on how the adversarial set (dataset, concept) itself is targeted of, or coverage of such. Such is to say, we also need the *precursor theory* of component treatments. Which is explained above, but such is also particularly extensible to explain also grokking (Power et al. [2022b], Liu et al. [2023]), benign overfitting (Bartlett et al. [2020], Muthukumar et al. [2021]), *neural collapse* (Papayan et al. [2020], Han et al. [2022]), optimization-based anomaly or phase-transition similar form (Cohen et al. [2021], Fort et al. [2019]), and the fuzzy, diverged front of explanations in deep learning phenomena, as criticized by Razeghi et al. [2025].

A note on polynomial neural networks

Because of the sensitive context of double descent, polynomial suffers the same fate of design, as so is many of the experiments. Its design implementation details must be explicit. That is, of which can be seen in the polynomial transition model. That said, the setting itself is rather more difficult to assess.

Originally, it is asked of the experimental team to set up **adversarial testing**, i.e. testing two mismatched polynomial models together, in a singular testing system. So, we will have (p_k, p_c) for $k \in [a, b]$

¹Conjunction space simply means that the full numerical encoding space of all activities involved

of complexity specification, in this case, the number of initialized degree of interest. As such, the experimental system is between two polynomials. But this turns out to have certain problem: *domain range*. The domain value, when admitted of the input $x \in \mathbb{R}$ or $\mathbf{x} \in \mathbb{R}^n$ of n -dimensional input sequence, is rather difficult to control, especially for $[a, b]$ that is nonstandard in $[-1, 1]$, and not of normalized unit. Usually, we understand the polynomial transition space as the modular $\mathbb{R} \rightarrow \mathbb{R}^h \rightarrow \mathbb{R}$ mapping for singular n -dimensional hidden field, the polynomial basis terms. However, if we are to upgrade it to multivariate polynomial, such changes dramatically, and so is the structure itself. Thereby, so is the behaviours on ranges. The justification for nonstandard, free monomial is that for increased value range, the pattern is not chaotic, but rather smoothed out by dominant terms, for example, the n -degree terms contribution in univariate case. Similar case happens when taken into consideration n -dimensional polynomial mapping on $\mathbb{R}^n \rightarrow \mathbb{R}^m$, but only this time the dominance is local to individual dimension itself. In general,

Conjecture 4.1.1. *The probability of dataset representing wrong n -univariate polynomial p_n , on finite arbitrary space $\mathcal{D} = [a, b]$, for $|\mathcal{D}| \rightarrow \infty$, increases, in presence of $\epsilon \sim \mathcal{N}(a, \sigma)$ noise.*

The main fundamental reason is because of the ill-conditioned Vandermonde basis, in which is typically used in polynomial expression. Elsewhere, on finite space $[a, b]_m$ on m samples, it is then conjectured that the chance of interpolation is much greater, and much more deviated from the actual concept shape if stray over large enough $[a, b]$, thus also influence the error rate and gradient-based learning form of different weights. The goal is then, to construct rich enough polynomial concept class, based in this new discrete, finite range and initialization goals, with the main goal is to find sufficient polynomial construction on arbitrary $[a, b] \in \mathbb{R}$ that y does not explode. This is fundamentally hard, and if unable to procure, we would have to choose to go for the normalized space on $[-1, 1]$. This is also for computational purpose, since float128, of which is the main data point form, is numerically unstable with large operations and the density of operations as such without sufficient safety control, in which it is very hard to reconcile.

One more point to note. What we have seen is based on the perception of typical machine learning setting that **monomial Vandermonde** basis is the polynomial representation space. In fact, this is not true, as there are many as Lagrange polynomial, Newtonian basis, Legendre polynomials, and different classes of orthogonal polynomials. However, without structural operational analysis, such addition and generalization is hard to tract of the structural droplet they introduce.

The following will list the amount or configuration that test should be taken into consideration, first with the guideline.

1. **Carefulness on numerical processing:** Introducing polynomial models meaning that for each component, it is reducibly equivalent to n linear components, for each component, with multiplicative combinator. So for example, a system of full n -degree model sequentially, would have $n! + 1$ linear equivalent mass size, including the bias on numerical space. Therefore, the value in practical calculation can be very large, even in $[-1, 1]$ as the pairwise normalized. Cares then must be taken to draw that out.
2. **Minimal implementation:** Use the least, but consistent code possible. Each code must be recommended to be designed modularly, such that you can snap it in others with only relevant knowledge as to *what it does* and *what is the typing allowed* of the unit itself, aside from utility systems.
3. **Loop collapse:** The order of loop, or changing test case is very necessary to be taken into account of. First one, is the **system test loop**, where the setting of the environment for operation, i.e. noise, concept allowed (dimension or description or class of types), number of samples, and so on. The second one is the **hypothesis-action loop**, of the empirical action of *train-test-validate-cross* partition

of the observational space $\mathcal{S}_n$ to mimic generalization in oracle setting, the hypothesis spaces, and so on. The final one is the **operational level**, where you specify the loss-landscape, learning metric, aggregated measures, loss function measure, actions on learning algorithms, the learning algorithm configurations (for example, learning rate, etc.), and so on.

Aside from such, each layer in the final point *procure measures and statistics*. For example, at system test loop, you have properties of the dataset itself. Distribution and shape or form of the model as the second loop, and so is the jump path, approaches, pathology of the reconstructed loss space. Such should be taken into consideration. The list of model is then:

1. Polynomials, function concept set: This is where we do not require the setting of two-polynomial model configuration together and to study them using a class of different elementary function, only fit the input-output space. For example, singular $\sin(\cdot)$ class that fits $\mathbb{R} \rightarrow \mathbb{R}$ is ideal, or else, or different pathological systems.
2. Polynomial, adversarial $[-1, 1]$ range: Construct both, only in $[-1, 1]$ range for evaluation, since higher range collapses into smoother values.
3. Polynomial, changing basis construction, $[-1, 1]$ range: Use modules, like `numpy.polynomial.legendre`, or `scipy.special.legendre`, and so on.
4. Polynomial, changing basis (plus monomial), $[-1, 1]$ range, input-dropout (filter): Filter out, in the form of hidden information from the hypothesis polynomial, and not the concept polynomial, to simulate ordered information mismatch.

One of the main examples would be to pitch different kind of models together, for example, targeting Legendre polynomial class on monomial going against Newtonian polynomials, or the Vandermonde basis going against other structural models as well. On the reduction of range, what we have done are conditioning upon the instability of large numerical values. This is only a temporary solution, and would almost change the geometry of the problem significantly. As such, methods aside from such mapping reduction can be suggested instead of flat collapse. Several choices are possible, including *affine normalization*, $x' = (x - \mu)/s$ for μ as the centre of the arbitrary range $[a, b]$ that was previously presented, and s as the characteristic custom scale parameter. The cost incurred here is reversible on affine transformation, preserving domain semantics. Standardization or z-score on $x' = (x - \bar{x})/\sigma$ rescaling, whitening (PCA) on $x' = \Lambda^{-1/2}U^\top(x - \bar{x})$ from SVD of the covariance on the dataset structure, Log/Box-Cox transformation, and many more to be specific. In the following up deployment of learning presentation, we would like to examine its behaviour on such space.

4.1.3 Preliminaries on observations - [Schaeffer et al. \[2023\]](#)

Conveniently, one of the few analytical sufficient experimental analysis that we can get, is from [Schaeffer et al. \[2023\]](#) paper on demystifying double descent. While the experimental setup is rather well-defined and interesting, it is also rather limited in scope, of which we want to dive deeper into it and further the analysis that was provided. The setting of the model in precondition was within the setup of a duplet correlation dataset,

$$\mathcal{D} = \{(\mathbf{x}_n, y_n)\}_{1 \leq n \leq N} \quad (4.2)$$

for N samples with variational vector x as input. The function is the classical functional type of $y = c(\mathbf{x})$ of the hidden concept, of which is assumed to utilize all variables in the setting of $\mathbf{x}$ itself. Thence, objective is to figure out for the model $f(x)$ given, such that

$$f(\mathbf{x}) = \mathbf{x} \cdot \boldsymbol{\beta} \approx c(\mathbf{x}) \quad (4.3)$$

In a sense, if the setting is ordinary linear regression, then such model equates to multilinear model, unless $|\mathbf{x}| = 1$, without bias b on the conditioned projection space. The experiment conducted in the test is rather simple, with the setting being,

$$|\mathbf{x}| = 1, \quad x \in [-1, 1], \quad y \in \mathbb{R}, \quad c(x) = 2x + \cos(25x) \quad (4.4)$$

and within such approximating $c(x)$ using a univariate polynomial on feature map ϕ_P of which maps any x to P -degree polynomial basis vector, and collapse it of such for the fitting model $f(x)$ to be now

$$f(x) \equiv \phi_P(x) \cdot \beta_P \approx c(x) \quad (4.5)$$

While the paper indeed introduced a larger and logical intuition concepts, plus analysis into how one might interpret such result using ordinary least square and the general regression problem, such is often poorly directed toward where the analysis itself is rather only additive than any. First, is the legible assumption of the generality of application on results, of the tested function class of cosine class functions with additive conditional bias drift on $2x$ term. The paper insufficiently so, does not tell us about how this visual explanation and analysis can help us in understand more about double descent aside from the optimization path analysis already did in several other analyses. The family of cosine equations itself imposes a rather non-trivial properties, it forces **periodicity** and value space bounding on dataset itself. That is, the choice of giving cosine class concept function morphs the moving optimization space to a more narrow space, on which there exists x -consistent spread and structures, i.e. a **concentrated space** class. During which, the shift can give the cosine class a linear upward drift, but not enough toward which to be called exploding, like in many cases of polynomial models or complex pollinated classes dynamics. Effectively, such is to constrain c into the real analytic on compact interval problem, or $c \in C^\omega([-1, 1])$ space, and of which spectral decays and heavy-tailed or sparse spectral structures are rare or nonexistent. Second, constraining the setup to $[-1, 1]$ directly interferes with the stability of the ill-conditioned Vandermonde basis polynomial space, of which setup a rather non-trivial optimization path under iterative algorithms. Thirdly, the experiment deliberately removes noise-conditioning setting, of which means that the setting assume of the synthetic analysis *full oracle access* on details, of which is a non-trivial choice.

Continuing forward, the three critical quantities are generally rather not explained. For example, it is mentioned that removing small singular values in the training features X prevents double descent, however such removal is non-trivial and requires justification on why this is a rather good choice, and whether a synthetic dataset can remove such and *where* would it be so. Preventing the test feature from varying in trailing singular modes prevents double descent, however with such singlar modes are there is unknown, and there are no justification - whether if such mode is detrimental to the actual dataset or not of the arbitrary concept classes. Similarly so, ensuring that the optimal model in the model class has zero residual prediction errors E is stated to prevent double descent, however, that is a rather metric-dependent optimization scheme, and might as well lead to the *overfitting regime*, where at this behaviour, one actually returns to the question of overfitting both test and train sets, and would have to rely on the third set to see if double descent appears there or not. Either way, the setting itself is interesting to analyse, and one would consider it interesting to expand on their analysis, especially to different finite-domain conditioned system and concepts.

4.2 Data generation

Dataset generation, coming off the analysis on the actual complexity of the concept, structural exchanges and thereof, requires at this stage deliberate elaboration on the method of synthesizing. Indeed, while typical dataset on synthesizer are often obtained by hard-coded artefact, plus random noise

$\epsilon \sim \mathcal{N}(0, 1)$, such is rather not desirable and certainly not the only type of noise in the toolbox of statistical analysis. Certainly so, this is illustrated in Figure 4.1, where there exists a rather rich family of noise to consider, that perhaps are non-trivial of analysis given the fact that different noise profile can be either irreducible, or by virtue contains structural details hidden from what we observed of the relationship itself. Thence, let us analyse how we can gain sufficient dataset corruption, noise channel, and so on,




Figure 4.1: The different variations of the noise profile that can be often seen when sampling from a simple regression model. By order, (1) represents mean-scale additive noise, (2) considers heteroskedastic multiplicative noise, (3) considers variance-scaled additive noise, (4) considers absolute additive noise (where scale independent of y), (5) considers heavy-tailed noise (Student- t modelling), and (6) is the standard canonical Gaussian noise (basically the default). All are under reduced dimension of the simplest linear model on feature space as possible for illustration.

given the fact that such exists of its proliferation. Which can be solved using a model of machining. This will be the precursor for the larger consideration in later chapters. Let us begin by the domain analysis of such. For irreducible and independent noise from concept structure of c , then we are allowed to add noise to any given observable components of the observation space. Specifically, given a concept under functional form, $c(x) \in \mathcal{C}$, the concept itself works by the following dimensional components,

$$\mathbb{R}^d \rightarrow \mathbb{R}^{w \times b \times \theta} \rightarrow \mathbb{R}^p, \quad \mathbb{R}^d \times \mathbb{R}^m \rightarrow \mathbb{R}^{w \times b \times \theta} \rightarrow [\mathbb{R}^p]^m \quad (4.6)$$

for the first case being singular pass, and the second one is m -size inference on c . Under the same chain, we can see that, for example, the noise can be proliferated by

$$\Sigma_1[\mathbb{R}^d \times \mathbb{R}^m] \rightarrow \Sigma_{2,3,4}[\mathbb{R}^{w \times b \times \theta}] \rightarrow \Sigma_5\{[\mathbb{R}^p]^m\} \quad (4.7)$$

for five distributive component entries on exertion. Now, the synthetic scenario has four cases, either where you only take the representative mass on $\mathbb{R}^{w \times b \times \theta}$ as black-box plus consider it inaccessible, or accesible, not black-box and you can condition specific breakage of noise profile into different structural packets, accessible but blackbox such that $\Sigma_{2,3,4}$ are chosen at will, randomly, or ignore such noise, which is the comparatively worst case under this ladder of transparency. Thus, it leaves us with trivial Σ_1 and




Figure 4.2: The different variations of the noise profile that can be often seen when sampling from a simple regression model, but for different controlling nob for strength of the noise, by σ in range of $[0.1, 0.2, 0.3, 0.4, 0.5, 0.6]$. Again, by order, (1) represents mean-scale additive noise, (2) considers heteroskedastic multiplicative noise, (3) considers variance-scaled additive noise, (4) considers absolute additive noise (where scale independent of y), (5) considers heavy-tailed noise (Student- t modelling), and (6) is the standard canonical Gaussian noise (basically the default). All are under reduced dimension of the simplest linear model on feature space as possible for illustration.

Σ_5 noise profile, which brings on itself several risks, including the fact that the model itself might have its noise interacting with the structure of the model itself. As during Figure 4.1 and Figure 4.2, the modification is often only for y , such is where it occurs that the $\mathcal{U}$ -measure, which will be motivated by the quantification of the dataset itself, predicts that the uncertainty profile differs by a series of margins. Specifically, this is illustrated by Figure ??, where the uncertainty aggregation on geometrical uncertainties are within $2\sigma^2$ since we fix only x -based corruption on random noise. We can see that each profile incurs different uncertainty profile, of which if to be also conditioned on x -corruption - which is plausible more or less in practice - and various range conditioning, such can be seen to happen as that more and more of the concept class would be reiterated into varying uncertainty class far from each other, thus raising the point about generality of noise channel class.

In general, noise are typically of consideration from two testing apparatus perspectives. First, is the *empirical laboratory consideration*, or supposedly full-control in designer's corruption mechanism and intervention on the observations given; the second one being *blind noise consideration*, in which noise is treated to be irreducible, of no origin and in which contributes only distortion on the dataset itself, without structural reason. As for such, how the dataset is interpreted depends on what version of the two extreme one take on the realization of the dataset itself, such is to say a way to analyze the dataset, within various statistical methods, or structural one.

Aside from the main statistical approach, of which can be considered in many different ways, most of-ten probabilistic-random variable approach to data extrapolation, the space encoding of the dataset itself is often very much interesting on itself, especially for a learner expression that would approximate such concept. Since we already used an illustration of the rather more complex notion of computational systems, we would like to use that in-effect of one of the notion of sample space analysis, i.e. the uncertainty




Figure 4.3: Noise uncertainty creation, aggregated through the metric $U(S'_n)$ of the diffused dataset. Here, we can see that for y -only corrupted noise, the uncertainty increased on bounding for $2\sigma^2$, of which diverges several noise profiles from one another. An interesting outlier is in the suppressing effect of mean-scale additive noise, which is not as expected.

metric U .

Polynomial regression is then one such analytical solution to function approximation task, using Weierstrass theorem as its basis. From such, we can also observe that polynomial regression, and hence polynomial interpolation, are correlated and can be reduced to each other. Polynomial regression can be reduced to interpolation via stating the criteria such that, denoting it as ℓ , then $\ell(f, p) = 0$, and interpolation can be reduced to regression via the condition that $p(x_i) - y_i \leq \epsilon$ for arbitrary $\epsilon > 0$. Hence, the regression formulae lie between the two extreme: the Weierstrass setting and the interpolation setting.

We then have to identify such extreme and what kind of extreme they represent. Fortunately, we know this as the problem of **transparency**, in the setting of approximation, which certainly correlates to the learning theory. According to such, the Stone-Weierstrass represents the total transparency case, where the function f is known to the approximator with polynomial class. Interpolation represents total black box, where the goal is only to interpolate the sample point by itself. In between such interpretation, there exists a fundamental assumption, hence the structural assumption.

Assumption 4.2.1 (Transparency). *For any bivariate pair (x, y) , we assume each pair depends on and originate from a based concept $y = f(x)$.*

Two kinds of knowledge is then available, which correlates to the degree of freedom the observation space can have. First is the knowledge of f . This knowledge of f is encoded directly by the number of sample point and their distribution in certain case. Accordingly, the worst case is as followed.

Conjecture 4.2.1. *If $\forall x$ there exists y corresponding to such, and $(x, y) \in \mathcal{S}$, then we say we have total transparency to the concept f . If there exists only finite x that we can have y corresponding to, then we say we have limited transparency to the concept f . If we assume no knowledge of f , then we say the setting is called the black-box setting.*




Figure 4.4: The different variations of the noise profile that can be often seen when sampling from a polynomial model, particularly constrained, 5th degree p_5 univariate polynomial on $[-1, 1]$ domain inverse range. By order, (1) represents mean-scale additive noise, (2) considers heteroskedastic multiplicative noise, (3) considers variance-scaled additive noise, (4) considers absolute additive noise (where scale independent of y), (5) considers heavy-tailed noise (Student- t modelling), and (6) is the standard canonical Gaussian noise (basically the default). We can see that each distributive pattern produces various different, often irreducible to Gaussian noise configurations.

Secondly, is the action on the dataset itself and what kind of concept is used to model the given structural observation. One of the main assumption that is usually observed is that, we presuppose a specific structures that after a finite amount of iteration, can return us back the desired results.

Assumption 4.2.2 (Fallacy). *We assume there exists every specific bivariate pair $(x_i, y_i) \in f[\mathcal{S}_n]$, such that $f(x_i) = y_i$.*

However, usually, this is not the case, as we have **noise** in addition from the dataset. The origin of the noise cannot be tracked. However, we can interpret what happens when such noise is introduced to the dataset. The noise transitions the transparency of the learned setting, on the field $\mathbb{R}$, by utilizing the dense property of the real space. Assuming we operate on $\mathbb{Z}$. Then, each discrete noise will introduce abnormality point such that no function on $\mathbb{Z}$ will be able to match, up to certain distance on $\mathbb{Z}$. This is because on such space, there exists no point x_1, x_2 , for example, that lies between $d(x_1, x_2) = 1$. If the noise is randomized on ± 1 , then there will exists no function that can approximate such disparity better than a random fluctuating function. However, on $\mathbb{R}$, the real number is inherently dense, such that there will always be certain point x_3 such that $x_3 \in [x_1, x_2]$, by the density of real number. By the distance metric on $\mathbb{R}$, we define $d(X, Y) = |X - Y|$ for distance between two data points. Noise intrinsically increases such distance, such that $d'(X, Y) > d(X, Y)$. Hence, the space in which density can appear increase, thus there then exists more fluctuation, and thus more uncertainty between such two points. This notion can be increased from local uncertainty to global uncertainty, hence the *noisy dataset*.

We then can measure or aggregate some properties of the dataset itself, including the **uncertainty** of the dataset. We define the simple uncertainty of the dataset as followed.

Definition 4.2.1 (Dataset uncertainty). Given a closed interval $[a, b]$, for the sample set $\mathcal{S}_n = (X, Y) \in [a, b]$, the uncertainty $\mathcal{U}(\mathcal{S}_n)$ is defined as

$$\mathcal{U}(\mathcal{S}_n) = \left[\mathbb{E}_{x_i \in X} (x_{i+1} - x_i)^2 + \mathbb{E}_{y_i \in Y} (y_{i+1} - y_i)^2 \right]^{1/2} \quad (4.8)$$

We assume X, Y are totally ordered sets.

Hence, the assumption of fallacy assumes $\mathcal{U}(\mathcal{S}_n) = 0$, though in reality it is often $\mathcal{U}(\mathcal{S}_n) > 0$. We would like to clarify, that such definition is currently reserved for only polynomial-like functional class; though, this can be extended pretty easily for any $\mathbb{R}^n \rightarrow \mathbb{R}^m$ functional concepts. Furthermore, it is *operational*, in the sense that it measures the numerical encoding inherent uncertainty, of which is also the space used of algorithms traversing in stochasticity. It affects cross-validation, gradient-based algorithm path finding, and roughness of the learning processes. Certain concerns must also be addressed, such as the definition, which assumes *sequential processing* and pair-wise measurement (which can also be extended to tri-fold measurement and so on). Indeed, modern deep learning use parallel batch processing, mini-batch sampling, or distributed learning across devices for optimization. Nevertheless, such includes sequential-based learning optimization, and so on, where mini-batch still inherently includes batch ordering; algorithms like RNNs and so on also are sequential. The computational cost per dataset, if not, is $\mathcal{O}(n)$ after initial order, or with mundane immediate measurement on the way that the dataset was configured at start, since the system and stochastic algorithms sequentially go through them.

Certain others aspects can be expected to be clarified. While indeed, this is order-dependent measure, and one might expect to go against it and say it is not real dataset property, but the key point is that dataset are measurements, captured and recorded into the encoded space that is used upon the learning regime (supervised or semi-supervised or so on), or the learning algorithms. In a classical algorithmic world, such must be done in discrete steps, and hence ordering strongly matters as well as the processing pipeline relevant to it. There exists natural ordering of data within noise ϵ , because if now, the projection on embedded space $\mathbb{R}^n \times \mathbb{R}^m$ of any given concept functional c holds no value, as the function itself, plus the distribution, with respect to the encoding space where the function mapping is defined already contains structural information encoded within 'value range' as the most basic form. A weaker posture we can adopt is that it is an operational uncertainty measure, given presentation order to a finite discrete processing machine, and is indeed *pre-training* diagnostic. There is also the issue in which X and Y might not be measurable or comparable. Indeed, since the range of value can jump intermittently between different concept class and form, but even then, it will only highlight the disparity between the contributing terms. For example, if we have three pairs, $(130, 1)$, $(150, 1)$ and $(225, 0)$ as the data, then the volatility is not that bad, and stable on Y , but *very high on X* , as there are missing information in large gap. From the machine-algorithmic, this matters. if we take the 130, 225 points as a case of shuffling, then the machine experiences a 95-to-1 volatility ratio. Nevertheless, one can attempt a normalized metric, i.e.

$$\mathcal{U}_{\text{norm}}(\mathcal{S}_n) = \sqrt{\mathbb{E}[(\Delta x_i)/r(X)]^2 + \mathbb{E}[(\Delta y_i)/r(Y)]^2} \quad (4.9)$$

for relative jumpiness across different problems with normalized morphism. With this, we can quantify how the dataset is behaving on the scale of the bivariate form typically seen. Noise intrinsically increases this distance, usually by the y -axis. This dataset uncertainty increases with multivariate dataset, for example, $(\mathbf{x}, y)$ of $\mathbb{R}^n \rightarrow \mathbb{R}$. We call the new metric as the *empirical uncertainty metric*. The basis of such argument on the concept of a tradeoff comes from Nakkiran [2019a], in which linear regression is affected, at least in such setting and consideration, such that dataset will both hurt and benefit the learning process.

Overall, however, this allows us to state of the following theorem:

Theorem 4.2.3 (Separation Encoding under Diffusion). *Fix the encoding space $\mathbb{D} = \mathbb{R}^q \times \mathbb{R}^p$ for a (q, p) conjunction space. Let $\mathcal{S}_n = \{(x_i, y_i)\}_{i=1}^n \in \mathbb{D}$ be a dataset of n ordered samples drawn from an oracle $EX(\theta)$. Let $\{\mathcal{N}\}$ denote the set of noise configurations $\eta = \{\eta_i\}_{i=1}^n$ where each $\eta_i \sim \mathcal{D}$ with $\mathbb{E}[\eta_i] = 0$ and $\mathbb{E}[\eta_i^2] = \sigma^2 < \infty$, independent across components. Define the diffused dataset $\mathcal{S}'_n = \{(x_i + \eta_i^{(x)}, y_i + \eta_i^{(y)})\}_{i=1}^n$. Then*

$$\mathbb{E}_{\eta \in \{\mathcal{N}\}} [\mathcal{U}(\mathcal{S}'_n)^2] = \mathcal{U}(\mathcal{S}_n)^2 + 4\sigma^2 \quad (4.10)$$

In particular, for any $\sigma^2 > 0$ and any zero-mean finite-variance $\mathcal{D}$,

$$\mathbb{E}_{\eta \in \{\mathcal{N}\}} [\mathcal{U}(\mathcal{S}'_n)^2] > \mathcal{U}(\mathcal{S}_n)^2 \quad (4.11)$$

Proof 1. Simple proof. Using normed metric, $\mathcal{U}_{\text{norm}}$. For noised data, that is,

$$\mathcal{S}'_n = \{(x_i + \eta_i^{(x)}, y_i + \eta_i^{(y)})\}_{i=1}^n \quad (4.12)$$

where the noise vector on encoding space is orthogonal, such that

$$\eta_i = (\eta_i^{(x)}, \eta_i^{(y)}) \sim \mathcal{D} \approx \mathcal{N}(a, \lambda \text{Cov}(\eta)) \quad (4.13)$$

where we can fix $\mathcal{D} = \mathcal{N}(a, \lambda \sigma^2)$ for $a = 0$ as normalized around the zero value. If we assume independent noise, then we have the covariance as

$$\lambda \text{Cov}(\eta) = \lambda \begin{pmatrix} \sigma_x^2 I_q & 0 \\ 0 & \sigma_y^2 I_p \end{pmatrix} \quad (4.14)$$

We can simplify slightly, as to have the case in consideration that $\sigma_x^2 = \sigma_y^2 = \sigma^2$, then the notation would effectively be such that $\lambda \sigma^2 I_{d+p}$ is enough. The uncertainty metric then should be expressed as

$$\begin{aligned} \mathbb{E}[\mathcal{U}(\mathcal{S}'_n)^2] &= \mathbb{E} \left[\mathbb{E}_{x_i \in X} (x_{i+1} - x_i)^2 + \mathbb{E}_{y_i \in Y} (y_{i+1} - y_i)^2 \right] \\ &= \mathbb{E}_i [(x'_{i+1} - x'_i)^2] + \mathbb{E}_i [(y'_{i+1} - y'_i)^2] \end{aligned} \quad (4.15)$$

for $\mathbb{E}_i$ here denotes the expectation over the uniform index i of each component of the sample space. Now, we substitute the following into the empirical average pair:

$$x'_{i+1} - x'_i = (x_{i+1} - x_i) + (\eta_{i+1}^{(x)} - \nu_i^{(x)}), \quad y'_{i+1} - y'_i = (y_{i+1} - y_i) + (\eta_{i+1}^{(y)} - \nu_i^{(y)}) \quad (4.16)$$

First, for x_i, x_{i+1} set, taking expectation over noise η gives

$$\begin{aligned} \mathbb{E}_\eta [(x'_{i+1} - x'_i)^2] &= \mathbb{E}_\eta [(x_{i+1} - x_i)^2 + 2(x_{i+1} - x_i)(\eta_{i+1}^{(x)} - \eta_i^{(x)}) + (\eta_{i+1}^{(x)} - \eta_i^{(x)})^2] \\ &= \mathbb{E}_\eta [(x_{i+1} - x_i)^2 + (\eta_{i+1}^{(x)} - \eta_i^{(x)})^2], \quad \mathbb{E}_\eta [2(x_{i+1} - x_i)(\eta_{i+1}^{(x)} - \eta_i^{(x)})] = 0 \\ &= \text{Var}(\eta_{i+1}^{(x)}) + \text{Var}(\eta_i^{(x)}) \\ &= 2\sigma^2 \end{aligned} \quad (4.17)$$

for x_i fixed and $\mathbb{E}[\eta_i^{(x)}] = 0$ that leads to the second zero equality, independent noise $\eta_{i+1}^{(x)}, \eta_i^{(x)}$ with variance σ^2 . Therefore, we have for each i pairing, that

$$\mathbb{E}_\eta [(x'_{i+1} - x'_i)^2] = (x_{i+1} - x_i)^2 + 2\sigma^2 \quad (4.18)$$

Averaging over i and apply the same argument to y -component, we gain identical form with doublecondition, such is,

$$\mathbb{E}[\mathcal{U}(\mathcal{S}'_n)^2] = \mathbb{E}_i [(x_{i+1} - x_i)^2] + \mathbb{E}_i [(y_{i+1} - y_i)^2] + 4\sigma^2 = \mathcal{U}(\mathcal{S}_n)^2 + 4\sigma^2 \quad (4.19)$$

Since $\sigma^2 > 0$, the expression gives $\mathcal{U}(\mathcal{S}_n)^2 + 4\sigma^2 > \mathcal{U}(\mathcal{S}_n)^2$ by Jensen inequality. The statement works under expectation in the squared sense directly, and for $\mathcal{U}$, we expect monotonicity on $\mathbb{R}_{\geq 0}$ for square terms reduces to the expression on non-squared condition. $\square$

Specifically, the unsquared version will give

$$\mathbb{E}[\mathcal{U}(\mathcal{S}'_n)] \leq \sqrt{\mathbb{E}[\mathcal{U}(\mathcal{S}'_n)^2]} = \sqrt{\mathcal{U}(\mathcal{S}_n)^2 + 4\sigma^2} \quad (4.20)$$

which again, requires the argument going through the squared form by condition. The condition that invalidates the theorem under such form would be for any non-zero mean noise, for $\mathbb{E}[\eta_i] = \mu \neq 0$, which does not collapse the contributing terms. The middle $2\mu \cdot \mathbb{E}_i[\cdot]$ term then is a telescoping sum proportional by $(x_n - x_1)/(n - 1)$, which can be either positive or negative depending heavily on trend in X space, so biased noise can actually **increase or decrease** uncertainty extrapolated depending on dataset structure. This is rather interesting, because it highlights a particular region where systemic bias (or so we called of such), interacts with dataset ordering in a non-trivial way, which is far harder to concentrate. In fact, the real form under arbitrary notion can be written as

$$\mathbb{E}_\eta[\mathcal{U}(\mathcal{S}'_n)^2] = \mathcal{U}(\mathcal{S}_n)^2 + 2(1 - \rho)(\sigma_x^2 + \sigma_y^2) + \text{Bias} \quad (4.21)$$

of the bias term arbitrary, and covariance $\rho\sigma^2$ for correlation of both η_i, η_{i+1} , which is rather interesting, and targets the case of sensor drift in observational spaces. Finally, the constraining partition is that if we have infinite variance on $\mathcal{D}$, then the finiteness of $\mathcal{U}(\mathcal{S}'_n)$ is not finite and the comparison is undefined. As such, while interesting, and indeed extensible to a wider variety of control space for dataset, we note the inherent understanding and restriction of such uncertainty measure.

Such type of experimental pre-pipeline investigation can also be very informative on the structure of the dataset in whole, particularly within the tension of **dataset identifiability** projection. Specifically, it is where we are to hinge upon the two perspective on absolute noise independency from structures, or structural noise, of which would be rather difficult to approach form. This can be raised of the problem as followed. In a specific dataset, between the *actual concept* and the observable spaces, lies the empirical and generalization setting. While it is not often mentioned in literatures, the switch to direct observable evaluation of statistical learning, induces certain limitations that would determine how any specific dynamical system operate under its domain, and of which the data points itself, i.e. here we call it as observational points of the observation space, be inferred of one another, or our dataset being a snapshot of a limited frame in which the dynamics are frozen in capture.

dataset
identifiability

Let us assume the following preliminary in discussion. We have two polars, on the mechanistic absolute, or *white box* knowledge, with full expressions on the underlying mechanics of a certain system concept. Explicitly, this can be written as $y = F(x; \theta)$ under standard morphism, for any kind of description θ . Extension of such notion can also include what we call as the mechanism, or $F(x; \theta, \Gamma)$ of which

complies both the parameters, and the formulations. The other polar is the *probabilistic domain*, of which then, we have the inference $\sim$ operator instead. Thus, we denote $y \sim \mathbb{P}(Y | X)$ in typical notation, in which the dependency is now Y on the basis of X . One can also do the reverse, as to generate the inverse conditions, but this is more probably general, if one switch to joint distribution (X, Y) under statistical schemes. Abstractly speaking, the dual structure contains what we call temporary, as the **system action** A_S . The notation then becomes

$$y A_S x \rightarrow f(x; \theta, \Gamma), \quad y A_S \rightarrow \mathcal{S}[\mathbb{P}(X, Y)] \quad (4.22)$$

For $\mathcal{S}$ here temporarily denotes the *sampling operations*, of observing randomly on the basis of said probability without any knowledge but the presumed parameters. If such process is with control parameter, from the optimization perspective, we have the notation as different to be $\mathcal{S}[\mathbb{P}(X, Y); \theta_P, \Gamma_P]$, where now Γ_P is the probabilistic process, i.e. distributions or other random structure, for example as stochastic processes as a finite-dimensional distributions plus regularity structure, for random variables θ_P . In general, we call it the **black-box domain**.

If we are to recover θ from θ_P , or at least on the similarity metric, typically or ML encoded by $\ell(\theta, \theta_P)$, where $\ell(\theta, \theta_P) \approx 0$, then we say that the system experiences a *white transition*. Note that we do not take into account of Γ and Γ_P , since they are inherently different in this white transition, but mechanistically recoverable; thus, *total white transition* means $\Gamma \approx \Gamma_P$ of the mechanistic patterns. If the reverse happens, where mechanistic system turn to probabilistic (X, Y) distribution on θ_P and structure $\Gamma \rightarrow \Gamma_P$.

Probability inherently, similar to mechanistic Γ , enforces causal or rather transformational relation (simply the dual $x \rightarrow y$ is transformational within f). There, the probability can be interpreted, for either Γ explicitly states that they are conditional, for X to be deterministic and explicitly stated, or joint distribution and so on. However, the interference in such recovering mode comes from observational space inconsistency. This is usually expressed of *noise*, where if we admit (x, y) of the mechanistic system, then for $x + \epsilon_x$ and $y + \epsilon_y$ we have instead a *noisy mechanistic system*. We would note that this is a larger class than pure mechanistic system, basically $x \rightarrow [\epsilon, f(x; \theta, \Gamma)]$ instead. Thus, for intrinsic noise of measurements or statistical information, probabilistic transition in recovery of mechanistic systems admits the recovery of the noisy system instead. Depends on the severity of ϵ , such system can be irrecoverable. If such approximation is impossible within certain margin, then probabilistic approximation remains in probable the only manageable system of approximation, rather than a quasi mechanistic recovery mode (we will prove this in a later theorem). If we to accept such view, then it means that in certain failure modes, you *cannot move a mechanistic (deterministic) system into probabilistic then back*, given enough disturbance and noise. This quantifies to the *information loss* system-wide, where we can create a prototypical functional, denoted $\mathcal{R} : F \rightarrow P$, where

$$\mathcal{R}[F](A) = \int_{\Theta} \mathbf{1}_{\{\mathbf{F}(\mathbf{x}; \theta, \Gamma) \in A\}} \mu(d\theta) \quad (4.23)$$

for such transition, of the probabilistic measure space $(\Omega, \mathcal{F}, \mathbb{P})$, for $\mu \in \mathcal{F}$ and $A \in (\mathcal{X}, \mathcal{Y})$ of the observational space. The implication here is that, under no information (information-agnostic), every noisy concept must meet certain limit on the basis and resultant structure itself. There exists, presently, no sufficient guarantee to return P to F , except the same operation in reversal, is forced on ideal setting, for infinite-range coverage, and approximately dense observational space. The limit is, inherently, enforced by arbitrary ϵ . It is to note that we do not claim that mechanistic systems are ontologically privileged; we claim that the mapping from mechanisms to observations is not generally invertible under disturbance. Thus, effectively, what we have done is simply the motivation for said approach, and the importance

of considering the entire dataset space rather than just aggregated on value-specific distribution under measurements.

4.3 Linear function class $h(u)$ - Lafon and Thomas [2024]

Linear-linear scenario learning is one particularly well-studied case of representing double descent or not, usually through simplicity in analysis (of the similar ease of proof compared to binary classification). Models in this setting can be said to share the same *parameter space* with no additional structure or scaling mask (like in monomial basis). The model that is representative of this preliminary observation is from Lafon and Thomas [2024], specifically, linear regression (function approximation) with Gaussian noise, and of univariate case. Consider the family class $(\mathcal{H}_p)_{p \in \llbracket 1, d \rrbracket}$ of linear functions $h : \mathbb{R}^d \mapsto \mathbb{R}$ where exactly p components are non-zero ($1 \leq p \leq d$), we have the following model.

Definition 4.3.1 (Lafon and Thomas [2024]). For $p \in \llbracket 1, d \rrbracket$, $\mathcal{H}_p$ is the set of functions $h : \mathbb{R}^d \mapsto \mathbb{R}$ of the form:

$$h(u) = u^T w, \quad \text{for } u \in \mathbb{R}^d$$

With $w \in \mathbb{R}^d$ having exactly p non-zero elements. The complexity here itself can be interpreted as the effective random initialization, trimming down data of hidden details instead - this is indicated by the amount of information it observes given d dimension, and p -hypothesis dimension.²

Using such model, they consider the setting of minimizing the customized mean square loss,

$$\min_{w \in \mathbb{R}^d} \frac{1}{2} \|Xw - Y\|^2 \rightarrow \min_{w \in \mathbb{R}^p} \frac{1}{2} \|X_{\sim p} w - y\|^2 \quad (4.24)$$

of the sub-problem on p cardinality, where again p is the dimensionality or independent degree of freedom. This results in the following theorem that 'mimic' double descent on this linear regression model.

Theorem 4.3.1 (Double descent on linear regression - Lafon and Thomas [2024]). Let $(x, \epsilon) \in \mathbb{R}^d \times \mathbb{R}$ independent random variables with $x \sim \mathcal{N}(0, I)$ and $\epsilon \sim \mathcal{N}(0, \sigma^2)$, and $w \in \mathbb{R}^d$. we assume that the response variable y is defined as $y = x^T w + \sigma \epsilon$. Let $(p, q) \in \llbracket 1, d \rrbracket^2$ such that $p + q = d$, $X_{\sim p}$ the randomly selected p columns sub-matrix of X . Defining $\hat{w} := \phi_p(\hat{w}), [\hat{w}]$ with $(\hat{w}) = X_{\sim p}^+ y$ and $[\hat{w}] = 0$. The risk of the predictor associated to $\hat{w}$ is:

$$\begin{aligned} & \mathbb{E}[(y - x^T \hat{w})^2] \\ &= \begin{cases} (\|w_{\sim q}\|^2 + \sigma^2)(1 + \frac{p}{n-p-1}) & p \leq n-2 \\ +\infty & n-1 \leq p \leq n+1 \\ \|w_{\sim p}\|^2(1 - \frac{n}{p}) & p \geq n+2 \\ + (\|w_{\sim q}\|^2 + \sigma^2)(1 + \frac{n}{p-n-1}) & \end{cases} \quad (4.25) \end{aligned}$$

The point $p = n$ is representative of the interpolation threshold from underparameterization to overparameterization.

We present two parts - first are those fundamentals predicted by the theorem, and second are verifications of the theorem with our own version of testing. The first experimental setting is then conducted

²According to the *manifold hypothesis*, there are chances where the actual concept can be interpreted and mimicked by seeing that it lies on a lower-dimension manifold inside its actual expressive dimensional space.

with fixed d , varying p for each varying n samples. Since this is a least square result, preliminary run is conducted using one-shot error-risk measure on least square closed-form approximation. Note that $p \leq d$, hence we are regarding the domain of the term underparameterized and equal parameterization setting. The result is presented in Figure 4.5




Figure 4.5: Theorem 4.3.1 behaviours on randomized setting. For this model, we have $d = 100$, $\sigma = 0.5$, and variational n . Full test cases are for $p = [1, 100]$, $n = \{12, 24, 36, 50, 76, 82, 90, 100, 106, 112, 121\}$ accordingly. The test function of the concept itself is the same linear model, but with different input only.

As we see from this particular experimental result, the theorem predicted a loose, left-peak absent, double descent curve estimation for a variety of test case with more and more samples in increment. There exist some fundamental conclusion that we have about the shift of the interpolation threshold upon to $p = n$, which is reflected in the piecewise prediction. However, this pattern breaks down for n higher than specific constant, for this case, starting from $n \approx 100$ to beyond. From here, at least from testing scenario, the test indicates that predictions suggest a breakdown of the double descent estimation. Beyond this, a simple semi-chaotic convergences is observed, from $n = 106, 112, 121$.

Interestingly, verifying this notion of easy enough to simplify it into what the statement holds under particular notion. Namely, as we find out, the above double-descent replicating formula can actually be realized by the parameter-wised error, comparing the parameter between the least-square solution and the actual concept. The formula for this is expressed by the exposed MSE on parameter function, for $\beta_{p,k}$ the coefficient value estimated for the k -th selected feature when the model is restricted for randomization of p features,

$$\text{MSE}_{\text{param}} = |\hat{w} - w|_2^2 = \sum_{j=1}^d (\hat{w}_j - w_j)^2 \quad (4.26)$$

where $\hat{w}_j$ is chosen such that it is equal 0 for $j \notin S$, otherwise $\beta_{p,k}$ for $S = \{i_1, \dots, i_p\}$ is the set of selected feature indices. Nevertheless, it is then found dubious why this is true, as for the original statement of

Lafon and Thomas [2024] indicates correlation to the prediction error of the result, $x^\top \hat{w}$ instead. The result of evaluating $\text{MSE}_{\text{param}}$ is presented in Figure 4.6.




Figure 4.6: Contrary to Theorem 4.3.1, this is the behaviours on randomized setting for standard parameter-wised mean square error measure (MSE-on-parameter). For this model, we have $d = 100$, $\sigma = 0.5$, and variational n . Full test cases are for $p = [1, 100]$, $n = \{12, 24, 36, 50, 76, 82, 90, 100, 106, 112, 121\}$ accordingly. The test function of the concept itself is the same linear model, but with different input only.

If, however, we turn our attention to the statement itself, of which refers to the prediction error MSE_{pred} , defined by the equation

$$\text{MSE}_{\text{pred}} = \frac{1}{n} \left\| y - X_p \hat{\beta} \right\|_2^2 = \frac{1}{n} \sum_{i=1}^n \left(y^{(i)} - x_p^{(i)T} \hat{\beta} \right)^2 \quad (4.27)$$

on h and c accordingly, a strong case of this happens to be not following the pattern indicted by the theorem. Rather than being the interpolation point of return, $p = n$ is now the convergence point of the prediction error instead. Test result is shown in Figure 4.7.

Thence, the theorem gives us two things – the interpolation point, in this setting, refers to the parameter error for their double descent behaviour. However, on the prediction results’ error measure, the interpolation point correlates to the point where the defined generalization risk in such case equals zero. These patterns hold up to certain point. However, the test case is only up to $d = p$. An unexplored test case for $p > d$ indicate the surprising fact according still, to the theorem: the generalization risk takes double descent as a region with disturbance in the middle, but bias-variance appears afterward. This is illustrated in Figure 4.8, where the same theorem is applied but in the domain which is usually called as *overparameterized region*.

Such can then be seen as a prediction that double descent in certain case, can be extrapolated to be the transition between the two curvature of bias variance, or the abnormal point on the error curve measure. Such must then be verified, both in the experiment’s correctness, and so on that lead to such result. Nevertheless, if true, it would mean that double descent indeed has another implication of itself.




Figure 4.7: Contrary to Theorem 4.3.1 and both figure 4.6, this is the behaviours on randomized setting for standard parameter-wised mean square error measure (MSE-on-prediction). Especially, we do not see double descent pattern in this setting, as it is. For this model, we have $d = 100$, $\sigma = 0.5$, and variational n . Full test cases are for $p = [1, 100]$, $n = \{12, 24, 36, 50, 76, 82, 90, 100, 106, 112, 121\}$ accordingly. The test function of the concept itself is the same linear model, but with different input only.




Figure 4.8: Similar experiment according to setting and result of Theorem 4.24, where p is in the range $[1, 5523]$. Here, we can see small double descent modelled exactly at the interpolation threshold, while the later region exhibit similar phenomenon as bias-variance tradeoff.

4.3.1 Extension

The current theory prompted for a more general treatment of the linear function class, as in-line with exploring reasonable existing alternatives in interpretation (Nakkiran [2019a]). We have much to desire of, for example, extension of prediction for stochastic or iterative algorithms, consideration on statistical formations, uncertainty metrics. Or furthermore, different kind of linear models, including multilinear

models and so on, with variational setting meaningful to the problem space.

Let us state some hypotheses about the previous experimental cases. First, is the lack of the beginning of the prediction on double descent of linear random cutoff features model. This can be explained informally as to be resulted from the uniformity of the operating space. Specifically, all parameters of the $\mathbb{R}^n$ parameter space, with respect to the dual space on such $f : \mathbb{R}^n \rightarrow \mathbb{R}$ projection is defined, are of the same order of magnitude. That is, for $f_{i,k}(x), f_{j,k}(x)$ of arbitrary parameter indexing on the set k of shared parameter choice, and $i \neq j$ being the distinct choice, then we say $f_{i,k}(x) = \mathcal{O}(f_{j,k}(x))$ of the growth of both function on value ranges. Approximately speaking, the choice of parameter does not dramatically growth to a point of no return, like in later examples of polynomial approximation model. Such special model of multi-variable linear projection of scalar output also reinforces the numeric growth of the numerical encoding expression on the hypothesis model. The behaviour of the MSE-on-parameter, the internal risk, is also explainable using the same notion, such is that for a black-box case approximation, such linear model does not allow for average parameter deviation per approximation to be oversaturated and skewed (for example, the degree of polynomial directly counteract this, by separating expressive complexity by growth degree). How such hypothesis can be evaluated, we will have to make it for the general extension experiment, which we will have to conduct.

Let us define the more general class of linear model, denoted $\mathcal{H}_L$, in the following structure of the linear model. We also note that the previous linear class $\mathcal{H}_p$ only solves for the problem of underparameterized, hidden parameters p' . Let $k \in \mathbb{N}$ denote the number of input vectors, and let $\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_k$ represent the input vectors where each $\mathbf{u}_i \in \mathbb{R}^q$ for $i \in \{1, 2, \dots, k\}$. Here, $q \in \mathbb{N}$ denotes the dimensionality of each input vector, with the convention that $q = 1$ corresponds to scalar inputs (i.e., $u_i \in \mathbb{R}$).

For a given dimension parameter $d \in \mathbb{N}$, we define $\mathcal{H}_L$ as the set of all functions $h : (\mathbb{R}^q)^k \rightarrow \mathbb{R}$ that can be expressed in the unified form:

$$h(\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_k) = \sum_{j=1}^d w_j \cdot \phi_j(\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_k) + b \quad (4.28)$$

where $\mathbf{w} = [w_1, w_2, \dots, w_d]^\top \in \mathbb{R}^d$ is the weight vector, $b \in \mathbb{R}$ is the bias term, $\phi_j : (\mathbb{R}^q)^k \rightarrow \mathbb{R}$ is a feature map for $j \in \{1, 2, \dots, d\}$ that transforms the input vectors into a scalar feature, we consider the following important special cases, distinguished by their choice of feature map ϕ_j . If we do not have such feature mapping of which preserves relative linearity, then we said it is the basic linear class dynamic.

Definition 4.3.2 (Linear class, $q = 1$). Consider class $h : (\mathbb{R}^{d+1})^q \rightarrow \mathbb{R} \in \mathcal{H}_{\text{Lin}}$ of all functionals of the form as Equation 4.28. Then, when $q = 1$ of scalar observables, i.e., $u_i \in \mathbb{R}$ for $i \in \{1, \dots, k\}$, then the function is simply as:

$$h(u_1, u_2, \dots, u_k) = \sum_{j=1}^k w_j \cdot u_j + b \quad (4.29)$$

for $\phi_j(u_1, \dots, u_k) = u_j$ for $j \in \{1, 2, \dots, k\}$, i.e. deactive feature map.

Such class can be extended rigorously, through a varied change in dynamics, specifically, considering the general feature representation on ϕ .

Definition 4.3.3 (General feature, kernel linear model). Consider class $\mathcal{H}_{\text{Lin}}$. For maximum flexibility, allow arbitrary feature maps $\phi : (\mathbb{R}^{d+1})^q \times \mathbb{R}^k$ for $k = \{1, 2, \dots\}$ as of hidden or feature mapping

processor internally. The functional form is then expressed by

$$h(\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_k) = \sum_{j=1}^d w_j \cdot \phi_j(\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_k) + b \quad (4.30)$$

for $j \leq d$ as the iterative form.

Common choices for such kernel mapping includes:

- **Polynomial features:** $\phi_j(\mathbf{u}) = u_{i_1} u_{i_2} \cdots u_{i_m}$ for various index combinations
- **Radial basis functions:** $\phi_j(\mathbf{u}_1, \mathbf{u}_2) = \exp(-\gamma \|\mathbf{u}_1 - \mathbf{u}_2\|^2)$
- **Fourier features:** $\phi_j(\mathbf{u}) = \sin(\boldsymbol{\omega}_j^\top \mathbf{u})$ or $\cos(\boldsymbol{\omega}_j^\top \mathbf{u})$
- **Neural network activations:** $\phi_j(\mathbf{u}) = \sigma(\mathbf{v}_j^\top \mathbf{u} + c_j)$ for activation σ , though we note that this is a very simplistic removal of structural basis class. Rather, it is performing shallow network in linear approximation on kernel system. The reason most of the time we do not use this is because of complexity in multi-kernel distortion.

From such, it basically hints at the generality of, at least being a *linear approximation* on various hypothesis spaces, while not actually being similar. Then, our setting of sparsity being synthetically induced introduces the basin of subclasses. For any $p \in \{1, 2, \dots, d\}$, we can then simply define the sparse subclass $\mathcal{H}_{L,p} \subseteq \mathcal{H}_L$ as the set of all functions $h \in \mathcal{H}_L$ where the weight vector $\mathbf{w} \in \mathbb{R}^d$ has exactly p non-zero entries. Formally, per Lafon and Thomas [2024],

$$\mathcal{H}_{L,p} = \{h \in \mathcal{H}_L : |\{j : w_j \neq 0\}| = p\} \quad (4.31)$$

The subset $S \subseteq \{1, 2, \dots, d\}$ of indices with $|S| = p$ corresponding to non-zero weights is assumed to be fixed throughout the learning process (i.e., the sparsity pattern is predetermined). This constraint yields:

$$h(\mathbf{u}_1, \dots, \mathbf{u}_k) = \sum_{j \in S} w_j \cdot \phi_j(\mathbf{u}_1, \dots, \mathbf{u}_k) + b \quad (4.32)$$

where only p features contribute to the final prediction. This effectively masks the data randomly, simulating the general observation-based information corruptions, without inducing specific bias like *dropout*, though potentially leads to strong perturbation in cases of meaningful information threshold for different function classes.

Baseline investigations

Preliminary model learning then will relies on theorem 4.3.1 for its initial prediction, using least square approximation results. The testing is adversarial and consistent, of which we have $c \in \mathcal{H}_L$ and $h \in \mathcal{H}_{l,p}$ for $p = d$, which can then further be modified. Because of their equal expressive capacity, respect to independent consideration of the parameter space $\mathbf{X}$, we define the complexity, which will be used later on, simply as

$$\mathbb{C}(h) = \frac{1}{\|\bar{\mathbf{x}} - \underline{\mathbf{x}}\|} \left[1 + p + q + \mathbb{E}_{i \in S} (w_i x_i) + b \right], \quad S = \{p_i\} \quad (4.33)$$

Functional prediction on the classical theorem is presented below, for $n = \{1, \dots, 241\}$ samples, and with $p \in [1, 200]$. We set the target degree to the maximal 200, since the theorem is not defined on the overparameterized region.




Figure 4.9: Extension general experimentation, using the prediction made of theorem 4.3.1 for verification. Here, $q = 1$ as the general basis case of verification.

A design of the above class will require, in general, a much less straightforward implementation than the least square method, at least when it comes to circumstances and case fielding. The previous experimental result was done using a straightforward, single-time calculated theorem. To verify such, we need to adopt the non-zero class in a way that is static of the entire training procedure, and so on, for larger and larger q . First, we discuss of the mismatch calculation and thereof taken from the iterative error segmentation. In our cases, we have h and c with *model mismatch*, that is, there might exist certain parameter that h has, but the concept c of its originality does not, and vice versa for the concept model of which the hypothesis does not have, or attain certain parameter of interest. This prompts us to make use of a type of **penalty** for such case. With respect to the mean square error on parameter, from the perspective of the supervisor, we have two ways to implement this, either of the truncate or padding mode. The truncate mode simply calculates parameter expression of those that can be compared. By default, for the learning problem to be consistent of the space, there must be at least one parameter that is not inhibited. Then,

$$\text{MSEP}_{\text{truncate}}(h, c) = \frac{1}{k} \sum_{i=1}^k (p_i - q_i)^2 \quad (4.34)$$

For $K = \min(n, m)$ of the overlap, or can be changed to be the set $k = \{(h_i, c_i)\}$ of which matches each other. Such can be changed to absolute error, if one does not want to inflate the error measure beyond its distance metric with respect to certain axis. The padding mode simply takes into account of all missing or

surplus information and of the optimization phase. By initialization, from the hypothesis side, all surplus parameters are padded with input 1. Hence, the measure is defined as

$$\text{MSEP}_{\text{pad}}(h, c) = \frac{1}{k} \sum_{i \in k} (p_i - q_i)^2 + \lambda \left(\frac{1}{n-k} \sum_{i \notin k} p_i^2 + \frac{1}{m-k} \sum_{i \notin k} q_i^2 \right) \quad (4.35)$$

The total error then is simply the above error, is simply calculated toward all layers possible or any architectural design parameters. This design will be also used in subsequent testing, including polynomial, in which the effectiveness is enhanced significantly.

4.4 Polynomial models

4.4.1 Structure

For mathematical analysis and experimental design, we would like to state the current main structure of a polynomial model. Specifically, in such example, we are concerned of the shallow polynomial neural network, of which is stated as such in [Arjevani et al. \[2025\]](#). Notice that it is defined on the normal Vandermonde basis, where an expression of such polynomial model under Newtonian or Lagrangian basis is not so clear-cut of an analysis.

Definition 4.4.1 (Shallow polynomial neural network). A shallow polynomial is a neural network of one layer that are functions $f_{\mathcal{W}}$, defined such as: $\mathbb{R}^d \rightarrow \mathbb{R}$ of the form

$$f_{\mathcal{W}}(x) = \sum_{i=1}^r \alpha_i (w_i \cdot x)^d, \quad (4.36)$$

$$\mathcal{W} = (\alpha, w_1, \dots, w_k) \in \mathbb{R}^r \times (\mathbb{R}^n)^k,$$

where $d \in \mathbb{N}$.

A modification is then available, in which we simply add $b(\cdot)$ for random initialization process. We note that this introduces certain flexibility to the model initial interpolation curvature.

Definition 4.4.2 (Polynomial neural network, general). The polynomial network based on Model 4.4.1 is modified as

$$f_{\Theta}(x) = \sum_{i=1}^d \theta_i (\mathbf{w}^{\top} x)^i + b, \quad (4.37)$$

where $\Theta = (\{\theta_i\}, \{w_k\}, b)$.

From this, we can state the model within a testing scenario listed as such, to fully capture the structural foundation of the related testing model class. For the notion of complexity, naively, we employ two kinds of complexity measure. First is the usual complexity measure via parameter counts. Second, is the expressive complexity on the basis of the linear transformation space. That is, an expressive complexity that take into account the fact that $\mathcal{O}(wx^n)$ will always be stronger than $\mathcal{O}(wx^{n-1})$ for $x > 1$. All of which can be contained in a singular expression, as seen in definition 4.4.3. However, we would have to analyse what is happening inside the metric itself. The model itself is illustrated in Figure 4.10.

In the equation for expressive complexity, the first term is the normalized linear parameters θ_i fixing and controlling the shift of the polynomial inlining parameter, the second term is the sum of weight-aware

Table 4.1: Controllable parameters and hyperparameters for polynomial class analysis

Information	Notation	Details
Hypothesis model	$h[\mathcal{W}_h, \mathbf{H}_h] \in \mathcal{H}$	Hypothesis model. For simplicity, it contains $\mathcal{W}$ as the configuration parameter set, and $\mathbf{H}$ the fixed hyperparameters
Concept model	$c[\mathcal{W}_c, \mathbf{H}_c] \in \mathcal{C}$	Concept or target model. Similar configuration as the hypothesis model
Complexity measure(s)	$\{\mathcal{C}(\cdot)\}$	Set of complexity measure of choice for any arbitrary model taken in.
Sample count	m	The number of sample available for hypothesis model of c , for $m = S $ dataset
Hypothesis dimension (hyperparameter)	$[d_h]$	An explicit hyperparameter contains the list of all in-testing polynomial degree for the hypothesis model.
Concept dimension (hyperparameter)	d_c	An explicit hyperparameter of the concept's true dimension. Theorized such that $d_c = d_{h,i}$ for arbitrary i will reduce to possible low-error problem.
Data noise	ϵ_S	Dataset-wide noise with relative strength to the value of the sample point. This is to reduce vanishing noise with higher-degree polynomial sample points.
Learning parameters	epoch, lr , bs	Two main parameters for gradient-descent, iterative-based algorithm - the epoch, the learning rate, and the batch size.

Notes: We output to $\{\mathbb{E}[\ell]\}$ as the expected value of the error measure on h, c . There are a few things possible, but under approximation setting, we use MSE-on-parameter and MSE-on-prediction for two different insights: internal representation and external representation.




Figure 4.10: Illustration of polynomial model in definition 4.4.2.

complexity with degree factor, and the degree factor over operating range $\|x_{max} - x_{min}\|$ of the dataset, for the degree complexity. The final term is the last weight-aware complexity for the dominating degree n of the polynomial function. Let us begin with the gauge on coefficient norm, i.e. the term $\|\alpha_d\| + 1$. Here, we set 1 as baseline, and the arbitrary norm on α there coefficient effect. The terms coefficient in this expression matters because, for a function $f : [a, b] \rightarrow \mathbb{R}$, the Lipschitz constant L for polynomial f goes by

$$L_p = \sup_{x \neq y} \frac{|f(x) - f(y)|}{|x - y|} \quad (4.38)$$

$$= \sup_{x \neq y} \frac{|f_\Theta(x) - f_\Theta(y)|}{|x - y|} \quad (4.39)$$

$$\leq \sup_{x \in [a, b]} \|f'_\Theta(x)\|_2 \quad (4.40)$$

$$= \sup_{x \in [a, b]} \left\| \sum_{i=1}^d i \theta_i (\mathbf{w}^\top x)^{i-1} \mathbf{w} \right\|_2 \quad (4.41)$$

$$\leq \|\mathbf{w}\|_2 \sup_{x \in [a, b]} \sum_{i=1}^d i |\theta_i| |\mathbf{w}^\top x|^{i-1} \quad (4.42)$$

$$\leq \|\mathbf{w}\|_2 \sum_{i=1}^d i |\theta_i| \max\{|\mathbf{w}^\top a|, |\mathbf{w}^\top b|\}^{i-1}. \quad (4.43)$$

Where we note that $f_\Theta(x) = \sum_{i=1}^d \theta_i (\mathbf{w}^\top x)^i + b$, and $\|\mathbf{w}\|_2$ is the Euclidean ℓ^2 norm. The magnitude of $f'_\Theta(x)$ depends directly on the coefficient size α_i , and also on how big is x on the domain. All of this mean for large expression and value range on f_Θ , the Lipschitz constant of that model is accordingly large. Defining α -norm that way is a cheap proxy for Lipschitz constant, since it gauges stability, and rate of change of the functional expression itself.

Definition 4.4.3 (Polynomial expressive complexity). For $x > 1$ or $\mathbf{x}_i > 1$ of $\mathbf{x} \in \mathbb{R}^n$ as input. Denote the polynomial class as $\mathcal{P}_n$ of their highest degree n , then we define the total expressive complexity of $p \in \mathcal{P}_n$ as followed.

$$\begin{aligned} \mathbb{C}_E(p) &= \mathbb{E} \left[\|\alpha\|_d + 1 \right. \\ &\quad + \frac{\sum_{i \in S \setminus \{p_n\}} w_i i}{\|p\| - p_n} \\ &\quad \left. + \frac{1}{\|\bar{x} - \underline{x}\|} \left(\frac{\sum_{i \in S} i}{\|p\|} + w_n n \right) \right], \end{aligned} \quad (4.44)$$

where $S = \{p_j\}$, $\|p\|$ denotes the amount of different degree separable terms and $\|x_{max} - x_{min}\|$ is the observable range of the dataset.

4.4.2 Scenario schema

The scenario schema will be as followed. The following set of parameters are required for controlling the testing case, aside from different system measures. For parameter error, we reuse equation from Definition 4.4.3 but with degree aware terms. This is covered in Table 5.1.

We note a few problems with both interpolation interpretation and polynomial regression of higher-degree terms. Firstly, is the intrinsic interpolation error and numerical, software issue error. Higher degree polynomial can be represented as solving for a wide Vandermonde system of n degrees for m sample points. Using monomial basis for interpolation, either for soft interpolation also exhibits instability and numerical garbage within the operations, as indicated by [Shen and Serkh \[2025\]](#). They can also be a potential source of double descent, as the majority of models tested around this range are not coded effectively for handling large data operations as such, and artifacts from such choices might be taken as the indication of double descent in some particular cases.

Aside from it, and aside from the fully-connected case, we have the particular testing scenario that can be expected of the ease of implementation. First, is the mundane implementation that is of gradually increasing degree mismatch. That is, for polynomial of h_p and c_p , one requires only d_h and d_c to figure out the structure as there will exist exactly $d + 1$ with the constant term, terms of which of finite many degrees. The second case is where for p chosen randomly, such that aside from such p entry, then the other degree or nodes of input are simply scattered dependent nodes, linear algebraically. That is, for a uniformly constructed d -degree network, choose at random p particular degree to be fixed, while the other of $d - p$ are assigned randomly of belonging to some $p_i \in \{p\}$ nodes given the random choice on p above. Linearly algebraically, scattered and dependent nodes then mean for example, node 4 might use node 35's degree, and vice versa (if they were both non-fixed). The dependency structure is deliberately randomized rather than organized. For overparameterized case, we simply design the hypothesis by the following line of justification. For a mapping $\mathbb{R}^d \rightarrow \mathbb{R}$, the mapping $\mathbb{R}^{d+n} \rightarrow \mathbb{R}$ for $n = 1$ requires additional information, since the source provision does not contain information to handle additional n dimensional input. Hence, we design it the following. For the hypothesis h , such unit of which does not fit to the original input space, the *internal input space* that h 'hypothesize' contains, for the other n entries, $q \leq d$ per unit input. That is, each input of n overparameterized case receives at random, q inputs can be received, for example, $w_{n1}x_j + w_{n2}x_k$ for k, j specific input numbering. By the simplicity argument, then $q = 1$ for the simplest case. However, this creates a skewed polynomial network, of which informally, the provided structural bias or additional information, is the fact that there exists overparameterization. To prevent this, instead, the final case considers the most general structure — for each unit, each one can take q inputs controlled as hyperparameter of the network, then the set $\{p\}$ of chosen degree influence their degree choice, and is uniform if the sampling process it is not included.

4.4.3 Procedure

The procedure for making the observation space, i.e. the most important one, will be as followed. Specify finite $|\mathbf{X}|_{\mathbb{E}}$ denotation for admissible value region (if multivariate polynomial) of the observables' intended context coverage. Second, initialize **random shutdown** choice of the Oracle with arbitrary structure, in which this case we simply use polynomial component. Generate portion of 1 for the true oracle without shutdown, and introduce a mixture as followed: for k shutdown sequence, more severe than another, assign choice of choosing such data into the final dataset by Poisson distribution on $\lambda^k e^{-\lambda} / (k!)$ of arbitrary $\lambda \in \mathbb{R}_+$. For fixture, we assume $\lambda = 0.2$. Then finally, sample on the generated shutdown corruption, and add intrinsic noise, for $\epsilon_{i \leq m} \sim \mathcal{N}(1, \sigma^2)$ tuneable for intrinsic noise change, and thus finish the initialization of the observation space.

4.5 Graph Neural Network (GNN) - Shi et al. [2024]

To confirm and observe particular insights on the problem, experimentally, we focus more on the results given by the analysis on a particular type of neural network, the graph neural network (GNN). However, preliminary experiments are also needed, for example, with various complexity and function class, though most of them are algorithm based on ERM, which typically can be called derivation of gradient descent.

According to major literature, for example, Shi et al. [2024], double descent is absent usually in GNN, for example, GCN networks. However, in fact, it is relatively unstable, as for certain papers and researches point out its existence and variation of it being ubiquitous to the network, some reports none of such occurrence in their experiments. Thereby, our first strategy would be to encounter this absence, and the way to force it to exhibit itself, if the phenomenon is perceived non-existent. In experiments by Shi et al. [2024], Buschjäger et al. [2020], we know that graph network are inherently sensitive to data configuration. The graph MPNN in Hamilton [2022], for example, depends heavily on the configuration of data to present its neighbourhood aggregation for each layer. However, we would like to first claim the dichotomy of bias-variance holds for it to be presented as the first interpolation point.

We would use, and utilizes a hybrid setting between MPNN, GCN and GAN (attention network), to configure out network in itself. However, heterogeneous network – either consists of only MPNN or GCN layers, is somewhat preferred for ease of analysis, and potency of experimental learning control. Because double descent lies between the concept of test error and training error, either supervised or semi-supervised would be used. According to Shi et al. [2024], semi-supervised setting gives better flexibility and overall 'range' of operation. Other than GNN, certain more abstract, 'vanilla' neural network formalism as specified in the above section treatment will also be conducted, in a supplementary manner.

4.5.1 Methodology

Because of the irregularity in the statement that GNN does not exhibit any phenomena that is related to double descent Shi et al. [2024], we would be investigating this phenomenon the most. This particularly means that a lot of our focus will be to analyse the GNN network by itself.

Quick introduction to graph theory

A **graph** G is a 2-tuple (V, E) where V is the set of **vertices**(or nodes) and E is the set of **edges**. The set $ne[n]$ stands for the neighbour of vertex n , while $co[n]$ is the set of edges that have n as vertex. Edge is often denoted by (u, v) for vertices u and v . We said that (u, v) **joins** u and v , and it can be directed or undirected. A graph is called a **directed graph** if all edges are direct or **undirected graph** if all edges are **undirected**. The **degree** of vertices v , denoted by $d(v)$, is t number of edges connected with v . The graph data can be loosely (not specifically in cases) as followed. Let $\mathcal{G} = (\mathcal{V}, \mathcal{E}, \mathcal{W})$ be a graph, where $\mathcal{V} = \{1, \dots, n\}$ is the set of nodes, $\mathcal{E} = \{1, \dots, M\} \subseteq \mathcal{N} \times \mathcal{N}$ is the set of edges, and $\mathcal{W} : \mathcal{E} \rightarrow \mathbb{R}$ is the edge's weight function. There is also the optional weight $\mathcal{W}' : \mathcal{V} \rightarrow \mathbb{R}$ for each vertex. In this case, it is for certain problem like TVP, where all cities have certain properties for the path. We say that a data sample $\mathbf{x}$ is a graph data, if its entries are related through the graph $\mathcal{G}$.

Our 2-tuple (V, E) and the collection of all such tuple forms a group of **simple graph** (undirected) and **directed graph**, where the only change is that $(u, v) \neq (v, u)$ for any given edge on a graph. From such, a typical setting, if we are to apply a machine learning setting on graph-theoretical problems, can be described as the following

Much information can be packed on a graph. Typically, the graph $\mathcal{G} = (\mathcal{V}, \mathcal{E}, \mathcal{M})$ will contains node

information and edge information. Edge information can be the path-only, or weight-specific, that is, there exists for an edge (u, v) a property $\phi_{uv} \in [0, 1]$ acting like a weight to gauge a walk on the graph. This can also be used for the *distance measure* of a graph, for such graph being represented with values per position, that is, each node/vertex has position, and similarly there exists arbitrary distance $|u, v|$ between u and v . Any space that encode a graph's information in such way is called an **embedding space**. For node information, connections to other nodes are one type of information; however, it also retains its own properties, typically inferred through the *node embedding space* which contains features of a node and the graph itself. There are a few restrictions of such embedding space, such as to define the node embedding space to be of the same shape n , to avoid 'edge cases' in the process.

The reason for choosing an encoding can be justified as followed. Typically, graphs can be classified and categorized into a specific type of data representation, called *non-Euclidean data*. More specifically, there exists a topological space encapsulating the graph, but exists no fundamental measure on such topological space housing it. For example, there exists no notion of a *discrete distance* on a graph, but only the induced notion of *graph distance* by counting the shortest path from one node to another in a specific graph. Because machine learning works on the assumption that the space of the system gives *measurable space*, a natural response is to encode the system into an encoding space of arbitrary meaning. One, for example, simple encoding is the degree map, where nodes are mapped into an $n \times n$ matrix of nodes and valued by the number of neighbour they have.

This type of encoding then creates, for each graph, a versatiel number of arbitrary properties space they can have, stacked on each other. Typically, for a hypothesis h to work on a graph $\mathcal{G}$, then it will have two main targets that act directly on the graph, as illustrated in Illustration 4.11. There are many ways




Figure 4.11: Illustrative decomposition of the graph on its simplexes edges and vertices: Each of the decomposed space $\mathcal{V}$ and $\mathcal{E}$ can be encoded separately to their own ordeal, for example, of the edge connection through incident matrix.

or aspect of a graph to be 'extracted' from the graph features itself. Classically, this is done using feature engineering (Hamilton [2022]) and hence the learner is also configured to utilize such features. Nowadays, it is mostly abstracted into an arbitrary representation space of the graph encoding, used in neural network architectures (Oono and Suzuki [2020], Lopushanskyy and Shi [2024], Scarselli et al. [2009], Hamilton [2022]). The only different between the classical encoding and the deep learning encoding is on the arbitrary properties of deep learning encoder, of which the transformed feature might not be well-defined feature engineering as the classical case.

This prompts two types of main learning targets: either to have the hypothesis h to learn the edge connections and properties, or to learn the vertices/node. They are, however, not separated, as to learn about which connections $e \in \mathcal{E}$ to make between two nodes requires information about the general graph's node itself, to predict which node would be connected under obscured information (removing edges). So, there are:

- **Edge reconstruction:** predict if there exists an edge at an arbitrary segment of the graph, for example, between $u, v \in \mathcal{V}$.
- **Edge classification:** An edge $e \in \mathcal{E}$ is endowed a label, and we will then try to classify it based on the actions, or properties it has on the graph.
- **Node classification:** Classify a node $v \in \mathcal{V}$ of certain class of labels.
- **Node regression:** Estimate values of certain nodes. In turns, it might also be used for ranking certain nodes based of specific criteria.

It is helpful to see that a graph can also be classified into two more main archetypes, based off their actions: either **static**, where the node and edges remains the same, or **dynamic**, where deletion of nodes and edges and vice versa are taken into account. The problem revolving a dynamically changing graph (as more prominent in real-life cases) is resolved using temporal methods for graph learning (Rossi et al. [2020]), however, it is much more difficult to analyse, and hence we would restrict ourselves in the case of static graph analysis.

We then state our problems in regard.

Definition 4.5.1 (Graph learning problem). Given a graph $\mathcal{G}(\mathcal{V}, \mathcal{E})$, for any type of graph-theoretical sub-type (multigraph, digraph, etc.) by m constraints. There then exist two topologies: the **graph-theoretical topology** of connections and nodes appearances, and the **encoded topology** of various measured properties for all $v \in \mathcal{V}, e \in \mathcal{E}$, denoted by (ϕ_G, ϕ_E) . A hypothesis $h \in \mathcal{H}$ is then tasked to learn its own encoding, or loosely speaking **interpretation** of the hypothesis about $\mathcal{G}$ based on the topological space given by $\mathcal{G}$, and use it to target or solve problems related to both ϕ_G and ϕ_E .

From this, we might derive various problems setting as it is, and considering the set $\{\phi\}$ of encoding intrinsic to the graph.

Graph Neural Network

We target the graph neural network structure (GNN), specifically a neural network implementation for solving graph data problems. A more detailed description can follow from Hamilton [2022], Scarselli et al. [2009]. Typically, graph neural network (Scarselli et al. [2009], Veličković [2023], Tanis et al. [2024], Lopushanskyy and Shi [2024]) follows the instruction flow of the *encoder-decoder* architecture. A GNN is formulated and structured by analysing a graph data system by conceptually apply a *neighbour-dependent* neural input arrangement on top of the data. That is, in principle, it reserves on its own a particular extractor and modifiers on the graph interpretation itself, through its layers.

Structurally, the description of a GNN is defined by the overall *message-passing neural network* (MPNN), defined by (??):

$$\mathbf{x}_i^{(k)} = \gamma^{(k)} \left(\mathbf{x}_i^{(k-1)}, \bigoplus_{j \in \mathcal{N}(i)} \phi^{(k)} \left(\mathbf{x}_i^{(k-1)}, \mathbf{x}_j^{(k-1)}, \mathbf{e}_{j,i} \right) \right), \quad (4.45)$$

for $\bigoplus$ the differentiable, permutation invariant function, usually called the aggregator, $\mathbf{x}_i^{(k-1)} \in \mathbb{R}^F$ the node features of node i in passing layer $k - 1$, $\mathbf{e}_{j,i} \in \mathbb{R}^D$ the optional edge features from node j to node i . Additionally, γ denotes the nonlinearity differentiable function (usually ReLU, or sigmoid) ϕ denote the MLP layer accompanied They are often called the **updater** and the **messenger**, respectively. A

simplified example of this structure in Scarselli et al. [2009], Hamilton [2022] as

$$\mathbf{x}_i^{(k)} = \sigma \left(\mathbf{W}_{\text{self}}^{(k)} \mathbf{x}_i^{(k-1)} + \mathbf{W}_{\text{neigh}}^{(k)} \sum_{v \in \mathcal{N}(i)} \mathbf{x}_v^{(k-1)} + \mathbf{b}^{(k)} \right) \quad (4.46)$$

where $\mathbf{W}_{\text{self}}^{(k)}$, $\mathbf{W}_{\text{neigh}}^{(k)}$ are trainable parameter matrices, and σ denotes an elementwise nonlinearity, and an optional bias term $\mathbf{b}^{(k)}$. Different flavours of the differentiable aggregator create the *graph convolutional networks* (GCNs), defined by:

$$\mathbf{x}_i^{(k)} = \sigma \left(\mathbf{W}^{(k)} \sum_{v \in \mathcal{N}(i) \cup \{i\}} \frac{\mathbf{x}_v}{\sqrt{|\mathcal{N}(i)| |\mathcal{N}(v)|}} \right) \quad (4.47)$$

A somewhat popular approach is to apply attentional layer and weights to facilitate neighbourhood attention, which is called graph attention network.




Figure 4.12: A conceptual illustration on the running flow of an n -layer GNN on particular structure of interest. Note that the data section itself has particular embedding structure on its own.

As noted by Hamilton [2022], node features are often required to be inputted into the GNN. Usually, this is of the form $\mathbf{x}_u$ for all $u \in \mathcal{V}$. The same can be said for edge in specific tasks required of them. However, if there are no sufficient node representation possible in the first glance, then we can use node/edge statistics as classical techniques.

A GNN is, when fitting into the framework of learning system and modeling, is a feature mask generation and adaptation mechanism of the modelling pipeline. What this means is that GNN *assumes* every data point has its own feature-embedded space, and according relationships. By this, we further mean that it creates an embedding space of the aggregator, or the GNN-embedding space within such. While input embedding space captures and represents the best possible interpretation of the input space, GNN-embedding space captures what is required of the aggregation properties that the GNN layers specify. Hence, we can assume relative fallacy in such embedded space. Those n -embedded space of data aggregated at the n th final layer, would then be fed into another straightforward - task-based and structured network, such as a pass to a single-layer sigmoid network, Hamming network, or generally an FCN.




Figure 4.13: Encoding space and the change of encoding inbetween a static graph (or snapshot-wise) GNN. The encoding space is warped toward arbitrary notions inside a GNN, to the point that after certain layers of processing, the encoding outputs different from expected encoding space (native to the model, not to the designer), though there might be universal notions preserved, like the degree of node represented in a different way.

4.5.2 Double descent and GNN

As of date, there has been no reports on double descent on GNN. This can be taken accounted of the complex and often difficult analysis for graph and the neural network adopted for such task. Because of this, we will have to attempt to force double descent appears instead of simply observing it. To do this, we need to specify both the task for evaluation for the error measure, and model complexity for the other axis. For such, we would often reserve to node-level or edge-level task, and not graph-level tasks.

Aside from layer-wised complexity, GNN also has two types of complexity: the size of the graph itself, and the size of the encoding in specification. Thence, the complexity would then at least be the tuple (m, n, l) for m -size of the original graph, n -depth of the GNN layers, and l layers of GNN. If n varies between layers, then we will have the index set instead. However, for simplicity, we would likely want to have constant n throughout all layer.

4.6 Support Vector Machine (SVM)

Chapter 5. Interpretation

There are a lot of problems regarding how we treat the problem, how we see it, and how we actively analyse it. Most of them lies in the space of statistical learning theory, since bias-variance on itself has variations and justification that is seemingly very natural in statistical theory, like the other version name *approximation-estimation error tradeoff*. All of them contribute somewhat to the current outlook, including operational factors and hidden assumptions underneath the framework of statistical learning, in which might only hold in special regions and not of a wider landscape, and non-applicable to richer constructions of semantic. Thus, we are prompted to look into both the **structural** problems, i.e. specifically on the reliance of bias-variance connection toward risk, for example, the fact that it is not general of a decomposition; the oblique definition of model structures and gaps between different discipline of model building; different operational definition of risks and hypothesis space, and the *operational problems*, of different modelling setting; different treatment in consideration of the learning scenarios; different conditions and mathematical structures enabled while also overclaim patterns seen within such space, and so on so forth. While we do not claim to solve or identify all of them, the baseline of this section is to **note down** such disparity. Following up such section would be the interpretation of double descent in its form, and inconclusively of what is the nature of it.

5.1 Structural and operational problems

There are several structural problems that require attention in this section. That said, it can be observed rather easily by our incapacity to express or explain recent structural failures or phenomena that appear of large system. Despite its prevalence in particular research interest, double descent itself has no definite statement to describe its particular behaviour. This is also true for a list of other phenomena, such as *grokking* (Power et al. [2022a]), *emergence behaviours* (Wei et al. [2022b]), typically discussed in large language models, which is understandable of large-scale specific ruling system), *implicit bias* (Soudry et al. [2024]), and *neural scaling law* (Kaplan et al. [2020], Wei et al. [2022b]). There is also *catastrophic forgetting* (van de Ven et al. [2024]) which is still fairly native of large systems; nevertheless, those phenomena will be of target later on (or discussed in later chapters). The fact remains that we only *partially understand* these particular results with conjectures and theories to support the facilitation of new insights, new analytical lens, and nothing much. Without a clear definition of what is even double descent, sometimes, we cannot proceed, lest to say even solving it. This is a problematic issue, since the current theory is so messy and not rigid that many even turns to exotic choices, like topological informatics learning theory, *Homotopy Theoretic and Categorical Models* (Manin and Marcolli [2024]), and several more. Without clear definitions, it remains ambiguous whether double descent constitutes a single phenomenon, a family of related behaviours, or an artefact of particular modelling and evaluation choices.

Thus, prompt several distinction classes that are needed to take into consideration. One is toward

the class between **small and large** models and what do they mean structurally or semantically; one is about **model description and richness** between them; another one is between **different metrics and evaluating spaces**, between different interface and *operational procedure*, i.e. between algorithm-agnostic and algorithm-dependent results or observations. They would be detrimental of the analysis, since the concept of bias-variance itself hedges on such gap between theoretical treatment and empirical observations. Operational problems are rather easier to spot aside from operational procedure. For example, the concept of bias and variance can be computed differently in different algorithmic settings, the risk profile can also be different, inherent design might introduce structural reduction to finite cases of the evaluation landscape, and different state of deformations of uncontrolled modules in randomized setting might deviate results or damp the observation down in noise. Such will have to be taken into consideration before moving into further sections focusing on experiments.

5.1.1 Bias-variance and the definitive form

Bias-variance tradeoff is not new in concept, their definitions on computation has been existing before the formal foundation of modern learning theory. However, it is not trivial on how it became the statement of overreach, much to the disparity observed between classical and deep neural network model. An analysis might be sufficed to analyse such disparity in treatment of the general working space. Mainly, models in classical settings are considered to be of **parameterized models**, and also the non-parameterized version also. With such in mind of the classical setting for which probing object is a structurally unspecified relation $x \mapsto y$, we can define the following definition as parameterization using $\Psi_\Gamma : \Theta \times X \rightarrow Y$ such that for functions $f \in \mathcal{F}_\Gamma$, $f(x) \approx \Psi_\Gamma(\theta, x)$.

Definition 5.1.1 (Parameterization). On arbitrary function space $\mathcal{F} : \{X\} \rightarrow \{Y\}$, of which a system admits its operation in such form, let $\mathcal{F}_\Gamma \subseteq Y^X$ denote a class of functions admitting representation under a fixed mathematical semantic Γ (e.g. algebraic, geometric, compositional, or dynamical semantics). Then such system is transformed to a similar **parameterized system** if and only if it admits the factorization form, or if there exist a parameter space Θ and a family of operators $\Psi_\Gamma : \Theta \times X \rightarrow Y$, where Γ is the mathematical semantic used, θ is the above-level replacement of direct relational connection into variable control of θ , and variations in y being controlled by θ , such that for functions $f \in \mathcal{F}_\Gamma$, $f(x) \approx \Psi_\Gamma(\theta, x)$.

Thus, it is in-spirit, a definition of parameterization of model that we can use. Most definition of parameterization differs from this form, not inherently a disadvantage, to presume independent choice of selecting a family of parameterization; undermine the structural form that are explicit in Γ ; using parameter as controllers that can be tuned of optimization, and often merge such description together with loss and training. The loss of Γ is rather troublesome, but they are often mediated in a weak form using similarity of elementary compositions, i.e. most structural forms in ML models are often elementarily composable, though there exists transcendental expressions of which can be formed during process and not inherent of the syntax itself. Concretely, Γ is the hypothesis about *how relations are connected of each other*, i.e. the grammar of admissible operators. When Γ is absent or mismatched, optimization is free to exploit the structural gaps by directly specify θ -only optimization on non-structured parameters, and can indeed imperfectly gauge the underlying hypothesis model structure itself. Two things can happen in this case: either optimization reduces empirical loss by fitting idiosyncrasies of the training set (overfitting to structure outside the true operator family or admissible approximation of the family to another external family); or optimization selects arbitrary parameter instantiations along many uninformative directions, producing unstable solutions with high variances. Those are potentially two failure mode on wrong operators and worse variance, and thus also correlate to the arbitrary notion of generalization, where this is particularly troublesome since some form of optimization setting assumes Γ_c of the concept $c \in \mathcal{C}$.

From such, we state the frame in which bias-variance is calculated of, with construction preconditioning specified by construction 5.1.1.1.

Construction 5.1.1.1 (Bias-variance structure preconditioning). We use the setup of a model M structured into $(\theta, I, \{f\})$, where $\{f\}$ is the set of descriptions using θ of the internal model's components, and I as the external receiver of connector. In effect, a model is structurally the set of all $\{I, (\theta, \{f\})\}$ with respect to the component form θ , and variation of the external connector I , often can be considered input, and if θ is reduced to variable, we call it a **parameterized model** that creates implicit linkage between I and its intended transformation from I , using θ . Furthermore:

1. Input space I and output space O , typically referred to as X, Y , measurable space taking values from the real field $\mathbb{R}$. Typical set of all feasible values of input and output forms can be specified of $(\mathbb{R}^d, \mathbb{R}^k)$ or for output to be a small embedded subset $Y = \{y_1, \dots, y_n\}$ being finite and discrete on $\mathbb{R}$.
2. Euclidean metric $\|\cdot\|$ is used for all real-valued components.
3. System relation between X, Y is the **function space** Y^X or a subset $\mathcal{F}_\Gamma \subset Y^X$, where Γ is the chosen semantic or representation class. Such system uses also the **parameter space** $\Theta \subseteq \mathbb{R}^p$ for finite p .
4. Hypothesis model, i.e. the learnable model admits the **parameterized form** if able, as Definition 5.1.1, or $\Psi_\Gamma : \Theta \times X \rightarrow Y$. A model instance is $\hat{f}_\theta(x) = \Psi_\Gamma(\theta, x)$. We denote all models as $h \in \mathcal{H}$ for $\mathcal{H}$ the class of all specific Γ and intended value space on $\theta_{\mathcal{H}}$.
5. We assume there exists an **objective oracle**, denoted either c or EX , of which belongs to a class $\mathcal{C}$ of all oracles fitting certain description (θ_c, Γ_c) . We assume this objective oracle is a part of the general oracle class $\mathcal{O} \supset \mathcal{C}$. Furthermore, we assume **knowledge** of this oracle is presented only by fixed instantiation of it, giving non-factorable description of linkage in the form of *dataset*, $D = \{(x_i, y_i)\}_{i=1}^n$ being finite and drawn from observations or probing of c . One more assumptions that can be removed: c exists and knowable of existence; in case this does not happen, we say the setting is *oracle-black boxed*. In case partial information are possible, we say that the setting is semi-transparent. The dataset cutoff implicitly creates $c' \in \mathcal{C}'$ of a proxy concept class, such that lowest condition remains $\mathcal{C} \cap_{\Gamma, \theta} \mathcal{C}' \neq \emptyset$, and so on of the inclusion structure.
6. We apply a **statistical interpretation** on the dataset of observables D to capture the observations' occurrence patterns replacing the internal Γ exposition, into probabilistic distribution P on $X \times Y$ of variations. Simplest assumption can be determined of being i.i.d. or identically and independently sampled, and thus observation are further assumed to come from instantiation of the **sampling action**.
7. Loss function $\ell : Y \times Y_c \rightarrow \mathbb{R}_{\geq 0}$ is defined on comparison space between h and c . On Euclidean space, use L_2 norm or square loss as measure. The semantic can change, thus we say the loss system contains $\{\ell_i\}$ of variational applied form, such that $\{\ell_i\}_{i=1}^3$ being $\ell_3(\ell_2(\ell_1(\cdot, \cdot)))$. Define **population (expected)** error named as risk, taking expected value over the entire space.
8. Operations making use of ℓ defines the **objective** for finding suitable h such that to reduce $\ell(y, y')$ of their representative observables.

One of the main question that should be asked of the modern landscape is that, would modern models even **admit bias-variance** in any reasonable way? We might as well also analyse in such fact, firstly with admission on computational complexity and decidability problem, and how bias-variance forces structure this way. As we have already seen of the setup, the construction itself is not neutral. Indeed, it asks

for many conditions simultaneously, further than the above set conditioning. Indeed, one might also extrapolate and see that under general setting, the bias itself is not total, in the sense that $\mathcal{B}(h, c)$ and $\mathcal{V}(h, c)$ are never enough or might measure overcapacitated domains, i.e. where two modes happen,

$$\mathcal{B}(h, c) + \mathcal{V}(h, c) < \{R\}, \quad \mathcal{B}(h, c) + \mathcal{V}(h, c) > \{R\} \quad (5.1)$$

of the bias-variance combination. Often, theories resolve this by collapsing it into ϵ , thus delegating it to noise or rather irreducible error. However, such is rather an unwarranted move, and often time it can also indicate **bias-variance artefact** instead than delegation to process-level fluctuation or inconsistency. The bias-variance notion itself is brittle, especially when its definition is explicitly defined in classical way through ℓ_2 norm class, and of which to be explored of how the divergence pattern happen to different settings, of which to violate the above construction. However, before again, moving on to identifying the structural justification and problem of bias-variance, we also would like to present another construction. This time, how *bias-variance is computed* in real system.

5.1.2 Computing bias and variance

The notion of bias is direct – that is, it does not depend that much on the "way of segregation" on which we ought to "test out generalization. Rather, it can be defined per specification, a measure of the model as a whole, throughout an algorithm operation.

In a loose fashion, the **bias** of any estimator/hypothesis model measures the **central tendency** of such model, to the true function of c . Of such, we have the first definition:

Definition 5.1.2 (Bias, I). Given a model $M[h, S]$, where $h \in \mathcal{H}$ and S is the associated data, the **bias** of h to the true concept $c \supset \{S\}$ is defined as the measure of estimation of the central tendency of the hypothesis to the true concept:

$$\text{Bias}(h_S) = \delta_{x \sim \mathcal{P}} \{ \mathbb{E}[h_S(x)], f(x) \} \quad (5.2)$$

where $\delta(\cdot, \cdot)$ is the difference measure, of certain measure space associated with the hypothesis and concept.

From such definition, we can change it, for example, to the more typical regression task, then

$$\text{Bias}(h_S) = \mathbb{E}_{x \sim \mathcal{P}} [h_S(x)] - f(x) \quad (5.3)$$

Notice that we are comparing the total function, that is perhaps every actionable data point, to the hypothesis. So, in fact, a better formulation given the practicality, under regression setting (this of course can be generalized later) should be:

$$\text{Bias}[h_S(x)] = (m^{-1}) \sum_{j=1, x_j \sim \mathcal{P}}^m \left\{ \left[(n^{-1}) \sum_{i=1}^n h_{S,i}(x_j) \right] - f(x_j) \right\} \quad (5.4)$$

where m, n is the total amount of sample point available for experiments (if we count letting the model run and operational as testing and experimenting), and the total times that sample is repeatedly experimented upon (or the total runs within the same sample set), respectively. There, the association index also lies accordingly, for example, x_j being the j -th points in the dataset, coresponding to "example counts", simply speak; while $h_{S,i}$ means the i -th iteration, of the sample S associated.

Within such formulation, we have the iterative definition of bias:

Definition 5.1.3 (Bias, Iterative). Suppose a model with the form $M(h, S)$, where $S \subset \mathbb{R}^{n+m}$, and a **logging set** containing previous iteration of the model in the procedure $\mathcal{T}(M, L)$, of such that $S \subset L$ and $L = \mathbb{R}^{m+p}$, where m is the index of the number of example, and p is the index of the number of ensemble (for each iteration in the past). Then, the bias at the t -th iteration is (assuming $p = t - 1$) is defined as

$$\text{Bias}[h_S(x), t] = (m^{-1}) \sum_{j=1, x_j \sim \mathcal{P}}^{p=t-1} \left\{ \left[(n^{-1}) \sum_{i=1}^m h_{S,i}(x_j) \right] - f(x_j) \right\} \quad (5.5)$$

In the more considerate approach, we would want to add non-negativity as a normal metric, hence we can square the term out as a way of "positive scaling", and then normalize it back. Coincidentally, this way of normalizing and treatment sounds, and feel similar to the notion of **standard deviation** in normal statistical analysis. This is expressible as

$$\text{Bias}[h_S(x), t] = (m^{-1}) \sqrt{\sum_{j=1, x_j \sim \mathcal{P}}^{p=t-1} \left\{ \left\| \left[(n^{-1}) \sum_{i=1}^m h_{S,i}(x_j) \right] - f(x_j) \right\|_2 \right\}^2} \quad (5.6)$$

where $\|\cdot\|_2$ is the ℓ_2 -norm, or the Euclidean standard norm that is normal distance on this metric. So, in one way or another, we can think of bias as the crude effective standard deviation of machine learning model, which aligns with what it was intended for statistical bias.

Overall, such definition enables us to inspect such bias with reasonable iterative behaviour. We should also note that, this can be generally generalized for the case of overfitting and underfitting. First, however, notice that in case of such fitness is observed, we have:

$$\text{Bias}[h(\mathcal{D})] = \frac{\text{Bias}[h(\mathcal{D}_{tr})] + \text{Bias}[h(\mathcal{D}_{te})]}{2} \quad (5.7)$$

This applies, generally, to others partitioning. However, the total score should also take into account of the **partition size** for each partitioning, especially in the case of uneven corss-validation testing, of which then such is prominent.

Definition 5.1.4 (Overfitting, Bias-term). Given the partition $\mathcal{D} \rightarrow \{\mathcal{D}_{tr}, \mathcal{D}_{te}\}$, then the model $M(h, \mathcal{D})$ is said to overfit if:

$$\lim_{|\mathcal{D}| \rightarrow \mathcal{D}_{max}} \text{ret}(M_{tr}, M_{te}) = \lim_{|\mathcal{D}| \rightarrow \mathcal{D}_{max}} \frac{\text{Bias}[h_{\mathcal{D}_{te}}(x)]}{\text{Bias}[h_{\mathcal{D}_{tr}}(x)]} = 0 \quad (5.8)$$

or, given an assumption:

$$\text{ret}(M_{tr}, M_{te}) \in \mathcal{O}(x^2) \quad (5.9)$$

In similar fashion, underfitting can also be defined

Definition 5.1.5 (Underfitting, Bias-term). Given the partition $\mathcal{D} \rightarrow \{\mathcal{D}_{tr}, \mathcal{D}_{te}\}$, then the model $M(h, \mathcal{D})$ is said to underfit if:

$$\lim_{|\mathcal{D}| \rightarrow \mathcal{D}_{max}} \text{ret}(M_{tr}, M_{te}) = \lim_{|\mathcal{D}| \rightarrow \mathcal{D}_{max}} \frac{\text{Bias}[h_{\mathcal{D}_{tr}}(x)]}{\text{Bias}[h_{\mathcal{D}_{te}}(x)]} = 0 \quad (5.10)$$

or, given an assumption:

$$\text{ret}(M_{te}, M_{tr}) \in \mathcal{O}(x^2) \quad (5.11)$$

by just changing the fractional term. As such, this *should be*, enough for our definition as of now. Note that, however, that bias term is really close to accuracy, in most sense. So it puts us at quite a position - how to analyse and realize the position of variance in such picture.

Variance

In a somewhat acceptable and identical manner, we can still, loosely define variance. But the spirit of the variance term is simple. Informally, it is a measure of fluctuation of a learner around its central tendency (again, expectation value), where, the fluctuations result from different sampling of the training set [B. Neal, 2019]. So how much of this is true? We first go for the formal definition of such:

Definition 5.1.6 (Variance, I). Given a model $M[h, S]$, where $h \in \mathcal{H}$ and S is the associated data, the **variance** of h to the true concept $c \supset \{S\}$ is defined as the measure of **fluctuations** of the hypothesis (learner) around the central tendency to the true concept:

$$\text{Var}(h_S) = \mathbb{E}_{x \sim \mathcal{P}} [(h_S(x) - \mathbb{E}[h_S(x)])^2] \quad (5.12)$$

Origin and usage of variance also comes from statistical theory, and was imported into machine much later because of a practical merge for regression analysis and parameterized optimization models. As such, the usage of bias and variance is practically different of such, and also, the general kind of *distinction* between two case. Specifically, variance asks of "normally, and now in learning theory - which **estimator** was used?", and "How many active estimator (and their classes) are there?". The deviation or sufficient difference between many realizations of the estimator is then used, and per each different instantiation of such, the variance is different for any one hypothesis compare to the average. For an algorithm on M -range hypothesis instances, thence the variance is not static, but rather a trajectory of both the average in comparison of the variance term, and the variance of the h_i hypothesis at configuration i in testing. Variance by then, operates on one higher level than bias, of which is per-instance evaluation threshold. In similar fashion to the bias notion, we can also define the variance term in a more iterative manner, of such to fits the task of monitoring such terms, with increment model iterative time.

First, notice that term now includes

$$(h_S(x) - \mathbb{E}[h_S(x)])^2$$

of which the moving term would be the second term concerned of the mean average of the hypothesis. In fact, we can move variance to a more generalized definition, that is,

Definition 5.1.7 (Variance, II). Given a model $M[h, S]$, where $h \in \mathcal{H}$ and S is the associated data, the **variance** of h to the true concept $c \supset \{S\}$ is defined as the measure of **fluctuations** of the hypothesis (learner) around the central tendency to the true concept:

$$\text{Var}(h_S) = \mathbb{E}_{x \sim \mathcal{P}} d_M(h_S, \langle h_S \rangle) \quad (5.13)$$

where the notation just shift to $\langle h_S \rangle = \mathbb{E}(h_S)$. The iterative form then should be concerned of the main decomposition component, hence the equating term should be quite different.

Definition 5.1.8 (Variance, Iterative, I). Suppose a model with the form $M(h, S)$, where $S \subset \mathbb{R}^{n+m}$, and a **logging set** containing previous iteration of the model in the procedure $\mathcal{T}(M, L)$, of such that $S \subset L$ and $L = \mathbb{R}^{m+p}$, where m is the index of the number of example, and p is the index of the number of

ensemble (for each iteration in the past). Then, the **variance** at the t -th iteration is (assuming $p = t - 1$) is defined as

$$\text{Var}(h_S, t) = \frac{1}{n} \sum_{j=1}^n \left[h_S(x_j) - \frac{1}{p} \left(\sum_{i=1}^p h_S(x_{i,j}) \right) \right]^2 \quad (5.14)$$

For $x_{i,j}$ is the i th iteration that the model is applied on such sample point, and j th is the index position of that sample point in the dataset.

The connection between variance and over/underfitting is not too apparent at first, but, let's us make an observation, because somehow, variance and bias still form the tradeoff termed *bias-variance tradeoff*, which indicate clearly that certain relation of the form must be included. .

bias-variance
tradeoff

On a side note, the variance term, again, can be combined, such as

$$\text{Var}(h(\mathcal{D})) = \frac{\text{Var}(h(\mathcal{D}_{tr})) + \text{Var}(h(\mathcal{D}_{te}))}{2} \quad (5.15)$$

for the usual case of two partitions.

5.1.3 The definitions of double descent


Double descent is often considered the direct opposition to the viewpoint of bias-variance control, and often not more. More fundamentally, the term **double descent** does not currently admit a formal definition beyond an empirical characterization of observed risk behaviour. While several works have suggested interpreting double descent as an instance of emergent behaviour, such interpretations necessarily remain conjectural, as emergence itself lacks a precise and operational definition within contemporary learning theory. As a result, appeals to emergence do not resolve the underlying indeterminacy of the phenomenon, but rather shift it to a different conceptual level.

Empirically, double descent has been reported in a variety of settings, while being absent in many others, even under closely related modelling and data conditions. This variability complicates attempts to treat double descent as a unified phenomenon and raises whether it constitutes a single mechanism, a collection of regime-dependent effects, or an artefact of specific modelling and evaluation choices. In the absence of necessary and sufficient conditions for its occurrence, theoretical accounts are forced to accommodate mutually incompatible behaviours, thereby undermining their explanatory specificity. Consequently, existing frameworks face a dual difficulty: they neither reconcile naturally with classical bias-variance analyses, nor align with the assumptions underlying earlier generalization results, such as uniform convergence over well-behaved hypothesis classes. Without a precise characterization of the object under study, it remains unclear how such theories can be systematically compared, falsified, or extended.

If we are to reframe double descent, then it would be rather better characterized if we to treat double descent not as a singular phenomenon, but rather a corpus of such, all of which containing different notions and observations than said mention, depends on various moving mechanism of the dynamical system in operations. There are at least, in records, three different flavours of double descent, starting from Belkin et al. [2019a] to Nakkiran et al. [2019], and Hastie et al. [2019]. Namely, they are *model-wise double descent*, *sample-wise double descent*, and *epoch-wise double descent*. By the ordering of the hyperparameter noted in the introductory chapter, we see that model-wise and sample-wise are of system hyperparameter, while epoch-wise case lies comfortably within the setup of the system, thus specific hyperparameter class. This will affect how factors and conditions can be taken into account, more than simply enough.

Model-wise double descent, Belkin et al. [2019a]

In Belkin et al. [2019a], when it was plotted of risk versus model capacity, which is then defined in different ways, for example, number of parameters, neural network width and depth, double descent occurs at the region where classical bias–variance expects monotone U-shape. Modern models instead, by results, shows a peak, which is then coined **interpolation threshold**, and then decrease again converging to an error $\gamma_{Te} \geq \gamma_{Tr}$ of the absolute convergence between the training procedure and the testing procedure. The term interpolation threshold under certain perspective, specifically statistical physics, is also coined as the model **phase transition**, of which typically reflects the dynamic mode of the model in-between such configuration of the system itself. Empirical evidence can be shown, from their own collection of models and the appearance of double descent in ResNet for CIFAR-10 dataset, and so on.



(a) Double descent risk curve for a fully connected neural network on MNIST. Training and test risks of a network with a single layer of H hidden units, learned on a subset of MNIST ($n = 4 \cdot 10^3$, $d = 784$, $K = 10$ classes). The number of parameters is $(d + 1) \cdot H + (H + 1) \cdot K$. The interpolation threshold (black dotted line) is observed at $n \cdot K$. Taken from Belkin et al. [2019a].

(b) Double descent risk curve for random forests on MNIST. The double descent risk curve is observed for random forests with increasing model complexity trained on a subset of MNIST ($n = 10^4$, 10 classes). Model complexity is controlled by the number of trees N_{tree} and the maximum number of leaves allowed for each tree $N_{\text{leaf}}^{\text{max}}$.

Figure 5.1: Double descent risk curves on MNIST for different model classes. Comparison of the phenomenon in (left) a fully connected neural network and (right) a random forest as model complexity increases. We use results from Belkin et al. [2019a], however, extension is rather natural.

Of itself, the problem with model-wise double descent is rather easy to see. Phenomenologically, complexity itself is used both in the practice of constructing models, in case of having the hyperparameters on defining the number of layers, the amount of weights, cumulative active number of available tuneable parameters, or total parameters taking altogether the hard-coded feature engineering. As such,

the other side contains notion of model complexity as per measure against, for example, VC-complexity or Rademacher complexities. There exists not a clean transformational analogue between structural complexity of the model, its dynamical complexity from the configuration itself, and the analytical measure thereof. Usually, we can use them approximately similar, such is for neural network being interpreted as being (m, n) -layers of n widths network, but otherwise such does not cleanly can represent each other. As such, the hedging critical point can be considered the point of **model complexity**, and from such we would be able to categorize them into various different notions, each perhaps will have different interpolation threshold, or exists none at all. Such is an open question by itself, and of which might also even comes to d' Ascoli et al. [2020] triple descent appearance.

The error rate also deserves another treatment. Under the same basis of the loss-function on natural Euclidean metric $\mathbb{R}$, there exists many distance metrics that can be taken on the loss space itself. For example, the shift into absolute error, and so on, might present different effects altogether, even if such interpolation threshold might be relatively similar. Whether certain type of error function actively hides double descent or its landscape is flat-forced or so, is often not known as the landscape itself is *second-order* in terms of determinacy, there exists not quite coherent way to see if such prediction on landscape mapping can be reached. That is, the observable risk curve is not a direct property of the model, but rather instead as we can speak, projection of a higher-dimensional optimization landscape.

Sample-wise double descent, Hastie et al. [2019], Nakkiran [2019a]

Epoch-wise double descent, Nakkiran et al. [2019]

5.1.4 Gradient descent and metric mismatch

Gradient descent is an effective approach to realize the case of **empirical risk minimization** into the formation of practical encodable algorithm, within certain structural constraints, or rather the reverse is true. For ERM to be applied, we require certain spaces of admissible structure:

5.2 Complexity notion and its pluralism

Classically, we can see that the problem of double descent is not inherently a problem that can be explained only by considering model complexity and the error rate. As such, it requires a framework that consider multiple factors together, and interpreting different hypothesis. This, of such, obviously includes the notion of complexity, which hinges as the higher-level double descent concept of itself. Nevertheless, classical theory does indeed point to certain factors or aspect that should be of concern. For example, we have seen plenty of analysis that indicates that *sample complexity* might be the biggest factor, considering the behaviour is intrinsically operable and perhaps observable in almost any model, except for certain model that is specialized of the data structure itself. Even by then, the error still can be obtained given large enough model disparity to the concept class.

Assume the working space is the real field $\mathbb{R}^n$. Let us recall a model $h \in \mathcal{H}$ of certain hypothesis class $\mathcal{H}$. Then, the **complexity** refers to the complexity of the hypothesis h itself; there is also the notion of class maximal complexity for $\mathcal{H}$. Complexity in this case refers to the mass of the model, or rather, to answer the following questions:

1. What is the effective expression range of a hypothesis?
2. What is the hypothesis structured in?
3. If the model is constructed by parameters, how many parameters are there?
4. How many components are there in specific model h , and how would another model $h' \in \mathcal{H}$ differs

from h ?

Those questions are not so easily answered. In fact, under the current learning theory landscape, such is often incoherent if taken individually, because the model structure are often incoherent the sense of exhibiting the **expressive plurality** problem, where the abbreviated $|c|$ and $|\mathcal{H}|$ size of the hypothesis model or hypothesis model space, respectively, are rather actually a set of all measures contracted on the model description. At most, there exists two types of model complexity, with the first being *operational complexity* or often known in discourse as expressive complexity (though such term is poorly used), and the second one being *structural complexity*, coming from the model configuration and static mass-operation conditioning (mass here means the collection of all objects in question of the model itself). Those are model-respective complexity. When we are to initialize testing scenario, such term is then also sample complexity, and so on. However, the majority of the plurality can be pointed to the insufficiency in expressing different classes of models together. Or rather, the lack thereof of structural equivalence beyond simple mathematical formulaic equivalence, which, if such is sufficient, the problem would be absolved.

Toward specificity, analysis often puts the hypothesis class $\mathcal{H}$ as a class of functions. So, there exists the class $\mathcal{H}_n$ of n -th degree polynomials, or else. Effective complexity in such case means the range of expression a function can give – for example, for n -th polynomial class. For the class $\mathcal{H}_n$ of polynomials of degree n on a closed interval $[a, b] \subset \mathbb{R}$, the Weierstrass approximation theorem guarantees that, as $n \rightarrow \infty$, any continuous function $f : [a, b] \rightarrow \mathbb{R}$ can be approximated arbitrarily closely in the uniform norm. Hence, the expressive complexity of $\mathcal{H}_n$ grows with n , but for fixed n only a finite-dimensional subset of $C([a, b])$ is exactly representable. This distinction highlights the difference between *approximation capacity* of a hypothesis class and the *exact representational capacity* captured by its parameterization. These expressions depend on the field or space the hypothesis class lives in, for example, if the hypothesis is configured to work in the set of all functions in $\mathbb{F} = \mathbb{R}$ of algebraic field, then we call it the algebraic family. There are various hypothesis classes in such case, and usually, not a lot of them can be analysed and represented in the same way, and hence also the analysis of the ‘size’ of the hypothesis, also between different hypothesis classes, prompted the usage of a more general notion, called **representation complexity**¹.

The above complexity (and so is in Table 5.1), aside from sample space-related measure, can also be applied onto the concept class’s members itself. We also take the class complexity in such case, as the argument that maximizes those measure: that is, the largest value of the measure, taken on the class as the class’s complexity.^{2,3,4,5,6}

The next complexity measures that can be taken of, comes from the learning process structure itself.

¹Representation complexity refers to the number of components needed to express h and generally, any hypothesis in $\mathcal{H}$. This is often the parameters count, since the parameters themselves are used in constructing the function in algebraic system. Thus, again, an n -th degree polynomial might have $2(n + 1)$ parameters, for $n + 1$ variational input (or variable), and $n + 1$ coefficients. Notice that we specifically at the beginning, use the notion of algebraic structure $\mathbb{R}$ because of it. If we restrict ourselves to the expression, for example, for $\mathbb{B} = \{0, 1\}$, then it can be expressed in various ways in logic theory, such as conjunctive normal form (CNF), disjunctive normal form (DNF), and else. This forbid the analysis based on the effective process itself, but rather the number of components required, hence taking them as complexity of construction.

²Expressive complexity measures the variety of input-output behaviours the model can implement.

³Structural complexity is decomposable into component-wise measures (e.g., per-layer, per-module) and useful for architectural comparisons.

⁴Sample complexity depends on hypothesis class, loss, and desired error/confidence. Further clarification must be made of a grounded definition for utilization.

⁵Sample space complexity infers the statistical measure over the sample space; affects learnability and effective sample size.

⁶We also note that the notion of optimization is usually hard to pin down analytically, often characterized empirically or by worst-case bounds.

Table 5.1: Descriptions of complexity measures applied to models and learning processes, from Hajek and Raginsky [2021], Sugiyama [2015], Mohri et al. [2012], Kearns and Vazirani [1994] and Vapnik [1999] with some custom notions.

Type	Notation	Purposes	Notes
Representation	$\mathcal{M}_r$	Complexity of encoding a model (size of representation, alphabet-dependent).	Model-level measure.
Expressive	$\mathcal{M}_e$	Complexity of the model's observable behaviour (operational expressivity of the mapping it implements).	Model-level measure.
Structural	$\mathcal{M}_s$	Complexity of the model's internal architecture (how components compose and interact).	Model-level measure.
Structural mass	$\mathcal{M}_\Sigma$	Aggregate representation size derived from representation complexity across the model.	Sum of $\mathcal{M}_r, \mathcal{M}_e, \mathcal{M}_s$
Sample	$\mathcal{M}_{p \in \mathcal{S}}, \mathcal{M}_p$	Number of samples required to reach a prescribed performance level (statistical sample-count).	Process-level measure.
Sample space	$\mathcal{M}_\mathcal{S}$	Intrinsic complexity of the data distribution / sample domain (richness, diversity, structured-ness).	Global process measure.
Optimization	$\mathcal{M}_\epsilon$	Complexity of the learning algorithm (convergence rate, computational/iteration cost, sensitivity) or worst-medium-best learning cost.	Process-level measure.

Notes: The notion of structural mass is convoluted. In certain setting, it is simply the parameter space $\theta \in \Theta$, but in some setting, there exists fundamental difference than identically weighted parameter space complexity, as seen in multi-basis polynomial models, not taking into account homogeneous deep learning networks. Optimization complexity coincides with PAC-learning bounds for polynomial runtime, and *sample complexity* sometimes can only be the dimension of the input/observation space.

Usually, in statistical learning theory (Hajek and Raginsky [2021], Mohri et al. [2012], Shalev-Shwartz and Ben-David [2014]), we often use the setting that is related to **supervised process**. This usually means the process can be configured in a feedback loop, or induced response toward the model – often come from the observational space, where the concept that is set to be the target of the process – then called as learning process – is conjured. For now, we reserve this to the classical learning theory and process, and will not inquire further of such direction.

The learning process assumes an operational space containing the structure, for example, $\mathbb{R}$ -encoding space of real values, and a set of **samples** living in such space, denoted $\mathcal{S}$. Those samples are often referred to as the **observation set** of a particular objective of learning. As mentioned above, $\mathcal{S}$ belongs to the concept c , assumed to be of certain concept class $\mathcal{C}$. $\mathcal{S}$ is obtained by using an observing process to create or generate the result. This process can be determined or considered in multiple way, however, generally can be considered by having a deterministic interpretation (using an explicit function $c(x)$) or by probabilistic approach (modelling a distribution $\mathcal{D}$). Hence, we gain another type of complexity, with regard to the process called **sample complexity** – usually defined of cardinality – the amount of sample needed for particular purpose, for particular objective. For the complexity category up to the structure of the learning setting by itself, we call it the **sample space complexity**.

One final complexity that is worth mentioning is then the **optimization complexity**, measure on the action of the learning algorithm. By such, it investigates the construction of an optimization algorithm solely dependent on its size of the algorithm, and other non-process related measures. Optimization complexity can also be related or expressed, in situation of time complexity, which is typical in computing system.

One might ask why we need to determine and categorize complexity measures and the need to generalize and differentiate of each one. Indeed, while in usual use cases, such complexity notion can be

considered sufficient in controlling and gauging the model itself. However, just as James et al. [2009] indicated in the section on deep learning double descent, such complexity might eventually, of some regions, inadequate to express and evaluate the complexity of the model in its operations any more. Such transition is unanswered, and while the founding of complexity notion is still fragmented, it is indeed difficult to not paying attention specifically to those notions themselves. Considering such of the learning framework and the more sophisticated goal that we push forth of reforming the learning theory, it is then natural to define such. That said, the list itself is not enough, and it would be rather difficult to think of, as without a clear structural scheme, we might as well cannot analyse the system holding such.

5.2.1 Current complexity measure

Within the current framework and theoretical understanding, certain conclusions can be taken for analysis. Supposedly, the different methods of defining complexity measures are presented. However, often time, they are organized in a fairly compressed notion, or fragmented under different perspectives. Let us first start with introducing the terms and problem on representation complexity.

One of the formal definition for representational complexity then works in assumption of the wider concept class $c \in \mathcal{C}$. However, it can also be applied onto the hypothesis class $\mathcal{H}$, and toward such bases the concept spaces into conjunction that can be compared of each other. To define it however, we need to define the concept of a representation scheme (Kearns and Vazirani [1994]), which defines under a particular sense, how the model expression is built.

Definition 5.2.1 (Representation scheme). For a given concept class $\mathcal{C}$, a **representation scheme** for such concept class is a function $\mathcal{R} : \Sigma^* \rightarrow \mathcal{C}$, where Σ is the finite alphabet of symbols. We call any configuration $\sigma \in \Sigma^*$ such that $\mathcal{R}(\sigma) = c$ a **representation** of c , under $\mathcal{R}$. For any given c , the representation scheme might not be unique.⁷

This can also be named differently to be the **encoding scheme**, lying around Layer III downward to Layer II in our ladder of functional. Specifically, we choose a type of string on $[a_1, a_2, \dots, a_n]$ finite for our descriptions, substituting each $a_i \rightarrow \rho$ for $\rho \in \Sigma$ of the alphabetical object used for representation. Under simplest circumstances, such correlates to binary representation on logical derivatives, $\{0, 1\}^n$ string with encoding. Note that this remains tangible only because it abstracts the *descriptive information* into the representation, discarding dynamical, operational details, existential concepts, 'volumetric' identifiability in exchange for parameter encoding that requires external readers to make sense of the string ordering and patterns. As such, representation of a specific concept or hypothesis classes can be different, as they rely on the parser and how everything works. Consider the learning problem of a Boolean string. The representation can take form as the **conjunctive normal form** (CNF) or **disjunctive normal form** (DNF), with their expressiveness on equal, however string length varies exponentially, as seen by Wegener [1987], Miltersen et al. [2005], Darwiche and Marquis [2002]. Since the learning algorithm has access only to the input-output behaviour of c , in the worst case, it must assume that the simplest possible mechanism is generating this behaviour. Within this assumption, we can now formulate the **representation size** or **representation complexity** of particular concept, denoted $size(c)$.

Definition 5.2.2 (Representation complexity). Given a concept $c \in \mathcal{C}$, the **representation complexity** of c is defined to be the size of the smallest representation of the concept c in the underlying representation

⁷Perhaps representation scheme can be further realized by applying a Borel σ -algebra on top, and put on it a perspective measure to handle the notion of representation size. For example, in the traditional example of axis-aligned rectangle, one such representation scheme could be the real number schema $\mathcal{R} : (\Sigma \cup \mathbb{R})^* \rightarrow \mathcal{C}$, which then can utilize the extended notion of a Lebesgue measure over $\mathbb{R}$ and its spaces.

scheme $\mathcal{R}$, that is,

$$size(c) = \min_{\mathcal{R}(\sigma)=c} \{size(\sigma)\}$$

for any string σ of the representation scheme $\mathcal{R}$ that still represents c .

Under such definition, the more "complex" the concept c is with the respective chosen representation scheme, the larger $size(c)$ is. This is also where the concept of the concept class $\mathcal{C}$ comes from – it comes from the *representation class induction* that we have in mind some fixed concept classes we study by their representation scheme. We note that, however, the definition works within the least-lower bound limit – we choose the infimum. In the worst-case scenario, the representation potential mismatch between h and c might be very large, hence the *effective representation* might not be the same. We also operate on the assumption that different representations are equal, even though practically, we do not have the exact form of c to be sure that problems can reduce from one representation to a smaller one – which forms the basis of the **manifold hypothesis**⁸. Hence, the following assumptions are needed.

Assumption 5.2.1. *For every concept $c \in \mathcal{C}$, there exists the representation family $\{\mathcal{R}\}$ of all representation with characteristic shared by each concept c . For each c , there exists one or more representation $\mathcal{R}$ such that its representation complexity can be defined. A weaker assumption assumes that for c of one or more representation scheme, the more complex representation scheme can always be reduced to a lower representation scheme of minimal complexity.*

However, the definition falls into several pitfalls regarding its potential usage. First, the definition itself is too restrictive. Kearns and Vazirani [1994] specifically defines it in a sense, on the system of binary representation scheme and a *Turing-theoretic accept state* – representing c means if there exists a Turing-like machine, then the output then must come out in the accept state of the membership of c . This interpretation is too general, and also too specific to be applied on any reasonable model in practice. Second of all, we do not know the representation scheme in practice too, hence making such measure fairly inconsistent within its own setting. Generally, to give a fairly nice bound or result on the representational complexity, we require the hypothesis and the concept to belong to the same alphabet space, as if there exists a substantial shift between them, then by no mean can h mimics c if they are too different and there exists no way for h to describe c . They are some of the serious issues with this complexity measure, though their theoretical significant is still present partially.

Optimization complexity

Optimization complexity refers to the runtime in solving the learning problem by itself. Taken for example, between polynomial and linear regression analytical-probable solution solving time, the complexity lies around $\mathcal{O}(np^2 + p^3)$ for the structural complexity defined by $p = d + 1$ parameter count (including bias) via ordinary least square (OLS). However, this measure itself is not enough, and is not relevant to our problem, partially because of its operational level usage in runtime, which is often attributed to the learning algorithm convergence bound itself. In practice, most analysis consider optimization complexity toward the slice of **time complexity**, where one deliberately take the reductio of the algorithmic process into the complexity where finite computation terminates under specific bounding. For example, certain analysis indicates, Abrahamsen et al. [2021] for neural network to be realizable on $\exists\mathbb{R}$ -complexity,

⁸The **manifold hypothesis** can be loosely expressed in the following statement by Lawrence Cayton

[...] the dimensionality of many data sets is only artificially high; though each data point consists of perhaps thousands of features, it may be described as a function of only a few underlying parameters. That is, the data points are actually samples from a low-dimensional manifold that is embedded in a high-dimensional space.

and so on. Such incurs two separate structure – either existentially of that complexity class, or the notion of time complexity on effective condition time complexity and structure shifts, which includes context and assumptions given to the model.

Still, we should give our own idea of such optimization complexity, called simply the ϵ -optimization complexity, rephrased from famous variations already established in the theory of optimization (hence, not a novel definition).

Definition 5.2.3 (Optimization complexity). The **optimization complexity**, or ϵ -optimization of a given algorithm $\mathcal{A}$ on $h \in \mathcal{H}$ is the worst-case computational complexity, such that $\hat{R}_S(h) \leq \epsilon$ for $\epsilon > 0$.

We can change this definition to the average-time complexity, though usually we often not opt for this. Arguably, setting a fairly stable point of the average-case would benefit the analysis, simply because of the randomness in argument. In the same time, we note that such is incomplete, and would be directed to most of the analysis when encountering real constructs, specifically in later chapters.

Sample complexity

The sample complexity can be defined in a fairly simple manner. Per usual, sample complexity, in context of learning systems, is defined as the number of training samples that it needs in order to successfully learn a target function. Certain definitions can be found, in for example, [Huang et al. \[2025\]](#), [Wegel et al. \[2025\]](#), [Mei et al. \[2022\]](#) or in a more specific context of active online learning, [Balcan et al. \[2010\]](#). However, such definition is usually too simplistic of the sample complexity itself, or compressed any reasonable information or properties are there of the sample set. Nevertheless, we should still recall the copulated definition,

Definition 5.2.4 (Sample complexity). A measure $\mathcal{M}_S$ gauges the *sample complexity* of a learning setting on sample set S , if, the learning algorithm $\mathcal{A}(h, c)$ can guarantee, for $M_S : \mathcal{P}(S) \rightarrow \mathbb{N}$, such that M_S outputs a particular value such that for S_i the i th ordered power subset of S , then

$$\mathbb{P}_{S_i} \left[\hat{R}(h) \leq \epsilon \right] \geq 1 - \delta \quad (5.16)$$

If, it can guarantee instead, such that,

$$\mathbb{P}_{S_i} \left[|\hat{R}(h) - R(h)| \leq \epsilon \right] \geq 1 - \delta \quad (5.17)$$

We call it the *generalization sample complexity* on empirical data.

Again, we emphasize the usage of this sample complexity term, however, either we *decompose it* (which is not possible), or we will need to define another notion for *sample complexity* of a given learning setting entirely. Furthermore, statistical metric of the sample set should also be considered. For example, a *statistical-aware sample complexity* can be defined by the (ϵ, μ) -sample complexity measure on Gaussian spread distribution, as followed.

Definition 5.2.5 ((ϵ, μ) -sample complexity). A measure $\mathcal{M}_S[\epsilon, \mu](h, c)$ gauges the *sample complexity* of a learning setting on sample set S , if, for the sample distribution on particular ϵ, μ , of a random permutation process, the learning algorithm $\mathcal{A}(h, c)$ can guarantee, for $M_S : \mathcal{P}(S) \rightarrow \mathbb{N}$, such that M_S outputs a particular value such that for S_i the i th *permuted power subset* of S , then

$$\mathbb{P}_{S_i} \left[\hat{R}(h) \leq \epsilon \right] \geq 1 - \delta \quad (5.18)$$

This definition sought to connect, in a sense, partially, the statistical analysis of the sample space to the complexity term with respect to the hypothesis h . Even though we use the statistical set (ϵ, μ) of normal distribution process, a different set can be used, and hence changing the definition accordingly.

Expressive complexity

Expressive complexity can subtly say itself to be the **effective representation size** of the model itself, though that would be rather conflating it to representation complexity. What is meant here often correlates to VC-theory, Rademacher theory, Depth-Width efficiency (DWE) in neural network, hence for such to mean the capacity in which the model can *change the representation scheme* to fit objects or to replicate certain amount of structural fixture. By then, let us investigate the proliferation of such concept on the baseline of deep learning, relying on works of [Hu et al. \[2021a\]](#).

Structural mass

The structural mass takes similar inspiration off the representation scheme and complexity, however, this time, it is defined in a usual parameter-based style. Basically, it is the total parameter that is used to specify the numerical encoding of a given model. Under normal setting, we would take this as the

5.3 The nitpicking set

Obviously, there are a lot of problems regarding how we treat the problem, how we see it, and how we actively analyse it. But not all fits into a clean analysis, and some of which are more institutional, or more rhetorical, and cannot be explained analytically. For them, we defer it to here of the analysis.

Most of them lies in the space of statistical learning theory, since bias-variance on itself has variations and justification that is seemingly very natural in statistical theory, like the other version name *approximation-estimation error tradeoff*. All of them contribute somewhat to the current outlook, and so far, it has been pretty bad, than what is being advocated of. Nevertheless, let us state them out, even if missing some particular details. Some of the problems might also intersect with the above sections, but not so much conclusive at will.

5.3.1 Authoritative statements

Despite its prevalence in particular research interest, double descent itself has no definite statement to describe its particular behaviour. This is also true for a list of other phenomena, such as *grokking* ([Power et al. \[2022a\]](#)), *emergence behaviours* ([Wei et al. \[2022b\]](#)), typically discussed in large language models, which is understandable of large-scale specific ruling system), *implicit bias* ([Soudry et al. \[2024\]](#)), and *neural scaling law* ([Kaplan et al. \[2020\]](#), [Wei et al. \[2022b\]](#)). There is also *catastrophic forgetting* ([van de Ven et al. \[2024\]](#)) which is still fairly native of large systems; nevertheless, those phenomena will be of target later on (or discussed in later chapters). The fact remains that we only *partially understand* these particular results with conjectures and theories to support the facilitation of new insights, new analytical lens, and nothing much. Without a clear definition of what is even double descent, sometimes, we cannot proceed, lest to say even solving it. This is particularly troublesome, since the current theory is so messy and not rigid that many even turns to exotic choices, like topological informatics learning theory, *Homotopy Theoretic and Categorical Models* ([Manin and Marcolli \[2024\]](#)), and several more. So this is definitely one of the main problem to pushback on.

5.3.2 Model complexity

The notion of model complexity can be said to be often not so well-defined. There are attempts by [Hu et al. \[2021a\]](#), [Luo et al. \[2024\]](#), [Barceló et al. \[2020\]](#), [Molnar et al. \[2020a\]](#), [Janik and Witaszczyk \[2021\]](#). The problem of model complexity is not so apparent until the inherent complex nature of models, typically in deep learning models (based on the MLP architecture) is observed to cannot be treated in similar way as classical models. In [Hu et al. \[2021a\]](#), most deep learning models are based on, and investigated of their complexity through measures like the expressive capacity, the effective capacity, and so on; aside from other teams and researchers investigated it by using computational complexity, or functional decomposition. One promising aspect of this is [Molnar et al. \[2020a\]](#), where we are able to decompose the functional form a specific model into variations of its complexity.

Nevertheless, it is reasonable to say there are no conclusive measure or definition of a model's complexity. Considering that double descent also appears outside the range of deep learning model, and is somehow consistent in classical models (for example in models tested by [Belkin et al. \[2019a\]](#)). This is particularly troublesome, as a theoretical analysis requires such definition, especially when the relation of interest is directly involved in such manner. This is a particular aspect needed to be resolved if there are to be progress made in this investigation. Additionally, the concept of overparameterization and underparameterization is also a problem, especially since its definition is not well-defined, hence become a problem in analysis by identifying the classification of the problem setting.

5.3.3 Model structure and the bounding inequality

In modern practice, the theoretical and formal treatment, rigorous consideration of different model structures and complexity is not realized. Partially, this is due to the bloom of machine learning and artificial intelligence from the early 2000s. Currently, there has been no conclusive theoretical definition or formulation about the structure of different models, typically can go under the name of mathematical modelling hypothesis, in a way that unify certain properties between architectures. Heuristics are also often employed in several newer models and widely-utilized structures typically seen in large-scale systems, leading to difficulty in formalization and factorization of affectants.

We also mentioned classical learning theory in accordance to our analysis. In particularly, we would like to draw out the fundamental results that have been driving learning theory, and what they say beforehand of the analysis. Would it just only that the old theory being misrepresented, misinterpreted, or there exists some fundamental restriction of the theory that lead to bias-variance being considerate of such setting only, with certain strong assumptions that lies within the theory itself, and can be extended or removed toward generality. Such is to say quite a bold claim, but worth to consider pursuing it.

In consideration of the PAC-ERM bound, of which restrict $R(h)$ and $\hat{R}(h)$ on some particular restriction, at least for ERM, they are fairly implicit in nature. We will have to reinterpret what they are saying to us, since it does not make much sense usually so. But the general pitfall that we have to be careful of usually lies in the nature of the bounds themselves. Specifically, their fame of being rather existential guarantees but no process-based analytic.

The first point of approach comes from either the *Probably Approximately Correct* framework and its variation (we would like to talk of the singular original framework) or the *Vapnik-Chervonenkis* framework of learning guarantees. Both of them serve what is there to be called a learning bounds within finite hypothesis class, or by reduced optimization problem with estimators. PAC bounds typically guarantee the upper, worst-case scenario bounds, while VC-dimension of its own definition is typically used for general lower bounds. Nevertheless, we have to know how they work to interpret and use them correctly.

PAC-learning criteria started by specifying a restriction bound in which, the probability that the generalization error can be achieved, for error margin $\epsilon > 0$, is less than a confidence point of $1 - \delta$ for $\delta > 0$. How they did this? Well, by bounding the error under specific m samples of the observation space such that the above criteria is guaranteed, given $m \geq \text{poly}(1/\epsilon, 1/\delta, n, \text{size}(c))$ where $\text{size}(c)$ is the representation space of the concept to be learnt from, and n is the input-dimension on itself. Basically, the guarantee under worst-case scenario, whether the generalization error evaluated of the algorithm $\mathcal{A}$ be specifically satisfying the criterion. This, is a *concept-aware* criteria – reflected by specifying the general copula, restricting the algorithm and consider the term of the *ambiance space*⁹ $\mathbb{A} = (\mathcal{C}, \mathbb{I}(n, m))$ upon the evaluator correspondence. The advantage of this method is well-known – it bounds, for example, for how many samples would it be possible to learn, under the assumption that the algorithm $\mathcal{A}$ can indeed, learn such that h^* of the final result is $\inf_{h \in \mathcal{H}} \hat{R}(h)$, within polynomial expression of m . However, the implication of PAC-learning, is large and fairly difficult to attend to, such is to say they also consider the consistent and inconsistent case already, which is good, but still fairly not so straightforward in particular expression.

For example, PAC-learning on ERM principle requires that, for some particular notion of the bounding method, such that the following bound will be true,

Proposition 5.3.1. *For any $h \in \mathcal{H}$, for any $\delta \in (0, 1)$,*

$$\mathbb{P}_S \left(R(h) > \hat{R}(h) + C \sqrt{\frac{\log(1/\delta)}{2n}} \right) \leq \delta \quad (5.19)$$

This bound works, under scrutiny assumption such that c and h effectively does not matter in such case, and only use the concentration inequalities on the estimated empirical error on loss function of a particular setting. This is by itself, not so rigorous, and not of the spirit of the PAC-learning boundary itself. The guarantee itself is, controversially, not even a good one. At least for our interest, it does not say anything about the algorithm, it does not say anything about the hypothesis internal mechanics (if the mechanistic model is exposed, like when the concept class is defined, usually, the hypothesis is supposed to work the same way, which is indeed, not the case). What we miss here, is a lot of those factors together. Nevertheless, the PAC-learning bound indicates a particular insight. We have to *formalize* the instance/observational space; we will have to formalize, of it, the sample m , the dimension of the encoding n , the representation system used by $\text{size}(c)$, the concept, and indeed, of the structure of the hypothesis space $\mathcal{H}$ of which gives detail to the hypothesis itself. Such is to say, we need to somehow make the PAC-criteria to be *structural-aware*, per meaning of the word.

For VC-theory and the concept of its dimension, its limit is again, well-known up to certain accordance. Classical VC-theory worked upon a fairly restricted, discrete framework of classification, usually binary classification (for SVM) for example, on *shatterable concept space*. The more, fairly advanced version of *fat-shattering* and multiclass-VC domain still fall short of the generality and more complex structure – in which case we mean the simpler the restriction on what the task is. Insights and ideas toward VC-theory is rather explicit of its own principle – fundamentally speaking, VC-theory conjure to the *combinatorial VC-dimension* of all possible solvable system, up to certain data points, of certain class, except of special cases. Such is to say that one of the aspect that Vapnik-Chervonenkis theory focus on is one particular *expressive complexity* of a specific model. This is also a problem that we have to solve.

⁹In general notion, the ambiance space is what we called the *object collection set*, containing the system's dynamical components and what to consider of itself. Basically, the S in the mathematical modelling setting

Classical Rademacher is also very hard to justify of its use here. Classical Rademacher analysis typically refers to the application of *Rademacher complexity*, a near-empirical theoretical definition of complexity class. One advantage that it poses, against VC and PAC-based structures, is that the ambient space of Rademacher complexity is not restricted toward particular class of problem, like binary labelling of the other theoretical analysis. Nevertheless, Rademacher complexity is unstable (leading to its dilemma of having very few closed-form results) and computationally difficult. Its bounds are also of the same fate, lest to say some of the results proceeded of the Rademacher complexity are often refuted in empirical experiments. Furthermore, strong assumptions and restricted bound cases still hinders its usefulness. And best of all, it does not have anything reasonable analysis for complex structure other than the plain estimators.

Those critiques against classical analysis is not alone of the author. Rather, there are indeed a lot of invoices on such matter toward new development of learning theory. Specifically so, [Zhang et al. \[2017\]](#) inquired of the fact that bounds are often vacuous in terms of deriving for deep neural network. Of the same paper is also the opinion that expressivity is not generalization, which is, obviously true from our theoretical preliminary. [Truong \[2025\]](#) discussed the instability and hard computation of the Rademacher bound, also of threshold behaviours. On the core topic of generalization (which a lot of people pursue for obvious reason), [Nagarajan and Kolter \[2021\]](#) argued against the current uniform convergence theory that seemingly fail to garner an understanding of such behaviour, and so is [Tunali \[2019\]](#). Solutions exist, yes, but either in the form of partial solution and variation of classical theory, like PAC-Bayes, some way of experimenting, like *stability-based analysis*, *margin/norm measures* or localized Rademacher complexity [Bartlett and Mendelson \[2005\]](#), [Bartlett et al. \[2017\]](#), [Bousquet and Elisseeff \[2002\]](#), [Neyshabur et al. \[2015\]](#), yet they are still as we said, *partial of its form*, and no solution has emerged, yet. We would likely want ourselves to not follow the same path again.

5.3.4 Theoretical insufficiency

The setting surrounding the learning theory, machine learning models is generally missing of its rigours and analytical properties. Most of the learning setting and model's training-testing setting are often narrowly stated in theoretical sense, and there are a lot of diffusing details that might or might not affect the model's perturbation without knowing. This in a certain way leads to the complexity in analysing different problem setting and experimental results, since the architecture used, the setting considered, model in questions, configuration (either custom or on system side) are not consistent overall.

Furthermore, a lot of terms, definitions are often hand-waived in papers or in discussions. This also led to the point that in [Nakkiran et al. \[2019\]](#), they have to somehow 'reinvent' another type of model complexity itself, albeit unsatisfying of a definition, is still a new kind of definition per insufficiently of such. While it is welcome to introduce new definitions, the fact is that it is subjectively inconclusive of terms, definitions and thereof leads to the perception and situation where multiple studies can be found, but all of them are defined on different ground.

Analysing statistical learning theory and overall landscape of statistical learning theory, and the learning theory as a whole, encountering new problems like double descent also revealed its weakness, and ultimately, perhaps one of the reason why it is ineffective against such problem. First, there are simply too many assumptions made, too many formulations made during said process. Furthermore, there are also unclear notions and concepts, of which make it even harder to analyse or fully formalize. Secondly, there are inherent conflicts and uncertainty within those theories, formulations and notions by itself. For example, in one sense, the No-free-lunch theorem is considered to be representative and true, whilst also simultaneously being considered the opposite of such. And amidst abnormality behaviours of the old

bias-variance formulation, we also find distinctive weakness in our theory, for example, the concerning difficulty in defining the notion of *model complexity* in various contextual ways. To analyse double descent, perhaps we also need a new theory or formulation to support it.¹⁰

Another problem with classical analytical solution is that it is very *loose*, almost too loose to even considered of such. But sometimes it is too *rigid* in a sense such that it simply does not work at all. Indeed, basing our works and foundational assumptions on top of approximation theory is not quite good at all, because they are very limited. Overall, we can say it is not equipped for handling similar situation. For example, almost all bounds that is subjected of classical learning theory only guarantee particular existential theorem, under very simplified and specific situation without the virtue of extending it to a more general setting.

This is particular event or rather pattern that we often see in learning theory development, or machine learning theory in general – the ‘**mathification**’ of a theory. Even though we need mathematics on either end, the approach is impractical or simply wrong, believing mathematics to be the singular thing that defines learning theory. This has many fallacies that can then be attributed to a lot of factors and whatnot, of several factors that plague this analysis even further than just the problem of double descent. For now,

1. Epistemological limits and interpretations (?) – We ultimately lack understanding of a lot of things. While those ‘theorems’ are very nice in learning theory, the real picture is that it is not real learning, for the word learning are not even defined, as such is to compare them to *human learning* on itself. Even by then, theorems are severely limited. Certain voices also concern of similar problem, including Lipton [2017], Doshi-Velez and Kim [2017], Molnar et al. [2020b] (Molnar et al. [2020b] shortly considers the misleading interpretation question instead), and philosophically, with several push-backs on structural anecdotes, Dreyfus [1965, 1972], Dreyfus and Dreyfus [1986], Suchman [1987], Brooks [1991], Searle [1980], McCarthy and Hayes [1969], Harnad [1990]. On the more modern side, of contemporary critiques, Pearl [2009], Marcus [2018], Sutton [2019]. Of the Chinese Room Argument, perhaps we can look into already copulated passages.
2. Applying wrongly, and is used to impress and not to explain anything (Lipton and Steinhardt [2018]) – mathematics is used to *impress* certain demographic of reviewers and readers, to provide a sense of rigours, to further enhancing the image of formal theory to the point that such theory, even if wrong, can be considered fairly correct by the sheer volume of practitioners believing in such. Such is to say the *mathiness* is turning things into ideology more than rigours itself, of which we can take a tangent to see in economic theory, one of the place to adopt a large portion of machine learning statistics, the pushback against such cursed devolution (Romer [2015], Syll [2024]).
3. Reproducibility and acute false claims (?) – In general, what we have done cannot be recreated, in certain way, and of certain too optimistic setting that is seemingly unrealistic – particularly in adoption toward practical means of actions.

¹⁰Furthermore, most of the general solution and bounds created by statistical learning theory is often in a very simplistic system. For example, if we are to utilize the Rademacher complexity measure (Mohri et al. [2012]), most of the time we will have to compute it through the growth function $\Pi_{\mathcal{H}}(m)$ for m points, on the standard finite hypothesis class $\mathcal{H}$ such that

$$\Pi_{\mathcal{H}}(m) = \max_{x_1, \dots, x_m \in \mathcal{X}} |\{(h(x_1), \dots, h(x_m)) \mid h \in \mathcal{H}\}| \quad (5.20)$$

Most of our problems resolve to binary classification, or rather, the problem of pattern recognizing discrete, reduced classification form. It is not so sure for now if all problems can be reduced to such way, so we cannot draw conclusive analysis that is not diluted of mathematical formulation for complex systems. That is not to count the computationally intensive operation required to calculate the supposedly classical measure, while not entirely of itself holds any substantial reasonable information about the internal dynamics of the model itself.

While most of those are not directly related to our problem – suffice to say the general narrative is so, it particularly indicates a very serious error in the theoretical development itself, and the scope of its reach has been perhaps confirmed at hand. Attempts to solve this problem might need to be an entirely new framework, of something that unify the theory at-hand, or by relinquished attendances to the current dilemma using old and trusted foundation, yet will inevitably fail to resolve further outcome of the solution since its onset is beyond the classical theory. Such is to say, we need an entire solution to cover this particular track. One is promising of them all, being the *Neural Tangent Kernel* (Jacot et al. [2018]), but other than that, nothing much came through. In a sense, it pushes forward the need for something else, yet also deliberately forcing constraints for it to not be something else.

Chapter 6. Intersectional insights

It is evident that we need another theory to explain what we are seeing, of plain sight. We still emphasize the need toward which the new theory can go, and the backward compatibility that is needed of such. Thus, this section will going through ground zero. We would not try to create the theory right away, but reformulate what has been seen first; this includes even the classical generalization and empirical optimization process, under different interpretations. Some of those results might not fit the formalism, or rather rigorously proven system, yet will ultimately form a logical framework on which developments can be made, and perhaps later on theorems to reaffirm the thesis. But the general idea is simple: *learning is not just approximation*, but an entirely paradoxically simple yet complex construct.

Specifically, in the first few sections, the fragmented insight would have to be considered, namely of the following dilemma:

1. We need a more general consideration between the environment and the agent, i.e. between the structures c required to extrapolate, and the modulus in which h explicitly exhibits such 'replica' trace.
2. Information destruction between encoding and modelling assumptions must be addressed, to accommodate the argument about cross-system information divergence or convergence, and so on. This includes also distribution shift and different kind of mechanistic fits.

Those are fragmented realization of the development of learning theory, especially with more experiments conducted ever since the developmental framework (Vapnik [1999], Valiant [1984] are now at least two decades old). The new theory must then, seed such understanding and factors into places, thus also hints at certain type of structural analysis that can conduct such approach. Prerequisite on such structural conditions, the firstly-analysed aspects of consideration includes:

1. To capture uncertainty naturally of mathematical modelling descriptions. That is, the mechanistic to probabilistic transition is inherently decomposable, and such transition from white-box and black-box can be captured by different configurations of space.

Before moving on, we emphasize the importance of realizing the gap between theory and practice. This is inherently serious, and if not for context, should have been included at the very top of the paper. Nevertheless, it should be emphasized here. Critically, this is not rare, and is inherently a complex matter, as seen in many tangent fields and analysis, of Semmelrock et al. [2023], Kapoor and Narayanan [2023] on reproducibility and experimental leakage, Kliemann and Sanders [2016], Tamassia and Vismara [1997] on the engineering reality of algorithms (or geometric algorithms), bgp [2024] goal on determining the

gap between theory and practices, [Revol and Theveny \[2014\]](#), [Collange et al. \[2015\]](#), [Diethelm \[2012\]](#) on numerical reproducibility limits and overall empirical safety, [Kaskenpalo and MacDonell \[2012\]](#) on theoretical-practice bridge. On the specific side of the theory, [Holk et al. \[2025\]](#) analysed such statistical guarantee between the theoretical and empirical mismatch, of which found certain problems regarding such gap, for example, the convergence rate in difussing controlled setting (statistically and system-address wise, stable); so is [Adcock and Dexter \[2021\]](#) on the gap between theory and practice in function approximation with deep neural networks; and so is also [Saxena et al. \[2025\]](#), but on the examination of AI mismatches, which usually is concerned of the *implementation details*. Interestingly, [Behrouz et al. \[2025\]](#) paper on nested learning exploits the line of argument on the mismatch, but however, is troublesome in their formulation and the disparity between theory and practice is not actually addressed, but simply metrics in which the new model is compared to the previous models. That said, it is clear, that any attempt to build theory will inherently attain certain amount of presumptions, implementation decisions, language, memory hierarchy, compiler behaviours, cache, practical limitation on hardware and platform, and the environment tested in which the model is captured of its results. Even when the theory can be correct, such can also be a factor, coupled with the side of uncertain given when learning process or any stochastic, long-runtime system run for a long time. Such also suggests us to consider the agenda of the *ergodic theory* in touch.

6.1 Interpreting the learning theory

Most of the paper will argue about different features of statistical learning theory ([Sterkenburg \[2024\]](#), [Vapnik \[1999\]](#), [Hajek and Raginsky \[2021\]](#)). It is then imperative to discuss the background setting of such. On one hand, when talking of the details on statistical learning theory and the understanding of the encoded language to describe the learning action, typically we are given the observations, or dataset of the form $\mathcal{S} = (\mathcal{X}, \mathcal{Y}) \subset \mathbb{R}^n \times \mathbb{R}^m$, of all 2-tuple pairs, assumed to be sampled or observed and governed by a distribution $\mathcal{D}$. On the other hand, there is also the push toward a more general and sophisticated view, in which might as well reveal deeper or richer structures on the theoretical framework to realize. Note that this contains reformulation, or unification, reinterpretation or patchworks of existing structures too, and not simply claimed to be entirely new.

6.1.1 Domain and process

In dataset, between the *actual concept* and the observable spaces, lies the empirical and generalization setting. While it is not often mentioned in literatures, the switch to direct observable evaluation of statistical learning, induces certain limitations.

Let us assume the following preliminary in discussion. We have two polars, on the mechanistic absolute, or *white box* knowledge, with full expressions on the underlying mechanics of a certain system concept. Explicitly, this can be written as $y = F(x; \theta)$ under standard morphism, for any kind of description θ . Extension of such notion can also include what we call as the mechanism, or $F(x; \theta, \Gamma)$ of which complies both the parameters, and the formulations. The other polar is the *probabilistic domain*, of which then, we have the inference $\sim$ operator instead. Thus, we denote $y \sim \mathbb{P}(Y | X)$ in typical notation, in which the dependency is now Y on the basis of X . One can also do the reverse, as to generate the inverse conditions, but this is more probably general, if one switch to joint distribution (X, Y) under statistical schemes. Abstractly speaking, the dual structure contains what we call temporary, as the *system action* A_S . The notation then becomes

$$y A_S x \rightarrow f(x; \theta, \Gamma), \quad y A_S \rightarrow \mathcal{S}[\mathbb{P}(X, Y)] \quad (6.1)$$

For S here temporarily denotes the *sampling operations*, of observing randomly on the basis of said probability without any knowledge but the presumed parameters. If such process is with control parameter, from the optimization perspective, we have the notation as different to be $S[\mathbb{P}(X, Y); \theta_P, \Gamma_P]$, where now Γ_P is the probabilistic process, i.e. distributions or other random structure, for example as stochastic processes as a finite-dimensional distributions plus regularity structure, for random variables θ_P . In general, we call it the **black-box** domain.

If we are to recover θ from θ_P , or at least on the similarity metric, typically or ML encoded by $\ell(\theta, \theta_P)$, where $\ell(\theta, \theta_P) \approx 0$, then we say that the system experiences a *white transition*. Note that we do not take into account of Γ and Γ_P , since they are inherently different in this white transition, but mechanistically recoverable; thus, *total white transition* means $\Gamma \approx \Gamma_P$ of the mechanistic patterns. If the reverse happens, where mechanistic system turn to probabilistic (X, Y) distribution on θ_P and structure $\Gamma \rightarrow \Gamma_P$.

Probability inherently, similar to mechanistic Γ , enforces causal or rather transformational relation (simply the dual $x \rightarrow y$ is transformational within f). There, the probability can be interpreted, for either Γ explicitly states that they are conditional, for X to be deterministic and explicitly stated, or joint distribution and so on. However, the interference in such recovering mode comes from observational space inconsistency. This is usually expressed of *noise*, where if we admit (x, y) of the mechanistic system, then for $x + \epsilon_x$ and $y + \epsilon_y$ we have instead a *noisy mechanistic system*. We would note that this is a larger class than pure mechanistic system, basically $x \rightarrow [\epsilon, f(x; \theta, \Gamma)]$ instead. Thus, for intrinsic noise of measurements or statistical information, probabilistic transition in recovery of mechanistic systems admits the recovery of the noisy system instead. Depends on the severity of ϵ , such system can be irrecoverable. If such approximation is impossible within certain margin, then probabilistic approximation remains in probable the only manageable system of approximation, rather than a quasi mechanistic recovery mode (we will prove this in a later theorem).

If we to accept such view, then it means that in certain failure modes, you *cannot move a mechanistic (deterministic) system into probabilistic then back*, given enough disturbance and noise. This quantifies to the *information loss* system-wide, where we can create a prototypical functional, denoted $\mathcal{R} : F \rightarrow P$, where

$$\mathcal{R}[F](A) = \int_{\Theta} \mathbf{1}_{\{F(x; \theta, \Gamma) \in A\}} \mu(d\theta) \quad (6.2)$$

for such transition, of the probabilistic measure space $(\Omega, \mathcal{F}, \mathbb{P})$, for $\mu \in \mathcal{F}$ and $A \in (\mathcal{X}, \mathcal{Y})$ of the observational space. The implication here is that, under no information (information-agnostic), every noisy concept must meet certain limit on the basis and resultant structure itself. There exists, presently, no sufficient guarantee to return P to F , except the same operation in reversal, is forced on ideal setting, for infinite-range coverage, and approximately dense observational space. The limit is, inherently, enforced by arbitrary ϵ .

It is to note that we do not claim that mechanistic systems are ontologically privileged; we claim that the mapping from mechanisms to observations is not generally invertible under disturbance.

6.2 The learning theoretic approach

Previously, we say that we have the dataset, or rather called as **observations** of a given system c , $\mathcal{S}_m$ of size m , such that $(x, \cdot) \sim \mathcal{D}$ of a distribution, optionally of also y to be sampled from such distribution. The dataset itself contains pairs, as

$$\mathcal{S} = \{(x_1, y_1), (x_2, y_2), \dots (x_n, y_n)\} \subset \mathcal{X} \times \mathcal{Y}$$

where the dataset is assumed to be i.i.d. sampled according to $\mathcal{D}$, which is unknown by the priori. Hence, we have the first set of assumption. This assumption, based on earlier insight, situated our finding or analysis into the range of black-box and probabilistic models.

Assumption 6.2.1. *The observational space is governed by a particular sampling distribution $\mathcal{D}$, for fixed y -response according to the concept $c \in \mathcal{C}$. The sampling process is further assumed to hold no further structure, hence is *identically and independently distributed* (i.i.d.).*

The learning problem is then formulated as followed. Given the machine learning model expressed a hypothesis h of the hypothesis class $\mathcal{H}$, the learning theory aims for creating a procedure to learn either elements of the concept class $\mathcal{C}$ of all concepts $c : \mathcal{X} \rightarrow \mathcal{Y}$, or the function class $\mathcal{F}$ of all functions $f : \mathcal{X} \rightarrow \mathcal{Y}$, which is usually set to $\{0, 1\}$ or $[0, 1]$. This distinction is trivial, hence, if it is clear, we will talk about the concept class $\mathcal{C}$ only as representative form. The learner $\mathcal{L}(h)$ consider the set of possible hypothesis $\mathcal{H}$, in which might not coincide with $\mathcal{C}$. It receives a partial image of sample $S = (x_1, \dots, x_2, \dots, x_n)$ drawn i.i.d. according to $\mathcal{D}$ as well as the label $(c(x_1), \dots, c(x_n))$. This constitutes the dataset $\mathcal{S}$, which are based on specific concept $c \in \mathcal{C}$ for the hypothesis to learn. The task is then to use (or *extract*) meaningful information to select a hypothesis $h_S \in \mathcal{H}$ that accurately mimic c , with marginal error Θ . The notation h_S stands for all hypothesis that can be inferred from the range of the dataset (the first argument, $\mathcal{X}$).

The marginal error Θ is considered of two parameters, the empirical error $\hat{R}(h)$ and the generalization error $R(h)$. We give the following definition of them.

Definition 6.2.1 (Empirical risk). Given a hypothesis $h \in \mathcal{H}$, a target concept $c \in \mathcal{C}$, and a sample $S = (x_1, \dots, x_m)$. For some particular $\epsilon > 0$, the **empirical error** or *empirical risk* of h is defined by

$$\hat{R}_S(h) = \mathbb{P}_{x \in S \sim \mathcal{D}} [\ell\{h(x), c(x)\} \geq \epsilon] = \frac{1}{m} \sum_{i=1}^m \ell\{h(x_i), c(x_i)\} \quad (6.3)$$

The generalization risk is then defined as the extension of the empirical risk, of which instead of being evaluated on the set $\mathcal{S}$, it is evaluated on arbitrary infinite density sample space. The notation for the risk calculation can also be R_S instead, of which the need is to classify the set of which it is evaluated from.

Definition 6.2.2 (Generalization risk). Given a hypothesis $h \in \mathcal{H}$, a target concept $c \in \mathcal{C}$, and an underlying distribution $\mathcal{D}$ on $\mathcal{X}$. For some particular $\epsilon > 0$, the generalization error or *risk* of h is defined by

$$\begin{aligned} R(h) &= \mathbb{P}_{x \sim \mathcal{D}} [\ell\{h(x), c(x)\} \geq \epsilon] \\ &= \mathbb{E}_{x \sim \mathcal{D}} [\ell\{h(x), c(x)\}] \\ &= \int_{x \in \mathcal{D}} \ell\{h(x), c(x)\} dP(x) \end{aligned} \quad (6.4)$$

In essence, the empirical risk measure the hypothesis to observation error, and the generalization risk captures the difference of the hypothesis to the actual concept. Notice that by doing this, we have implicitly assumed that the observations and the concept is not the same thing, under the majority of scenarios. For the observations to be *approximately equal* or sufficiently captures the concept, there then would have to be various criteria to fulfil. This can be illustrated as Illustration 6.1. For fixed $\mathcal{H}$, for fixed and sufficiently large $\mathcal{S}$, and no observation (data) errors, the empirical risk is the generalization




Figure 6.1: **Illustrative dynamic of the learning problem.** For an incremental hypothesis space sequence $\mathcal{H}_1, \mathcal{H}_2, \mathcal{H}_3$, we aim to obtain the procedure $\mathcal{T}$ that would either reach the *generalization solution* h , or the *empirical solution* h^* for a particular hypothesis class, with respect to the concept c and its observed concept c' . Model selection hence dictates how an algorithm or procedure than choose the best possible hypothesis to approximate the generalization solution from the empirical solution set. Do notice that here we explicitly state that the data would create a **proxy concept** that overlaps with the concept class and of some arbitrary ‘distance’ from the true concept by $d(c, c')$. In case the two classes overlap, there still exists the arbitrary distance. Also, we can also observe the intuitive notion of increasing hypothesis class – while it indeed can help in getting closer to the concept class, the somewhat intrinsic property of current learning theory being the relative random procedure class make it probabilistically unstable.

risk. These two measures between h and c constitute the learning problem, which can also be separated into both cases – either empirical learning or generalization learning, one to minimize $\hat{R}(h)$, and one to minimize $R(h)$. Here, we emphasize such hierarchy by defining the generalization goal as the **nested criteria** for empirical risks minimization.

Definition 6.2.3 (Empirical learning problem). We present the formal form of the empirical learning. Suppose we have a **target**, $c \in \mathcal{C}$, where $\mathcal{C}$ is an arbitrary concept class that captures targets of the same type. Suppose we are provided a set of observations $\mathcal{S}$. The problem is to use certain algorithm $\mathcal{A}$ using $\mathcal{S}$, to obtain a hypothesis h^* for a fixed $\mathcal{H}$ such that:

$$R(h^*) = \min_{h \in \mathcal{H}} \hat{R}(h) = \min_{h \in \mathcal{H}} \mathbb{E}_{x \sim \mathcal{D}, x \in \mathcal{S}} \ell\{h(x), c(x)\} \quad (6.5)$$

The generalization best h is also called in literature as the *Bayes model*, such that it reduces the Bayes risk that is the infimum of the generalization risk, $R^* = \inf_{h \in \mathcal{H}} R(h)$. The hypothesis h^* is often called the *empirical best*, for it being the minimal, finite hypothesis of the lowest loss evaluation on the entire observation space $\mathcal{S}$. There exists no certified assumption regarding whether h^* aligns with the minimal generalization error.

Definition 6.2.4 (Generalization learning problem). We present the formal form of the generalization learning problem. Suppose we have a **target**, $c \in \mathcal{C}$, where $\mathcal{C}$ is an arbitrary concept class that captures targets of the same type. Suppose we are provided a set of observations $\mathcal{S}$. Supposed we have an algorithm $\mathcal{A}$ that for fixed hypothesis space $\mathcal{H}$, equation 6.5 holds true. The problem is to use certain algorithm $\mathcal{A}'$ such that, under limited availability, to obtain h , satisfies:

$$R(h) = \min_{h \in \mathcal{H}} R(h) \leq \{\epsilon\}, \quad \epsilon > 0 \quad (6.6)$$

For a set of risk bounds ϵ . If the setting is deterministic, then there exists $\epsilon = 0$ accordingly.

We can show that normally, the problem of empirical learning and generalization learning is not the same. That is, if one tries to solve the empirical learning problem to certain measure, then they will fail the generalization problem. This can be illustrated by figure 2.1. This problem between the disparity of the empirical and generalization target is what we often called as the *Allegory of the cave*. This is the phenomena where we interpolate too much of the empirical space, when turning to the generalization space, the model instead is marginally significantly inaccurate. Here, we risk being dependent ourselves on false hypothesis obtained by empirical data acquisition, but also have to aim for the general form of the target by itself, which is often assumed unobtainable or impossible to know of prior belief. Hence, either that we take a skewed representation of reality through empirical data, or we diverge from such in a manner that gets closer to the ground truth's concept. At least in this setting, it is then observed as above that empirical learning problem is a subproblem toward generalization problem – one have to solve the empirical learning problem before attempting the generalization one. We also want to introduce the term **data shape** and **concept shape** for data $\mathcal{S} \in \mathcal{C}$ and the shape of $\mathcal{C}$ by itself, to reflect this difference.

Further progression on such description for the process of generalization goal can be, realized in certain form as the global optimization objective in which empiricism can obtain from reasonable set of assumptions and conditions. The following algorithmic framing indeed aims to focus and extradite such sense of operational goal, where we directly attempt such. We note that the differentiable field ∇ itself is a placeholder on similar-typed object, so as for us to not commit to any closure systems prematurely.

Optimization problem 2: Empirical-generalization learning dynamics (non-proxy)

Data: $(x_i, y_i) \in \mathcal{S}, \Gamma \supset \mathcal{S}$, objective $\ell(y, f(x)), \nabla(\cdot)$ differentiable gradient field on ℓ .

Result: Expected optimized model to the minimal error on the objective $\ell(\cdot, \cdot)$, dependent on $\mathcal{S}$.

begin

$\Gamma \supset \mathcal{S}$, the general set;

$R(\cdot), \hat{R}(\cdot) \leftarrow \ell$ as objective functional;

for Fixed $\mathcal{H}$, with $h \in \mathcal{H}$, **for** $\epsilon > 0$ **do**

 Find $h \rightarrow h'$, using algorithm $\mathcal{A}$ for minimization objective, such that

$$\mathcal{A} \left(\min_{h \in \mathcal{H}} \hat{R}_{\mathcal{S}}(h) + R(h) \right) \leq \epsilon \quad (6.7)$$

for procedure $\mathcal{A}(\cdot, \cdot)$ of structural constraints.;

end

end

return Optimized model h on arbitrary objective space (ℓ, ∇) .

These two stages of learning procedure comes of as the intrinsic design of machine learning setting. As opposed to interpolation, where empirical learning problem would have to be put into first priority to achieve the best possible point-through model, machine learning leans on approximation more and a separated criterion that is only visible in the learning setting – the fact that we are using $\mathcal{C}'$ to extract $\mathcal{C}$. This disparity is exactly the **generalization problem**, of which is then also the goal of machine learning, to approximate the general solution using the observed solution. One then can ask what is the difference between the two notions, and when we can guarantee that both will ‘converge’ to the same point. The process of **model selection** is then the procedure that will optimize $h \in \mathcal{H}$ to a given target fixture, such as the empirical best.

One of the major assumption or observation from the point of the hypothesis class, is that the hypothesis cannot know the generalization error. Thereby, the more efficient and often used procedure/algorithm in training, is then the *empirical risk minimization* (ERM) method, minimizing only the empirical risk to the empirical best; then, certain measures to ‘generalize’ this heuristically is apprehended on the hypothesis. The process of model selection is also the place that we get the concept of **bias-variance tradeoff**, and its longstanding form to be a foundational metric on model selection.

In general, to fit also the former theories, in the classical learning setting, we can dissect the problem into what constitute the following cleaner version of the setting, includes both the generalization learning and empirical learning problem. Again, under ideal condition, and ideal setup, the learning setting will have empirical learning and generalization learning to be on the same track. Later on, we would formalize what exactly on the same track means.

Setting 6.2.1. Given an op-space $\mathcal{X} \times \mathcal{Y}$, define the concept class $\mathcal{C}$ of all $c : \mathcal{P}(\mathcal{X}) \rightarrow \mathcal{P}(\mathcal{Y})$ of all power set for both, where $\mathcal{X} \sim \mathcal{D}$ for some distribution $\mathcal{D}$ of the probabilistic assumption. The learning problem consists of defining a hypothesis class $\mathcal{H}$, for h being an arbitrary hypothesis, either with specific architecture or not, and a learning algorithm $\mathcal{A}$, of a **learner** such that the following is acquired:

- Set $\epsilon, \delta > 0$ be the error bound and the confidence measure (that is – the probability of failure) of h , for ℓ the evaluative measure of $\{h(x), c(x)\}$, define $\hat{R}(h)$ being the empirical risk of h on the observation set $\mathcal{S}$ available. Then, the *empirical learning problem* is tasked to find the hypothesis such that $\hat{R}(h) = \min_{h_0 \in \mathcal{H}} \hat{R}(h_0) \leq \epsilon$ for some arbitrary h_0 , with confidence $1 - \delta$. This h is called the *empirical best*.
- Similarly, define $R(h)$ the generalization error, assumed outside $\mathcal{S}$. Then, the *generalization learning problem* is tasked to find the hypothesis such that $\hat{R}(h) = \min_{h_0 \in \mathcal{H}} R(h_0) \leq \epsilon$ for some arbitrary h_0 , with confidence $1 - \delta$. This h is called the Bayes model.
- The generalization setting requires choosing the optimal point for both h^* of the empirical best, and h_B for Bayes hypothesis.

To choose and to compare between the empirical best, and the generalization best, we need certain notions to compare the two, algorithm to choose and evaluate, and metric to determine how comparison should be performed, with both measurable. At best, we need to find the correlation between the subconcept and the actual target concept itself.

6.2.1 Overfitting and underfitting

Following the above framework of statistical probabilistic learning theory, one understanding about the concept of underfitting and overfitting (similar to James et al. [2009]) can be reached, using the notion of *partitioning* of observation space. Such work is typically done because of **practical condition**, in which the following constraints is applied.

Assumption 6.2.2. No knowledge of $c \in \mathcal{C}$ is provided. Consequentially, $R(h)$ cannot be accessed for any particular h ¹.

¹Few points need to be said about this assumption. This assumption represents the *worst-case possible* of a system dynamic. From the perspective of a designer, inherent leakage of information can be provided in the design of a model. In such, case we are operating in the mode of a *resource-constrained, closed* game, in which the model with further assumptions of fixed class, operates on. The problem space itself is inherently large and penalize any additional structural changes. If one is to occur, we usually mention it as *implicit bias*, but another name that would be more fitting would be *structural attractor*, after the notion of gradient attractor in chaotic systems.

Without such knowledge, $R(h)$ cannot be computed, accordingly, and so under circumstances, it can only satisfy the empirical risk notion if one is to procedurally ignore generalization possibilities. We specify such approach as *heuristic*, as it is only an approximation, and for better definition, we define it as followed.

Definition 6.2.5 (Heuristic empirical risk). Given a dataset $\mathcal{S}$ of size m , supposedly of the concept $c \in \mathcal{C}$, and the hypothesis $h \in \mathcal{H}$, then the **heuristic empirical risk** take a partitioning $k < m$ with $k/m \geq 1/2$, hence defining the heuristic empirical loss on k , denoted $\hat{R}_k(h)$.

Definition 6.2.6 (Heuristical generalization risk). Given a dataset $\mathcal{S}$ of size m , supposedly of the concept $c \in \mathcal{C}$, and the hypothesis $h \in \mathcal{H}$. Fixed $\hat{R}_k(h)$, for the final optimization of algorithm $\mathcal{A}$, then the **heuristic generalization risk** take the remaining 2-partition of portion $p = k - m$, hence defining the heuristic generalization loss on p , denoted $\hat{R}_p(h_{\mathcal{A}})$.

The two heuristic definition however, can be noted to not be equal to the true generalization risk - empirical risk pair in reality. Nevertheless, one can then define overfitting and underfitting in such manner.

Definition 6.2.7 (Underfitting). A model h is called **underfitting** if for a constant $\epsilon > 0$ chosen for the setting, then $\hat{R}_k(h_{\mathcal{A}}) > \epsilon$, $\hat{R}_p(h_{\alpha}) > \epsilon$ over the entire observation space.

Usually, underfitting is defined to be such that test error is smaller than training error partition, however, we do not think that is a good definition. Especially since, at least in this setting, the empirical error under algorithm $\mathcal{A}$ is reduced to the representative smallest error that $\mathcal{A}$ can achieve on that particular model, within particular setting of h , then come the heuristic generalization error. Considering also that $k/m \geq 1/2$, that is, at least the training partition is half of the observation set, then only under very specific case will $\hat{R}_p(h_{\alpha})$ be smaller than $\hat{R}_k(h)$, statistically speaking. Thence this line of thought also reinforces the idea that both of them will be larger than the threshold ϵ . Usually, heuristically, this threshold of error is 0.3, normalized. Next, we define the notion of overfitting.

Definition 6.2.8 (Overfitting). A model h is called **overfitting** if there exists a specific constant $\epsilon > 0$ chosen for the setting, such that $\hat{R}_k(h) < \epsilon$, $\hat{R}_p(h_{\alpha}) > \epsilon$ over the entire observation space.

Namely, this mean that the static learning makes the generalization error worse. This draws the analogy of the student - studying explicitly for the past paper tests, yet unaware of the larger margin of questions in the large set of the test paper pool, leading to subpar performance. Similarly, overfitting is theorized to be fitting too fit to the reduced observation space - hence when it comes to the generalization observation space, it fails within such heuristic. The notion works on the mask of the risks, $\hat{R}_k$ and $\hat{R}_p(h_{\alpha})$.

In practice, Assumption 6.2.2 means we cannot have $c \in \mathcal{C}$, or its probabilistic interpretation $\mathcal{D}$ such that $x \sim \mathcal{D}$, or in the extreme non-structured case, $(x, y) \sim \mathcal{D}$. In such case, the typical practical fix relies on the partitioning of the available partitioning space, that is, the empirical-generalization proxy. For large enough singleton count of $|\{(X_k, Y_k)\}|$, we assume that $\hat{R}_{\text{Gen}}(h) \approx R(h)$. Thus, the algorithm can be stated in a form as Algorithm 3.

The assumption above of the generalization set to the generalization truth error is apparent in the usual system theorem, Theorem 6.2.3.

Optimization problem 3: Empirical-generalization learning dynamics (proxy)**Data:** $(x_i, y_i) \in \mathcal{S}$, $y_i = \{\pm 1\}$, hinge objective $\ell(y, f(x))$, $\nabla(\cdot)$ differentiable gradient field on ℓ .**Result:** Expected optimized model to the minimal error on the objective $\ell(\cdot, \cdot)$, dependent on $\mathcal{S}$.**begin** $((X_1, Y_1), \dots)_n \leftarrow \text{Partition}_n(\mathcal{S});$ $\{(X_j, Y_j)\} \leftarrow \text{Train};$ $\{(X_k, Y_k)\} \leftarrow \text{Gen for } k \neq j;$ **for** Fixed $\mathcal{H}$, with $h \in \mathcal{H}$, for $\epsilon > 0$ **do** (Optimistically) achieve $h \rightarrow h'$, using algorithm $\mathcal{A}$ for minimization objective, such that

$$\mathcal{A} \left(\min_{h \in \mathcal{H}} \hat{R}_{\text{Train}}(h) + \hat{R}_{\text{Gen}}(h) \right) \leq \epsilon \quad (6.8)$$

 for procedure $\mathcal{A}(\cdot, \cdot)$ of structural constraints.; **end****end****return** Optimized model $h_{\mathcal{A}}$ on arbitrary objective space (ℓ, ∇) .

Theorem 6.2.3 (Mohri et al. [2012]). For fixed $\mathcal{H}$, for fixed and sufficiently large $\mathcal{S}$, and no observation errors, the empirical risk is the generalization risk:

$$\mathbb{E}_{S \sim \mathcal{D}^m} [\hat{R}_S(h)] = R(h) \quad (6.9)$$

Proof. By linearity of the expectation, and the fact that we assume the sample is given i.i.d., we can write:

$$\begin{aligned} \mathbb{E}_{S \sim \mathcal{D}^m} [\hat{R}_S(h)] &= \frac{1}{m} \sum_{i=1}^m \mathbb{E}_{S \sim \mathcal{D}^m} [\mathbb{1}_{h(x_i) \neq c(x_i)}] \\ &= \frac{1}{m} \sum_{i=1}^m \mathbb{E}_{S \sim \mathcal{D}^m} [\mathbb{1}_{h(x) \neq c(x)}], \end{aligned} \quad (6.10)$$

for any x in sample S . Thus, we result in

$$\begin{aligned} \mathbb{E}_{S \sim \mathcal{D}^m} [\hat{R}_S(h)] &= \mathbb{E}_{S \sim \mathcal{D}^m} [\mathbb{1}_{h(x) \neq c(x)}] \\ &= \mathbb{E}_{x \sim \mathcal{D}} [\mathbb{1}_{h(x) \neq c(x)}] \\ &= R(h). \end{aligned} \quad (6.11)$$

which yields the desired result. $\square$

This theorem can be transferred to a more generalized conjecture, of which fit the generalized notion of the empirical proxy and so on.

Conjecture 6.2.1. *Supposed $\mathcal{H}, \mathcal{C}$ are fixed, for $\mathcal{S}$ is the empirical c^* for $c \in \mathcal{C}$ the true concept. Then, for algorithm $\mathcal{A}(\cdot, \cdot)$ enforcing minimization apparatus, and the specific objective on the second argument, we have:*

$$\begin{aligned} & \mathcal{A}_{j \rightarrow \infty, \mathcal{S}_i \neq \mathcal{S}_k} \left(\min_{h \in \mathcal{H}} \hat{R}_{\text{Train}}(h) + \hat{R}_{\text{Gen}}(h), [\ell] \right) \\ & \approx \mathcal{A} \left(\min_{h \in \mathcal{H}} \hat{R}_{\mathcal{S}}(h) + R(h), [\ell] \right) + \epsilon_{\mathcal{S}} \end{aligned} \quad (6.12)$$

For $\mathcal{S}_i \neq \mathcal{S}_j$ such that there are only distinct samples, j the observation set size, $[\ell]$ the set of structural or empirical objectives enforced on the problem landscape, and $\epsilon_{\mathcal{S}}$ is the empirical minimum that the algorithm result can achieve assuming sufficient ‘perfect interpolation’.

By the end of this section, we would also want to note that we deliberately choose an expression basis of $\mathcal{A}$ via the couple tuple $(\min, [\ell])$ as a choice for *algorithm-agnostic class* of algorithm satisfies the dynamic only, not just implementations. The algorithm, in certain sense, depends much more on structural assumptions, of which is not accessible via observations but via design, and intrinsic information specified in the observation class itself. This is a requirement for fine-tuning analysis in the long-run, as there exists a myriad of different approaches. Indeed, the classical theory of learning simply does not have much to offer in terms of a unified perspective, in which important factors are analysed of its realization. Specifically speaking of such, classical theory does do worse in many parts, for example, the aspect of transparency given the model, supervisor and the concept, of which their roles are not fully understood or formulated correctly, both in controlled setting and extension to replication of real life examples. Thereby, the problem requires a newer theory of learning to come forward. A new learning theory must extend beyond what constitute the basic learning model, such as the idealized PAC-learning, and statistical learning setting in general. More specifically, we want to design a learning theory that is more akin to **modelling theory** of its internalization, but also attain the learning of machine learning theory.

Within the constraints of such theoretical framing, we can then construct a different notion in which the generalization and the empirical learning problem can possibly relate to each other, at least from the statistical framework. Let us consider the problem of mathematical modelling in such sense. Beside the model class construction, which will be discussed in further sections, the entire idea of this structuralization is to refine the framework of which learning and so on can be defined itself.

6.3 System-wide structuralization

From which we can draw out of the above discussions, concurrent theories are inherently unstable and with many variations, making them ill-suited for instant application on analysis of different models without losing major information about the structural details and expressions. Hence, we propose a novel analysis on the basis of the axiomatization of the learning setting, under various structural assumptions, which lies in the implementation sections. For now, we defer to the insight in which justify such movements.

6.3.1 The idea of neural network formalism

Previously, we have been discovering, examining past results and testing on models of different types, generally classical models. While decisive, each modelling setting requires different analysis and is often not so streamline, typically fragmented as per different functional schemes, probabilistic schemes, optimization and algorithmic-based choice of designs, and so on. They have barely any unified view or general patterns for extension, thus making analysis inherently difficult. Within such, our interest, and discovery,

posit that the formalism or the idea of neural network (accordance to designs and patterns of McCulloch and Pitts [1943], Hebb [1949], Rosenblatt [1958b], Widrow and Hoff [1960b], Minsky and Papert [1969], Fukushima [1980], Kohonen [1982], Hopfield [1982], Werbos [1974], Rumelhart et al. [1986], Cybenko [1989], Hornik [1991], Hochreiter and Schmidhuber [1997], LeCun et al. [1998], Bishop [1995], Haykin [1994]) and some of its modern understanding (as of Goodfellow et al. [2016], Zhang [2023], Hinton et al. [2006], LeCun et al. [2015], Goodfellow et al. [2014], Srivastava et al. [2014b], Szegedy et al. [2014], Simonyan and Zisserman [2014], Vaswani et al. [2017b], He et al. [2016], Devlin et al. [2018], Brown et al. [2020b]) practiced and exhibited, to be, in its general form, the correct but ill-defined and poorly developed framework. Thus is our goal in this section to layout such task, and the development that will be described in this section.

Basic understanding

In modern form, a neural network (Goodfellow et al. [2016], Zhang [2023]) N is typically specified by the dual $(L, [S])$ where L is the number of processing layer of a neuron, and $[S]$ is the list of neurons configuration per layer – for example, the number of neurons per layer, and so on. For such network, the definition can be simple as followed

Definition 6.3.1 (Standard multilayer network, Zhang [2023]). We define a K -layer fully-connected deep neural network with real-valued output. Let $m^{(0)} = d$ and $m^{(K)} = 1$ for m the width of the network, and d the shape of the input vector. We then recursively define:

$$x_j^{(0)} = x_j \quad (j = 1, \dots, m^{(0)}), \quad (6.13)$$

$$x_j^{(k)} = h \left(\sum_{j'=1}^{m^{(k-1)}} \theta_{j,j'}^{(k)} x_{j'}^{(k-1)} + b_j^{(k)} \right) \quad (j = 1, \dots, m^{(k)}), \quad (6.14)$$

$$f(x) = x_1^{(K)} = \sum_{j=1}^{m^{(K-1)}} u_j x_j^{(K-1)} \quad (6.15)$$

where $k = 1, 2, \dots, K-1$, for the model parameters can be represented by $w = \{[u_j, \theta_{j,j'}^{(k)}, b_j^{(k)}] : j, j', k\}$ with $m^{(k)}$ being the number of hidden units at layer k ; $\theta \in \mathbb{R}^m$ the weight of the neuron.

Here, each neuron is denoted by their expression as $f(\mathbf{W}^\top \mathbf{x} + b)$ for the activation function f . We can simplify a lot of those structures above into the neural network schemes. For example, the linear function class case can be simplified to a single $(2, [n, 1])$ neural network, such as:

$$x_j^{(0)} = x_j \quad (j = 1, \dots, n) \quad (6.16)$$

$$x^{(1)} = \sum_{j'=1}^n w_{j'} x_{j'}, \quad (6.17)$$

without the activation function h . The polynomial neural network can be represented by simply a shallow, $[3, p]$ neural network. More concretely, if we take the *input channel* as a (pre-)layer, then

$$x_j^{(0)} = x_j \quad (j = 1, \dots, m^{(0)}), \quad (6.18)$$

$$x_j^{(1)} = h \left(\sum_{j'=1}^{m^{(0)}} \theta_{j,j'}^{(1)} x_{j'}^{(0)} \right) + b_j^{(k)} \quad (j = 1, \dots, m^{(1)}), \quad h(x_j) = x_j^j \quad (6.19)$$

$$f(x) = x_1^{(2)} = \sum_{j=1}^{m^{(1)}} u_j x_j^{(1)} \quad (6.20)$$

for a neural network of shape $(3, [n, n, 1])$, and of polynomial custom scaling function. For support vector machine, this is even harder to do, since its definition varies between different notion of the model itself. However, it is interesting because we can treat it as an arbitrary generation of a neural network structure. There are papers about formalizing SVM in the sense of neural network, including making it to be an infinitely-wide shallow neural network, for example, [Chen et al. \[2021\]](#), [Wang et al. \[2019\]](#), [Zhang and Wang \[2014\]](#), however, the structure that is formed from the SVM learner can be expressed as followed:

Setting 6.3.1 (Support vector machine). For any given optimization problem on input space $x_i \in \mathbb{R}^d$, and label $y_i = \{\pm 1\}$, the shallow neural network resultant from the training procedure is always of the form:

$$x_j^{(0)} = x_j \quad (j = 1, \dots, m^{(0)}), \quad (6.21)$$

$$x_j^{(1)} = \sum_{j' \in m^{(1)}} \theta^{(1)} x_{j'}^{(0)} K(x_{j'}^{(0)}, x) \quad (6.22)$$

$$f(x) = \text{sign}(x_1^{(2)}) \quad (6.23)$$

for a specific kernel $K(x, x')$.

Those three fundamental conversions to neural network raise interesting insight into the neural network formalism. For linear and polynomial shallow neural network, we design the network and then provide it with a training scenario, evaluated of the proxy generalization. SVM however can be considered an adaptive data construction, where the neural network is constructed from the dataset or observation set itself, using minimization on RKHS and with constraints. This prompts exactly as such, on one side the pre-optimization construction, and one is the post-construction optimization of a neural network architecture. This interpretation is surprisingly useful when talking about neural network's unit-like structure. Hence, from such evidence prompts us to analyse a neural network, not in only the sense of managing its layers, but also by its unit components together.

This way of analysing unit-wise gives us much more interpretability, and also potential for a newer kind of network specification. As far as we are concerned of, a neural network can represent an active *mathematical modelling system*. This then indicates, for example in adversarial learning case between the target neural network tailored to specific concept, and the learning hypothesis network, then the process of learning itself can be loosely interpreted as a concept of mathematical modelling representation learning, even in the case of black-box model, and not just simply phenomenological modelling of statistical data. To do this, we will examine the structure from the perspective of a **generator**.

The SVM generator

Support vector machine, first formally originated from [Vapnik \[1999\]](#), is a model inherently, specifically defined for pattern recognition task, and binary classification via an *optimal separating hyperplane*. There are two variants for SVM, namely, for linear and nonlinear hyperplane, but the main general aspect of SVM is its reversal of the direction on model building – instead, it aims to minimize the structure of a model, using the data-dependent structure, or *support vectors*.

The SV machine implements the following two precursor ideas: It maps the input vectors x , supposed of the setting, into a high-dimensional feature space Z through some nonlinear mapping, chosen a priori. In this space, an optimal separating hyperplane is constructed. By statistical learning theory, Vapnik

restricted the function class (as for infinite hypothesis it is impossible to learn) to the class of hyperplanes by

$$\langle \mathbf{w} \cdot \mathbf{x} \rangle + b = 0; \quad \mathbf{w} \in \mathbb{R}^n, b \in \mathbb{R} \quad (6.24)$$

which $\mathbf{w}$ is the controlling weight, $\mathbf{x}$ is the input space to the space of all binary $\{-1, +1\}$ category. Hence, this basically divide the input space into two: one part containing vectors of the class -1 and the others being $+1$. If there exists such a thing, then it is said to be *linearly separable*. To find the class of a particular vector $\mathbf{x}$, we use the following decision function

$$f(\mathbf{x}) = \text{sgn}[\langle \mathbf{w} \cdot \mathbf{x} \rangle + b] \quad (6.25)$$

This is called the **hyperplane classifier** class². The above form of finding final classification can then be presented in dual form, which then depends only on dot products between vectors, as

$$f(\mathbf{x}) = \text{sgn} \left(\sum_{i=1}^{\ell} y_i \alpha_i \langle \mathbf{x} \cdot \mathbf{x}_i \rangle + b \right) \quad (6.26)$$

where $\alpha_i \in \mathbb{R}$ is a real-valued variable that can be viewed as a *measure* of how much informational value $\mathbf{x}_i$ has (for $\mathbf{x}_i$ the support vectors we get from training)³. The second idea is of the *kernel method*. This particular method, useful and well-known of the time SVM was created, gain a more prominent position among other techniques, including neural network. Specifically, this method is characterized by the mapping of the input vectors into a richer (usually high-dimensional) feature space where they satisfy *linear separable* criteria. This prompted the possibility to solve nonlinear problem through such mapping, and indeed yields a nonlinear decision surface in the input space, which is linear in the feature space. A **kernel** (function) is then a function $k(\mathbf{x}, \mathbf{y})$ that given two vectors in input space, return the dot product of their images in feature space such that $k(\mathbf{x}, \mathbf{y}) = \langle \phi(\mathbf{x}), \phi(\mathbf{y}) \rangle$ for the nonlinear mapping ϕ . The general form of SVM is then usually cited as

$$f(\mathbf{x}) = \text{sgn} \left(\sum_{i=1}^{\ell} y_i \alpha_i k(\mathbf{x} \cdot \mathbf{x}_i) \right) \quad \ell \leq n \quad (6.27)$$

for any particular final categorization. One of the major problem for analysing SVM, is in the properties of it potentially having infinitely many parameters, depends on the size of the dataset. This makes the curve of both bias-variance and double descent not applicable in such regard – infinitely many parameters' complexity, yet still finite complexity class⁴.

With respect to the observable space, SVM and neural network, we can call the construction scheme of SVM a type of *backward construction* from data $\mathcal{S}$. Or rather, the following algorithm represents SVM construction problem:

Such a factory is reversed inherently, and with its formulation constrained in the form $\mathbf{w}^\top \phi(\mathbf{x}) + b$, of which is equivalent to a non-linear shallow neural network with infinite width (Wang et al. [2017],

²As can be understood simply from such, there exists many hyperplanes that can correctly classifies the classes. It has been then shown that the hyperplane that guarantees the best generalization problem performances is the one with the maximal margin of separation between two classes (?).

³The *support vector* of a particular hyperplane $\mathbf{y}$ is the collections of all points that lies closest to the hyperplane, which is specified as condition for maximization on both side of maximal margin vector machine.

⁴From such observation, it would seem imperative that the notion of complexity based solely on the parameter counts of certain model construction is ill-equipped for certain flavours of machine learning model, in this case support vector machine (in the same type is the Gaussian Mixture Model). No double descent has been observed by empirical reports record per the author's knowledge.

Factory 4: Support Vector Machine – backward construction.**Data:** $(x_i, y_i) \in \mathcal{S}$, $y_i = \{\pm 1\}$, hinge objective $\ell(y, f(x))$, $\nabla(\cdot)$ differentiable gradient field on ℓ .**Result:** Expected optimized model to the minimal error on the objective $\ell(\cdot, \cdot)$, dependent on $\mathcal{S}$.**begin** $d \leftarrow \dim |\mathcal{S}|;$ $n \leftarrow |\mathcal{S}|;$ $\ell(y, f(x)) = \max(0, 1 - yf(x));$ **Kernalization** $\phi : \mathbf{x} \rightarrow \mathbf{x}';$ $f(\mathbf{x}) \leftarrow \mathbf{w}^\top \phi(\mathbf{x}) + b;$ **Objective**

Attain the minimization

$$\min_{w, b \in \mathbb{R}} \frac{1}{2} \|w\|^2 + C \sum_{1 \leq i \leq n} \ell(y_i, f(x_i)) \quad (6.28)$$

Ideally for $y_i(\mathbf{w}^\top \phi(\mathbf{x}_i) + b) \geq 1$, $\forall i, y_i \in \mathcal{S}$, penalized by ℓ ;**end****return** Optimized model $\text{SVM}(f)$ on objective ℓ of ∇ field.

Chen et al. [2021]), but is not a full-fledged neural network in technicality, hence it is only a functional equivalence. That is, such result only indicates that for $N \rightarrow \infty$ of the number of neurons of a single-hidden layer neural network, $\sum_{j=1}^N a_j(w_j^\top x + b_j)$, then

$$\text{NN}(x) \approx \text{SVM}(x) = \sum_i^\ell \alpha_i k\langle x, x_i \rangle, \ell \leq n \quad (6.29)$$

The support vector, ℓ list, is inherently implicit of the factory itself. On the other hand, typical deep learning or neural network constructions are rather *forward-constructed*, that is, the model is in a sense, prefabricated then tested of its efficacy, and thus the effective performance of the model is limited by iterations of its restricted frame. We can easily visualize the factory of which neural network are constructed, different of the SVM factory case:

The difference is deliberate.

6.3.2 Attractor-based insights

The above reformulation in this chapter, giving insights of the learning formulation typically observed, arguably still attains a lot of assumptions, of which we settle for the next section. Optimization procedure 2 and procedure 3 define the precursor tradeoff in the problem space, or rather, the heuristic-generalization tradeoff. We can state it as the following to clarify such problem in plain sight.

Theorem 6.3.1 (Heuristic-generalization tradeoff, model selection). *An algorithm $\mathcal{A}$ on the system $(\mathcal{H}, \mathcal{C}, \mathcal{A})$, for the expected empirical heuristic $\hat{R}_{\text{Train}}(h)$ and generalization heuristic $\hat{R}_{\text{Gen}}(h)$ on $\mathcal{S}$, on the measure landscape of $\mathcal{L}[\ell]$ for $\hat{R}(\cdot)$, of which $\mathcal{A}$ runs separately on $\mathcal{A}'(\text{Train}, h)$ and $\mathcal{A}'(\text{Gen}, h)$ for the above heuristic, there exists*

$$\left\{ \arg \min_{h \in \mathcal{H}} \hat{R}_{\text{Train}}(h), \arg \min_{h \in \mathcal{H}} \hat{R}_{\text{Gen}}(h) \right\} \in \mathcal{L}[\ell] \quad (6.31)$$

Factory 5: Neural Network – forward construction class**Data:** $(x_i, y_i) \in \mathcal{S}$, loss $\ell(y, f(x))$, $\nabla(\cdot)$ differentiable gradient field on ℓ .**Result:** Optimized model on objective $\ell(\cdot, \cdot)$, dependent on prefabricated $[\mathcal{H}]$.**begin** $d \leftarrow \dim |\mathcal{S}|;$ $n \leftarrow |\mathcal{S}|;$ Fix architecture $[\mathcal{H}] : \{L, \{d_l\}_{l=1}^L, \sigma\};$ Initialize $\theta_0 \sim p(\theta)$ on $[\mathcal{H}]$; $f(\mathbf{x}) \leftarrow \sigma_L(W_L \cdots \sigma_1(W_1 \mathbf{x} + b_1) \cdots + b_L);$ **Objective**

Attain the minimization

$$\min_{\theta \in [\mathcal{H}]} \frac{1}{n} \sum_{i=1}^n \ell(y_i, f_{\theta}(x_i)) + \lambda \Omega(\theta) \quad (6.30)$$

via iterative ∇ -field update $\theta_{t+1} \leftarrow \theta_t - \eta_t \nabla_{\theta} \ell;$ until $t \leq N$ or $\ell[h_{M,N}] \leq \epsilon;$ **end****return** Optimized model $\text{NN}(f_{\theta})$ on objective ℓ of ∇ field.

as convex attractor on the measure landscape. Then, algorithm $\mathcal{A}$ solves the generalization if for the transition $\mathcal{T} : \mathcal{L}_{\text{Train}}[\ell] \rightarrow \mathcal{L}_{\text{Gen}}[\ell]$ for separated loss-landscape, it satisfies for $\mathcal{A}(\mathcal{S}, h) \rightarrow h'$, that

$$h' = \arg \min_{h' \in \mathcal{H}} \left(\hat{R}_{\text{Train}}(h) + \hat{R}_{\text{Gen}}(h) \right) \quad (6.32)$$

This tradeoff relies on the precursor concept of the **attractor** landscape, of which is observable within the objective-space expansion from $\mathcal{L}_{\text{Train}}[\ell]$ to $\mathcal{L}_{\text{Gen}}[\ell]$ of the heuristic space. The partition $\{(X_j, Y_j)\}$ and $\{(X_k, Y_k)\}$ covers the practical notion of cross-validation, though the cross-validating procedure itself works much different, as it covers multiple re-shuffling, i.e. unordered re-partition of the default permutating groups. This is the theorem of purpose in which we are to prove. Such requires further clarification, and thus form the basis for this particular len.

6.4 Reattaching the neural perspective

Since the previous analysis of the learning model, we can then notice a few disadvantages and a few very subtle points that we omitted. It is apparent that the analysis is somewhat entangled, with various situations, interpretations, setting and assumptions together, and with a non-unified structure of models in question. The system in which the model is considered and analysed is also missing in essence. There are a lot of ambiguities around what is considered and what is not, even with probabilistic nature and the deterministic aspect. Terms like **latent spaces** and many do not capture the entire dynamic or meaning of the system. It is then we introduce our novel analytical framework, namely the **Unified Unitary Construct** framework. It is however, can be said otherwise that the learning theory is rather not so novel, but approaching the problem is a more distinct way of the already existing framework.

6.5 Motivation

The predecessors of neural network, for operations and purposes, take in the form of a specific application of function constructs, finite singleton algorithms. For such purposes, a specific classical model m takes the form of a single structure, with complex construction, often offering fixed architecture of its representation. Neural network is an attempt, for its philosophy, to change the aforementioned philosophy of constructing machine model by considering compute unit [Rosenblatt \[1958a\]](#).

Motivated by earlier analysis and observations, it is apparent that the analysis is fragmented, with various situations, interpretations, setting and assumptions together, and with a non-unified structure of models in question. The system in which the model is considered and analysed is also ineffective when considering larger structures or more ‘realistic’ consideration aside from very simplified, strong system itself. There are a lot of ambiguities around what is considered and what is not, even with probabilistic nature and the deterministic aspect. Terms like **latent spaces** and many do not capture the entire dynamic or meaning of the system. It is then we introduce our novel analytical framework, namely the **Unified Unitary Construct** framework. It is however, can be said otherwise that the learning theory is rather not so novel, but approaching the problem is a more distinct way of the already existing framework.

Recalled the classical learning framework, for h . We know and construct the principle of the system by settle on $h : \mathcal{X} \rightarrow \mathcal{Y}$, for any arbitrary defined in-out 2-tuple. Hence, the overall construction of a given model can be abstracted into layers of responsive results. This in particular draws some advantages, for example, can be easily approximated or expressed in the black box model style, or empiricism approach. However, it also comes with the drawback that we have seen, in the sense that the interpretation of such black box model is very narrow, even after we add different kind of flavours, for example, the assumptions of the Gaussian noise. This is majorly because of applying and addition of different components without a general framework, leading to a lot of messiness when looking directly into the theory itself. Specifically, for example, is the fact that there exists two kinds of learning analysis, one considers the information-theoretic approach of information exchange, while the other is simply the classical approach given in such the intrinsic noise error given by the observation set itself, keeping the classical theory intact. How to formalise this is a big hurdle to incorporate each other in, without considering the wider framework – one of the main problem with the speciality of learning theory. Nevertheless, at least accordance to general theory of the current date, formalization in to unit-like processing unit as [neural network components](#) is one thing that should be considered for the learning framework. This is also one of the main point for the word **Unified** in our theory.

6.5.1 Neurons perspective

The formalization to be [atoms](#) comes from the biological inspiration of exactly the biological neurons in human and species’ brain. However, the mechanism and how they operate is a bit more interesting.

For the biological more advanced equivalent, the brain consists of a very large number of highly connected elements, called [neurons](#). The simplified structure of such neuron consists of three components: the [dendrite](#), which is tree-like receptive network that carry signals (electrical) to the cell body; the [cell body](#) effectively processes these signals, usually as sums and threshold response; and the [axon](#), a single long fibre that carries signal from the cell body to other neurons. The point of contact between an axon and a dendrite of another cell is called a synapse. The neural network, hence, is established by the synapses and various arrangements of the neuron.

This configuration of biological neuron is particularly powerful. Not only that they can process information allowing someone to write this text, but also the fact that of billions or neuron, they parallelly

work at the same time. Furthermore, the intricacy lies further ahead than the above formation – the entire structure is of too much complexity, for the brain itself. However, a principle can be seen – they are contained of *building blocks* together, in which case is neuron. And **artificial networks** also take this approach.

Note 6.5.1. From such, we can conclude that, we have to have some classification that requires the smallest component neuron to be the *atoms* that everything else is formed upon.

One of the most fundamental thing of the structure of the *neural network formalism*, hence, is the notion of a *unit neuron*. Specifically, a neuron x is defined, and designed to be a discrete operating unit on its own. The word *operating* here means that it has all the facility requires for an input-output process – the most simple one includes the input receiver, the processing formulae, and the output transmitter.

The neural network formalism is best expressed by the following principle. This theorem is by design, and thus our design is to fix the foundational minimization to a specific neuronal basis module. Such is to say, we state the principle for *inductive recursion* of neuronal system collapse.

Theorem 6.5.1 (Fundamental theorem of neural formalism). *Let $\mathcal{S}$ be an automated operational system with internal composition system, admitting*

1. *nontrivial input-output interaction,*
2. *internal closure under composition, and*
3. *irreducibility with respect to its operational primitives.*

Then any minimal operational component of $\mathcal{S}$ must admit a representation whose cardinality and arbitrary (analytic, topological, etc.) structures are equivalent to that of the minimally defined neural structure $\mathfrak{N}(\mathcal{N}, \mathcal{C})$.

We begin by examining the model of a typical neuron. Note that, this is the simplest one, but also is the fundamental block. We have the following.

Definition 6.5.1 (Neuron). Given the neural network formalism. Define a construct x . Then x is called a **neuron** if it belongs component-wise to the class $\mathcal{N}$ of all neurons, which is minimally expressed by $\mathcal{N}(n_i, n_o, M(\dots))$, where n_i is the input handler, n_o is the output handler, and $M(\dots)$ is the internal system.

From this, we then come up with the definition of the minimally defined neuron – the single most iteration of the above definition.

Definition 6.5.2 (Minimally defined neuron). Given the neural network formalism. Then, for $x \in \mathcal{N}$, we call a neuron **minimally defined** if it belongs to the class $\mathcal{N}_0 \subset \mathcal{N}$ of $\mathcal{N}_0(n_i[\mathbf{w}], n_o[\mathbf{w}'], M(f, b))$ for $f : n_i[\mathbf{w}] \times b \rightarrow \mathcal{O}$ is the internal function of M . $\mathcal{O}$ is the domain of n_o in which it receives the value, and in cases, there exists no $\mathbf{w}'$ configured for the output channel.

The notation $n_i[\mathbf{w}], n_o[\mathbf{w}']$ inherently indicate the notion of *control* of the neuron on the two gate of input and output. More specifically, a given neuron, aside from the upper-level expression of $\mathcal{N}(\dots)$, can also be realized by its *parameters*, or rather, its configuration of the neuron itself. By that, for $x \in \mathcal{N}$, the *configuration parameters* for the minimally defined neuron can be taken in the form $\mathcal{C}(\mathbf{w}_i, \mathbf{w}_o, \mathbf{w}_M)$, here we use instead the pairing notation. The letter $\mathbf{w}$ used here is historical by certain accounts: the original idea comes from the term **weight**, in which the neuron can conceptually influence the input, for example, the signal received or information received, by certain amount, either *downplay* it or *signify* it.

We clearly clarify the need for both $\mathcal{N}$ and $\mathcal{C}$ here. We know, that we want the neuron to be an *operating unit*. This means that for it to be fully realized, it needs to also operate, and hence, to be 'observed'. Hence, you need both the description of the parameters which defines it – taken the interpretation where each parameter and configuration is one specific building block on its own, then the block of all those parameters form the shape of the neuron. Similarly, the *operations* on such neuron assure the interaction and working mechanism of all those components together, inside the neuron, by itself. Hence, to fully, minimally express a minimalized neuron, you need to have both $\mathcal{N}$ and $\mathcal{C}$ as its minimality requirement. We then revise it to the following definition.

Definition 6.5.3 (Classical neuron class). Given the *neural network formalism*. Then, we define a **neuron** A to be *minimally defined* if it is equivalent to any unit x of the class $\mathfrak{N}_m(\mathcal{N}, \mathcal{C})$, called a *neuron class*, where the 2-tuple expanded to $\mathcal{N}(n_i, n_o, M)$ and $\mathcal{C}(\mathbf{w}_i, \mathbf{w}_o, \mathbf{w}_M)$.




Figure 6.2: Minimal neuron structure

With this, we formalize the abstract neuron into two descriptions. Note that, however, the *minimally defined* neuron has a parameterized exposure protocol to the construction. In practice, for any neuron $a \in \mathfrak{N}$, the tuple $(\mathcal{N}, \mathcal{C})$ can be different, not to say complex. One of the strong points of this construction of the neural class is that it is *recursive*. The following proposition might help aid in understanding why it is recursive.

Proposition 6.5.2 (Compound construction). *An abstract neuron A can be dissected into three most important compartments (I, O, M), where I, O is respectively the in-out interface. Hence, a cluster of neurons $\{A_i\}_{i \leq n}$ for such $|\{A_i\}| \geq 3$ can be compartmentalized into a new neuron unit, that is, $\{A_i\} \in \mathfrak{N}$, if and only if $|\{A_i\}| \geq 3$.*

Proof. This mostly comes of as definition (2.1). Notice that for any structure of a neuron to be minimally defined, the tuple $\mathcal{N}$ must minimally contain (n_i, n_o, M) . We want them to be discretely defined, then a cluster of neuron will have to have at least one neuron such that to satisfy the requirement of the 3-tuple component. If $n = |\{A_i\}| = 1$, it reduces to a singular neuron; for $n = 2$, at least one compartment do not have any component, hence minimally we need 3 neurons to effectively be arranged as a neuron unit. $\square$

Hence, we can recursively define neuron, either going up or down, if they satisfy such condition. However, usually, we will only go up, as we will have to define in detail the set $\{A_{M,i}\}$ of all possible *minimally defined* neuron.⁵

There are several assumptions and properties accompanied by this type of neuron structure that we would like to note.

1. This neuron structure works *sequentially*. A minimally defined neuron, such as in a practical setting, will have $\mathcal{C}$ to support sequential operation. What this means is present in the definition of the minimal neuron, such that $\mathcal{C} = (\mathbf{w}_i, \mathbf{w}_o, \mathbf{w}_M)$, and the operation follows $i \rightarrow M \rightarrow o$ for an input-process-output session.
2. By extension of the point above, the neuron order in a cluster will also be *fixed*. Usually, this follows a single dimension, or, for example, if there exists a singular span directional vector $\vec{dr}$ that indicates the direction of operation in the cluster.
3. We suppose the *independence* of any components inside a neuron. This is meant for both *operational independence* for non-connective and resultant actions case. If subcomponents are connected in a causal form and so on, then such independence is violated, but not otherwise.
4. The neuron is expressed in a *parameterized manner*. That is, the behaviour of the neuron can be expressed, controlled, and designed by modifying and constructing a set of non-realized parameters. If this set of parameters is finite, we say that it is *closed parameterization*. If not—i.e., there can be infinitely many parameters—we call it *open parameterization*.
5. This neuron structure is called the *forgetful neuron*, simply because it works, in principle and in actuality, as a processing unit, in the minimally defined case. We will soon touch upon this definition for the detailed structure of the minimally defined neuron.

When we structure our neuron in this way, it is natural to ask whether the neural unit can be expressed similarly to a *finite automaton* or not. But aside from such an analogy, our minimal neuron is already very powerful. Certain clusters of such neurons, embedded with structure $\mathcal{L}$ of *layer-order clustering* (which can be thought of as a more formal notion for a layered neural network), can approximate any continuous representation scheme $c : \mathcal{X} \rightarrow [\cdot, \Sigma]$, given an arbitrary closed interval and to an arbitrary degree of accuracy. We often refer to this as the *universal approximation theorem* for the minimally defined neuron system, and it will serve as one of the focuses of the future analysis.

Analysis

The above construction of the fundamental neuron fundamental create the component model of the minimum processing unit which will be used in various collections of system of related mean. In mathematical modelling context, the neuron model is *descriptive*, meaning that it aims specifically not only for the input-output protocol, but also the unit neuron itself. To do this, we recall any given concrete, functional structure under mathematical formalism is given of the basis, on which abstraction and quantification (without physical realization). Hence, our model here aligns well with such mathematical formalism, in the sense that they end up with abstraction without the impromptu need to signify the inner physical realization – or existence itself. And doing this leaves us with formalizing the *minimally defined neuron* of the class $\mathfrak{N}(\mathcal{N}, \mathcal{C})$ into specific requirement.

The inspiration for the concept of a neuron is a biological one: it stems from the study of the brain or any given processing organs of species with specialized region to be called brain. Considering such, we

⁵At this point, I am considering using *group theory* to aid such set compartment.

found out that the brain mostly consists of the fundamental component called biological neuron and its supporting structures – other resource matters of the inner brain that helps to make it function. Without physical realization depends on many things, but most of the time, it is expressed by the proposition that is to remove the physical attachment from any given object of interest, giving it a more quantified look. For a neuron, it can be treated in such way by formalize the neuron component without its supporting structure, replacing electrical signal and other innate signalling mechanism with abstract 'data flow' and 'operations', without much worry about the actual physical realization of chemical reaction and channel – as we do not need such.

Back there, we all said of the neuron class. What does this mean? For starter, this section on the neuron is the relatively minimal construction we can get it onto. Recall that we say mathematical formalism is to put objects in quantifiable notions, and abstractions of the given subject; also, to express them in a language (or framework, depends on your choice of lexical sense) that support such. Hence, one of the first thing we would like to do, is to formalize the neuron into specific *description scheme*; that is, we would like to treat our objects – the neurons – specifically in certain descriptions that captures the essence of the neural unit, without much ambiguity, and an overhead view of the neuron. The tuple $(\mathcal{N}, \mathcal{C})$ did exactly that – it tells us that the neuron is expressed by that separation of descriptions – in the language of mathematical quantification; for then, $\mathcal{N}$ stands for the quantitative *parameters* that characterize the resources, component scale, or the **mass** of a given neuron, which $\mathcal{C}$ stands for the operations that governs the neuron's internal dynamics. One feasible assumption here is that the neuron does not exhibit any of its operation, that is, interfering with the macroscopic world outside itself, aside from the input-output formalism. Those two descriptions, one governs the mass, one govern the internal mechanics that confounds the subject's behaviours, is what the mathematical description do – and it works very well as it is.

In such formalism, $\mathcal{N}$ is typically easier to specify than $\mathcal{C}$ – which is trivial since the internal mechanics is harder to define and construct, for there exists many relationship and connection between components in the same system. And, considering that for any working system, especially for an interactive one, which inherently dependent on the scale t , or given any reference point of operation to be referred upon, there can be either infinite configuration, or infinite type of ordering of the system itself. We may come to the analysis on *invariant structure* of the space of configuration, but otherwise, it is uncommon to have invariant structure presented in the majority inside a typical system, except fitting of certain criteria.

Overall, the neuron formalism of capturing them into a neural class also helps to figure out the *macroscopic configuration* and the *microscopic configuration* of any given neuron. Under a certain instance of the neuron class $\mathfrak{N}$, iterations of individual neurons can be vastly different, depends on their representation scheme $\mathcal{R}$ for individual sub-description of subcomponent – for example, if there exist the component called the *loss function*, then the macroscopic behaviour will only be configured up to certain accuracy, the rest is left for the detail microscopic setting. Hence, we can partition $\mathcal{N}$ into

$$\mathcal{N} \supset \mathcal{N}_{\mathcal{H}}(\dots) \times \mathcal{N}_{\mathcal{P}}(\dots), \quad \mathcal{N}_M \supset \mathcal{N}_{\mathcal{H},M}(n_i, n_o, M) \times \mathcal{N}_{\mathcal{P},M}(n_i, n_o, M)$$

of which here we use the word *hyperparameter* for $\mathcal{H}$ a bit different from its usage of literature, and $\mathcal{P}$ for microscopic *parameters*, under the proposition that the properties can be configured using parameters as abstraction. Similar notion goes for $\mathcal{C}$, though it is less apparent.

Therefore, the minimally defined neuron, or what we can take as the **fundamental model of neural processing unit** works well, and provides us with the first iteration of the constructing brick for our larger, wider implementation. Though our aim is to advance this structure, first, we must take a look to the formalization and propositional-nation of the rudimentary concept of the neuron, even of the

fundamental, minimal neuron. And this will happen at the end of this chapter. Until then, examining the maximal usage of the current structure will be more than enough as it aligns with historical literature's neuron and neural network formalism. Following sections will see how far we can go with this.

6.5.2 Why we call it minimal

It seems a bit weird why we call our construction on its own a *minimal construction*. To see this, though, we have to go back to the definition of a neuron, and the historical viewpoint on the constructing principle. The need for structuralization and complexity measure is impending on the theory itself, for example, the fact that there exists no current singular class of measure for any specific structures of feedforward networks, as seen of [Hu et al. \[2021b\]](#). Often, the minimal construction in typical layered network is the neuron unit only, of which is simply understood as a function with changing activation function. However, this is a limited approach, as it only satisfy mathematical elegance or simplicity of form, which it does well on standard systems, without targeting specific automaton-reliant and actual ergodic properties. For example, take the polynomial neuron. By certain structural minimization as we defined of the fundamental theorem

6.5.3 Preliminary Formalisation

The structure of our theory on the neural formalism is influenced by the object-abstracted treatment of mathematically embedded structures, and the unit-wise principle of particular neuron. Before we meet ourselves into the notion of epistemic circularity⁶ problem, we might as well clarify a few prerequisites for such structure to exhibit.

First, we indict on the fundamental encoding environment that any object can take. The main point of any structure here is that there exists fundamentally the encoding space of two types. First is the object's cardinality space, denoted $\Gamma = (\mathbb{N}, F)$ for any given categorization F . Second is the encoding primitive of the field $\mathbb{R}$ for generality – in general any field is alright, and they define the analytic structure of the system itself. Any extension, for example, the $\mathbb{R}$ -algebra of complex number $\mathbb{C}$ is then the primitive field's extension.⁷ Any type of data or system can then be decomposed to such, with additional structure on top of such primitive. Such is then called the *primitive framework*.

Definition 6.5.4 (Primitive framework). Let us define the primitive framework $\mathcal{P}_0$ of the dual $(\Gamma, \mathbb{R})$ where $\Gamma = (\mathbb{N}, F)$ is the cardinality encoding space, and $\mathbb{R}$ is the base field primitive of the analytical encoding. An object $X \in \mathcal{P}_0$ admits a dual representation,

$$X \mapsto (\gamma(X), \rho(X)) \quad (6.33)$$

for $\gamma : \text{Obj} \rightarrow \Gamma$ of cardinality encoding, and $\rho : \text{Obj} \rightarrow \mathcal{E}(\mathbb{R})$ for the field extension of $\mathbb{R}$ category, or the category of all $\mathbb{R}$ -algebras. For $\mathcal{E}(\mathbb{R})$ without extension, then $\mathcal{E}(\mathbb{R}) \cong \{\mathbb{R}\}$ of all $\mathbb{R}$ -algebra.

For a structure that is supposed to be unit-wise constructed like neural network and the like, one of the main principle is the principle of abstraction. With this, come the idea of layer. Specifically, a *layer*

⁶Also called bootstrapping, where to understand or define certain notion, requires the knowledge of the object wishes to be defined itself – for example, defining the size of a finite natural number set, using the elements of the number set itself.

⁷This corresponds to a variety of idea. For example, see ? for the idea of discrete and continuous representations and processing, and ? for the same idea, but used on computer graphic process where it is constructed as discrete indexing structure (hash tables) with continuous neural fields. Cardinality space takes inspiration from combinatorial species, sheaf theory of global-local set, and Grothendieck's approach to algebraic topology. Extension is the idea from field extension itself, for structures that can be extended. The general idea of this section relies on interpretation of category theory.

separates abstraction in terms of subspace. Let us take an example of such kind for clarification. Let us define the primitive framework $\mathcal{P}_0$ as now the layer L_0 of this layering scheme. Then, we define the layer $\mathcal{L}_1$ of all unit-wised neuron-like units taking over the representation scheme on $\mathcal{L}_0$. We then define the structure of the neuron class $\mathcal{N}_1$ upon such as followed.

Definition 6.5.5 (Base neuron class). We define the base neuron class $\mathcal{N}_1 \equiv L_1$ as followed. For any $\mathcal{U} \in L_1$ for $\mathcal{U}$ as unit-wise construction over L_0 , then every $\mathcal{U}$ satisfies the input signature $\mathcal{I}_{sig} : \mathcal{E}'(\mathbb{R}) \rightarrow L_0$, output signature $\mathcal{O}_{sig} : L_0 \rightarrow \mathcal{E}'(\mathbb{R})$, and $\mathcal{C}_{ext}$ as extensible construction over L_0 , for $\mathcal{E}'(\mathbb{R})$ particular extension on analytical encoding of L_0 . The **type** template of a neural unit $\mathcal{U} \in L_0 \equiv \mathcal{N}_1$ is then defined as $\mathcal{U} = \langle \mathcal{I}_{sig}, \mathcal{C}_{ext}, \mathcal{O}_{sig} \rangle$ where $\mathcal{I}_{sig}$ and $\mathcal{O}_{sig}$ are invariant as type.

outer typing Those definitions directly link it to type theory, while such development is perhaps more complex. In general, this specifies the *type* of a particular unit using the invariant analytical typing in the general environment linking to it. This is the *outer typing* of a given model construct, of which defers the type of which interaction between the environment, or the *global space state* happens and what is received of such. While the ambience space can be forgiving, such cannot be said for the typing of given structure, since it must have the correct typing for identification. The extension is then is fairly simple, but the more important notion is the equivalence of the variant $\mathcal{C}_{ext}$ extensible unit of the inner structure itself. We will develop this more in the later sections, but the insight can be treated as such.

6.6 Polynomial insights

Polynomial space is one of those famous example that is often cited of the bias-variance tradeoff, for example, in Goodfellow et al. [2016], Sugiyama [2015]. Indeed, it is considered, usually, a very effective and high-volatility construct class. Let us reiterate what we know about polynomial class.

There exists two type of elementary polynomial classification, based on their (i) variable count or (ii) its domain range. We first define the univariate polynomial function.

Definition 6.6.1 (Univariate polynomial). A function of a single variable t is a polynomial on its domain if we can put it in the form

$$a_n t^n + a_{n-1} t^{n-1} + \cdots + a_1 t + a_0 t^0 \quad (6.34)$$

where $\{a_i\}$ are constants, and of a singular variable t .

From this, we can extend to polynomial of multiple variables. In general, a function $f(x, y)$ is a *multivariate polynomial* in x, y if and only if it can be represented as a finite sum of terms of the form $c x^k y^m$, where c is a coefficient and k and m are nonnegative integers. The number $k + m$ is called the degree of the term, and the degree of the polynomial $f(x, y)$ is equal to the highest degree of its own terms. We can then generalize multivariate polynomial and univariate polynomial in the same kind.

Definition 6.6.2 (Polynomial, multivariate generalization). Let K be a field. We then denote $K[x_1, \dots, x_n]$ the space of all polynomials in n variables, $x_1, \dots, x_n$. A **multivariate monomial** is then an expression $x_1^{\alpha_1} \cdots x_n^{\alpha_n}$ where $\alpha_1, \dots, \alpha_n$ are integers. The degree of such multivariate polynomial is the largest degree of terms of any monomial term.

Effectively, we can see that multivariate polynomial can take on many forms, and also unconventional forms. In order to capture any reasonable details, most of the time focus are on symmetric polynomials $f(x, y)$ in which $f(x, y) = f(y, x)$, or homogeneous polynomials in which all the terms are of the same degree.

Given such, operationally, polynomial is used for expressing other functions, due to their usefulness because of such simplicity and versatility. In fact, applications ranges, even to a polynomial basis on n -dimensional space, $(p_1, \dots, p_n) \in K[x, \dots, x_n]$; their growth rate being polynomial also added more to their strength of approximating fast increasing function, usually expressed and used in setting, for example, polynomial growth class $\mathcal{O}(p(x))$ in computer science. Two applications stand out for such properties: function approximation of arbitrary criteria ℓ , or polynomial interpolation on arbitrary points. Those two are interconnected, though their objective and object of concern is much different.

We can define the application of analytical polynomial on arbitrary function by using the criteria on approximation. Usually, this is expressed in the theory of approximation, and requires a guarantee that polynomial can indeed approximate to arbitrary accuracy, of arbitrary function on specific interval. This is the statement of Weierstrass theorem.

Theorem 6.6.1 (Weierstrass Approximation Theorem). *Let $f \in C[a, b]$ of all continuously differentiable functions on the closed interval $[a, b]$. Then, for every $\epsilon > 0$, there is a polynomial p such that $\|f - p\| < \epsilon$, for $\|f - p\| = \sum_{i=1}^{\infty} |f(x_i) - p(x_i)|$ ⁸*

Proof. This is one of the many standard proofs for WAT (Schep [2007] to be specific). We begin by extending f to a bounded uniformly continuous function on $\mathbb{R}$, defining $f(x)$ as to be $f(a)(x - a + 1)$ on the domain $[a - 1, a)$, for $f(x) = -f(b)(x - b - 1)$ on the domain $(b, b + 1]$, and $f(x) = 0$ on $\mathbb{R} \setminus [a - 1, b + 1]$. There exists, in particular, $R > 0$ such that $f(x) = 0$ for $|x| > R$. Let $\epsilon > 0$ and M such that $|f(x)| \leq M$ for all $|x|$. let $\epsilon > 0$ and M such that $|f(x)| \leq M$ for all x . Then by the above theorem there exists $h_0 > 0$ such that for all $x \in \mathbb{R}$, we have

$$|S_{h_0}f(x) - f(x)| < \frac{\epsilon}{2} \quad (6.35)$$

Since we also have $f(u) = 0$ for $|u| > R$, we can write

$$S_{h_0}f(x) = \frac{1}{h_0\sqrt{\pi}} \int_{-R}^R f(u) \exp - \left(\frac{u-x}{h_0} \right)^2 du \quad (6.36)$$

On $\left[-\frac{2R}{h_0}, \frac{2R}{h_0}\right]$ the power series of e^{-v^2} converges uniformly, so there exists N such that

$$\left| \frac{1}{h_0\sqrt{\pi}} e^{-\left(\frac{u-x}{h_0}\right)^2} - \frac{1}{h_0\sqrt{\pi}} \sum_{k=0}^N \frac{(-1)^k}{k!} \left(\frac{u-x}{h_0} \right)^{2k} \right| < \frac{\epsilon}{4RM}. \quad (6.37)$$

for all $|x| \leq R$ and all $|u| \leq R$, since in that case $|u - x| \leq 2R$. This implies that

$$\left| S_{h_0}f(x) - \frac{1}{h_0\sqrt{\pi}} \int_{-R}^R f(u) \sum_{k=0}^N \frac{(-1)^k}{k!} \left(\frac{u-x}{h_0} \right)^{2k} du \right| < \frac{\epsilon}{2}. \quad (6.38)$$

for all $|x| \leq R$. If we put

$$P(x) = \frac{1}{h_0\sqrt{\pi}} \int_{-R}^R f(u) \sum_{k=0}^N \frac{(-1)^k}{k!} \left(\frac{u-x}{h_0} \right)^{2k} du, \quad (6.39)$$

⁸We again note the disparity between $h(x) - c(x)$ or $c(x) - h(x)$ for clarity. In essence, $c(x) - h(c)$ is looking from the perspective of the concept itself. While $h(x) - c(x)$ is comparison on difference from the perspective of the hypothesis.

then $P(x)$ is a polynomial in x of degree at most $2N$ such that $|S_{h_0}f(x) - P(x)| < \varepsilon/2$ for all $|x| \leq R$. This implies that $|f(x) - P(x)| < \varepsilon$ for all $x \in [a, b]$. $\square$

Here, ϵ is defined to be the *global error* of the polynomial with respect to the arbitrary function, and with uniform approximation. Another variation works on the squared error, such that $\|f - p\| = \sum_{i=1}^{\infty} (f(x_i) - p(x_i))^2$, or the mean absolute

$$\|f - p\| = \lim_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n (f(x_i) - p(x_i))$$

Weierstrass Approximation Theorem provides tighter bound for the mean absolute error, but not the square error, or the mean square case. Nevertheless, they are natural measure on the uniform metric space $C(X)$ with measure $d(f, g) = \sup_{x \in X} |f(x) - g(x)|$. If said metric also use the L^p metric instead, which is defined by

$$d_p(f, g) = \left(\int_X |f(x) - g(x)|^p dx \right)^{1/p} \quad (6.40)$$

we can also see that it is weaker than the bound given by Weierstrass theorem. Generalization of Weierstrass theorem can be made by Stone-Weierstrass theorem instead, of which is state in real space as followed, though generally Weierstrass theorem is sufficed.

Theorem 6.6.2 (Stone-Weierstrass Theorem (real version)). *Let S be a locally compact Hausdorff space. Let A be the subset of bounded continuous real functions on X , or of subalgebras $\mathcal{C}(X, \mathbb{R})$ which contains a non-zero constant function, then A is uniformly dense in $\mathcal{C}(X, \mathbb{R})$ if and only if it separates points.*

In the rigorous proof, or so is the proof for Theorem 6.6.1, the only objective is to establish the existence in which such polynomial might exist, but talk nothing effectively about how to obtain such polynomial, or the feasibility of such model. Indeed, in Bernstein's constructive proof (Sury [2019], Muresan [2023]), the proof utilized the construction of the monomial-basis spanning sub-family of the Bernstein polynomial, and show that converges uniformly to f arbitrarily defined on the mapping space. Such is a sufficient proof, in which Bernstein's family is dense in $C[0, 1]$ and thus specify the polynomial structures itself; but overall, there exists no necessity proof on such matter. Such a necessity theorem does not exist, because WAT is usually expressed as *far too weak* to enforce such generality. A heavier version must be used, for example, Korovkin's on the sequence of positive linear operators, but otherwise this is inherently lacking. Expressively, this is the same for the dynamics and how different families of polynomial will emerge or moves under dynamical process and true approximation. However, the theorem is still useful in its sense of the universality, thus we will still use it, but *only* so.

The setting of polynomial interpolation is rather different from the setting of function approximation. The setting of interpolation gives an n samples set S , of which in this case we restrict to $\mathbb{R} \rightarrow \mathbb{R}$. Then, we are required to fit exactly the entire dataset, such that the following setting is sufficed.

Definition 6.6.3 (Interpolation). A function class $\mathcal{F}$ is said to **interpolate** a distinct dataset $S_n = (x_n, y_n)$ of n samples, if there exists $f \in \mathcal{F}$ such that $\mathbb{E} \|S_i - f(x_i)\| = 0$. Such f is then called the **interpolating function** of S_n .

Under such definition, the question is then if the polynomial class $\mathcal{P}_n$ of any given polynomial with the highest degree n is sufficed to approximate such dataset. The answer can be found using unisolvence theorem.

Theorem 6.6.3 (Unisolvence theorem). *Given $n + 1$ bivariate data points $(x_0, y_0), (x_1, y_1), \dots, (x_n, y_n) \in \mathcal{S}_n \subset \mathbb{R}^n$, with all x_i distinct, there exists a unique polynomial $p_n(x)$ of degree at most n such that $p(x_i) = y_i$ for $i = 0, 1, \dots, n$.*

In determining such problem, we can also say that there exists a unique n -saddle point function represented here as polynomial, such that to fit n points perfectly by the definition of the saddle point. The process in obtaining such formula is difficult, however, and finding the perfect interpolation polynomial would be computationally significant, and erroneous. The following definition specifies the concept of such saddle point.

Definition 6.6.4 (Saddle point, Nie et al. [2021]). Let $X \subseteq \mathbb{R}^n$, $Y \subseteq \mathbb{R}^m$ be two sets (for dimension $n, m > 0$) and let $F(x, y)$ be a continuous function in $(x, y) \in X \times Y$. A pair $(x^*, y^*) \in X \times Y$ is said to be a saddle point of $F(x, y)$ over $X \times Y$ if

$$F(x^*, y) \leq F(x^*, y^*) \leq F(x, y^*), \quad \forall x \in X, \forall y \in Y \quad (6.41)$$

Effectively speaking, each interpolation degree of the polynomial interpolation scheme admits a specific (x^*, y^*) of which that the convexity is absolved. Thus, the reason why you require n saddle points for n points, or d -degree polynomial for $d + 1$ points. The landscape of such interpolation is "wiggling", as indicted, since for any two interpolation point, the macroscopic details would be similar as Figure 6.3, of such extend for any local monotonicity on n saddle points, or basically a secant line.

x	a	x_i^*	x_{i+1}^*	b
$\mathcal{D}_1[p]$	—	0	+	0
$p_1(x)$	$p(a)$	y_i^*	y_{i+1}^*	$p(b)$
$\mathcal{D}_2[p]$	+	0	—	0
$f_2(x)$	$p(a)$	y_i^*	y_{i+1}^*	$p(b)$

Figure 6.3: Table of monotonicity for the saddle point landscape between $[a, b]$, for both cases of the landscape.

Polynomial regression is then one such analytical solution to function approximation task, using Weierstrass theorem as its basis. From such, we can also observe that polynomial regression, and hence polynomial interpolation, are correlated and can be reduced to each other. Polynomial regression can be reduced to interpolation via stating the criteria such that, denoting it as ℓ , then $\ell(f, p) = 0$, and interpolation can be reduced to regression via the condition that $p(x_i) - y_i \leq \epsilon$ for arbitrary $\epsilon > 0$. Hence, the regression formulae lie between the two extreme: the Weierstrass setting and the interpolation setting.

We then have to identify such extreme and what kind of extreme they represent. Fortunately, we know this as the problem of **transparency**, in the setting of approximation, which certainly correlates to the learning theory. According to such, the Stone-Weierstrass represents the total transparency case, where the function f is known to the approximator with polynomial class. Interpolation represents total black box, where the goal is only to interpolate the sample point by itself. In between such interpretation, there exists a fundamental assumption, hence the structural assumption.

Assumption 6.6.4 (Transparency). *For any bivariate pair (x, y) , we assume each pair depends on and originate from a based concept $y = f(x)$.*

Two kinds of knowledge is then available, which correlates to the degree of freedom the observation space can have. First is the knowledge of f . This knowledge of f is encoded directly by the number of sample point and their distribution in certain case. Accordingly, the worst case is as followed.

Conjecture 6.6.1. *If $\forall x$ there exists y corresponding to such, and $(x, y) \in \mathcal{S}$, then we say we have total transparency to the concept f . If there exists only finite x that we can have y corresponding to, then we say we have limited transparency to the concept f . If we assume no knowledge of f , then we say the setting is called the black-box setting.*

Secondly is the action on the dataset itself. One of the main assumption that is usually observed is that

Assumption 6.6.5 (Fallacy). *We assume there exists every specific bivariate pair $(x_i, y_i) \in f[\mathcal{S}_n]$, such that $f(x_i) = y_i$.*

However, usually, this is not the case, as we have **noise** in addition from the dataset. The origin of the noise cannot be tracked. However, we can interpret what happens when such noise is introduced to the dataset. The noise transitions the transparency of the learned setting, on the field $\mathbb{R}$, by utilizing the dense property of the real space. Assuming we operate on $\mathbb{Z}$. Then, each discrete noise will introduce abnormality point such that no function on $\mathbb{Z}$ will be able to match, up to certain distance on $\mathbb{Z}$. This is because on such space, there exists no point x_1, x_2 , for example, that lies between $d(x_1, x_2) = 1$. If the noise is randomized on ± 1 , then there will exists no function that can approximate such disparity better than a random fluctuating function. However, on $\mathbb{R}$, the real number is inherently dense, such that there will always be certain point x_3 such that $x_3 \in [x_1, x_2]$, by the density of real number. By the distance metric on $\mathbb{R}$, we define $d(X, Y) = |X - Y|$ for distance between two data points. Noise intrinsically increases such distance, such that $d'(X, Y) > d(X, Y)$. Hence, the space in which density can appear increase, thus there then exists more fluctuation, and thus more uncertainty between such two points. This notion can be increased from local uncertainty to global uncertainty, hence the *noisy dataset*.

We then can measure or aggregate some properties of the dataset itself, including the **uncertainty** of the dataset. We define the simple uncertainty of the dataset as followed.

Definition 6.6.5 (Dataset uncertainty). Given a closed interval $[a, b]$, for the sample set $\mathcal{S}_n = (X, Y) \in [a, b]$, the uncertainty $\mathcal{U}(\mathcal{S}_n)$ is defined as

$$\mathcal{U}(\mathcal{S}_n) = \left[\mathbb{E}_{x_i \in X} (x_{i+1} - x_i)^2 + \mathbb{E}_{y_i \in Y} (y_{i+1} - y_i)^2 \right]^{1/2} \quad (6.42)$$

We assume X, Y are totally ordered sets.

Hence, the assumption of fallacy assumes $\mathcal{U}(\mathcal{S}_n) = 0$, though in reality it is often $\mathcal{U}(\mathcal{S}_n) > 0$. We would like to clarify, that such definition is currently reserved for only polynomial-like functional class; though, this can be extended pretty easily for any $\mathbb{R}^n \rightarrow \mathbb{R}^m$ functional concepts. Furthermore, it is *operational*, in the sense that it measures the numerical encoding inherent uncertainty, of which is also the space used of algorithms traversing in stochasticity. It affects cross-validation, gradient-based algorithm path finding, and roughness of the learning processes. Certain concerns must also be addressed, such as

the definition, which assumes *sequential processing* and pair-wise measurement (which can also be extended to tri-fold measurement and so on). Indeed, modern deep learning use parallel batch processing, mini-batch sampling, or distributed learning across devices for optimization. Nevertheless, such includes sequential-based learning optimization, and so on, where mini-batch still inherently includes batch ordering; algorithms like RNNs and so on also are sequential. The computational cost per dataset, if not, is $\mathcal{O}(n)$ after initial order, or with mundane immediate measurement on the way that the dataset was configured at start, since the system and stochastic algorithms sequentially go through them.

Certain others aspects can be expected to be clarified. While indeed, this is order-dependent measure, and one might expect to go against it and say it is not real dataset property, but the key point is that dataset are measurements, captured and recorded into the encoded space that is used upon the learning regime (supervised or semi-supervised or so on), or the learning algorithms. In a classical algorithmic world, such must be done in discrete steps, and hence ordering strongly matters as well as the processing pipeline relevant to it. There exists natural ordering of data within noise ϵ , because if now, the projection on embedded space $\mathbb{R}^n \times \mathbb{R}^m$ of any given concept functional c holds no value, as the function itself, plus the distribution, with respect to the encoding space where the function mapping is defined already contains structural information encoded within 'value range' as the most basic form. A weaker posture we can adopt is that it is an operational uncertainty measure, given presentation order to a finite discrete processing machine, and is indeed *pre-training* diagnostic. There is also the issue in which X and Y might not be measurable or comparable. Indeed, since the range of value can jump intermittently between different concept class and form, but even then, it will only highlight the disparity between the contributing terms. For example, if we have three pairs, $(130, 1)$, $(150, 1)$ and $(225, 0)$ as the data, then the volatility is not that bad, and stable on Y , but *very high on X* , as there are missing information in large gap. From the machine-algorithmic, this matters. if we take the 130, 225 points as a case of shuffling, then the machine experiences a 95-to-1 volatility ratio. Nevertheless, one can attempt a normalized metric, i.e.

$$\mathcal{U}_{\text{norm}}(\mathcal{S}_n) = \sqrt{\mathbb{E}[(\Delta x_i)/r(X)]^2 + \mathbb{E}[(\Delta y_i)/r(Y)]^2} \quad (6.43)$$

for relative jumpiness accorss different problems with normalized morphism.

With this, we can quantify how the dataset is behaving on the scale of the bivariate form typically seen. Noise intrinsically increases this distance, usually by the y -axis. This dataset uncertainty increases with multivariate dataset, for example, $(\mathbf{x}, y)$ of $\mathbb{R}^n \rightarrow \mathbb{R}$. We call the new metric as the *empirical uncertainty metric*. The basis of such argument on the concept of a tradeoff comes from Nakkiran [2019a], in which linear regression is affected, at least in such setting and consideration, such that dataset will both hurt and benefit the learning process.

Overall, however, this allows us to state of the following theorem:

Theorem 6.6.6 (Separation Encoding under Diffusion). *Fix the encoding space $\mathbb{D} = \mathbb{R}^q \times \mathbb{R}^p$ for a (q, p) conjunction space. Let $\mathcal{S}_n = \{(x_i, y_i)\}_{i=1}^n \in \mathbb{D}$ be a dataset of n ordered samples drawn from an oracle $EX(\theta)$. Let $\{\mathcal{N}\}$ denote the set of noise configurations $\eta = \{\eta_i\}_{i=1}^n$ where each $\eta_i \sim \mathcal{D}$ with $\mathbb{E}[\eta_i] = 0$ and $\mathbb{E}[\eta_i^2] = \sigma^2 < \infty$, independent across components. Define the diffused dataset $\mathcal{S}'_n = \{(x_i + \eta_i^{(x)}, y_i + \eta_i^{(y)})\}_{i=1}^n$. Then*

$$\mathbb{E}_{\eta \in \{\mathcal{N}\}} [\mathcal{U}(\mathcal{S}'_n)^2] = \mathcal{U}(\mathcal{S}_n)^2 + 4\sigma^2 \quad (6.44)$$

In particular, for any $\sigma^2 > 0$ and any zero-mean finite-variance $\mathcal{D}$,

$$\mathbb{E}_{\eta \in \{\mathcal{N}\}} [\mathcal{U}(\mathcal{S}'_n)^2] > \mathcal{U}(\mathcal{S}_n)^2 \quad (6.45)$$

Proof 1. Simple proof. Using normed metric, $\mathcal{U}_{\text{norm}}$. For noised data, that is,

$$\mathcal{S}'_n = \{(x_i + \eta_i^{(x)}, y_i + \eta_i^{(y)})\}_{i=1}^n \quad (6.46)$$

where the noise vector on encoding space is orthogonal, such that

$$\eta_i = (\eta_i^{(x)}, \eta_i^{(y)}) \sim \mathcal{D} \approx \mathcal{N}(a, \lambda \text{Cov}(\eta)) \quad (6.47)$$

where we can fix $\mathcal{D} = \mathcal{N}(a, \lambda \sigma^2)$ for $a = 0$ as normalized around the zero value. If we assume independent noise, then we have the covariance as

$$\lambda \text{Cov}(\eta) = \lambda \begin{pmatrix} \sigma_x^2 I_q & 0 \\ 0 & \sigma_y^2 I_p \end{pmatrix} \quad (6.48)$$

We can simplify slightly, as to have the case in consideration that $\sigma_x^2 = \sigma_y^2 = \sigma^2$, then the notation would effectively be such that $\lambda \sigma^2 I_{d+p}$ is enough. The uncertainty metric then should be expressed as

$$\begin{aligned} \mathbb{E}[\mathcal{U}(\mathcal{S}'_n)^2] &= \mathbb{E} \left[\mathbb{E}_{x_i \in X} (x_{i+1} - x_i)^2 + \mathbb{E}_{y_i \in Y} (y_{i+1} - y_i)^2 \right] \\ &= \mathbb{E}_i [(x'_{i+1} - x'_i)^2] + \mathbb{E}_i [(y'_{i+1} - y'_i)^2] \end{aligned} \quad (6.49)$$

for $\mathbb{E}_i$ here denotes the expectation over the uniform index i of each component of the sample space. Now, we substitute the following into the empirical average pair:

$$x'_{i+1} - x'_i = (x_{i+1} - x_i) + (\eta_{i+1}^{(x)} - \nu_i^{(x)}), \quad y'_{i+1} - y'_i = (y_{i+1} - y_i) + (\eta_{i+1}^{(y)} - \nu_i^{(y)}) \quad (6.50)$$

First, for x_i, x_{i+1} set, taking expectation over noise η gives

$$\begin{aligned} \mathbb{E}_\eta [(x'_{i+1} - x'_i)^2] &= \mathbb{E}_\eta [(x_{i+1} - x_i)^2 + 2(x_{i+1} - x_i)(\eta_{i+1}^{(x)} - \eta_i^{(x)}) + (\eta_{i+1}^{(x)} - \eta_i^{(x)})^2] \\ &= \mathbb{E}_\eta [(x_{i+1} - x_i)^2 + (\eta_{i+1}^{(x)} - \eta_i^{(x)})^2], \quad \mathbb{E}_\eta [2(x_{i+1} - x_i)(\eta_{i+1}^{(x)} - \eta_i^{(x)})] = 0 \\ &= \text{Var}(\eta_{i+1}^{(x)}) + \text{Var}(\eta_i^{(x)}) \\ &= 2\sigma^2 \end{aligned} \quad (6.51)$$

for x_i fixed and $\mathbb{E}[\eta_i^{(x)}] = 0$ that leads to the second zero equality, independent noise $\eta_{i+1}^{(x)}, \eta_i^{(x)}$ with variance σ^2 . Therefore, we have for each i pairing, that

$$\mathbb{E}_\eta [(x'_{i+1} - x'_i)^2] = (x_{i+1} - x_i)^2 + 2\sigma^2 \quad (6.52)$$

Averaging over i and apply the same argument to y -component, we gain identical form with doublecondition, such is,

$$\mathbb{E}[\mathcal{U}(\mathcal{S}'_n)^2] = \mathbb{E}_i [(x_{i+1} - x_i)^2] + \mathbb{E}_i [(y_{i+1} - y_i)^2] + 4\sigma^2 = \mathcal{U}(\mathcal{S}_n)^2 + 4\sigma^2 \quad (6.53)$$

Since $\sigma^2 > 0$, the expression gives $\mathcal{U}(\mathcal{S}_n)^2 + 4\sigma^2 > \mathcal{U}(\mathcal{S}_n)^2$ by Jensen inequality. The statement works under expectation in the squared sense directly, and for $\mathcal{U}$, we expect monotonicity on $\mathbb{R}_{\geq 0}$ for square terms reduces to the expression on non-squared condition. $\square$

Specifically, the unsquared version will give

$$\mathbb{E}[\mathcal{U}(\mathcal{S}'_n)] \leq \sqrt{\mathbb{E}[\mathcal{U}(\mathcal{S}'_n)^2]} = \sqrt{\mathcal{U}(\mathcal{S}_n)^2 + 4\sigma^2} \quad (6.54)$$

which again, requires the argument going through the squared form by condition. The condition that invalidates the theorem under such form would be for any non-zero mean noise, for $\mathbb{E}[\eta_i] = \mu \neq 0$, which does not collapse the contributing terms. The middle $2\mu \cdot \mathbb{E}_i[\cdot]$ term then is a telescoping sum proportional by $(x_n - x_1)/(n - 1)$, which can be either positive or negative depending heavily on trend in X space, so biased noise can actually **increase or decrease** uncertainty extrapolated depending on dataset structure. This is rather interesting, because it highlights a particular region where systemic bias (or so we called of such), interacts with dataset ordering in a non-trivial way, which is far harder to concentrate. In fact, the real form under arbitrary notion can be written as

$$\mathbb{E}_\eta[\mathcal{U}(\mathcal{S}'_n)^2] = \mathcal{U}(\mathcal{S}_n)^2 + 2(1 - \rho)(\sigma_x^2 + \sigma_y^2) + \text{Bias} \quad (6.55)$$

of the bias term arbitrary, and covariance $\rho\sigma^2$ for correlation of both η_i, η_{i+1} , which is rather interesting, and targets the case of sensor drift in observational spaces. Finally, the constraining partition is that if we have infinite variance on $\mathcal{D}$, then the finiteness of $\mathcal{U}(\mathcal{S}'_n)$ is not finite and the comparison is undefined. As such, while interesting, and indeed extensible to a wider variety of control space for dataset, we note the inherent understanding and restriction of such uncertainty measure.

Thirdly, is the treatment on such information. Even though information is known, such as from the assumption of transparency and the fallacy conjecture, it relies on the method in which the model uses that make use of such information. Polynomial interpolation effectively removes information about transparency, and thus the interpolation setting is a kind of **black-box absolute setting**. This setting is *extremely brittle*, as to see that if y_i is slightly wrong due to measurement noise, the interpolant oscillates wildly, which is explained by Runge phenomenon (Runge [1901], Corless and Sevyeri [2018], Boyd [2009]). Probability then arises from such setting naturally, though this particular degree of probabilistic setting must be quantified, which comes from the process that the model simulates upon. Regression is inherently not probabilistic, but treatment about the probabilistic assumption, for example, such that (X, Y) inherently as random variables. We can then separate this into three parts – the data-level probabilistic layer (data uncertainty), the model-level layer, and the algorithm-layer that determines the solution solver within such model. We can define that as followed.

Conjecture 6.6.2 (Probabilistic layer). *Given a learning setting with dataset $\mathcal{S}$ of bivariate pairs, $h \in \mathcal{H}$, $c \in \mathcal{C}$ as the hypothesis and concept, we have the following conceptual probabilistic setting:*⁹

1. *Data-level probability: Sampling assumptions (observation sample points) such that there exists the distribution $\mathcal{D} \sim (X, Y)$, intrinsic data uncertainty $\mathcal{U}(\mathcal{S}_n)$.*
2. *Model-level probability: Structure of the dataset is controlled by the concept structure $\{\theta\} \in c$, c is modelled as a target sampler $c \equiv \mathcal{D}$, knowledge about the structure of c , Bayesian prior-posterior assumptions, $\mathbb{P}(Y | X, \{\theta\})$. For Bayesian polynomial regression, the model-level probability is formulated as*

$$\mathbb{P}(Y | X) = \int \mathbb{P}(Y | X, \theta) \mathbb{P}(\theta) d\theta, \quad \mathbb{P}(\theta_j) = \mathcal{N}(0, \sigma^2) \quad (6.56)$$

of which assumes aleatoric uncertainty and epistemic uncertainty from the weight.

⁹Currently, we have no way to quantify the ‘probabilistic degree’ of the system setting itself

3. *Algorithm-level probability: Stochastic method, decision conditional (dropout, MCMC), ensembles criteria, Monte-Carlo initialization, variational inference.*

Dynamically, in this paper, we would like to investigate this particular junction between uncertainties, to see what happens when the model transparency, the treatment of data, and so on, that exhibits the observations that we would receive from the model itself. The experimental details will infer such result.

Generally, we would be entitled to see particular trajectory specifically designed around the point of interpolating match, or $n = m$ of n degree of freedom, technically, and m data points. After this point, there exist two possibility of dynamic – either the model dilutes, in which case result in *double descent*, or the model fluctuate strongly afterward, leading to *bias-variance emergence* again later on. This conclusion however, would have to be analysed within the current framework of the possible interjection to its fluctuation strength. For example, in the monomial basis $\{x^i\}^n$, any given fluctuation contributes to the entire polynomial. In such case, fluctuation of the other $[0, n - k]$ of arbitrary k and large enough n would be unaccounted for by the scale. Conversely, restrict them to $[1, 0]$ range will reveal something about the system, however will destroy particular properties of the prediction, which is not recommended. Nevertheless, we must conduct the test before giving in any reasonable conclusion.

Previously, we said that we did indeed focus on the interpolation problem. Now, assuming we have given the interpolating algorithm $\mathcal{A}$ to approximate the interpolation, what will this look like? Firstly, solving this is *very computationally expensive*, because it is a dense, $(n + 1) \times (n + 1)$ linear system, and even more so if we consider the overparameterized outcome. However, this can be reduced further down if we attempt to solve them using convergence bound, but will be still very expensive. That and incremental soft interpolation is somewhat out of the question since it will be kind of the same. Aside from the insight, we would still have to know what to do. Then there exists two ways – we have some version of this to optimize of simply evaluating by gradient descent optimization, or, we can use somewhat analytical fit for the case of soft interpolation modification.

A theory on such problem can be obtained of its mechanics, as will be a fairly intuitive way in consideration of the theory of interpolation, as well as the technique of interpolation by itself. For a fixed-range interpolation threshold problem, a hard-encoded interpolation solution requires n -degree polynomial, of which has $n + 1$ degree of fitness, to fit $n + 1$ distinct points. Suppose all points are distinct of the set S which is finite, then this is the case. If we switch to the problem of interpolation but induce a more forgiving bound, there exists a possibility that it induces a relative *density averaging path* for the points, especially if we are considering the case where $p > n$ for p saddle points or degree of fitness, and n the number of point to interpolate. In a more nuanced formulation, we have

$$\text{Flatten} \left(\begin{bmatrix} 1 & x_0 & x_0^2 & \cdots & x_0^p \\ 1 & x_1 & x_1^2 & \cdots & x_1^p \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_p & x_p^2 & \cdots & x_p^p \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_n & x_n^2 & \cdots & x_n^p \end{bmatrix}_{n \times p} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ \vdots \\ a_p \end{bmatrix} - \begin{bmatrix} f(x_0) \\ f(x_1) \\ f(x_2) \\ f(x_3) \\ \vdots \\ f(x_n) \end{bmatrix} \right) \leq \epsilon \quad (6.57)$$

or equivalently, $\text{Flatten}(\mathbf{X}_{n \times p} \mathbf{w}_n - \mathbf{Y}_n) \leq \epsilon$ where $\text{Flatten}(\cdot)$ is the compression of vector to scalar, or so $\text{Flatten} : \mathbb{R}^n \rightarrow \mathbb{R}$. This essentially is the problem of solving n points approximation, toward using only p saddle points available. If $\epsilon = 0$, it is well known that interpolation is impossible. However, if $\epsilon \neq 0$, there are possibilities, depends on the shape of the dataset, that there will exist a data-dependent

property or more that indict a particular *density preference* toward those saddle points – or rather, the act of density averaging. This is reflected rather implicitly in the structure itself. In machine learning, the error measure is usually the mean average, such is to say that the criterion now would look more like $\mathbb{E}[\text{Flatten}(\mathbf{X}_{n \times p} \mathbf{w}_n - \mathbf{Y}_n)] \leq \epsilon$. Formally, we have the following hypothesis

Hypothesis 6.6.1 (Polynomial double descent). *For a polynomial $\{a_i x^i\}_n$ on Vandermonde basis, approximating on the set $\{(x_k, y_k)\}_m$, then*

1. *For $n < m$, we expect the polynomial form factor to be *inherently sparse*, such that for an iterative approximation, for each p_i saddle point of the polynomial expansion form, there exists an ϵ -ball B_ϵ for $\epsilon > 0$ with at least one data point satisfying $(x_k, y_k) \in B_\epsilon$.*
2. *For $n = m$, we expect the polynomial to be able to fit all points, up to hard interpolation threshold such that $\hat{R}_S \approx 0$.*
3. *For $n > m$, we expect the polynomial form factor to be *locally dense* then turning to *globally dense*, over random ensembles. That is, for every point (x_k, y_k) there exists a β -ball B_β such that there is at least one saddle point of the polynomial belong to the ball.*

The shift that will be hypothesized to appear here is expected to be from sparse in the polynomial term sense, out to globally dense with enough model complexity, over an iterative process. One also main point of clarity is, at certain point right now, we haven't discovered the effect of error deviation, yet. For such, we theorize the following.

Conjecture 6.6.3 (Double descent). *Double descent happens, for a polynomial class model $h_n \in \mathcal{H}_p$, is upper bounded to certain error deviation $\epsilon > 0$ from the mean only of the generalization (test) dataset verification.*




Figure 6.4: **Intuition of Hypothesis 6.6.1.** Effectively, in the overparameterized range ($p > n$), the problem turns into soft interpolation, where the polynomial is possible to be dense enough such that relative error remains internally low. This can also happen if the dataset itself is intrinsically large.

That is, the disparity between the test and training dataset helps in such regard. If the two dataset is partitioned in the sense that their deviation is practically the same (which can be done using synthetic generation for a control group), then this will amount to not much of it. After all, it will be the consequence of the algorithm and error landscape interpretation used.

For this, we propose another speculative observation beside such, as to consider adjourned transform of the Vandermonde basis to the Newtonian or similar basis, in which the effects and motion can be easily observed, as least easy. This is because aside from mapping down to $[-1, 1]$, normal polynomial on monomial basis gives suboptimal result because of intrinsic high-degree term domination. This will become more relevant later on in the analysis.

Chapter 7. Modelling neural network

We have been

7.1 Multilayer perceptron

Neural architecture, and furthermore of neural network, presents an interesting topic of analysis, in which provide a rather rich experimental opportunity for any given admission. As such, we would like to take on testing neural network in a more complex analysis, but with normalized or simplified case basis, for example, analysing the *semantic* and features of the model evolution. Before that can happen, we would like to cover our convention and model intervention. As seen of [Cybenko \[1989\]](#), [Hornik \[1991\]](#), neural network can theoretically capture and represent a very wide class of functional encoding in closed, finite regime, furthermore with even Turing computability ([Siegelmann and Sontag \[1991\]](#)). Of such, trainable neural architecture like neural network layering, offer NP-complete or $\exists\mathbb{R}$ -complete training complexity, and the perceived power of more depth, layering as seen of [Eldan and Shamir \[2016\]](#), [Telgarsky \[2016\]](#). The question of double descent though, lies in the domain of the problem upon *reaching such approximation goal* instead. This in question, includes computational complexity. If we are to treat that the above presumption on complete training complexity is right, the problem is rather hinging on a dilemma, of which we defer to previous sections on decidability.

Suppose bias-variance applies, and we do not have double descent in question, thus position ourselves in the stead of [Belkin et al. \[2019a\]](#) when they first discovered double descent. However, what we would be doing is task admission-free, and consider a finite setting on approximation. Then, for any measure of arbitrary model complexity $\mathcal{M}(\mathcal{N})$ on network $\mathcal{N}$, results from the neural network construction, we create a set of $\{\mathcal{N}_i\}$ with variational complexity. The task is now to see the behaviours of this, from the viewpoint of structural analysis on different construction of network $\mathcal{N}$, and evolutionary dynamics. Since there are many ways to construct $\mathcal{M}$, what we would be doing is to connect error decision template to different model complexity measure. The level of changing model complexity would be one-level above training regimen and empirical test generalization.

7.2 Experiments

7.2.1 Experimental A - Conditioned MLP

Let us try an experiment on the matter. We fix the structure of the neural architecture to the typical layered neural network per accordance of our analysis on multilayer perceptron. Thus, for a network $\mathcal{N}_{m,\{n\}}$, specify by m layers of n_i , $i < m$ number of per-layer nodes (we implicitly assume homogeneity of layer construction, thus variation difference only matters for layer themselves), the construct $\mathcal{N}$ is single-directed, forward-only in operation, and consider modification post-evaluation on training or

optimization to be *post-operation* of state transformation. That is, the network itself does not evolve, but is externally modified and hence exists no state nor phase transition structure. We can then proceed with cutoff in the same way of polynomial network, but this time much more better. We classify three specific set of components:

1. Input components that feeds the system, we denote such special structure or ambiance context as $n_i^{\rightarrow}$ for ‘in channel’ feed.
2. Structural components of the neuron to be separated into three parts: weighted transceiver T_{kj} of the j neuron unit of layer k , mechanistic interpretation $M_{kj}(T_{kj})$ that contains, in the most plain singular form the activation function, and P_{kj} as the **outer pushout** (i.e. can be said as output) of the neuron component. We admit such modular structure for all components.
3. Ordering of neuronal unit. Here, we fix this to be constructed via **layer order**, such that there exists the acyclic graph from layer L_1 to L_2 up to L_m , not taking into account the input channel. There, there exists no cross-component connection between components of the same layer by acyclic condition, and the order of computation is the same as the directed acyclic graph – all computation are independent of another in-layer, and proceed only afterward when all layer-specific operations are completed. While constraints might matter, typical multilayer perceptron is *fully-connected*, that is, processed information from L_1 is received in entirety of every component in L_2 , and so on. A looser condition suggests the usage of n th order information transceiver. That is, for a node $N_{kj}(T, M, P)$ of neural unit from layer k , then at best connection is guaranteed such that transformation of layer L_{k-j} will be accessible to the of at best n th layer processing connection, assuming the network exists sparse connections. Another condition for the entire network is that all connection must be closed, that is if it starts, it must end of the last layer.

Simply said, we decompose the network out on itself into multiple components. About the transceiver concept, we can formalize such insight as followed of a construction diagnosis.

Definition 7.2.1 (Order- n transceiver). A neural architecture $\mathcal{N}_{m,\{n\}}$ admits an order n transceiver structure, iff for every neuronal unit N_{kj} , there exists a subset of layers

$$\mathcal{A}_{kj} \subset \{L_{k-1}, \dots, L_{k-n}\} \quad (7.1)$$

such that for all $L_\ell \in \mathcal{A}_{kj}$, the output pushouts $P_{\ell,*}$ are admissible inputs to T_{kj} , and no admissible inputs originate from $L_{<k-n}$ to be guaranteed on randomized model initialization, as to mean as determined solely by the architectural connectivity, independent of parameter realization under random initialization. Inversely, for a potentially sparsely configured, layer-based network definition, then the **inverse order n** condition is applied for $\mathcal{N}_{m,\{n\}}$ iff it can guarantee that

$$\mathcal{A}_{kj} \subset \{L_1, \dots, L_n\} \quad (7.2)$$

is admissible to arbitrary layer $L_{>n}$.

The condition at the final end of the neural architecture varies, but usually, approximation requires it to reduce to one. If we are to test multi-to-multi feature mapping, for example, $\mathbb{R}^q \rightarrow \mathbb{R}^l$, this would be helpful, however, the choice of canonical example is such that $l = 1$. Order- n transceiver also pushes another structural preference, that is, equipping traffic restriction such that each component only has g channels count to connect to the next layer aside from closure or beginning input sections of the network. Implicitly, this has justification on a kind of model complexity – the *highest admissible channel count*, and

enumerative total channel available of the network. That said, it does not mean the network can, or should connect all channels. In fact, we might as well in the future devise a connection-aware optimization scheme, though still external, will gauge network wiring configuration altogether, almost in the sense of inverse-dropout, though it can also be said to be dropout if we impose removal condition. The illustration of the network is shown in Figure 7.1. That said, for experiment, we have a few initialization setting that




Figure 7.1: **Diagrammatic illustration of the decomposed neural architecture.** In case of non-sparsity, all the graph would have to be fully-connected, and the order of the network is m . The canonical construction on $l = 1$ imposes final compression to (T_F, M_F, P_F) , which is then the admissible form used for evaluation of prediction on based input setting.

can be used. Explicitly, we do not use conjectured structure above of learning behaviours on wiring change, or order complexity. We do it the **old-fashioned way**, but with component-awareness. So, for now, we provide procedure:

Definition 7.2.2 (Standard oracle procedure). Construct an oracle that can be admitted of a form, either in full or best approximation of the decomposed neural architecture (component-based distinguishability). Choose finite oracle context semantic by the size of context per choice n , specify $|\mathbf{X}|_{\mathbb{E}} \in D$ of the domain of interest (if $n = 1$, then $D \subset \mathbb{R}$). Next, **corrupt** the oracle by explicitly setting up $\xi - 1$ -scenarios of removing $\{N_{qh}\}_{i \leq k}$ set of neuron components or equivalent approximated representation of the oracle as the **1st-order corruption**, or corrupting each if not more of (T, M, P) subcomponents, hence $\{(T_{qh}, M_{qh}, P_{qh})\}_{j \leq k}$ for **2nd-order corruption**, and intrinsic noisy interference $\epsilon_{rkj} \sim \mathcal{D}$ for $r = \{1, 2, 3\}$ of the three subcomponents; such can be combined into ϵ_{kj} of the neuron unit for simplicity. Canonical choice for ϵ is $\epsilon_{kj} \sim \mathcal{N}(1, \sigma^2)$. Then, define a Poisson distribution measure $\text{Poisson}(\lambda, k)$. Define a deterministic intensity profile, $\{\lambda_Z \geq 0\}$ of gauging the notion of corruption level in this ledger, thus $\lambda_Z = a - bZ$ for arbitrary $a, b \in \mathbb{Z}$. Then, sample of respective ξ -scenario, the amount of data from each configuration you require, using Poisson, for

$$n_i \sim \text{Poisson}_S(\lambda_i), \quad (7.3)$$

from the distribution specified (can be changed, on consideration). After this, we have the total size of the dataset by $N = \sum n_i$. For each of the ξ -category or corruption model, sample the mixture by counts specified above. Finally, add intrinsic 3rd-order noise to complete the set. The dataset now contains $\xi - 1$ scenario on corruption, and 1 on non-corrupted, noise-only scenario. For fully-corrupted set, change to ξ scenario without direct sampling from the oracle.

From this the setting is deliberately simple. We consider constructing the oracle similar to the procedure, then construct a suitable (m, n) -size neural network, first being (1) a family of **multilayer perceptron**, fully-connected with varying shape and size, and then, (2) a family of non-fully connected neural network, or similar functional approximation model type, with different shape and size. Semantic and structural comparison is then available for analysis from the experimenter side. In sparse or near-sparse connection, initialize as normal and impose no g -cap (it is a theorized notion for later expansion on experiment).

Chapter 8. Statistical physics

It is not surprising that the modern landscape, insofar of analysis and deep learning architectures, inspired a particular cohort of researchers with ideas on using **statistical physics** to infer on a learning theoretic system as for it to be a *statistical physics-induced complex system analysis*, i.e. a complex multi-degree of freedom system that statistical physics can take a hold on. The other direction is to apply artificial neural network to simulate a complex interacting systems, and thus verifying a version of mathematical descriptions on it, such as the **duality** of statistical physical descriptions. The list of principle illustrative works on this is not lacklustre in the least, such as

1. [Ferrari et al. \[2024\]](#): Machine learning and optimization-based approaches to duality in statistical physics, i.e. where two mathematical descriptions can describe the same system. The paper used artificial neural network, which benefits from arbitrary component encoding, thus allow for representation on arbitrariness. Similar analysis is used of ML-assisted system in [Hou and You \[2023\]](#), for the renormalization group.
2. [Barbier et al. \[2025\]](#): Works on the statistical physics of deep learning, and sought to analyse deep neural network similar to statistical system, and thus derive an optimal learning scheme on such. Works similar to this can be of comparison, toward **statistical physics interpretation**, such as [Pei \[2025\]](#), [Faulkner and Livingstone \[2024\]](#), or specialized case, most often as **Graph Neural Network (GNN)** like in the case of [Duranthon and Zdeborová \[2025\]](#). An interesting phenomenon or observation on deep learning system, and marginally classical learning systems, is **phase transition**, which is explored in [Ziyin and Ueda \[2022\]](#) for example, and [Ziyin and Ueda \[2023\]](#) for a recent concept of first-order and second-order understanding of deep neural network interpretation for phase transitions.

Utilization of such statistical physics structures onto deep learning includes [Wainwright and Jordan \[2008\]](#) of the Ising or Markov Random Field models, Sherrington-Kirkpatrick (SK) spin glass in [Sherrington and Kirkpatrick \[1975\]](#) for nonconvex loss landscape in ML (as per direct interpretation language); similar purpose is presented by using **p-spin or spherical p-spin models** ([Crisanti and Sommers \[1992\]](#)). Specifically on neural network and deep learning, [Hopfield \[1982\]](#), [Ackley et al. \[1985b\]](#), [Hinton \[2002\]](#) and [Gardner \[1988\]](#) utilized many statistical physics analogues onto neural networks, though in Gardner case it is a **perceptron optimal-capacity model**. In modern time, it is the **cavity model** on high-dimensional understanding, as in [Donoho et al. \[2009\]](#), [Jacot et al. \[2018\]](#) Neural Tangent Kernel (NTK) with statistical inspirations, Langevin mechanics of stochastic dynamics analogue as in [Mandt et al. \[2017\]](#). Energy-based models, as in [Ho et al. \[2020\]](#), [Song and Ermon \[2020\]](#) and entropy-SGD of [Chaudhari et al. \[2016\]](#), mean-field analysis ([Mei et al. \[2018\]](#)), and a rather modern **free-energy interpretation** of modern ML, in [Kingma and Welling \[2013\]](#). It is fair to say that statistical physics abstraction itself gives plenty applications and inspirations on the current modern ML landscape.

However, it is our subjective opinion that is to validate, that statistical physics application, and most of its concurrent analysis are inherently messy, disorganized, and of intrusive mapping from statistical physics onto deep learning or classical systems. We notice a central pattern that typically, statistical analogues are applied onto the machine learning or dynamic model system, without cares and consideration for a direct *formal and conceptual mapping* between the two theories. For example, it is not sure how well all different axiomatic presumption of statistical physics, statistical field theory works on machine learning system, or simple as old classical models with human-aided designs in the 1970s–1980s, or symbolic branches of machine learnable methods. Concepts like temperatures and entropy are typically presented as fact, or rather a list of intrinsic properties gave rise to a ‘natural measure’, yet is not so applicable to a wider degree of the deep learning landscape. Interpretation exploded during the rise of deep learning models, as Goodfellow et al. [2016] also introduced some in his book, since neural network produces a corridor for *isolated component structures*, treated similar to statistical physics without structural difference between the two. As such, there exists the decoupling, that leads to Hinton’s both ‘failure’ and ‘success’ of the Boltzmann neural network (Ackley et al. [1985b], Hinton [2012]), being a failure since it requires to express neural network in a niche, toy-liked, explicitly redesigned, narrow structure, and closed end-to-end loop for equilibrium evolution in-system. Such is often considered to be utilizing a representation scheme, this time, neural network, with fixed descriptions on encoding choice, rather than neural network in general.

Thus, we lay our goal to identify this gap, clarify concepts and deliver a closest possible working the problem, of applying statistical physics scheme onto the ontologically different space of objects. Within reasons, we also utilize such onto the set of **double descent problem**, in which way interpretations make sense and historical patterns or theories, and some experiment to test the hypothesis.

8.1 Note for team writing member

Please skip this page, right of the `clearpage` command right behind this. It is for organization during paper writing, and will be removed afterward.

8.1.1 Main structure

The following outlines the organizational structure in proposal specifically for the statistical physics paper. There are two parts, few claims that we posit in such discussion. Those are the main central point, but can also be mutated when further evolutions are possible.

1. We posit that statistical physics is a quite formal and general method, yet unrefined, and often is applied incorrectly in the framework of perspective of machine learning. There are certain situations (or rather frequently) that statistical physics do describe a deeper interpretation of a dynamical system, but on the other hand, forces on machine learning the **physical perspective** and description, and thus enables false claims, restricted claims, setup that is impossible, of assumptions that never satisfy the theory itself, to force everything into a physical system instead of machine learning, then claiming the reverse is the truth.
2. We claim that to understand statistical physics interpretation to machine learning or system, we need more than apply superficially, that is, to simply do the one-one mapping of this framework’s concept to another.
3. We will highlight such mismatch in the analysis, both in details of interpretation, and both formal way of mathematical expression, proofs, or so on, concretely.
4. We will propose a way to look at such, particularly through the problem of double descent, where

other theories or statistical physics interpretation fails, and the analogous similarity to other previous theories and attempts historically.

5. We then bring up forth a new theory, or rather, a new abstract framework to explain it. Couple with that are some toy experiments, to explain of the system we want to highlight.

Thus, the following sectioning is adopted as the formal form for the document. This includes sections on the introduction of statistical theory, the argument and concurrent (present theory) analysis, then reformulation, then transition to a new theory with experiments.

1. Introduction to Statistical Physics: Simply the introduction of the discussion about the main position, as to what is statistical physics, and the morphism to machine learning, or the learning theory.
 - a) Reasons of statistical physics. This is to dissect *what*, and why statistical physics is even applied, and what does it mean or how to specifically connect to machine learning. This includes the analysis of different *abstraction* level that must be made to apply certain parameters or so on of ML, to fit statistical physics (for example, what is required to *simplify* or *add details* into it such that the notion of temperature is borderline acceptable in ML algorithms?).
 - b) Axioms of statistical physics. What are the axioms and formulation that statistical physics use, from assuming properties of multi-particle interpretation, the dynamic allowed, what kind of dynamics internally and externally, etc., such that statistical physics is expressed most generally? And what can be of the application and mismatch of between?
 - c) Assumptions of statistical physics. Explicit assumptions list, and incompatibility explanation, derivation, analysis. For example, it can be like that independent of components system i.e. the molecular behaviours, which is often not true in an integrated ML system.
 - d) Relevant concepts that it uses in formation, assumptions, already established notions.
2. Problem and settings. Here, it is simple as to *identify the errors*, speaking out loud of the problems, the problems of people solving the problems, the fallacy that make everything fails, what can be retained, what actually works, its limitation and so on.
3. Identification: Strictly identify the problem solutions and identification on "what exactly the problem can be categorized", proposing hypothesis, solutions, there of.
4. Reformulation: Reformulate it using a new theory. Chooses two mainline approach. One on the **predictive patterns**, and one on the **structural formation**, i.e. not just predicting *the peak or rise* in double descent, but to identify the structural reason why that happens, and so on.

All of this works in conjunction with double descent as a medium of illustration. Specifically, the *problem and setting* part would have to utilize in-context illustration, which then would be the double descent experimental cases. So is *identification* and reformulation. The **reformulation** part is however much more general, but use double descent as a toy problem for test bed.

8.1.2 Notices

There are a few points that must be addressed. This will help and will be included in a much more sophisticated writing style prefix than before, ensuring that details will be added now and not later where necessary.

1. **Results are written as if they are already settled theory.** The entire thing contains intermediate reasoning, failed paths, test cases, and scope limits are erased (actually provided later but still, it

is *not finalized but pushed as if is*). You need **expendable structure** instead of **finished fragmented endpoints**.

2. The current working of the statistical physics treatment below is *suboptimal*. You jump to the conclusion too fast, and whenever submissions are in, it is in the state of already finalized semi-manuscript with conclusion. Elaborate on details, testing out cases, and so on, and make incremental gains as cumulative on a main document, not simply ‘separated documents’ that cannot be used in tangent with other team members.
3. Hiding too many stuffs. For example, even though it is said that there is isomorphism, you **must** clarify what is isomorphism here, why it is here, what is it, what isomorphic to what. All concepts are the same. Phase transition? If you list it out or seriously consider it part of the theory, include it or explain in great details.
4. Even though you stated, for example, the **ergodic theorem** or axioms or principle, you must *explicitly explain them*, with **full setting** of what object to consider, what is the central lining of the axioms, what it asserts, why is it there, what situation makes it to not *work* any more, and vice versa, add many more of the full details as to then justify the content itself.
5. Identify **actual notational system** and how they mean, as ontology and notations must be clarified of for one symbol contains many contexts and so on, especially with a long manuscript. Even Langevin dynamics will have to be clarified as to **why it is defined like that**, and thus why machine learning people use it in the theory. Exploit many topics, for example, Hamiltonian of a system but now ask ‘since we are working in machine learning, why don’t we test how and who the heck did the Hamiltonian of a ML system, or are there any *analogous* notion in here?’ so on and so forth. Equilibrium assumption is very *structural*, and requires many knowledges beforehand, so **list them out carefully and elaborate**.

Overall, this section misses details. It misses the basic foundational sections, explanation, completeness of details in a sensitive setting that we **interpret** and *unify* a fragmented connection between two fields, of which one assumes higher authority (statistical physics) than the other (learning theory). If we keep doing this, everything will fail and misrepresentation will happen, overclaiming will happen, something that seems to be important and big will turn out to be a slight cheeky representation and nothing more. And all the good ideas might just be not developed in blindness of the other mediocre but ‘nice-looking’ and hand-wainable.

This is a setback and must be addressed. The concept itself is highly valuable in the followed idea templates, and frankly to say genuinely good, but the elaboration and formation of such is lacking seriously in several fronts. Also, the **styling** is also too ‘short’ and is too dense in the wrong way of dense – in the sense of compacting many technical jargons but confusing in many parts. For example, ‘energy landscape as rugged, multi-valleyed, dominated by metastability’ contains the concept of a **rugged affine n -hyperspace in a $n + 1$ space**, convex peak or valley and **why is it even there** on the metric of the system, metastability but does not clarify whether it is a **global metastability** from chaotic distribution or something else. Spin-glass theory is also non-trivial, but is sighted without prior blanket. This standard **must be met** without question. Another thing to finally add as to why that kind of density is dangerous is that, from that simple sentence, you have several shots onto your leg:

1. What metric defines distance in the n -space? What is that space and the structures provided? Does it have inner product? Why the heck inner product is even in there? Suddenly now a configuration space is given pseudometric – why?

2. What notion of convexity is being used? What if it is not convex? What is the reference point (usually convexity is defined by 'probing two points' intuitively speaking)? What is alternative than convex? Why is convex there? Is the theory restricted to this convexity or that convexity and if it is non-convex in general, is it useless? Is convexity even a principle subject that defines the theory?
3. Is ruggedness quantified (e.g. Hessian spectra, barrier scaling)? What is rugged? Under what context? Is it emergent or inherent of a system up to certain threshold?
4. Is metastability dynamical (algorithm-dependent) or thermodynamic? What is being thermodynamic here means? Is it about a thermodynamic interpretation of a componential deep neural network structure? Is the algorithm being agnostic? Is the algorithm of different nature (physical operational algorithm or simply numerical or automata-kind one)?
5. What is the embedding space, what is encoded, how it is encoded, what is encoded? Even if information theoretic ideas are imported in interpretation from us researchers, **what does it mean for a system to have information?** Is it intrinsic, or observers' interpretation only?

Too many vector of analysis and thus too many fall-line that makes any further attempt to collapse under its own weight. We are not in the stage of building ideas any more in a fragmented way. There are guidelines which I provided, and there are structural in which we agreed upon. Ideas go in ideas, but structural must stay loyal to the structure and not to avoidance of details. The manuscript is already 150 pages long. There is nothing to shy from *keeping it short either* as the nature of this document. This is for anyone, so do **myself** on this section that I keep in my mind:

You can't expect people to understand what you write when you yourself cannot even invoke it whenever asked. Don't delegate meaning to those you subconsciously know, because you also know consciously you don't. In foundational, unifying, or interpretive work, the opposite is true. Every unstated assumption **becomes a liability**, every borrowed term **must be re-earned**, and every failure mode **must be named**.

Keep on with this and it would be alright. Otherwise, foundational work will find its foundation in another corner of a dusty library and nothing more. If they cannot meet this standard, they are free to leave and fragments of what they have done will still be recognized in the manuscript as proof of their involvement, no question asked. However, this is imperative. They must adhere to this and so am I, also free to be criticized when wrong. Science does not await for anyone to be soft-talked into obscurity. When I asked of you this project being foundational, I said to beware of the contract you must sign before you commit. Whether you now feel it too harsh of your own or not, the judgement is yours. But if you still work, this will be the judgement I give before you improve on this, for the betterment of both our work and ourselves.

8.1.3 In detail systematic projection

Now then, for the main portions, we should go with the following list, though it can be modified any given time. If you find something lacking, please insert it in this frame on details more than the above conceptual schematic, providing it is a good addition to the fold. In general, there is

1. Introduction to statistical physics. Specifically, this part will grapple about the *foundational criteria of statistical physics*, i.e. what is the system of laws, how objects are presented. This includes the following topic to include in, not per specific order but relative conceptual order:

- What is statistical physics? Why is it necessary?
 - What is the foundational framework that precedes it? I.e., for example, statistical physics of course goes before quantum but still applicable in quantum system with modification, so it effectively **switch register** in between. So, let's say we go to the foundation on how Newtonian foundation, theoretical mechanics and the needs of statistical method influences such foundation and thus the assumptions that follow.
 - The tools and perspective that statistical physics uses for such problem and so on. Especially on, for example, the **simulated scale problem** that statistical physics seems to work around, because of the requirement to *estimate large and multi-body* systems.
 - What is the tradeoff in forming that system, what is assumed during the process of 'statisticalized' the system and a new second-layer ordering of axiomatic set, for example? We need to systematically identify this, by proposing your own questions, interpretations, aspects.
 - Now, after we are done, move on to the typical and standard analysis, layout and rigorous analysis of statistical mechanics in actions, the principles, the theorems, the system analysis and observations in the form of physical theory that statistical theory has made. Going for this into the heart of the theory itself, especially those that would often be used in ML, like replica methods, and why it is there, is a crucial problem.
 - Following the previous step, ask – why you use *replica method* and so on, but not the others' analysis? This is because theory in itself compounds from analysis and aspects, or descriptions and encoding in-field, but this is where you deliberately **map the theory to another place**, which means that this particular relationship is rather 'toxic' – the other might not have what you previously identify as *natural evolution in solution of methods*. For example, even the action of formalizing Newtonian mechanics to Lagrangian requires the presumption and properties by uniformity of system description – meta-theoretical duality, so to speak.
2. The connection to learning theory, and the modelling schema in general. This includes the following problem sections:
 - Current methods and categorization of them (with references as always) from statistical physics and was used in ML. Similarly, analogous concepts from statistical physics to ML, and the reverse relationship if any in the way back.
 - Present results and historical researches, caveats in great details, analysis of the results themselves, author's own observations, the scope limitation, the setting self-constraints, and so on.
 - Analysis of each fundamental case – what is used, what is the **object of concept** in statistical physics, and **object under transformation** of understanding in learning theory or computational model used, i.e. what is the core objects and so on that are mapped to *what object over there* in ML, TML, learning theory and so on.
 - Amplification on detail research of drawbacks, i.e. this is the section to limit the scope and range of applicability of those above concepts and researches, nitpick on them, and so on. The goal is to reframe their reach here, and to provide windows of opportunity resulted from such act.
 - Replication of old experiments and applications, and explore their specific edge cases.
 3. Double descent on learning models and the statistical compounding explanation. This section is rather straightforward – find out what is going on with double descent and **the statistical physical**

explanation already presented. This is crucial as to connect back to earlier sections on *why it works* and when it does not, and to analyse the specific double-descent specific categories here. That said, focus on the conclusion or methods used on the notion of **model complexity**, **model structure**, **error landscape and model evaluation**, then **algorithmic performance and operational awareness** (i.e. the difference between deterministic and probabilistic plus operational systems are that they are volatile enough for us to justify actual operational features), **algorithm complexity** both in the time-complexity and structural complexity, agnosticity of algorithm and operational processes of models, hidden assumptions and reduction of scalability, for example, in certain texts, they specifically *constraints* the appearance of double descent to a very niche portion of the domain – if they generalize it further than it should be – that is our cue. Additionally, if able, the structural changes and pathological pathways explored by those models.

4. The novel development. I will leave this to you and the collaborators else in your own judgement. However, as said, there are two types of theoretical development in question – the *predictive framework* and the *structural framework*. The first one succeed in *the prediction is inherent of this domain and system*, while failing at to point out specifically why and which components break, while the other one *can explain well the structural reasons that lead to this particular phenomenon*, but phenomenologically performs worse do to the nature of structural conformation.

The general inquiry of this is to develop them in rigorous detail, for the repeated usage of this particular word as-is. We do not require being super-formal and mathematical, but the right way of being rigorous – as to hide nothing, unpack everything, precise language or can be often generating as the topic itself is largely uncharted domain, and mathematically *effective* – we don't want performative mathematical overformalism. The writing style fixture has been included in earlier notes, so we should not bring it here. However, if to add, then the mainline argument is to "don't overclaim where shouldn't, over-commit where trivial, under-communicate and describe where obscure, attack persons where should be targeted toward their works".

8.2 Approach A - Isomorphism

Isomorphism here, for the sake of clarify means the **structure preservation mapping** between two systems of abstract objects. The categorical error here is to assume isomorphism between two fields, and mapping them onto the general line of basis, and is a very famous *failure point* in the entire thesis of this section.

8.2.1 From Analogy to Isomorphism

IN here, we critically examine the intersection of statistical physics and machine learning. We posit that while statistical physics provides a rigorous framework for analyzing complex systems with many degrees of freedom, its application in machine learning has often been hindered by superficial analogies rather than mathematical isomorphisms. To truly leverage this framework, one must move beyond metaphorical comparisons and establish a dictionary of strict correspondence between physical quantities and learning dynamics.

Reasons for Statistical Physics in Machine Learning

The primary motivation for applying statistical physics to machine learning lies in the “curse of dimensionality” and the emergent phenomena observed in deep neural networks. Just as thermodynamics describes the macroscopic behaviour of 10^{23} particles without tracking individual trajectories, statistical learning theory seeks to describe the generalization performance of models with billions of parameters.

However, a critical mismatch often occurs in the abstraction level.

Common approaches often equate the “learning rate” in Stochastic Gradient Descent (SGD) to “temperature” in Langevin dynamics. We argue that this abstraction is insufficient and often misleading. In equilibrium physics, temperature is an isotropic scalar field. In contrast, the noise introduced by SGD is inherently anisotropic and state-dependent, scaled by the gradient covariance matrix. Therefore, a direct application of equilibrium concepts (like canonical distributions) without accounting for this structural mismatch leads to incorrect predictions about the system’s stationary state.

To be more precise, let us consider the standard Langevin dynamics:

$$dW = -\nabla L(W) dt + \sqrt{2T} dB_t \quad (8.1)$$

where T is the temperature and dB_t is a Wiener process. The stationary distribution is the Boltzmann distribution $\rho(W) \propto \exp(-L(W)/T)$. However, SGD dynamics are fundamentally different:

$$W_{t+1} = W_t - \eta \nabla L_{\mathcal{B}}(W_t) \quad (8.2)$$

where $\mathcal{B}$ is a mini-batch. The noise covariance is:

$$\Sigma(W) = \mathbb{E}_{\mathcal{B}} [(\nabla L_{\mathcal{B}} - \nabla L)(\nabla L_{\mathcal{B}} - \nabla L)^{\top}] \quad (8.3)$$

This covariance is **state-dependent** and **anisotropic**, fundamentally different from isotropic thermal noise.

Axioms and Formulation

Statistical physics is built upon fundamental axioms, most notably the *Ergodic Hypothesis*, which assumes that a system evolves over time to explore all accessible microstates with a probability proportional to their Boltzmann weight.

The Mismatch. Deep learning dynamics, particularly under SGD, are fundamentally non-ergodic on practical timescales. The system does not explore the entire weight space but is confined to specific trajectories determined by the initialization and the landscape geometry.

The Isomorphism. To resolve this, we propose a strict isomorphism:

- **Microstate:** A specific configuration of weights $W \in \mathbb{R}^P$.
- **Hamiltonian (Energy):** The Loss function $L(W, \mathcal{D})$.
- **Quenched Disorder:** The dataset $\mathcal{D}$, which remains fixed during the training dynamics, analogous to the random interactions J_{ij} in spin-glass models.

This mapping situates machine learning not in the realm of simple ideal gases, but squarely within the physics of *disordered systems*, where the energy landscape is rugged, multi-valleyed, and dominated by metastability. The disorder average over datasets corresponds to the quenched average in spin-glass theory:

$$\overline{\langle O \rangle} = \int d\mathcal{D} P(\mathcal{D}) \langle O \rangle_{\mathcal{D}} \quad (8.4)$$

where $\langle \cdot \rangle_{\mathcal{D}}$ denotes the thermal average for fixed dataset $\mathcal{D}$.

Assumptions and Incompatibilities

Standard statistical mechanical treatments (e.g., the Replica Method) assume the system reaches thermodynamic equilibrium. This implies that the final solution is independent of the training path. However, the phenomenon of double descent—and specifically epoch-wise double descent—demonstrates that the “solution” is a transient state of a non-equilibrium process.

Assumption 8.2.1 (Equilibrium Assumption). *The system reaches a stationary distribution $\rho^*(W)$ independent of initial conditions and training trajectory.*

This assumption is violated in practical deep learning for several reasons:

1. **Finite training time:** Training is stopped before equilibration.
2. **Non-convex landscape:** Multiple local minima prevent ergodic exploration.
3. **Time-dependent dynamics:** Learning rate schedules break time-translation invariance.

Forcing an equilibrium assumption onto a dynamic process like SGD is the fundamental error we aim to correct in subsequent sections.

8.2.2 Limitations of Equilibrium Theories

Here, we identify the specific failures of current theoretical frameworks when addressing the Double Descent (DD) phenomenon. While significant progress has been made, a comprehensive theory remains elusive due to a reliance on static analysis.

The Failure of Classical Bias-Variance

As established in Theorem 3.3.2, the classical U-shaped curve predicts that as model complexity increases beyond the optimal point, the variance term dominates and test error monotonically increases. Formally, for the decomposition:

$$\mathbb{E}[\mathcal{M}(f, y)] = \mathcal{B}(f, y) + \mathcal{V}(f, y) + \sigma_{\text{noise}}^2 \quad (8.5)$$

classical theory predicts $\mathcal{V}(f, y) \rightarrow \infty$ as model complexity $\rightarrow \infty$. However, empirical evidence from [Belkin et al. \[2019a\]](#), [Nakkiran et al. \[2019\]](#) shows a second descent in test error for over-parameterized models ($P > N$).

The Limitation of RMT and Replica Theory

Recent applications of Random Matrix Theory (RMT) and Replica Theory have successfully derived the “peak” of the double descent curve at the interpolation threshold ($P \approx N$). For linear regression with random features, the expected test error takes the form:

$$R_{\text{test}} = \sigma^2 \frac{P}{N - P - 1} + \text{bias terms}, \quad P < N \quad (8.6)$$

which diverges as $P \rightarrow N$. However, these successes are limited to “static” or “shallow” models. They treat the solution as a global minimum (min-norm solution) found in a convex landscape:

$$W^* = \arg \min_W \|W\|^2 \quad \text{s.t.} \quad XW = Y \quad (8.7)$$

The Identification of Errors. The critical error in these approaches is the neglect of *feature learning dynamics*. Deep neural networks do not merely fit fixed features; they evolve their representation space. Furthermore, the existence of *epoch-wise double descent*—where the test error exhibits non-monotonic behavior over time even with fixed model size—proves that DD is a dynamical phenomenon. Equilibrium theories, by definition, cannot explain transient dynamical effects.

8.3 Identification: The Dynamic-Landscape Interaction

We categorize the Double Descent problem not merely as a statistical anomaly, but as a complex interaction between *multi-scale dynamics* and a *critical landscape*.

The Gap Between Statics and Dynamics

There exists a “gap” between the static description of the loss landscape and the dynamic trajectory of the optimizer. We identify three distinct regimes that must be unified:

Under-parameterized Regime ($P \ll N$): Dominated by simple signal learning. The landscape is relatively smooth with a unique global minimum.

Critical Regime ($P \approx N$) — **Jamming**: At $P \approx N$, the system undergoes a “Jamming Transition”. The landscape becomes rugged with sharp minima, corresponding to satisfying a precise number of constraints equal to the degrees of freedom. The condition number of the Hessian diverges:

$$\kappa(H) = \frac{\lambda_{\max}}{\lambda_{\min}} \rightarrow \infty \quad \text{as} \quad P \rightarrow N \quad (8.8)$$

Over-parameterized Regime ($P \gg N$): The system becomes “floppy” or unjammed. There exists a continuous manifold of zero-loss solutions, and the geometry of this manifold determines generalization.

Hypothesis: Double Descent as Phase Transition Navigation

Hypothesis 8.3.1 (Dynamic Origin of Double Descent). *Double Descent is the result of the SGD algorithm navigating through phase transitions in the loss landscape. Specifically:*

1. The “peak” at $P \approx N$ is a trapping event in the jammed phase (sharp minima).
2. The second descent is a relaxation process driven by the specific noise structure of SGD, which allows the system to escape sharp minima and settle into flat, generalizing minima.

The key insight is that SGD noise is not merely a nuisance but an *implicit regularizer* that performs geometric selection among the manifold of solutions.

Reformulation: Non-Equilibrium Statistical Physics Framework

To address the identified gaps, we reformulate the learning problem using a Non-Equilibrium Statistical Physics (NESP) framework. We move from describing static probability distributions to modeling the evolution of the probability density $\rho(W, t)$.

The Fokker-Planck Description

The evolution of the probability density over weight space is governed by the Fokker-Planck equation:

$$\frac{\partial \rho(W, t)}{\partial t} = \nabla \cdot (\rho \nabla L) + \nabla \cdot (D(W) \nabla \rho) \quad (8.9)$$

Crucially, the diffusion tensor $D(W)$ is not constant (as in thermal equilibrium) but is **state-dependent**, defined by the covariance of the stochastic gradients:

$$D(W) = \frac{\eta}{2B} \Sigma(W) \quad (8.10)$$

where η is the learning rate, B is the batch size, and $\Sigma(W)$ is the gradient noise covariance.

Noise-Driven Entropic Selection

This reformulation leads to a new mechanism for generalization. The SGD noise is anisotropic and approximately proportional to the Hessian (curvature) of the loss landscape:

$$\Sigma(W) \approx \frac{1}{N} \sum_{i=1}^N \nabla \ell_i \nabla \ell_i^\top \sim H(W) \quad \text{near minima} \quad (8.11)$$

This leads to curvature-dependent diffusion:

- **In Sharp Minima** (associated with the interpolation peak/jamming): The curvature is high, so $D(W)$ is large. This creates a strong effective “repulsive force” that drives the system out of these non-generalizing states.
- **In Flat Minima** (associated with the second descent): The curvature is low, so $D(W)$ is small. These states are dynamically stable under SGD.




Figure 8.1: **NESP Mechanism of Double Descent.** At the interpolation threshold (Sharp Minima), high curvature induces strong anisotropic SGD noise (red wavy arrows), driving the system to escape. The system then relaxes to flat minima where low noise ensures stability. This process explains the second descent in test error.

The Mechanism of Double Descent

This framework reinterprets the Double Descent curve as a dynamical timeline:

1. **Descent 1 (Signal Learning):** Fast learning of dominant signal components (low-frequency modes). The model captures the primary structure in the data.

2. **Ascent (Peak/Jamming)**: The system attempts to fit high-frequency noise but encounters the “Jamming transition” at the interpolation threshold. The landscape is rugged, and the trajectory is temporarily trapped in sharp, overfitting minima.
3. **Descent 2 (Relaxation)**: In the over-parameterized regime, the “entropic selection” mechanism of SGD dominates. The landscape-dependent noise “anneals” the system, driving it out of sharp minima and into flat minima. This *implicit regularization* via non-equilibrium dynamics is the physical origin of the second descent.

Mathematically, the escape rate from a sharp minimum with curvature λ scales as:

$$\Gamma_{\text{escape}} \sim \exp\left(-\frac{\Delta L}{D_{\text{eff}}}\right) \sim \exp\left(-\frac{\Delta L \cdot B}{\eta \lambda}\right) \quad (8.12)$$

where ΔL is the barrier height. For sharp minima with large λ , the escape rate is high, explaining the instability of overfitting solutions.

Predictions and Implications

The NESP framework makes several testable predictions:

1. **Batch size dependence**: Larger batch sizes reduce noise ($D \propto 1/B$), making sharp minima more stable and potentially suppressing the second descent.
2. **Learning rate dependence**: Higher learning rates increase noise ($D \propto \eta$), accelerating escape from sharp minima and enhancing the second descent.
3. **Epoch-wise double descent**: The non-monotonic behavior over training time is a natural consequence of the system first falling into sharp minima (fast convergence) then slowly escaping (relaxation).

These predictions align with empirical observations in [Nakkiran et al. \[2019\]](#) and provide a unified explanation for both model-wise and epoch-wise double descent phenomena.

8.3.1 Experimental Verification: A Toy Model Proposal

The theoretical framework developed in preceding sections—positing that Double Descent arises from a Non-Equilibrium phase transition between Jammed and Relaxed states, mediated by anisotropic SGD noise—demands empirical validation. However, the complexity of modern deep neural networks obscures the fundamental mechanisms we seek to isolate. Therefore, we propose a series of *Gedanken-experimente* utilizing minimal, analytically tractable toy models that preserve the essential physics while eliminating confounding factors.

The Teacher-Student Protocol

We propose the **Teacher-Student framework** as our primary experimental testbed. This paradigm, extensively employed in the statistical physics of learning [Engel and Van den Broeck \[2001\]](#), [Advani et al. \[2013\]](#), provides the crucial advantage of a known ground truth, enabling precise measurement of generalization error without resort to heuristic train-test splits.

Setup. Consider a *Teacher network* $f^*(x) = \sigma(W^*x)$ with fixed, randomly drawn weights $W^* \in \mathbb{R}^{P^* \times d}$, where $\sigma(\cdot)$ denotes a nonlinearity (e.g., ReLU or tanh). The Teacher generates labels:

$$y = f^*(x) + \xi, \quad \xi \sim \mathcal{N}(0, \sigma_\xi^2) \quad (8.13)$$

where ξ represents irreducible label noise—a critical ingredient for inducing the interpolation peak.

The *Student network* $\hat{f}(x; W) = \sigma(Wx)$ with $W \in \mathbb{R}^{P \times d}$ is trained on N samples $\{(x_i, y_i)\}_{i=1}^N$ drawn i.i.d. from the Teacher’s distribution. By varying the Student’s capacity P while holding N fixed, we sweep across the interpolation threshold $P/N \approx 1$.

The Critical Control Parameter. The ratio $\alpha = P/N$ serves as our control parameter, analogous to the “packing fraction” in jamming physics. Our NESP framework predicts:

- $\alpha \ll 1$: Under-parameterized regime. Classical bias dominated behavior.
- $\alpha \approx 1$: Critical/Jammed regime. Interpolation threshold.
- $\alpha \gg 1$: Over-parameterized regime. Entropic selection dominates.

8.3.2 Physical Observables: Beyond Test Error

To validate the NESP mechanism—rather than merely confirming the existence of Double Descent—we must measure *physical observables* that directly probe the landscape geometry and noise structure. We propose three such observables:

Observable I: The Hessian Trace (Sharpness)

The trace of the Hessian, $\text{Tr}(H) = \text{Tr}(\nabla^2 L)$, provides a scalar summary of the local curvature at the solution W_{SGD}^* . This quantity is computationally accessible via Hutchinson’s trace estimator:

$$\text{Tr}(H) \approx \frac{1}{K} \sum_{k=1}^K v_k^\top H v_k, \quad v_k \sim \mathcal{N}(0, I) \quad (8.14)$$

where the Hessian-vector products Hv can be computed efficiently via automatic differentiation.

Hypothesis 8.3.2 (Sharpness Peak at Jamming). *The Hessian trace, measured at the end of training, exhibits a pronounced maximum at the interpolation threshold:*

$$\text{Tr}(H)|_{\alpha \approx 1} \gg \text{Tr}(H)|_{\alpha \ll 1} \quad \text{and} \quad \text{Tr}(H)|_{\alpha \approx 1} \gg \text{Tr}(H)|_{\alpha \gg 1} \quad (8.15)$$

Furthermore, in the over-parameterized regime, we predict a monotonic decay:

$$\frac{d}{d\alpha} \text{Tr}(H) < 0 \quad \text{for} \quad \alpha > 1 \quad (8.16)$$

This decay directly evidences the system’s migration from sharp (Jammed) to flat (Relaxed) minima as capacity increases.

Observable II: Noise-Hessian Alignment

The central claim of our NESP framework is that SGD noise is *not* thermal (isotropic) but rather *geometric* (aligned with the Hessian). To test this, we propose measuring the alignment between the noise covariance matrix $\Sigma(W)$ and the Hessian $H(W)$:

$$\mathcal{A}(\Sigma, H) = \frac{\text{Tr}(\Sigma H)}{\|\Sigma\|_F \|H\|_F} \quad (8.17)$$

where $\|\cdot\|_F$ denotes the Frobenius norm. For thermal noise, $\Sigma = \sigma^2 I$ and thus $\mathcal{A} \propto \text{Tr}(H)/\|H\|_F$, which is independent of the noise structure. For geometry-dependent noise, we expect strong positive correlation.

Hypothesis 8.3.3 (Anisotropic Noise Alignment). *The alignment $\mathcal{A}(\Sigma, H)$ is significantly elevated compared to the null hypothesis of isotropic noise:*

$$\mathcal{A}(\Sigma, H) \gg \mathcal{A}(\sigma^2 I, H) \quad \text{across all } \alpha \quad (8.18)$$

Moreover, this alignment should *strengthen* near the interpolation threshold, where the curvature structure is most pronounced. This would provide direct evidence that SGD performs geometric, not thermal, exploration of the loss landscape.

Observable III: Eigenspectrum Evolution

The full eigenspectrum of the Hessian, $\{\lambda_i\}_{i=1}^P$, encodes richer information than the trace alone. We propose tracking the spectral density $\rho(\lambda) = (1/P) \sum_i \delta(\lambda - \lambda_i)$ across the α -sweep.

Hypothesis 8.3.4 (Spectral Signature of Phase Transition). *At the Jamming transition ($\alpha \approx 1$), the spectral density exhibits:*

1. *A hard edge at $\lambda = 0$, corresponding to the vanishing of the null space as constraints saturate degrees of freedom.*
2. *Outlier eigenvalues at large λ , corresponding to sharp directions associated with fitting noise.*

In the Relaxed phase ($\alpha \gg 1$), the spectrum should exhibit:

1. *An extensive null space (bulk of eigenvalues at $\lambda \approx 0$), corresponding to flat directions in the solution manifold.*
2. *Suppression of outliers, indicating the system has escaped sharp directions.*

8.3.3 Experimental Protocol

To empirically validate the Entropic Selection mechanism, we propose the following experimental protocol:

Predicted Signatures. If our NESP framework is correct, the experimental results should exhibit the following correlated signatures:

1. $R_{\text{test}}(\alpha)$ displays the characteristic Double Descent curve.
2. $\text{Tr}(H)(\alpha)$ peaks at $\alpha \approx 1$ and decays for $\alpha > 1$, *anti-correlated* with the second descent.
3. $\mathcal{A}(\Sigma, H)$ remains elevated throughout, peaking near $\alpha \approx 1$.
4. The spectral density transitions from hard-edge (Jammed) to bulk-zero (Relaxed) as α crosses unity.

The simultaneous observation of these four signatures would constitute strong evidence for the NESP interpretation of Double Descent as a dynamical phase transition driven by geometry-dependent noise.

Experimental Protocol 6: NESP Validation via Teacher-Student Model

Data: Teacher weights W^* , noise level σ_ξ , sample sizes N , capacity range $P \in [P_{\min}, P_{\max}]$.

Result: Measurements of $R_{\text{test}}(\alpha)$, $\text{Tr}(H)(\alpha)$, $\mathcal{A}(\Sigma, H)(\alpha)$, $\rho(\lambda; \alpha)$.

begin

for $P \in \{P_{\min}, \dots, P_{\max}\}$ **do**

 Initialize Student network with random weights W_0 ;

 Train via SGD: $W_{t+1} = W_t - \eta \nabla L_{\mathcal{B}}(W_t)$;

 ; // At convergence:

 Compute $R_{\text{test}}(W_T) = \mathbb{E}_x[(f^*(x) - \hat{f}(x; W_T))^2]$;

 Compute $\text{Tr}(H)$ via Hutchinson estimator;

 Estimate $\Sigma(W_T)$ from mini-batch gradient variance;

 Compute alignment $\mathcal{A}(\Sigma, H)$;

 (Optional) Compute top- k Hessian eigenvalues via Lanczos iteration;

end

end

return Observables as functions of $\alpha = P/N$.

Alternative Minimal Model: Two-Layer Linear Networks

For maximal analytical tractability, we also propose the **two-layer linear network**:

$$\hat{f}(x; U, V) = UV^\top x, \quad U \in \mathbb{R}^{d_{\text{out}} \times k}, \quad V \in \mathbb{R}^{d_{\text{in}} \times k} \quad (8.19)$$

Despite being linear in the input-output map (equivalent to $W = UV^\top$), this model exhibits *nonlinear dynamics* under gradient descent due to the factorization. Recent work Advani and Saxe [2017b], Saxe et al. [2014] has shown that this model captures essential features of deep learning dynamics, including staged learning and critical slowdowns.

The advantage of this model is that the Hessian and noise covariance can be computed *analytically* in certain regimes, enabling precise theoretical predictions against which experimental measurements can be compared. The NESP predictions for this model are:

- The “effective temperature” $T_{\text{eff}} = \eta \cdot \text{Tr}(\Sigma)/P$ should exhibit non-monotonic behavior, peaking at the interpolation threshold.
- The condition number $\kappa(H) = \lambda_{\max}/\lambda_{\min}$ should diverge as $\alpha \rightarrow 1$, signaling the Jamming transition.

8.3.4 Extension

The idea itself offers many extensions for deployments.

Synthesis: From Equilibrium to Non-Equilibrium

In this work, we have undertaken a critical re-examination of the theoretical foundations underlying the Double Descent phenomenon. Our central thesis is that the apparent “anomaly” of Double Descent—its violation of the classical bias-variance tradeoff—is not a failure of statistical physics *per se*, but rather a failure of **equilibrium** statistical physics inappropriately applied to an inherently **non-equilibrium** dynamical process.

We have argued that the dominant theoretical approaches—Random Matrix Theory, Replica Methods, and classical bias-variance decomposition—share a common, often implicit, assumption: that the learning algorithm converges to a stationary state characterized by a Boltzmann-like distribution over weight space. This equilibrium assumption, while mathematically convenient, fundamentally mischaracterizes the nature of Stochastic Gradient Descent.

Our reformulation, grounded in Non-Equilibrium Statistical Physics (NESP), makes three key departures from the equilibrium paradigm:

1. **From static distributions to dynamical evolution:** We model the probability density $\rho(W, t)$ via the Fokker-Planck equation, treating the “solution” as a transient state rather than an equilibrium configuration.
2. **From isotropic to anisotropic noise:** We recognize that SGD noise is not thermal but geometric—its covariance $\Sigma(W)$ is state-dependent and aligned with the loss landscape curvature $H(W)$.
3. **From energy minimization to entropic selection:** We posit that generalization arises not from finding the global minimum, but from the *dynamical stability* of flat minima under geometry-dependent noise. Sharp minima, despite potentially lower loss, are rendered unstable by their own curvature.

Bridging the Gap: The Play and The Stage

A persistent tension in theoretical machine learning has been the disconnect between two perspectives:

- **The Stage (Landscape):** Static analysis of the loss surface—its minima, saddle points, and global structure.
- **The Play (Dynamics):** The trajectory of the optimizer through weight space—its convergence, oscillations, and implicit biases.

Equilibrium theories privilege the Stage: they characterize the set of possible solutions and assume the optimizer will find them. But they cannot explain *which* solution is selected, nor *how* the selection depends on hyperparameters like learning rate and batch size.

The NESP framework bridges this gap by coupling the Stage to the Play. The landscape geometry (Hessian) determines the noise structure (diffusion tensor), which in turn shapes the dynamical trajectory. Double Descent emerges naturally from this coupling: the interpolation peak corresponds to a *Jamming transition* where landscape constraints saturate available degrees of freedom, and the second descent corresponds to *entropic relaxation* as the system escapes sharp minima driven by its own noise.

This perspective offers a resolution to the longstanding puzzle: *Why do over-parameterized models generalize?* Our answer: not because they find better minima (they often do not), but because SGD dynamics are *geometrically biased* toward flat regions of the solution manifold. Generalization is an emergent property of non-equilibrium dynamics, not a static property of the loss landscape.

Moreover, if the NESP framework, is formally validated, it opens several exciting avenues for future research, as followed.

Grokking as Extreme Relaxation

The phenomenon of “Grokking”—wherein models exhibit delayed generalization long after achieving zero training error Power et al. [2022b]—presents a striking challenge to equilibrium theories. From the NESP perspective, we speculate that Grokking represents an *extreme* instance of the Jamming-Relaxation transition.

Specifically, certain data structures (e.g., modular arithmetic) may induce loss landscapes where the Jammed phase is *metastable* over extended timescales. The system achieves interpolation (zero training loss) but remains trapped in sharp, non-generalizing minima. Generalization occurs only after sufficient time for the noise-driven annealing process to escape these metastable states.

This interpretation predicts that Grokking should be:

- **Accelerated** by increasing learning rate or decreasing batch size (higher effective noise).
- **Suppressed** by explicit regularization that pre-selects flat minima.
- **Accompanied** by a monotonic decrease in Hessian trace during the “grokking” phase.

Transformers and Large Language Models

The training dynamics of Large Language Models (LLMs) represent perhaps the most consequential—and least understood—application of deep learning. The NESP framework suggests several speculative connections:

Temperature in Sampling vs. Training. LLMs employ a “temperature” parameter T during inference to control the sharpness of the output distribution. Intriguingly, our framework identifies an “effective temperature” $T_{\text{eff}} = \eta \cdot \text{Tr}(\Sigma)/(2B)$ that governs training dynamics. We pose the question: *Is there a principled relationship between these two temperatures?*

One speculative hypothesis: models trained with higher T_{eff} (larger learning rate, smaller batch) may naturally produce sharper internal representations, requiring lower inference temperature for coherent outputs. This would establish a non-trivial link between training dynamics and inference behavior.

Scaling Laws and Phase Transitions. Empirical “scaling laws” Kaplan et al. [2020] describe how LLM performance improves with model size, data, and compute. From the NESP perspective, these smooth power-law scalings may obscure underlying *phase transitions* in the training dynamics. As model capacity crosses critical thresholds relative to data complexity, the system may undergo Jamming-Relaxation transitions analogous to Double Descent.

Understanding these transitions could inform more efficient training schedules: rather than smooth scaling of learning rate, one might employ *punctuated* schedules that exploit phase boundaries—high noise to escape Jammed phases, low noise to stabilize in Relaxed phases.

Toward a Unified Theory of Implicit Regularization

The NESP framework provides a candidate mechanism for “implicit regularization”—the empirical observation that SGD finds generalizing solutions without explicit regularization terms. Our entropic selection mechanism offers a *dynamical* explanation: SGD implicitly regularizes by destabilizing sharp minima.

A unified theory would connect this mechanism to other forms of implicit bias:

- **Minimum-norm solutions** in linear regression.
- **Max-margin classifiers** in linearly separable problems.
- **Low-rank bias** in matrix factorization.

We conjecture that all these biases arise from the same underlying principle: the geometry-dependent noise of SGD preferentially stabilizes solutions with specific structural properties (low norm, large margin, low rank) because these properties correlate with flat loss landscape geometry.

Index

Bayes model, 24
Chernoff's inequality, 6
Cramer condition, 6
dataset identifiability, 97
empirical risk, 21
ERC, 42
generalization risk, 21
growth function, 46
halfspace, 49
hypothesis class, 20
Markov's inequality, 2
mechanistic layer, 16
model-free learning, 27
model-specific learning, 27
noise, 38
operationalisation, 16
outer typing, 158
phenomenological layer, 16
Rademacher complexity, 43
Rademacher variables, 42
representation scheme, 29
representation size, 30
Sauer's lemma, 49
shattering, 48
temporal difference, 78
transparency, 16
VC-dimension, 48

To verify the theorem itself, we test it with a standard setting without the theorem results. This is reflected in Figure 4.5. The result is somewhat consistent with the theorem and its latter parameter error and prediction error results, however, this time, it correlates exactly as the prediction error instead. Additionally, peculiarity occurs in the non-double descent region, with much less volatility than in the prediction of the theorem.

```

Data:  $n, p, d, \sigma, r$ , fixed range  $[a, b]$ ,  $u \in [0, 1]$ 
Result: Threefold error measure - MSE on parameters  $p, d$ , MSE on prediction  $Y', Y$ , and theoretical risks measure by prescription.

begin
  Randomizer( $\cdot$ )  $\leftarrow$  PCG64( $n \leq 100$ );                                Randomizing pseudo-random engine
   $\sigma \leftarrow 0.5, d \leftarrow 150$ ;                                    Common presetting for parameters
   $nl \leftarrow \{12, \dots, 280\}$ ;                                       Common full-list  $n$  count of datapoints
   $p \leftarrow 150$ ;                                                       Underparameterization to equal
   $\mathbf{p} \leftarrow \text{range}(1, p)$ ;
   $[a, b] \leftarrow [-100, 100]$ ;
   $w_c \sim \text{PCG64.Uniform}([1, d], 150)$ ;
   $b_c \sim \text{PCG64.Uniform}([1, 10])$ ;
   $w_p \sim \text{PCG64.Uniform}([1, p], p)$ ;
   $b_p \sim \text{PCG64.Uniform}([1, 10])$ ;
   $X \sim \text{PCG64.Uniform}(\text{range} \leftarrow [a, b], \text{size} \leftarrow [n, d])$ ;       $n$  samples of  $d$  dimensions.
   $\text{concept}(\cdot, \cdot, \cdot) \leftarrow (d, w_c, b_c)$ ;
   $Y \leftarrow \text{concept}(X, w_c, b_c)$ ;
   $\sigma_{\text{scaled}} \leftarrow \sigma \cdot \mathbb{E}[|Y|]$ ;
   $Y \leftarrow Y + \text{PCG64} \dots \text{Normal}(\mu \leftarrow 0, \sigma \leftarrow \sigma_{\text{scaled}}, \text{size} \leftarrow |Y|)$ ;      Noisy featuring
  for  $n \in nl$  do
     $w_p^* \leftarrow \text{PCG64.Permutation}(d, p)$ ;
    for  $i \leftarrow 1$  to  $|w_p|$  do
      if  $i \neq w_p^*$  then
         $w_p[i] \leftarrow 0$ , fixed zero point.
      end
    end
    end
     $\text{model}(\cdot, \cdot, \cdot) \leftarrow (d, w_p, b_p)$ ;
     $Y \leftarrow \text{model}(X, w_p, b_p)$ ;
     $\text{approx} \leftarrow \text{LeastSquare}(\text{model}, Y', X, Y)$ , fixed  $w_p$ .
  end
   $\text{MSE}_{\text{param}}(\cdot, \cdot, \cdot) \leftarrow \text{MSE}(Y, \{\text{approx}\}, \text{model})$ ;
   $\text{MSE}_{\text{pred}}(\cdot, \cdot) = \text{MSE}(Y, Y') \hat{R}_{\text{Thm}} = \text{TR}(w_c, \sigma, p, d, n)$ ;
end
return  $\text{MSE}_{\text{param}}, \text{MSE}_{\text{pred}}, \hat{R}_{\text{Thm}}$ 

```

We use the Algorithm 7 for verification, including three modes of error measure, which will then be calculated accordingly. For gradient descent, we refer to another experiment. For such algorithm, we notice that the operational space of the model is $\mathbb{R}^n \rightarrow \mathbb{R}$. The masking of parameters then is not well-received, since the output space is far more skewed - single dimension resultant with no scaling difference




Figure 1: Risk curves for varying sample size n (dimension $d = 100$, noise $\sigma = 0.5$). Double descent is determined in similar sequence, even though the interpolation theoretical $p = n$ is actually shifted over to the right from the actual interpolation observables, in parameterization. Afterward, correlation fails, and the theoretical p is ineffective, similar to previous theorem-based test run. Sample set size range is $n \in \{12, 24, 36, 50, 76, 82, 90, 100, 106, 112, 121, 150, 231, 250, 256, 280\}$. One additional remark is that the error landscape is relatively thin in either side, making it abnormal in comparison to actual bias-variance curvature.

because the model in question is supposed to be linear. Regarding such, it makes sense to realize that the model itself can still attain relatively stable error, because its feature compounds on the scaling parameters

To test the stability of the model, we run for a stable section of the permutations done on p . The result is outlined in 4.5. Interestingly, while the interpolation threshold theorized being $p = n$ does not hold, the pattern of double descent without head descent still fundamentally exists, albeit in a much more constrained region, as can see from $n = 12$ to $n = 150$. The pattern from then onward follows a reduced trend, without the chaotic behaviours seen. The result is shown in Figure 1.

Remark

At least from this experiment, overall, results gathered using Theorem 4.3.1 and variations of it suggests inconsistent results. The supposed interpolation threshold has different interpretation and indication in different scenario, including both parameter-wise error and prediction-only error measure. Furthermore, verification also shows inconsistency with the interpolation threshold, albeit given the fact that the margin of error from the highest peak to the interpolation point is consistent up to certain n . Breakdown of the result from theorem, as well as double descent behaviour starts with high n counts, as well as interesting outlook of bias-variance reappearance for $p \sim 5000$, far more than what is tested in normal testing cases. Such result suggests further experiments to validate the observations, as well as statistical analysis of the entire system.

Specifically, the interpretation of the result is not perfect. Under such setting, then for the concept of the same class, so $d = 150$, any concept of $p > d$ would result in the metric taking zero value of the ‘truncated region’ where d is literally flat on the $d + 1$ onward space. Hence, the error being added on top of such error measure is the component distance of p to the flat sub $p - d$ space itself – which explains the

curve going upward, as more and more risks are configured. Nevertheless, the amount of data provided, at least in terms of this model, does indeed have non-trivial effect on the model itself, and the least square solution of best fit which gives the potential minimal, hence the empirical best. Considering the results, we can see regions where the theorem ‘breaks’ of its bias-variance pattern. Further experiment to verify such claim and reconfiguration is needed. Worryingly, in this particular case, there exists no saddle point that was theorized by bias-variance tradeoff, which must then be explained accordingly. Furthermore, we also notice the troublesome fact that this particular pattern works on *same operational space* models, that is, linear models approximating linear models. This setting is fairly limited, and would not generalize well. However, we also do not know what kind of measure to take in such case. Conflicting notions and reports aid into the confusion of such topic more than can be normally observed, and the anomalous behaviours in some researches (for example, in [Shi et al. \[2024\]](#), we received that in GNN, there exists no trace of double descent) makes it even harder. As such, both the status of bias-variance being false, albeit in a given range only, and double descent in masking certain behaviours of the modelling process, have taken a mysterious stance in the modern machine learning community.

While discovering the double descent phenomena, it also exposes theoretical concepts that are not organized or formalized, portions of the theory and empirical observations that are within conflict of each others, issues with definitions, theorems, and interpretation of them. Those problems would then ultimately hinder further research in such direction, as well as broader theoretical requirement of the theory of machine learning, and specifically to a larger and more complex system of deep learning.

Observing gradient descent decomposition

While the interpretation of what the bias-variance tradeoff, the usual training-versus-test curve, and statistical learning consideration, another interesting aspect to consider (as well for double descent), is how bias-variance decomposition acts on gradient descent and its components. This claim is perhaps reinforced by the work of [Adlam and Pennington \[2020\]](#) which indicates that different bias-variance decompositions effects differently on the total system mechanism.

The main term of a gradient descent algorithm is the gradient of the loss function $\mathcal{L}(h(x), y)$. For the hypothesis h , the number of instances supplied denoted by k has three main cases: $k = 1$ for stochastic descent, $k \in (1, m)$ for batch gradient descent, and $k = m$ for standard gradient descent, for a given space of m samples. The gradient in such case is calculated using the empirical risk on m size, that is:

$$\mathcal{L}_k(h(x), y) = \hat{R}_{S[k]}(h) = \frac{1}{k} \sum_{i=1}^k \ell(h(x_i), y_i) \quad (20)$$

We use the classical bias-variance tradeoff at first. Since the gradient is a linear operator, we have:

$$\begin{aligned} \nabla_k(\mathcal{L}) &= \nabla_k \left((\mathbb{E}\hat{y}(\mathbf{x}) - \mathbb{E}y(\mathbf{x}))^2 + \mathbb{V}[\hat{y}(\mathbf{x})] + \mathbb{V}[y(\mathbf{x})] \right) \\ &= \nabla_k (\mathbb{E}[\hat{y}(\mathbf{x})] - \mathbb{E}[y(\mathbf{x})])^2 + \nabla_k \mathbb{V}[\hat{y}(\mathbf{x})] + \nabla_k \mathbb{V}[y(\mathbf{x})] \\ &= \nabla_k (\mathbb{E}[\hat{y}(\mathbf{x})] - \mathbb{E}[y(\mathbf{x})])^2 + \nabla_k \mathbb{V}[y(\mathbf{x})] \end{aligned} \quad (21)$$

This effectively decompose the gradient into three parts that influence the overall gradient calculation. Because $\nabla_k \mathbb{V}[y(\mathbf{x})]$ calculates the gradient of irreducible error, which is supposed to be constant of the sample space, it is then equal 0. Suppose a sample $\mathbf{x}_k$ of size km , of partition k . Hence, we have gradient




Figure 2: (a-e) **The different bias-variance decompositions.** (f-j) Corresponding theoretical predictions for $\gamma = 0$, $\phi = 1/16$ and $\sigma = \tanh$ with $\text{SNR} = 100$ as the model capacity varies across the interpolation threshold (dashed red). (a,f) The semi-classical decomposition of [Hastie et al. \[2019\]](#), [Mei and Montanari \[2019a\]](#) has a nonmonotonic and divergent bias term, conflicting with standard definitions of the bias. (b,g) The decomposition of [Neal et al. \[2018\]](#) utilizing the law of total variance interprets the diverging term V_D^C as “variance due to optimization”. (c,h) An alternative application of the law of total variance suggests the opposite, *i.e.* the diverging term V_P^C comes from “variance due to sampling”. (d,i) A bivariate symmetric decomposition of the variance resolves this ambiguity and shows that the diverging term is actually V_{PD} , *i.e.* “the variance explained by the parameters and data together beyond what they explain individually.” (e,j) A trivariate symmetric decomposition reveals that the divergence comes from two terms, V_{PX} and $V_{PX\epsilon}$ (outlined in dashed red), and shows that label noise exacerbates but does not cause double descent. Since $V_\epsilon = V_{P\epsilon} = 0$, they are not shown in (j). Taken from [Adlam and Pennington \[2020\]](#)

descent based on k th run over the entire dataset, permuted. Then, the loss function calculates in the form

$$\begin{aligned} \nabla_{k,n}(\mathcal{L}) = & \nabla_{k,n} \left[\frac{1}{n} \sum_{i \leq n} \hat{y}(x_{kn+i}) - \frac{1}{n} \sum_{i \leq n} y(x_{kn+i}) \right]^2 \\ & + \nabla_{k,n} \left[\frac{1}{n} \sum_{i=1}^n y(x_i)^2 - \left(\frac{1}{n} \sum_{i=1}^n y(x_i) \right)^2 \right] \end{aligned} \quad (22)$$

The indices is now changed to (k, n) for $m/k = n$ as the i th partitioned, ordered run, thereby $kn + i$ runs for the index. If we assume a simplified structure θ_t of parameters $y(x_i)$ depends on, we can reduce it

further.

$$\begin{aligned}
& \nabla_{k,n}(\mathcal{L}) \\
&= 2 \left(\frac{1}{n} \sum_{i=1}^n \hat{y}(x_{kn+i}) - \frac{1}{n} \sum_{i=1}^n y(x_{kn+i}) \right) \\
&\times \left(\frac{1}{n} \sum_{i=1}^n \nabla_k \hat{y}(x_{kn+i}) - \frac{1}{n} \sum_{i=1}^n \nabla_k y(x_{kn+i}) \right) \\
&+ \frac{2}{n} \sum_{i=1}^n y(x_i) \nabla_k y(x_i) - \frac{2}{n^2} \left(\sum_{i=1}^n y(x_i) \right) \left(\sum_{i=1}^n \nabla_k y(x_i) \right) \\
&= 2 \left[(\mathbb{E}_n[\hat{Y}] - \mathbb{E}_n[Y]) (\mathbb{E}_n[\hat{G}_k] - \mathbb{E}_n[G_k]) \right. \\
&\quad \left. + \mathbb{E}_n[Y G_k] - \mathbb{E}_n[Y] \mathbb{E}_n[G_k] \right] \\
&= 2 \left[(\mathbb{E}_n[\hat{Y}] - \mathbb{E}_n[Y]) (\mathbb{E}_n[\nabla_k \hat{Y}] - \mathbb{E}_n[\nabla_k Y]) \right. \\
&\quad \left. + \text{Cov}_n(Y, \nabla_k Y) \right] \\
&= 2 (\mathbb{E}_n[\hat{Y}] - \mathbb{E}_n[Y]) (\mathbb{E}_n[\nabla_k \hat{Y}] - \mathbb{E}_n[\nabla_k Y]) \\
&\quad + 2 \text{Cov}_n(Y, \nabla_k Y)
\end{aligned}$$

for the expectation over Z_i , and covariance,

$$\mathbb{E}_n[Z] = \frac{1}{n} \sum_{i=1}^n Z_i, \quad \text{Cov}_n(U, V) = \mathbb{E}_n[UV] - \mathbb{E}_n[U] \mathbb{E}_n[V]. \quad (23)$$

Here, we can see two terms multiplication, for the first one being the correlation between the expected error $\mathbb{E}_n[\hat{Y}] - \mathbb{E}_n[Y]$ and the gradient $\mathbb{E}_n[\nabla_k \hat{Y}] - \mathbb{E}_n[\nabla_k Y]$. This measures the correlation between the expected error, and the gradient measure. The second term is exactly the empirical covariance between the target Y and the true gradient component. Since we are taking a fixed concept case, we can simply disregard the truth gradient hence $\mathbb{E}_n[\nabla_k Y]$, for which the term is reduced to $2 (\mathbb{E}_n[\hat{Y}] - \mathbb{E}_n[Y]) (\mathbb{E}_n[\nabla_k \hat{Y}])$. Under this same assumption applied to the covariance term, we obtain

$$\begin{aligned}
\nabla_{k,n}(\mathcal{L}) &= \frac{2}{n} \sum_{i=1}^n (\hat{Y}_i - Y_i) \partial_k \hat{Y}_i \\
&= 2 \left((\mathbb{E}_n[\hat{Y}] - \mathbb{E}_n[Y]) \mathbb{E}_n[\nabla_k \hat{Y}] \right)
\end{aligned} \quad (24)$$

Looking into some particular algebraic representation, we can write this simply as

$$\nabla_{k,n} \mathcal{L} = 2 \mathbb{E}_n \left[(\hat{Y} - Y) \nabla_k \hat{Y} \right] \quad (25)$$

$$= 2 \left(\underbrace{(\mathbb{E}_n[\hat{Y}] - \mathbb{E}_n[Y])}_{[\mathbb{B}(\hat{Y}, Y)]^{1/2}} \mathbb{E}_n[\nabla_k \hat{Y}] \right) \quad (26)$$




Figure 3: An illustration of the (supervised) statistical process. Phase III contains two parts: First is the evaluation $\nabla(h, c)$ according to the data $\mathcal{D}$, and second is the Update process to re-align c to the actual target.

$$+ \underbrace{\mathbb{E}_n[\hat{Y} \nabla_k \hat{Y}] - \mathbb{E}_n[\hat{Y}] \mathbb{E}_n[\nabla_k \hat{Y}]}_{\text{Cov}_n(\hat{Y}, \nabla_k \hat{Y})} \quad (27)$$

$$- \underbrace{\mathbb{E}_n[Y \nabla_k \hat{Y}] - \mathbb{E}_n[Y] \mathbb{E}_n[\nabla_k \hat{Y}]}_{\text{Cov}_n(Y, \nabla_k \hat{Y})} \quad (28)$$

The first term can be attributed to self-correct of bias-error gradient correlation, while the two covariance term suggests correlation between model randomness and its gradient, and the last one deal with the truth-gradient correlation. What those gradients mean will have to be tested in some preliminary models. Generally, we should not expect the result to be quite effective. Similar decomposition can be conducted for the other decomposition, as suggested, however, the gradient effect is particularly interesting, and will have to be resolved.

Details of supervised learning structure

In a typical machine learning setting, we obtain assume the following figure 3. This includes the ground, input space (the setting of which the observations are observed), the concept $c \in \mathcal{C}$, the hypothesis $c \in \mathcal{H}$, and a supervisor ∇ that ‘correct’ the behaviours of the hypothesis to match with c . This is called a supervised learning setting, in which the learning part is governed by the supervisor and the observations from the concept.

Phase 1.

Begin with the construction and initialization of the feature space $\mathcal{X}^\infty$. By the assumption of a probabilistic process, $\mathcal{X}_m$ or simply $\mathcal{X}$ representing the dataset is assumed to be sampled, or appeared from the distribution $\mathcal{D}(p, \cdot)$ for p an arbitrary probabilistic system.

The *ground space*, or in literature generally called **feature space**, is created. This results in the first component of the tuple D , being $\mathcal{X} \subseteq \mathbb{R}^{m \times k}$, where $|\mathcal{X}| = m$, but $\text{size}(x) = k$, for $x \in \mathcal{X}$. We assume there exists the random variable x of a given distribution $\mathcal{D}$, such that the set $\mathcal{X} \sim \mathcal{D}$ of all observation is given by an underlying distribution. A well-known assumption in this case is that $x \in \mathcal{X}$ has no dependent components. That is, for example, there does not exist any function $\psi_x : \{x_i\} \rightarrow \{x_j\}$, for

$i \neq j$. (Thought, this only helps in Bayesian priori analysis). Furthermore, by specifying that $\mathcal{X}$ belongs to a specific distribution, we argue of the assumption that $x \in \mathcal{X}$ is *independently and identically distributed* (i.i.d.), that is, for a given random sampling interpretation $S \subset \mathcal{X} \sim \mathcal{D}$, all data is sampled with the events space governed by the distribution $\mathcal{D}$ for every instance, and the existence of x_i to x_j for $i \neq j$ is none.

The role of the distribution configuration is important. If $\mathcal{D}$ is uniform, that is, $D \sim \mathcal{U}(a, b)$, then we can disregard the aspect of $\mathcal{X}$ with random variables. For $[a, b]$ to be $(-\infty, +\infty)$, the problem changed to a *regression problem*. If $\mathcal{D}$ is some other distribution function, then the problem of *representative data* and *population size* becomes apparent, which requires statistical analysis to be taken into consideration.

Phase 2.

For the constructed feature space $\mathcal{X}$, let $EX(c, \mathcal{D})$ be a procedure (or *oracle*) acting on $\mathcal{C}$ and $\mathcal{X}$ to output $\langle x, c(x) \rangle$. The hypothesis then contains a similar procedure $EX_{\mathcal{H}}(h, \mathcal{D})$, such that to output $\langle x, h(x) \rangle$. We call this the *inference phase*.

The feature space $\mathcal{X}$ is provided to h . We assume c is processed of the same process, resulted in $\mathcal{Y}_c$. Then, $h(\mathcal{X})$ outputs the set of hypothesis' target $\mathcal{Y}_h$. Thereby, we are approximating the *process* c itself. We can approach this in two main ways¹: either *deterministic* or *probabilistic*². We define such as followed. For the constructed feature space $\mathcal{X}$, let $EX(c, \mathcal{D})$ be a procedure (or *oracle*) acting on $\mathcal{C}$ and $\mathcal{X}$ to output $\langle x, c(x) \rangle$. The hypothesis then contains a similar procedure $EX_{\mathcal{H}}(h, \mathcal{D})$, such that to output $\langle x, h(x) \rangle$. We call this the *inference phase*.

Definition .0.1 (Deterministic – *discriminative modelling*). A **deterministic model** h_d assumes its internal space as followed. For any $h \in \mathcal{H}_d$ of deterministic model, h is characterized by the mapping $h_d : \mathcal{X} \rightarrow \mathcal{Y}_h$, such that $h \in C^n$ of n -differential space.

Similarly, if the interpretation of the model is *probabilistic*, we have the following definition of probabilistic-based model.

Definition .0.2 (Probabilistic – *generative modelling*). A **probabilistic model** h_p assumes its internal space as followed. For any $h \in \mathcal{H}_p$ of all probabilistic hypothesis, h is characterized by the probability distribution $\Gamma(p_X(X))$, where Γ is the discrete decision process outside the probabilistic interpretation.³

Notice that the process and the learning setting is probabilistic, however, the interpretation of the hypothesis h can be either deterministic or probabilistic, or anything else. This puts the flexibility into the setting of the model, and permits us to consider various representation of the model structure. Furthermore, the hypothesis for the learning process is evaluated relative to the same probabilistic setting, in which the training takes place, and we allow hypotheses that are only approximation of the target concept Sugiyama [2015], Kearns and Vazirani [1994].

The similarity between both of the two general types is that the model is characterized by their **parameters**, both **model-specific parameters** and **hyperparameters**.

¹Note that, in some aspect, this is similar to the usage of mathematical simulation of certain dynamical system, or *any* system that has certain patterns or relations observable.

²The model itself can be entirely deterministic, while the learning process is not. in fact, one of the very first assumption and axiomatic view of the learning problem is that learning occurs in a *probabilistic setting*.

³The discrete decision process Γ is very simple to argue. For example, given a Naive Bayes model with $p(C_k | D[i])$ of the dataset, the probability distribution is regarded as the argument

$$Y^* = \arg \max_{c_k \in C} P_{\theta}(c_k | x) = \arg \max_{c_k \in C} P_{\theta}(x | c_k) \cdot P(c_k)$$

Phase 3.

There exists an evaluator (or *supervisor*) $\nabla(h, c)$ that evaluates specific *loss framework* $\ell\{h(x), c(x)\}$ based on available information, and a *update modifier* $U(\ell, \mathcal{A})$ of certain objective to algorithm $\mathcal{A}$.

In this step of the procedural setting, the *supervisor* includes a loss function, ℓ , which takes discrete arguments, and ‘binarily’ evaluate the discrete result between h and c . This loss function is supposed to have a global minimum, or at least a certain region of minimal approximation. We have the following definition.

Definition .0.3 (Loss function). A loss function is a non-negative function $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow [0, +\infty)$.

In general, the following is supposed of the loss function:

Conjecture .0.1 (Loss function convexity). For any supervisor ∇ of a learner framework $\mathcal{L}(\mathcal{H}, \mathcal{A})$, the loss function ℓ is assumed to be a *p-converging function*, or as a *convex function*. That is, for $\ell : \mathbb{R}^n \rightarrow \mathbb{R}$, then ℓ is convex on its domain, or is locally convex on $[a, b] \subset \mathbb{R}^n$, if, for $x, y \in \mathbb{R}^n$ of such interval, and for $\lambda \in [0, 1]$, it satisfies

$$f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y) \quad (29)$$

Under a single structure, that is, for example, a class of mapping $\mathbb{R}^n \rightarrow \mathbb{R}$, there can exist many loss functions to be considered. Take for example the setting of regression analysis. Then, ℓ can be either the absolute error, $|h(x) - c(x)|$, or the L^2 norm, $\|h(x), c(x)\|_2$. Different loss function provides different path and landscape, though they still share somewhat similar interpolation patterns, or either some might be subjected to, under the consideration of a loss landscape, local minima than others. Furthermore, the form and representation of loss functions is not discrete – but rather based of its interpretation and the supposition of the underlying notion required for the loss function to operate in. For example, if the setting is *generative*, the loss function would take the form of a *divergence measure* between two probability distribution. Notice however, that the action potential provided by the model in their operation disappears, or rather, is not considered, when using the loss function.

Using such loss function, for the case of no noises or interference uncertainty, if the goal is to minimize the empirical risk on such dataset, then we call this the method of *empirical risk minimization*, or ERM. This is born out of the consideration that in practice, the oracle $EX(c, \mathcal{D})$, and the actual generalization error $R(h)$ is not obtainable.

Appendix D - Bias-variance decomposition derivation in Geman et al. [1992]

Proof. We present the original derivation of bias-variance tradeoff and its analysis from Geman et al. [1992]. Geman’s paper is, by his own word, concerned of *parametric models*, i.e. hypothesis class without strong assumption of parameters defined, though the definition of bias-variance tradeoff afterward and its justification is made in the sense of parametric model. We partially clarify⁴ this with the following definition as customary.

Definition .0.4 (Parameterization). A **parametric model** is one that can be parameterized by a finite number of parameters. In general, for a hypothesis class $\mathcal{H}$ of all parametric hypothesis is then expressed as:

$$\mathcal{H} = \{f(x; \theta) : \theta \in \Theta \subset \mathbb{R}^d\} \quad (30)$$

⁴Generally, there is no definite definition on what would be considered non-parametric. Though, an example might be drawn from the (vanilla) neural network and Gaussian Process Regression (GPR)

where Θ is called the *parameter space*. A **nonparametric model** is one which cannot be parameterized by a fix number of parameters.

Suppose of a training dataset $\mathcal{D}$ of N 2-tuples $(\mathbf{x}_i, y_i)$, that is:

$$\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\} \quad |\mathcal{D}| = N, \mathbf{x} \in \mathbb{R}^d, y \in \mathbb{R} \quad (31)$$

The regression problem is to construct a function $f(\mathbf{x}) \rightarrow y$ based on $\mathcal{D}$, which is denoted $f(\mathbf{x}; \mathcal{D})$ to show this dependency. Then,

$$\begin{aligned} & \mathbb{E}[(y - f(\mathbf{x}; \mathcal{D}))^2 \mid \mathbf{x}, \mathcal{D}] \\ &= \mathbb{E}\left[\left((y - \mathbb{E}[y \mid \mathbf{x}] + (\mathbb{E}[y \mid \mathbf{x}] - f(\mathbf{x}; \mathcal{D})))\right)^2 \mid \mathbf{x}, \mathcal{D}\right] \\ &= \mathbb{E}\left[(y - \mathbb{E}[y \mid \mathbf{x}])^2 \mid \mathbf{x}, \mathcal{D}\right] + \left(\mathbb{E}[y \mid \mathbf{x}] - f(\mathbf{x}; \mathcal{D})\right)^2 \\ &\quad + 2\mathbb{E}\left[(y - \mathbb{E}[y \mid \mathbf{x}]) \mid \mathbf{x}, \mathcal{D}\right] \cdot (\mathbb{E}[y \mid \mathbf{x}] - f(\mathbf{x}; \mathcal{D})) \\ &= \mathbb{E}\left[(y - \mathbb{E}[y \mid \mathbf{x}]) \mid \mathbf{x}, \mathcal{D}\right] + (\mathbb{E}[y \mid \mathbf{x}] - f(\mathbf{x}; \mathcal{D}))^2 \\ &\geq \mathbb{E}\left[(y - \mathbb{E}[y \mid \mathbf{x}])^2 \mid \mathbf{x}, \mathcal{D}\right] \end{aligned} \quad (32)$$

Hence, we decompose it to:

$$\begin{aligned} & \mathbb{E}\left[\left((y - f(\mathbf{x}; \mathcal{D}))\right)^2 \mid \mathbf{x}, \mathcal{D}\right] \\ &= \mathbb{E}\left[(y - \mathbb{E}[y \mid \mathbf{x}])^2 \mid \mathbf{x}, \mathcal{D}\right] + (f(\mathbf{x}; \mathcal{D}) - \mathbb{E}[y \mid \mathbf{x}])^2 \end{aligned} \quad (33)$$

Where $\mathbb{E}\left[(y - \mathbb{E}[y \mid \mathbf{x}])^2 \mid \mathbf{x}, \mathcal{D}\right]$ is regarded to be a relative constant of variance on y given $\mathbf{x}$. Taking the expectation on $\mathcal{D}$, we gain:

$$\begin{aligned} & \mathbb{E}_{\mathcal{D}}\left\{\mathbb{E}\left[\left((y - f(\mathbf{x}; \mathcal{D}))\right)^2 \mid \mathbf{x}, \mathcal{D}\right]\right\} \\ &= \mathbb{E}_{\mathcal{D}}\left\{\mathbb{E}\left[(y - \mathbb{E}[y \mid \mathbf{x}])^2 \mid \mathbf{x}, \mathcal{D}\right] + (f(\mathbf{x}; \mathcal{D}) - \mathbb{E}[y \mid \mathbf{x}])^2\right\} \\ &= \mathbb{E}_{\mathcal{D}}\left\{\mathbb{E}\left[(y - \mathbb{E}[y \mid \mathbf{x}])^2 \mid \mathbf{x}, \mathcal{D}\right]\right\} + \\ &\quad \mathbb{E}_{\mathcal{D}}\left\{(f(\mathbf{x}; \mathcal{D}) - \mathbb{E}[y \mid \mathbf{x}])^2\right\} \end{aligned} \quad (34)$$

The second term is of importance, and is further decomposed to:

$$\begin{aligned}
& \mathbb{E}_{\mathcal{D}} \left\{ (f(\mathbf{x}; \mathcal{D}) - \mathbb{E}[y | \mathbf{x}])^2 \right\} \\
&= \mathbb{E}_{\mathcal{D}} \left\{ (f(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})] + \mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})] - \mathbb{E}[y | \mathbf{x}])^2 \right\} \\
&= \mathbb{E}_{\mathcal{D}} \left[(f(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})])^2 \right] \\
&\quad + \mathbb{E}_{\mathcal{D}} \left[(\mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})] - \mathbb{E}[y | \mathbf{x}])^2 \right] \\
&\quad + 2\mathbb{E}_{\mathcal{D}} \left[(f(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})]) \cdot (\mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})] - \mathbb{E}[y | \mathbf{x}]) \right] \\
&= \underbrace{\{\mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})] - \mathbb{E}[y | \mathbf{x}]\}}_{\text{bias term}} \\
&\quad + \underbrace{\mathbb{E}_{\mathcal{D}} \left\{ (f(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[f(\mathbf{x}; \mathcal{D})])^2 \right\}}_{\text{variance term}}
\end{aligned} \tag{35}$$

which yields the desired bias-variance tradeoff. $\square$

Practical consideration

While we have been talking in the learning setting pretty loosely on the side of the data itself, it is the data that gives us the problem setting. For example, if the data is a pair (X, Y) of input output representation, then we know that our working space would be $\mathbb{R}^n \times \mathbb{R}^m$ or any of its subset, given the fact that data must be encoded into numerical representations, even in discrete or logical case (in such case, typically, 0 – 1 pair is enough). In more complex and complicated setting, for example, on the structure of a graph and its embedded representation, it is much more difficult to see the *form of dataset* to be, hence, also the working space.

Another thing with data is the *availability of the data* and the *distribution of the data* also matters. In some cases, our data consists of only skewed or partitioned sections of the actual concept range, at least up to the observable criterion of the concept (some concept does not exist in certain range, at least in a numerical example, and some other reasons). In the other case, the availability of data, meaning the density of data, as well as added *observational effect* also alters the learning setting, simply because learning theory relies on, at least from our glance, simple statistical, observational learning from observation space, without the added structures. An actual, practical concern however, lies in the range of **practical empirical-generalization** setting.

Observational space partitioning

Previously, we established that at least in said formal setting, generalization errors cannot be known beforehand. Thereby, all of our operations, procedures, processes has to be conducted using the observed data, or the sample space $\mathcal{S}$. A solution to utilize that is to consider the generalization region as the region of *new, unseen observations* that while still is of the same concept c , but it is indeed, unseen if the model is never provided of such instances. Hence, it prompts the utilization of limited resources by then to partition the sample space into various subspaces, that would be then fed into the learning process as observational space. This is simply called **sample set partitioning**. Denoted the number of partition (discrete) as k , then for $k = 2$, then this is called the **train-test partition** (or dataset). This contains a dataset $\mathbf{x}_1$ and $\mathbf{x}_2$ of the

same original sample set, such that $x_1/x_2 \geq 1$ (one is bigger than another, usually). The ratio does indeed affect the overall running procedure, accordingly⁵.

Based on the above conclusion, the empirical error and generalized error can be re-calculated to fit the general scheme. For $|S| = m$ and partition $k + p = m$ for train and test partition respectively,

$$\hat{R}_k(h) = \frac{1}{k} \sum_{i=1}^k \ell\{h(x_i), c(x_i)\}, \quad (36)$$

$$\hat{R}_p(h_{\mathcal{A}}) = \frac{1}{k} \sum_{i=1}^p \ell\{h_{\mathcal{A}}(x_i), c(x_i)\} \quad (37)$$

where $h_{\mathcal{A}}$ is the result from the learning algorithm (procedure), for $h^* \neq h_{\mathcal{A}}$. Then, the empirical learning seeks to optimize the first term, that is, for the objective

$$\arg \min_{h \in \mathcal{H}} \hat{R}_k(h) = \arg \min_{h \in \mathcal{H}} \frac{1}{k} \sum_{i=1}^k \ell\{h(x_i), c(x_i)\} \quad (38)$$

and the generalization learning as:

$$\begin{aligned} & \arg \min_{h \in \mathcal{H}} \hat{R}_p(h_{\mathcal{A}}) \\ &= \hat{R}_k(h_{\mathcal{A}}) + \arg \min_{h \in \mathcal{H}} \hat{R}_p(h_{\mathcal{A}}) \\ &= \arg \min_{h \in \mathcal{H}} \left[\frac{1}{k} \sum_{i=1}^k \ell\{h_{\mathcal{A}}(x_i), c(x_i)\} + \frac{1}{k} \sum_{i=1}^p \ell\{h_{\mathcal{A}}(x_i), c(x_i)\} \right] \\ &= \arg \min_{h \in \mathcal{H}} \left[\epsilon + \frac{1}{k} \sum_{i=1}^p \ell\{h_{\mathcal{A}}(x_i), c(x_i)\} \right] \end{aligned} \quad (39)$$

By the definition order, the generalization error can only be calculated post- $\mathcal{A}$. Hence, we often called the empirical, or ERM-generalizer case to be dynamic, and the solution for ERM is tested statically on generalization set. This partition of order of computation though, brings up the problem of updating factors and others contributing potential that will damage or diffuse the measure.

For cross-validation and other randomize partition, the formulation requires more care in the notion by itself. The structural risk minimization scheme dilemma can be understood from the above consideration naturally. We have effectively defined the two proxies, **heuristic empirical risk** and **heuristic generalization risk**.

⁵If $k > 3$ and above, we call it the n th partitioning. If we further consider the choice of observational spaces being provided, then if $k > 3$, and subspaces are chosen randomly, it is called **cross-validating partition**. The effect of them are fairly well-understood, and we might as well refer to authoritative sources on such issue.

Bibliography

- Bias-variance tradeoffs in program analysis | Proceedings of the 41st ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, a. URL <https://dl.acm.org/doi/10.1145/2535838.2535853>.
- [PDF] A Unifeid Bias-Variance Decomposition and its Applications | Semantic Scholar, b. URL <https://www.semanticscholar.org/paper/A-Unifeid-Bias-Variance-Decomposition-and-its-Domingos/e1ed9d24db5e8f7ab326aeb797e965a94f5ad6d3>.
- [PDF] A Unifeid Bias-Variance Decomposition and its Applications | Semantic Scholar. URL <https://www.semanticscholar.org/paper/A-Unifeid-Bias-Variance-Decomposition-and-its-Domingos/e1ed9d24db5e8f7ab326aeb797e965a94f5ad6d3>.
- Bridging the gap between practice and theory in deep learning. In *ICLR 2024 Workshop*, Vienna, Austria, May 2024. International Conference on Learning Representations. URL <https://sites.google.com/view/bgpt-iclr24>. Workshop held on May 11, 2024.
- Mikkel Abrahamsen, Linda Kleist, and Tillmann Miltzow. Training neural networks is $\exists\mathbb{R}$ -complete. In *Proceedings of NeurIPS 2021*, 2021. URL <https://arxiv.org/abs/2102.09798>. arXiv:2102.09798.
- Panos Achlioptas. Stochastic Gradient Descent in Theory and Practice. *Lecture note, Stanford's AI*.
- D. H. Ackley, G. E. Hinton, and T. J. Sejnowski. A learning algorithm for boltzmann machines. *Cognitive Science*, 9(1):147–169, 1985a. doi: 10.1016/0364-0213(85)90014-6. URL [https://doi.org/10.1016/0364-0213\(85\)90014-6](https://doi.org/10.1016/0364-0213(85)90014-6).
- D. H. Ackley, G. E. Hinton, and T. J. Sejnowski. A learning algorithm for boltzmann machines. *Cognitive Science*, 9: 147–169, 1985b.
- Ben Adcock and Nick Dexter. The gap between theory and practice in function approximation with deep neural networks, 2021. URL <https://arxiv.org/abs/2001.07523>.
- Ben Adlam and Jeffrey Pennington. Understanding double descent requires a fine-grained bias-variance decomposition, 2020. URL <https://arxiv.org/abs/2011.03321>.
- Madhu S. Advani and Andrew M. Saxe. High-dimensional dynamics of generalization error in neural networks. 2017a. URL <https://arxiv.org/abs/1710.03667>. arXiv:1710.03667; (published versions/derivatives exist).
- Madhu S. Advani and Andrew M. Saxe. High-dimensional dynamics of generalization error in neural networks, 2017b. URL <https://arxiv.org/abs/1710.03667>.
- Madhu S. Advani, Subhaneil Lahiri, and Surya Ganguli. Statistical mechanics of complex neural systems and high dimensional data. *Journal of Statistical Mechanics: Theory and Experiment*, 2013. URL <https://arxiv.org/abs/1301.7115>. arXiv:1301.7115.

- Oskar Allerbo. Changing the kernel during training leads to double descent in kernel regression, 2025. URL <https://arxiv.org/abs/2311.01762>.
- Pierre Alquier. User-friendly introduction to pac-bayes bounds. *Foundations and Trends in Machine Learning*, 17(2): 174–303, 2024. ISSN 1935-8245. doi: 10.1561/22000000100. URL <http://dx.doi.org/10.1561/22000000100>.
- Shun-ichi Amari. Neural learning in structured parameter spaces – natural riemannian gradient. In *Advances in Neural Information Processing Systems 9 (NIPS’96)*, pages 287–293, 1996. Also cited as “Natural Riemannian Gradient”.
- Shun-ichi Amari. Natural gradient works efficiently in learning. *Neural Computation*, 10(2):251–276, 1998. doi: 10.1162/089976698300017746.
- D. J. Amit, H. Gutfreund, and H. Sompolinsky. Spin-glass models of neural networks. *Physical Review A*, 32(2): 1007–1018, 1985. doi: 10.1103/PhysRevA.32.1007. URL <https://doi.org/10.1103/PhysRevA.32.1007>.
- Mouin Ben Ammar, David Brellmann, Arturo Mendoza, Antoine Manzanera, and Gianni Franchi. Double descent meets out-of-distribution detection: Theoretical insights and empirical analysis on the role of model complexity, 2025. URL <https://arxiv.org/abs/2411.02184>.
- Yossi Arjevani, Joan Bruna, Joe Kileel, Elzbieta Polak, and Matthew Trager. Geometry and optimization of shallow polynomial networks, 2025. URL <https://arxiv.org/abs/2501.06074>.
- Sanjeev Arora, Simon Du, Wei Hu, Zhiyuan Li, and Ruslan Salakhutdinov. On exact computation with an infinitely wide neural net. *Advances in Neural Information Processing Systems (NeurIPS)*, 32, 2019.
- Kai Arulkumaran, Marc Peter Deisenroth, Miles Brundage, and Anil Anthony Bharath. A brief survey of deep reinforcement learning. *IEEE Signal Processing Magazine / arXiv*, 2017. URL <https://arxiv.org/abs/1708.05866>. Useful survey and roadmap for deep-RL research.
- Yasaman Bahri, Jonathan Kadmon, Jeffrey Pennington, Sam S. Schoenholz, Jascha Sohl-Dickstein, and Surya Ganguli. Statistical mechanics of deep learning. *Annual Review of Condensed Matter Physics*, 11:501–528, 2020. doi: 10.1146/annurev-conmatphys-031119-050745. URL <https://doi.org/10.1146/annurev-conmatphys-031119-050745>.
- Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, et al. Constitutional ai: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073*, 2022. URL <https://arxiv.org/abs/2212.08073>. Anthropic’s approach to AI alignment using AI-generated feedback.
- Maria-Florina Balcan, Steve Hanneke, and Jennifer Wortman Vaughan. The true sample complexity of active learning. *Machine Learning*, 80(2–3):111–139, 2010. doi: 10.1007/s10994-010-5174-y.
- Jean Barbier, Francesco Camilli, Minh-Toan Nguyen, Mauro Pastore, and Rudy Skerk. Statistical physics of deep learning: Optimal learning of a multi-layer perceptron near interpolation, 2025. URL <https://arxiv.org/abs/2510.24616>.
- Pablo Barceló, Mikaël Monet, Jorge Pérez, and Bernardo Subercaseaux. Model Interpretability through the Lens of Computational Complexity, 2020.
- Pablo Barceló, Mikaël Monet, Jorge Pérez, and Bernardo Subercaseaux. Model interpretability through the lens of computational complexity, 2020. URL <https://arxiv.org/abs/2010.12265>.
- Peter L. Bartlett and Shahar Mendelson. Rademacher and gaussian complexities: Risk bounds and structural results. *Journal of Machine Learning Research*, 3:463–482, 2002. URL <https://www.jmlr.org/papers/volume3/bartlett02a/bartlett02a.pdf>. Introduces data-dependent complexity measures that often yield tighter generalization bounds than VC dimension.

- Peter L. Bartlett and Shahar Mendelson. Empirical minimization. In *Probability Theory and Related Fields*, volume 135, pages 311–334, 2005. doi: 10.1007/s00440-005-0460-0.
- Peter L. Bartlett, Dylan J. Foster, and Matus Telgarsky. Spectrally-normalized margin bounds for neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, pages 6241–6250, 2017.
- Peter L. Bartlett, Philip M. Long, Gábor Lugosi, and Alexander Tsigler. Benign overfitting in linear regression. *Proceedings of the National Academy of Sciences*, 117(48):30063–30070, 2020. doi: 10.1073/pnas.1907378117.
- Peter L. Bartlett, Andrea Montanari, and Alexander Rakhlin. Deep learning: a statistical viewpoint. *Acta Numerica*, 30:87–201, 2021.
- Ali Behrouz, Meisam Razaviyayn, Peilin Zhong, and Vahab Mirrokni. Nested learning: The illusion of deep learning architectures. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025. URL <https://openreview.net/forum?id=nbMeRvNb7A>.
- Mikhail Belkin, Siyuan Ma, and Soumik Mandal. To understand deep learning we need to understand kernel learning, 2018. URL <https://arxiv.org/abs/1802.01396>.
- Mikhail Belkin, Daniel Hsu, Siyuan Ma, and Soumik Mandal. Reconciling modern machine learning practice and the bias-variance trade-off. *Proc. Natl. Acad. Sci. U.S.A.*, 116(32):15849–15854, August 2019a. ISSN 0027-8424, 1091-6490. doi: 10.1073/pnas.1903070116. URL <http://arxiv.org/abs/1812.11118>. arXiv:1812.11118 [cs, stat].
- Mikhail Belkin, Daniel Hsu, Siyuan Ma, and Soumik Mandal. Reconciling modern machine learning practice and the bias-variance trade-off. *Proc. Natl. Acad. Sci. U.S.A.*, 116(32):15849–15854, August 2019b. ISSN 0027-8424, 1091-6490. doi: 10.1073/pnas.1903070116. URL <http://arxiv.org/abs/1812.11118>. arXiv:1812.11118 [cs, stat].
- Mikhail Belkin, Daniel Hsu, and Ji Xu. Two models of double descent for weak features. *SIAM Journal on Mathematics of Data Science*, 2(4):1167–1180, January 2020. ISSN 2577-0187. doi: 10.1137/20m1336072. URL <http://dx.doi.org/10.1137/20m1336072>.
- Richard Bellman. Dynamic programming. *Princeton University Press*, 1957. Foundational book on dynamic programming.
- Dimitri P. Bertsekas. *Dynamic Programming and Optimal Control*. Athena Scientific, vols. i & ii, 4th ed. (selected vols.) edition, 2017. Fundamental reference on DP and optimal control.
- Dimitri P. Bertsekas and John N. Tsitsiklis. *Neuro-Dynamic Programming*. Athena Scientific, Belmont, MA, 1996.
- Christopher M. Bishop. *Neural Networks for Pattern Recognition*. Oxford University Press, 1995.
- Avrim L. Blum and Ronald L. Rivest. Training a 3-node neural network is np-complete. *Neural Networks*, 5(1): 117–127, 1992.
- Olivier Bousquet and Andre Elisseeff. Stability and generalization. In *Journal of Machine Learning Research*, volume 2, pages 499–526, 2002. One of the foundational works on algorithmic stability capturing generalization independent of hypothesis space complexity.
- Olivier Bousquet, Steve Hanneke, Shay Moran, Ramon van Handel, and Amir Yehudayoff. A theory of universal learning, 2020. URL <https://arxiv.org/abs/2011.04483>.
- JP. Boyd. Divergence (runge phenomenon) for least-squares polynomial approximation on an equispaced grid and mock-chebyshev subset interpolation. *Applied Mathematics and Computation*, 213(1):292–303, 2009. doi: 10.1016/j.amc.2008.12.087. Study of divergence behavior and mitigation by Mock-Chebyshev grid.

- Ronald J. Brachman and Hector J. Levesque. *Knowledge Representation and Reasoning*. Morgan Kaufmann, 2004.
- David Brellmann, Eloise Berthier, David Filliat, and Goran Frehse. On double descent in reinforcement learning with LSTD and random features. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=9RIbNmx984>.
- Michael M. Bronstein, Joan Bruna, Yann LeCun, Arthur Szlam, and Pierre Vandergheynst. Geometric deep learning: Going beyond euclidean data. *IEEE Signal Processing Magazine*, 34(4):18–42, July 2017. ISSN 1558-0792. doi: 10.1109/msp.2017.2693418. URL <http://dx.doi.org/10.1109/MSP.2017.2693418>.
- Michael M. Bronstein, Joan Bruna, Taco Cohen, and Petar Veličković. Geometric deep learning: Grids, groups, graphs, geodesics, and gauges, 2021. URL <https://arxiv.org/abs/2104.13478>.
- Rodney A. Brooks. Intelligence without representation. *Artificial Intelligence*, 47:139–159, 1991. doi: 10.1016/0004-3702(91)90053-M. URL <https://people.csail.mit.edu/brooks/papers/representation.pdf>.
- Gavin Brown and Riccardo Ali. Bias/variance is not the same as approximation/estimation. *Transactions on Machine Learning Research*, 2024. ISSN 2835-8856. URL <https://openreview.net/forum?id=4TnFbv16hK>.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. In *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901. Neural Information Processing Systems, 2020a. URL <https://arxiv.org/abs/2005.14165>. NeurIPS 2020. Introduced GPT-3 with 175B parameters demonstrating few-shot learning.
- Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. *arXiv preprint*, 2020b. URL <https://arxiv.org/abs/2005.14165>.
- Sebastian Buschjäger, Lukas Pfahler, and Katharina Morik. Generalized Negative Correlation Learning for Deep Ensembling, December 2020. URL <http://arxiv.org/abs/2011.02952>. arXiv:2011.02952 [cs, stat].
- Augustin-Louis Cauchy. General method for the resolution of simultaneous systems of equations. *Comp. Rend. Sci. Paris*, 25:536–538, 1847.
- Nicolò Cesa-Bianchi and Gábor Lugosi. *Prediction, Learning, and Games*. Cambridge University Press, 2006. URL <https://www.cambridge.org/core/books/prediction-learning-and-games/8C6D3E2F7D9D1C7F7A98C9F1C9C5E2F2>.
- Pratik Chaudhari, Anna Choromanska, Stefano Soatto, and et al. Entropy-sgd: Biasing gradient descent into wide valleys. 2016. arXiv:1611.01838; JSTAT 2019.
- Yilan Chen, Wei Huang, Lam M. Nguyen, and Tsui-Wei Weng. On the equivalence between neural network and support vector machine. *arXiv preprint arXiv:2111.06063*, 2021. URL <https://arxiv.org/abs/2111.06063>.
- Vladimir Cherkassky and Eng Hock Lee. To understand double descent, we need to understand vc theory. *Neural Networks*, 169:242–256, 2024. doi: 10.1016/j.neunet.2023.10.014. URL <https://doi.org/10.1016/j.neunet.2023.10.014>.
- Lenaic Chizat and Francis Bach. On the global convergence of gradient descent for over-parameterized models using optimal transport. *Advances in Neural Information Processing Systems (NeurIPS)*, 31, 2018.

- Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, et al. Palm: Scaling language modeling with pathways. *arXiv preprint arXiv:2204.02311*, 2022. URL <https://arxiv.org/abs/2204.02311>. Google’s 540B parameter model demonstrating breakthrough capabilities.
- William F. Clocksin and Christopher S. Mellish. *Programming in Prolog: Using the ISO Standard*. Springer, 4th edition, 2003.
- Jeremy Cohen, Tomaso Poggio, Alexander Rakhlin, and Ohad Shamir. Gradient descent on neural networks typically occurs at the edge of stability. *arXiv preprint*, 2021.
- Caroline Collange, David Defour, Stef Graillat, and Roman Iakymchuk. Numerical reproducibility for the parallel reduction on multi- and many-core architectures. *Parallel Computing*, 49:83–97, 2015. ISSN 0167-8191. doi: <https://doi.org/10.1016/j.parco.2015.09.001>. URL <https://www.sciencedirect.com/science/article/pii/S0167819115001155>.
- Robert M. Corless and Leili Rafiee Sevyeri. The runge example for interpolation and wilkinson’s examples for rootfinding. *arXiv preprint*, 0(arXiv:1804.08561), 2018. URL <https://arxiv.org/abs/1804.08561>. Modern analysis of Runge’s phenomenon and conditioning.
- A. Crisanti and H.-J. Sommers. The spherical p-spin interaction spin glass model: the statics. *Zeitschrift für Physik B*, 87:341–354, 1992.
- Alicia Curth, Alan Jeffares, and Mihaela van der Schaar. A u-turn on double descent: Rethinking parameter counting in statistical learning, 2023. URL <https://arxiv.org/abs/2310.18988>.
- G. Cybenko. Approximation by superpositions of a sigmoidal function. *Mathematics of Control, Signals and Systems*, 2:303–314, 1989. doi: 10.1007/BF02551274.
- Stéphane d’Ascoli, Levent Sagun, and Giulio Biroli. Triple descent and the two kinds of overfitting: where & why do they appear? In *Advances in Neural Information Processing Systems*, volume 33, pages 3058–3069. Curran Associates, Inc., 2020. URL <https://proceedings.neurips.cc/paper/2020/hash/1fd09c5f59a8ff35d499c0ee25a1d47e-Abstract.html>.
- Amit Daniely and Gal Katzhendler. Approximate description length, covering numbers, and vc dimension. *arXiv*, 2209.12882, 2022.
- Adnan Darwiche and Pierre Marquis. A Knowledge Compilation Map. *Journal of Artificial Intelligence Research*, 17: 229–264, 2002.
- Xander Davies, Lauro Langosco, and David Krueger. Unifying Grokking and Double Descent, March 2023. URL <http://arxiv.org/abs/2303.06173>. arXiv:2303.06173 [cs].
- Bruno de Finetti. *Theory of Probability: A Critical Introductory Treatment*. John Wiley & Sons, 1990. Originally published in Italian, 1974.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint*, 2018. URL <https://arxiv.org/abs/1810.04805>.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186. Association for Computational Linguistics, 2019. URL <https://arxiv.org/abs/1810.04805>. NAACL 2019. Introduced bidirectional pre-training for language representations.

- Kai Diethelm. The limits of reproducibility in numerical simulation. *Computing in Science and Engineering*, 14(1):64–72, 2012. doi: 10.1109/MCSE.2011.21. URL https://opus4.kobv.de/opus4-fhws/frontdoor/deliver/index/docId/2063/file/Diethelm_Limits_reproducibility_numerical_sim.pdf.
- Pedro Domingos. A unified bias-variance decomposition and its applications. In *Proceedings of the 17th International Conference on Machine Learning (ICML-2000)*, pages 231–238. Morgan Kaufmann, 2000a. URL <https://homes.cs.washington.edu/~pedrod/papers/mlc00a.pdf>.
- Pedro Domingos. A unified bias-variance decomposition for zero-one and squared loss. In *Proceedings of the 17th National Conference on Artificial Intelligence (AAAI-2000)*, pages 564–569. AAAI Press, 2000b. URL <https://homes.cs.washington.edu/~pedrod/papers/aaai00.pdf>.
- Pedro M. Domingos. A unified bias-variance decomposition for zero-one and squared loss. In *AAAI/IAAI*, 2000c. URL <https://api.semanticscholar.org/CorpusID:2063488>.
- Pedro M. Domingos. A Unifeid Bias-Variance Decomposition and its Applications. In *Semantic Scholar*, June 2000d. URL <https://www.semanticscholar.org/paper/A-Unifeid-Bias-Variance-Decomposition-and-its-Domingos/e1ed9d24db5e8f7ab326aeb797e965a94f5ad6d3>.
- David L. Donoho, Arian Maleki, and Andrea Montanari. Message passing algorithms for compressed sensing. 2009. arXiv:0907.3574; Proc. Natl. Acad. Sci. 2009.
- Finale Doshi-Velez and Been Kim. Towards a rigorous science of interpretable machine learning. *arXiv preprint arXiv:1702.08608*, 2017. URL <https://arxiv.org/abs/1702.08608>.
- Hubert L. Dreyfus. Alchemy and artificial intelligence. Technical Report P-3244, RAND Corporation, 1965. URL <https://www.rand.org/pubs/papers/P3244.html>.
- Hubert L. Dreyfus. *What Computers Can't Do: A Critique of Artificial Reason*. Harper & Row, 1972. ISBN 0060110821.
- Hubert L. Dreyfus and Stuart E. Dreyfus. *Mind Over Machine: The Power of Human Intuition and Expertise in the Era of the Computer*. Free Press, 1986. ISBN 0029080606.
- O. Duranthon and L. Zdeborová. Statistical physics analysis of graph neural networks: Approaching optimality in the contextual stochastic block model. *Physical Review X*, 15(4), November 2025. ISSN 2160-3308. doi: 10.1103/lfxj-hbsk. URL <http://dx.doi.org/10.1103/lfxj-hbsk>.
- George Casella E. L. Lehmann. *Theory of Point Estimation*. Springer Texts in Statistics. Springer-Verlag, New York, 1998. ISBN 978-0-387-98502-2. doi: 10.1007/b98854. URL <http://link.springer.com/10.1007/b98854>.
- Ronen Eldan and Ohad Shamir. The power of depth for feedforward neural networks. In *Proceedings of the 29th Annual Conference on Learning Theory (COLT)*, volume 49 of *Proceedings of Machine Learning Research*, pages 907–940, 2016. URL <https://proceedings.mlr.press/v49/eldan16.html>.
- C. Elliot. *On De Finetti's Philosophy of Probability*. PhD thesis, Tilburg University, June 2019. URL <https://research.tilburguniversity.edu/en/publications/on-de-finettis-philosophy-of-probability>. Available from Tilburg University repository.
- A. Engel and C. Van den Broeck. *Statistical Mechanics of Learning*. Cambridge University Press, Cambridge, 2001. ISBN 978-0521773072.
- Damien Ernst, Pierre Geurts, and Louis Wehenkel. Tree-backup and batch-mode reinforcement learning. *Journal of Machine Learning Research*, 2005. Classic work on batch/offline RL techniques.

- Molavi et al. Model Complexity, Expectations, and Asset Prices | The Review of Economic Studies | Oxford Academic. URL <https://academic.oup.com/restud/article-abstract/91/4/2462/7222145?redirectedFrom=fulltext&login=false>.
- Michael F. Faulkner and Samuel Livingstone. Sampling algorithms in statistical physics: A guide for statistics and machine learning. *Statistical Science*, 39(1), February 2024. ISSN 0883-4237. doi: 10.1214/23-sts893. URL <http://dx.doi.org/10.1214/23-STs893>.
- Andrea E. V. Ferrari, Prateek Gupta, and Nabil Iqbal. Machine learning and optimization-based approaches to duality in statistical physics, 2024. URL <https://arxiv.org/abs/2411.04838>.
- Gino Del Ferraro, Chuang Wang, Dani Martí, and Marc Mézard. Cavity method: Message passing from a physics perspective. 2014. URL <https://arxiv.org/abs/1409.3048>.
- R. Fletcher and C. M. Reeves. Function minimization by conjugate gradients. *The Computer Journal*, 7(2):149–154, 1964. doi: 10.1093/comjnl/7.2.149.
- Stanislav Fort, Adam Scherlis, and Yann Dauphin. Deep learning versus kernel learning: An empirical study of loss landscape geometry. *arXiv preprint*, 2019.
- Jerome Friedman, Trevor Hastie, Holger Höfling, and Robert Tibshirani. Pathwise coordinate optimization. *The Annals of Applied Statistics*, 1(2):302–332, 2007. doi: 10.1214/07-AOAS131.
- Justin Fu, Avi Kumar, Ofir Nachum, George Tucker, and Sergey Levine. D4rl: Datasets for deep data-driven reinforcement learning. *arXiv preprint arXiv:2004.07219*, 2020. Benchmark dataset for offline RL.
- Kuniyiko Fukushima. Neocognitron: A self-organizing neural network model for a mechanism of pattern recognition unaffected by shift in position. *Biological Cybernetics*, 36:193–202, 1980. doi: 10.1007/BF00344251.
- E. Gardner and B. Derrida. Optimal storage properties of neural network models. *Journal of Physics A: Mathematical and General*, 21(1):271–284, 1988. doi: 10.1088/0305-4470/21/1/031. URL <https://doi.org/10.1088/0305-4470/21/1/031>.
- Elizabeth Gardner. The space of interactions in neural network models. *Journal of Physics A: Mathematical and General*, 21:257–270, 1988.
- Stuart Geman, Elie Bienenstock, and Rene Doursat. Neural networks and the bias/variance dilemma. *Neural Computation*, 4(1):1–58, 1992. doi: 10.1162/neco.1992.4.1.1.
- Joseph C. Giarratano and Gary D. Riley. *Expert Systems: Principles and Programming*. Thomson Learning, 4th edition, 1998. ISBN 9780534957866.
- Ian Goodfellow, Yoshua Bengio, Aaron Courville, and Yoshua Bengio. *Deep learning*, volume 1. MIT Press, 2016.
- Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial nets. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2014. URL <https://arxiv.org/abs/1406.2661>.
- Suriya Gunasekar, Blake Woodworth, and Nathan Srebro. Mirrorless mirror descent: A natural derivation of mirror descent. In *Proceedings of the 24th International Conference on Artificial Intelligence and Statistics (AISTATS 2021)*, volume 130, pages 2305–2313. PMLR, 2021.
- Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, and Sergey Levine. Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor. *International Conference on Machine Learning (ICML)*, 2018. SAC — popular off-policy, entropy-regularised RL algorithm.

- Bruce Hajek and Maxim Raginsky. *Statistical Learning Theory*, volume 1. 2021. URL <https://maxim.ece.illinois.edu/teaching/SLT/>.
- William L. Hamilton. Graph Representation Learning. *Synthesis Lectures on Artificial Intelligence and Machine Learning*, 14(3):1–159, 2022.
- X. Y. Han, Vardan Papyan, and David L. Donoho. Neural collapse: A review. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2022. doi: 10.1109/TPAMI.2022.3158673.
- Stevan Harnad. The symbol grounding problem. *Physica D: Nonlinear Phenomena*, 42:335–346, 1990. doi: 10.1016/0167-2789(90)90087-6.
- Ouns El Harzli, Bernardo Cuenca Grau, Guillermo Valle-Perez, and Ard A. Louis. Double-descent curves in neural networks: a new perspective using gaussian processes, 2023. URL <https://arxiv.org/abs/2102.07238>.
- Trevor Hastie, Andrea Montanari, Saharon Rosset, and Ryan J Tibshirani. Surprises in high-dimensional ridgeless least squares interpolation. *arXiv preprint arXiv:1903.08560*, 2019.
- Simon S. Haykin. *Neural Networks: A Comprehensive Foundation*. Macmillan, 1994.
- Elad Hazan. Introduction to online convex optimization. *Foundations and Trends in Optimization*, 2(3–4):157–325, 2016. URL <https://www.cs.princeton.edu/~ehazan/papers/OCO.pdf>.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 770–778, 2016. URL <https://arxiv.org/abs/1512.03385>.
- D. O. Hebb. *The Organization of Behavior: A Neuropsychological Theory*. John Wiley & Sons, New York, 1949.
- Rebekka Heckel and Murat Yilmaz. Early stopping in deep networks: Double descent and how to eliminate it. *arXiv preprint*, 2020. URL <https://arxiv.org/abs/2007.10099>. discusses epoch-wise double descent and mitigation strategies.
- Thomas Hellström, Virginia Dignum, and Suna Bensch. Bias in Machine Learning – What is it Good for?, September 2020. URL <http://arxiv.org/abs/2004.00686>. arXiv:2004.00686 [cs].
- Magnus R. Hestenes and Eduard Stiefel. Methods of conjugate gradients for solving linear systems. *Journal of Research of the National Bureau of Standards*, 49(6):409–436, 1952. doi: 10.6028/jres.049.044.
- G. E. Hinton. A practical guide to training restricted boltzmann machines. In G. Montavon, G. B. Orr, and K.-R. Müller, editors, *Neural Networks: Tricks of the Trade (2nd ed.)*, pages 599–619. Springer, 2012. doi: 10.1007/978-3-642-35289-8_32. URL https://doi.org/10.1007/978-3-642-35289-8_32.
- Geoffrey E. Hinton. Training products of experts by minimizing contrastive divergence. *Neural Computation*, 14(8):1771–1800, 2002.
- Geoffrey E. Hinton, Simon Osindero, and Yee-Whye Teh. A fast learning algorithm for deep belief nets. *Neural Computation*, 18:1527–1554, 2006. URL <https://www.cs.toronto.edu/~fritz/absps/ncfast.pdf>.
- Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. 2020. arXiv:2006.11239.
- Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural Computation*, 9(8):1735–1780, 1997. doi: 10.1162/neco.1997.9.8.1735.

- Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, et al. Training compute-optimal large language models. *arXiv preprint arXiv:2203.15556*, 2022. URL <https://arxiv.org/abs/2203.15556>. Introduced Chinchilla and revised scaling laws for compute-optimal training.
- Asbjørn Holk, Claudia Strauch, and Lukas Trottnner. Statistical guarantees for denoising reflected diffusion models, 2025. URL <https://arxiv.org/abs/2411.01563>.
- J. J. Hopfield. Neural networks and physical systems with emergent collective computational abilities. *Proceedings of the National Academy of Sciences*, 79(8):2554–2558, 1982. doi: 10.1073/pnas.79.8.2554.
- Kurt Hornik. Approximation capabilities of multilayer feedforward networks. *Neural Networks*, 4(2):251–257, 1991.
- Wanda Hou and Yi-Zhuang You. Machine learning renormalization group for statistical physics, 2023. URL <https://arxiv.org/abs/2306.11054>.
- Erdong Hu. An overview of double descent and overparameterization. *independent*, (0), 2021.
- Xia Hu, Lingyang Chu, Jian Pei, Weiqing Liu, and Jiang Bian. Model complexity of deep learning: A survey, 2021a. URL <https://arxiv.org/abs/2103.05127>.
- Xia Hu, Lingyang Chu, Jian Pei, Weiqing Liu, and Jiang Bian. Model Complexity of Deep Learning: A Survey, August 2021b. URL <http://arxiv.org/abs/2103.05127>. arXiv:2103.05127 [cs].
- Baihe Huang, Shanda Li, Tianhao Wu, Yiming Yang, Ameet Talwalkar, Kannan Ramchandran, Michael I. Jordan, and Jiantao Jiao. Sample complexity and representation ability of test-time scaling paradigms, 2025. URL <https://arxiv.org/abs/2506.05295>.
- Arthur Jacot, Francois Gabriel, and Clement Hongler. Neural tangent kernel: Convergence and generalization in neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018. Kernel view of wide-network behavior.
- Gareth James, Trevor Hastie, Robert Tibshirani, and Jerome H. Friedman. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer Series in Statistics. Springer, 2nd edition, 2009.
- Romuald A. Janik and Przemek Witaszczyk. Complexity for deep neural networks and other characteristics of deep feature representations, 2021. URL <https://arxiv.org/abs/2006.04791>.
- Leslie P. Kaelbling, Michael L. Littman, and Andrew W. Moore. *Reinforcement Learning: A Survey*. 1996. Classic early survey (also useful historically).
- Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models, 2020. URL <https://arxiv.org/abs/2001.08361>.
- Sayash Kapoor and Arvind Narayanan. Leakage and the reproducibility crisis in machine-learning-based science. *Patterns*, 4(9):100804, 2023. ISSN 2666-3899. doi: <https://doi.org/10.1016/j.patter.2023.100804>. URL <https://www.sciencedirect.com/science/article/pii/S2666389923001599>.
- Petteri Kaskenpalo and Stephen G. MacDonell. Valuing evaluation: methodologies to bridge research and practice. In *Proceedings of the 2nd international workshop on Evidential assessment of software technologies*, ESEM '12, pages 27–32. ACM, September 2012. doi: 10.1145/2372233.2372242. URL <http://dx.doi.org/10.1145/2372233.2372242>.
- Michael J. Kearns and Umesh V. Vazirani. *An introduction to computational learning theory*. MIT Press, Cambridge, MA, USA, 1994. ISBN 0262111934.

- Mohammad Emtiyaz Khan and H° avarid Rue. The Bayesian Learning Rule, June 2024. URL <http://arxiv.org/abs/2107.04562>. arXiv:2107.04562 [cs, stat] version: 4.
- Diederik P. Kingma and Max Welling. Auto-encoding variational bayes. 2013. arXiv:1312.6114.
- Lasse Kliemann and Peter Sanders, editors. *Algorithm Engineering: Selected Results and Surveys*, volume 9220 of *Lecture Notes in Computer Science*. Springer, Cham, 2016. ISBN 978-3-319-49487-6. doi: 10.1007/978-3-319-49487-6. URL <https://doi.org/10.1007/978-3-319-49487-6>.
- Teuvo Kohonen. Self-organized formation of topologically correct feature maps. *Biological Cybernetics*, 43:59–69, 1982. doi: 10.1007/BF00337288.
- V. R. Konda and J. N. Tsitsiklis. Actor-critic algorithms. In *Advances in Neural Information Processing Systems (NIPS)*, pages 1008–1014, 2000.
- Ilya Kostrikov, Seyedkamyar Ghasemipour, and Sergey Levine. Offline reinforcement learning with implicit q-learning. *arXiv preprint arXiv:2110.06169*, 2021. Representative of modern offline RL developments.
- Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton. Imagenet classification with deep convolutional neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2012. URL <https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>.
- Marc Lafon and Alexandre Thomas. Understanding the Double Descent Phenomenon in Deep Learning, March 2024. URL <http://arxiv.org/abs/2403.10459>. arXiv:2403.10459 [cs, stat].
- Tor Lattimore and Csaba Szepesvári. *Bandit Algorithms*. Cambridge University Press, 2020. Covers bandits; useful foundational material for exploration.
- Yann LeCun, Leon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998. URL <http://yann.lecun.com/exdb/publis/pdf/lecun-01a.pdf>.
- Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. Deep learning. *Nature*, 521:436–444, 2015. URL <https://www.cs.toronto.edu/~hinton/absps/NatureDeepReview.pdf>.
- Eng Hock Lee and Vladimir Cherkassky. Vc theoretical explanation of double descent, 2022. URL <https://arxiv.org/abs/2205.15549>.
- Jaehoon Lee, Lechao Xiao, Samuel S. Schoenholz, Yasaman Bahri, Roman Novak, Jascha Sohl-Dickstein, and Jeffrey Pennington. Wide neural networks of any depth evolve as linear models under gradient descent. In *Advances in Neural Information Processing Systems (NeurIPS)*, pages 8572–8583, 2019.
- Qianyi Li and Haim Sompolsky. Statistical mechanics of deep linear neural networks: The backpropagating kernel renormalization. *Physical Review X*, 11(3):031059, 2021. doi: 10.1103/PhysRevX.11.031059. URL <https://doi.org/10.1103/PhysRevX.11.031059>.
- Xinyue Li and Rishi Sonthalia. Least squares regression can exhibit under-parameterized double descent. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 25510–25560. Curran Associates, Inc., 2024. URL https://proceedings.neurips.cc/paper_files/paper/2024/file/2d43f7a61b57f83619f82c971e4bddc0-Paper-Conference.pdf.
- Timothy P. Lillicrap, Jonathan J. Hunt, Alexandre Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver, and Daan Wierstra. Continuous control with deep reinforcement learning. *arXiv preprint arXiv:1509.02971*, 2015. URL <https://arxiv.org/abs/1509.02971>. DDPG — deterministic policy gradient for continuous control.

- Zachary C. Lipton. The mythos of model interpretability, 2017. URL <https://arxiv.org/abs/1606.03490>.
- Zachary C. Lipton. The mythos of model interpretability. *Communications of the ACM*, 61(10):36–43, 2018. doi: 10.1145/3233231. URL <https://arxiv.org/abs/1606.03490>.
- Zachary C. Lipton and Jacob Steinhardt. Troubling Trends in Machine Learning Scholarship, 2018.
- Nick Littlestone and Manfred K. Warmuth. The weighted majority algorithm. *Information and Computation*, 108(2): 212–261, 1994. URL <https://www.cs.cmu.edu/~avrim/ML94.pdf>.
- Chris Yuhao Liu and Jeffrey Flanigan. Understanding the role of optimization in double descent, 2023. URL <https://arxiv.org/abs/2312.03951>.
- Fanghui Liu, Zhenyu Liao, and Johan A. K. Suykens. Kernel regression in high dimensions: Refined analysis beyond double descent, 2021. URL <https://arxiv.org/abs/2010.02681>.
- Ziming Liu, Eric J. Michaud, and Max Tegmark. Grokking explained: A statistical phenomenon. *arXiv preprint*, 2023.
- Dmytro Lopushanskyy and Borun Shi. Graph Neural Networks on Graph Databases, 2024.
- P. J. F. Lucas. *Principles of Expert Systems*. University Lecture Notes / PDF, 1989. URL <https://www.cs.ru.nl/P.Lucas/proe.pdf>. PDF available online.
- Scott M. Lundberg and Su-In Lee. A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems (NeurIPS)*, 30, 2017. URL <https://arxiv.org/abs/1705.07874>.
- Jing Luo, Huiyuan Wang, and Weiran Huang. Investigating the impact of model complexity in large language models, 2024. URL <https://arxiv.org/abs/2410.00699>.
- Yifan Luo and Bin Dong. Double descent of discrepancy: A task-, data-, and model-agnostic phenomenon, 2023. URL <https://arxiv.org/abs/2305.15907>.
- Zhi-Quan Luo and Paul Tseng. On the convergence of the coordinate descent method for convex differentiable minimization. *Journal of Optimization Theory and Applications*, 72(1):7–35, 1992. doi: 10.1007/BF00939948.
- Duy V. N. Luong, Panos Pappas, Daniel Rueckert, and Berc Rustem. A weighted mirror descent algorithm for nonsmooth convex optimization problem. *Journal of Optimization Theory and Applications*, 170(3):900–915, 2016. doi: 10.1007/s10957-016-0963-5.
- Stephan Mandt, Matthew D. Hoffman, and David M. Blei. Stochastic gradient descent as approximate bayesian inference. *Journal of Machine Learning Research*, 2017. arXiv:1704.04289.
- Yuri Manin and Matilde Marcolli. Homotopy theoretic and categorical models of neural information networks. *Compositionality*, Volume 6 (2024), September 2024. ISSN 2631-4444. doi: 10.46298/compositionality-6-4. URL <http://dx.doi.org/10.46298/compositionality-6-4>.
- Gary Marcus. Deep learning: A critical appraisal. *arXiv preprint arXiv:1801.00631*, 2018. URL <https://arxiv.org/abs/1801.00631>.
- John McCarthy and Patrick J. Hayes. Some philosophical problems from the standpoint of artificial intelligence. In B. Meltzer and D. Michie, editors, *Machine Intelligence 4*, pages 463–502. Edinburgh University Press, 1969.
- Warren S. McCulloch and Walter Pitts. A logical calculus of the ideas immanent in nervous activity. *Bulletin of Mathematical Biophysics*, 5:115–133, 1943.

- Ninareh Mehrabi, Fred Morstatter, Nripsuta Saxena, Kristina Lerman, and Aram Galstyan. A Survey on Bias and Fairness in Machine Learning, January 2022. URL <http://arxiv.org/abs/1908.09635>. arXiv:1908.09635 [cs].
- Pankaj Mehta and David J. Schwab. An exact mapping between the variational renormalization group and deep learning. 2014. URL <https://arxiv.org/abs/1410.3831>. arXiv:1410.3831.
- Shibin Mei, Chenglong Zhao, Shengchao Yuan, and Bingbing Ni. Towards Bridging Sample Complexity and Model Capacity. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36 of *AAAI'22 Technical Tracks*, pages 1972–1980, June 2022. doi: 10.1609/aaai.v36i2.20092.
- Song Mei and Andrea Montanari. The generalization error of random features regression: Precise asymptotics and double descent curve. *arXiv preprint arXiv:1908.05355*, 2019a.
- Song Mei and Andrea Montanari. The generalization error of random features regression: Precise asymptotics and double descent curve. *arXiv preprint*, 2019b. URL <https://arxiv.org/abs/1908.05355>.
- Song Mei and Andrea Montanari. The generalization error of random features regression: Precise asymptotics and double descent curve, 2020. URL <https://arxiv.org/abs/1908.05355>.
- Song Mei, Andrea Montanari, and Phan-Minh Nguyen. A mean field view of the landscape of two-layer neural networks. *Proc. Natl. Acad. Sci. USA*, 2018. arXiv:1804.06561.
- Jan Mendling, Benoît Depaire, and Henrik Leopold. Theory and practice of algorithm engineering, 2021. URL <https://arxiv.org/abs/2107.10675>.
- M. Mezard, G. Parisi, and M. A. Virasoro. *Spin Glass Theory and Beyond*. World Scientific, Singapore, 1987. ISBN 978-9971501533.
- Marc Mezard and Andrea Montanari. *Information, Physics, and Computation*. Oxford University Press, Oxford, 2009. ISBN 978-0199233210.
- Peter B. Miltersen, Jaikumar Radhakrishnan, and Ingo Wegener. On Converting CNF to DNF. *Theoretical Computer Science*, 347(1–2):325–335, 2005.
- Marvin Minsky and Seymour Papert. *Perceptrons: An Introduction to Computational Geometry*. MIT Press, 1969.
- Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, et al. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533, 2015. doi: 10.1038/nature14236. URL <https://www.nature.com/articles/nature14236>. DQN — Deep Q-Network; landmark deep-RL result.
- Volodymyr Mnih, Adrià Puigdomènech Badia, Mehdi Mirza, Alex Graves, Tim Lillicrap, Timothy Harley, David Silver, and Koray Kavukcuoglu. Asynchronous methods for deep reinforcement learning. *Proceedings of the 33rd International Conference on Machine Learning (ICML)*, 2016. A3C — efficient parallel optimisation for deep RL.
- Mehryar Mohri, Afshin Rostamizadeh, and Ameet Talwalkar. *Foundations of Machine Learning*. The MIT Press, 2012. ISBN 026201825X.
- Pooya Molavi, Alireza Tahbaz-Salehi, and Andrea Vedolin. Model Complexity, Expectations, and Asset Prices. *The Review of Economic Studies*, 91(4):2462–2507, July 2024. ISSN 0034-6527. doi: 10.1093/restud/rdad073. URL <https://doi.org/10.1093/restud/rdad073>.
- Christoph Molnar. Interpretable machine learning. *Online book*, 2020. URL <https://christophm.github.io/interpretable-ml-book/>. Free, continuously updated reference book.

- Christoph Molnar, Giuseppe Casalicchio, and Bernd Bischl. *Quantifying Model Complexity via Functional Decomposition for Better Post-hoc Interpretability*, pages 193–204. Springer International Publishing, 2020a. ISBN 9783030438234. doi: 10.1007/978-3-030-43823-4_17. URL http://dx.doi.org/10.1007/978-3-030-43823-4_17.
- Christoph Molnar, Gunnar König, Julia Herbinger, Timo Freiesleben, Susanne Dandl, Christian A. Scholbeck, Giuseppe Casalicchio, Moritz Grosse-Wentrup, and Bernd Bischl. General Pitfalls of Model-Agnostic Interpretation Methods for Machine Learning Models. *arXiv preprint arXiv:2007.04131*, 2020b.
- Mircea Muresan. Bernstein approximation and beyond: Proofs by means of elementary probability theory. *arXiv preprint arXiv:2307.11533*, 2023.
- Vidya Muthukumar, Kailas Vodrahalli, Vignesh Subramanian, and Anant Sahai. Understanding deep learning requires rethinking generalization. *Journal of the ACM*, 68(4), 2021. doi: 10.1145/3446776.
- Vaishnavh Nagarajan and J. Zico Kolter. Uniform convergence may be unable to explain generalization in deep learning, 2021.
- Preetum Nakkiran. More data can hurt for linear regression: Sample-wise double descent, 2019a. URL <https://arxiv.org/abs/1912.07242>.
- Preetum Nakkiran. More data can hurt for linear regression: Sample-wise double descent. *arXiv:1912.07242*, 2019b. URL <https://arxiv.org/abs/1912.07242>. *arXiv preprint*.
- Preetum Nakkiran, Gal Kaplun, Yamini Bansal, Tristan Yang, Boaz Barak, and Ilya Sutskever. Deep Double Descent: Where Bigger Models and More Data Hurt, December 2019. URL <http://arxiv.org/abs/1912.02292>. *arXiv:1912.02292 [cs, stat]*.
- Preetum Nakkiran, Prayaag Venkat, Sham Kakade, and Tengyu Ma. Optimal regularization can mitigate double descent (openreview entry). OpenReview forum submission, 2020. URL <https://openreview.net/forum?id=7R7fAoUygoa>.
- Preetum Nakkiran, Prayaag Venkat, Sham Kakade, and Tengyu Ma. Optimal regularization can mitigate double descent, 2021. URL <https://arxiv.org/abs/2003.01897>.
- Balas K. Natarajan. On learning sets and functions. *Machine Learning*, 4:287–309, 1989. Generalizes VC dimension to multiclass classification.
- Robert F. Nau. De finetti was right: Probability does not exist. *Theory and Decision*, 51(2/4):89–124, 2001. doi: 10.1023/A:1015525808214.
- Brady Neal. On the bias-variance tradeoff: Textbooks need an update, 2019. URL <https://arxiv.org/abs/1912.08286>.
- Brady Neal, Sarthak Mittal, Aristide Baratin, Vinayak Tania, Matthew Scicluna, Simon Lacoste-Julien, and Ioannis Mitliagkas. A modern take on the bias-variance tradeoff in neural networks. *arXiv preprint arXiv:1810.08591*, 2018.
- Arkadii S. Nemirovski and David B. Yudin. *Problem Complexity and Method Efficiency in Optimization*. Wiley-Interscience Series in Discrete Mathematics. John Wiley and Sons, Chichester, UK, 1983. ISBN 9780471103455.
- Behnam Neyshabur, Ryota Tomioka, and Nathan Srebro. Norm-based capacity control in neural networks. In *Conference on Learning Theory (COLT)*, pages 1376–1401, 2015.
- Behnam Neyshabur, Srinadh Bhojanapalli, David McAllester, and Nathan Srebro. Exploring generalization in deep learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, pages 5947–5956, 2017.

- Jiawang Nie, Zhi Yang, and Guoyin Zhou. The saddle point problem of polynomials. *Foundations of Computational Mathematics*, 22(5):1133–1169, 2021. doi: 10.1007/s10208-021-09526-8. URL <https://link.springer.com/article/10.1007/s10208-021-09526-8>.
- Peter Norvig. *Paradigms of Artificial Intelligence Programming*. Morgan Kaufmann, 1992.
- Amanda Olmin and Fredrik Lindsten. Towards understanding epoch-wise double descent in two-layer linear neural networks, 2024. URL <https://arxiv.org/abs/2407.09845>.
- Kenta Oono and Taiji Suzuki. Graph neural networks exponentially lose expressive power for node classification. In *International Conference on Learning Representations*, 2020. URL <https://openreview.net/forum?id=S1ld02EFPp>.
- Urszula Ostaszewska and Krzysztof Zajkowski. Cramér transform and t-entropy. *Positivity*, 18(2):347–358, June 2014. ISSN 1572-9281. doi: 10.1007/s11117-013-0247-3. URL <http://dx.doi.org/10.1007/s11117-013-0247-3>.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744, 2022. URL <https://arxiv.org/abs/2203.02155>. Introduced RLHF for aligning language models with human preferences.
- Liam Paninski. Statistics 4107: Intro to Math Stat (fall 2005), 2005. URL <https://sites.stat.columbia.edu/liam/teaching/4107-fall105/>.
- Vardan Papayan, X. Y. Han, and David L. Donoho. Prevalence of neural collapse during the terminal phase of deep learning training. *Proceedings of the National Academy of Sciences*, 117(40):24652–24663, 2020. doi: 10.1073/pnas.2015509117.
- Judea Pearl. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, 2nd edition, 2009. ISBN 9780521895606. doi: 10.1017/CBO9780511803161.
- Zongrui Pei. Statistical physics for artificial neural networks, 2025. URL <https://arxiv.org/abs/2512.06518>.
- Mohammad Pezeshki, Amartya Mitra, Yoshua Bengio, and Guillaume Lajoie. Multi-scale feature learning dynamics: Insights for double descent, 2021. URL <https://arxiv.org/abs/2112.03215>.
- David Pfau. A generalized bias-variance decomposition for bregman divergences. Technical report, 2013.
- Jr. Philip C. Jackson. *Introduction to Artificial Intelligence*. Dover Publications, Mineola, NY, 2nd edition, 1985. ISBN 9780486843070. URL https://archive.org/details/Introduction_To_Artificial_Intelligence_2nd_Ed_Philip_C_Jackson_Jr. Scanned archived copy available at Archive.org.
- David Pollard. Convergence of stochastic processes. In *Springer Series in Statistics*. Springer, 1984. Contains definitions and results for pseudo-dimension of real-valued function classes.
- Alethea Power, Yuri Burda, Harri Edwards, Igor Babuschkin, and Vedant Misra. Grokking: Generalization beyond overfitting on small algorithmic datasets, 2022a. URL <https://arxiv.org/abs/2201.02177>.
- Alethea Power, Yuri Burda, Harri Edwards, Greg Baker, and Igor Babuschkin. Grokking: Generalization beyond overfitting on small algorithmic datasets. In *International Conference on Learning Representations (ICLR)*, 2022b. URL <https://arxiv.org/abs/2201.02177>.
- Victor Quéto and Enzo Tartaglione. Can we avoid Double Descent in Deep Neural Networks? In *2023 IEEE International Conference on Image Processing (ICIP)*, pages 1625–1629, October 2023a. doi: 10.1109/ICIP49359.2023.10222624. URL <http://arxiv.org/abs/2302.13259>. arXiv:2302.13259 [cs].

- Victor Quétu and Enzo Tartaglione. Can we avoid Double Descent in Deep Neural Networks? In *2023 IEEE International Conference on Image Processing (ICIP)*, pages 1625–1629, October 2023b. doi: 10.1109/ICIP49359.2023.10222624. URL <http://arxiv.org/abs/2302.13259>. arXiv:2302.13259 [cs].
- Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. Improving language understanding by generative pre-training. *OpenAI blog*, 2018. URL https://cdn.openai.com/research-covers/language-unsupervised/language_understanding_paper.pdf. The original GPT paper introducing generative pre-training.
- Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9, 2019. URL https://d4mucfpksywv.cloudfront.net/better-language-models/language_models_are_unsupervised_multitask_learners.pdf. Introduced GPT-2, demonstrating zero-shot task transfer.
- Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140):1–67, 2020. URL <https://arxiv.org/abs/1910.10683>. Introduced T5, treating every NLP task as a text-to-text problem.
- Garvesh Raskutti and Sayan Mukherjee. The information geometry of mirror descent. *arXiv Preprint*, 2013. arXiv:1310.7780.
- M. Rattray, D. Saad, and Shun-ichi Amari. Natural gradient descent for on-line learning. *Physical Review Letters*, 81: 5461–5464, 1998. doi: 10.1103/PhysRevLett.81.5461.
- Yasaman Razeghi et al. Not all explanations for deep learning phenomena are created equal. *arXiv preprint*, 2025.
- Nathalie Revol and Philippe Theveny. Numerical reproducibility and parallel computations: Issues for interval algorithms. *IEEE Transactions on Computers*, 63(8):1915–1924, August 2014. ISSN 2326-3814. doi: 10.1109/tc.2014.2322593. URL <http://dx.doi.org/10.1109/TC.2014.2322593>.
- Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. "why should i trust you?": Explaining the predictions of any classifier. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1135–1144, 2016. doi: 10.1145/2939672.2939778. URL <https://arxiv.org/abs/1602.04938>.
- Zohar Ringel, Noa Rubin, Edo Mor, Moritz Helias, and Inbar Seroussi. Applications of statistical field theory in deep learning, 2025. URL <https://arxiv.org/abs/2502.18553>.
- Herbert Robbins and Sutton Monroe. A stochastic approximation method. *Annals of Mathematical Statistics*, 22(3): 400–407, 1951. doi: 10.1214/aoms/1177729586.
- Daniel A. Roberts, Sho Yaida, and Boris Hanin. The principles of deep learning theory. Technical report, 2021. URL <https://arxiv.org/abs/2106.10165>. arXiv:2106.10165; expanded into a Cambridge Univ. Press book.
- Daniel A. Roberts, Sho Yaida, and Boris Hanin. *The Principles of Deep Learning Theory: An Effective Theory Approach to Understanding Neural Networks*. Cambridge University Press, 2022.
- Paul M. Romer. Mathiness in the Theory of Economic Growth. *American Economic Review*, 105(5):89–93, 2015.
- Frank Rosenblatt. The perceptron: a probabilistic model for information storage and organization in the brain. *Psychological review*, 65 6:386–408, 1958a. URL <https://api.semanticscholar.org/CorpusID:12781225>.
- Frank Rosenblatt. The perceptron: A probabilistic model for information storage and organization in the brain. *Psychological Review*, 65(6):386–408, 1958b. doi: 10.1037/h0042519.

- Emanuele Rossi, Ben Chamberlain, Fabrizio Frasca, Davide Eynard, Federico Monti, and Michael Bronstein. Temporal graph networks for deep learning on dynamic graphs, 2020. URL <https://arxiv.org/abs/2006.10637>.
- Sebastian Ruder. An overview of gradient descent optimization algorithms, June 2017. URL <http://arxiv.org/abs/1609.04747>. arXiv:1609.04747 [cs].
- David E. Rumelhart, Geoffrey E. Hinton, and Ronald J. Williams. Learning representations by back-propagating errors. *Nature*, 323:533–536, 1986. doi: 10.1038/323533a0.
- Carl Runge. Über empirische funktionen und die interpolation zwischen äquidistanten ordinaten. *Zeitschrift für Mathematik und Physik*, 46:224–243, 1901. The original Runge example showing divergence with equispaced interpolation.
- Daniel Russo and James Zou. Controlling bias in adaptive data analysis using information theory. In *Proceedings of AISTATS*, 2016. Introduced info-theoretic generalization bounds.
- Sarah Sachs, Tim van Erven, Liam Hodgkinson, Rajiv Khanna, and Umut Simsekli. Generalization guarantees via algorithm-dependent rademacher complexity. *arXiv*, 2307.02501, 2023. Focuses on algorithm- and data-dependent complexity for modern learning, bridging classic measures with practical generalization.
- N. Sauer. On the density of families of sets. *Journal of Combinatorial Theory, Series A*, 13(1):145–147, 1972. doi: 10.1016/0097-3165(72)90033-X.
- Andrew M. Saxe, James L. McClelland, and Surya Ganguli. Exact solutions to the nonlinear dynamics of learning in deep linear neural networks. *arXiv preprint arXiv:1312.6120*, 2014. ICLR 2014.
- Devansh Saxena, Ji-Youn Jung, Jodi Forlizzi, Kenneth Holstein, and John Zimmerman. Ai mismatches: Identifying potential algorithmic harms before ai development, 2025. URL <https://arxiv.org/abs/2502.18682>.
- Franco Scarselli, Marco Gori, Ah Chung Tsoi, Markus Hagenbuchner, and Gabriele Monfardini. The graph neural network model. *IEEE Transactions on Neural Networks*, 20(1):61–80, 2009. doi: 10.1109/TNN.2008.2005605.
- Rylan Schaeffer, Mikail Khona, Zachary Robertson, Akhilan Boopathy, Kateryna Pistunova, Jason W. Rocks, Ila Rani Fiete, and Oluwasanmi Koyejo. Double Descent Demystified: Identifying, Interpreting and Ablating the Sources of a Deep Learning Puzzle, March 2023. URL <http://arxiv.org/abs/2303.14151>. arXiv:2303.14151 [cs, stat].
- Anton R. Schep. Weierstrass’ proof of the weierstrass approximation theorem. Technical report / preprint, University of South Carolina, 2007. URL <https://people.math.sc.edu/schep/weierstrass.pdf>. PDF, May 1.
- John Schulman, Sergey Levine, Philipp Moritz, Michael Jordan, and Pieter Abbeel. Trust region policy optimization. *Proceedings of the 32nd International Conference on Machine Learning (ICML)*, 2015. TRPO; policy optimisation with guaranteed monotonic improvement.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017. URL <https://arxiv.org/abs/1707.06347>. PPO — simple, widely-used policy-gradient algorithm.
- John R. Searle. Minds, brains, and programs. *Behavioral and Brain Sciences*, 3(3):417–457, 1980.
- Harald Semmelrock, Simone Kopeinik, Dieter Theiler, Tony Ross-Hellauer, and Dominik Kowald. Reproducibility in machine learning-driven research, 2023. URL <https://arxiv.org/abs/2307.10320>.
- H. S. Seung, H. Sompolinsky, and N. Tishby. Statistical mechanics of learning from examples. *Physical Review A*, 45(8):6056–6091, 1992. doi: 10.1103/PhysRevA.45.6056. URL <https://doi.org/10.1103/PhysRevA.45.6056>.

- Shai Shalev-Shwartz. Online learning and online convex optimization. *Foundations and Trends in Machine Learning*, 4(2):107–194, 2012. URL <https://www.cs.huji.ac.il/~shais/papers/OLsurvey.pdf>.
- Shai Shalev-Shwartz and Shai Ben-David. *Understanding Machine Learning: From Theory to Algorithms*. Cambridge University Press, USA, 2014. ISBN 1107057132.
- Shai Shalev-Shwartz, Ohad Shamir, Nathan Srebro, and Karthik Sridharan. Learnability, stability and uniform convergence. *Journal of Machine Learning Research*, 11(90):2635–2670, 2010. URL <http://jmlr.org/papers/v11/shalev-shwartz10a.html>.
- Rahul Sharma and Alex Aiken. Bias-variance tradeoffs in program analysis. In *Proceedings of the 41st ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL ’14*, pages 127–137, New York, NY, USA, 2014. Association for Computing Machinery. ISBN 978-1-4503-2544-8. doi: 10.1145/2535838.2535853. URL <https://doi.org/10.1145/2535838.2535853>.
- Saharon Shelah. A combinatorial problem; stability and order for models and theories in infinitary languages. *Pacific Journal of Mathematics*, 41(1):247–261, 1972.
- Zewen Shen and Kirill Serkh. On polynomial interpolation in the monomial basis. *SIAM Journal on Numerical Analysis*, 63(2):469–494, March 2025. ISSN 1095-7170. doi: 10.1137/23m1623215. URL <http://dx.doi.org/10.1137/23M1623215>.
- D. Sherrington and S. Kirkpatrick. Solvable model of a spin-glass. *Physical Review Letters*, 35:1792–1796, 1975.
- Cheng Shi, Liming Pan, Hong Hu, and Ivan Dokmanić. Homophily modulates double descent generalization in graph convolution networks, 2024. URL <https://arxiv.org/abs/2212.13069>.
- Hava T. Siegelmann and Eduardo D. Sontag. Turing computability with neural nets. *Applied Mathematics Letters*, 4(6):77–80, 1991. doi: 10.1016/0893-9659(91)90080-F. URL <https://www.sontaglab.org/FTPDIR/aml-turing.pdf>.
- David Silver, Aja Huang, Chris J. Maddison, Arthur Guez, Laurent Sifre, George van den Driessche, Julian Schrittwieser, Ioannis Antonoglou, et al. Mastering the game of go with deep neural networks and tree search. *Nature*, 529:484–489, 2016. doi: 10.1038/nature16961. URL <https://www.nature.com/articles/nature16961>. AlphaGo — combines supervised learning + reinforcement learning + MCTS.
- Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition. *arXiv preprint*, 2014. URL <https://arxiv.org/abs/1409.1556>.
- Jascha Sohl-Dickstein. The natural gradient by analogy to signal whitening, and recipes and tricks for its use. *arXiv Preprint*, 2012. arXiv:1205.1828.
- Yang Song and Stefano Ermon. Score-based generative modeling through stochastic differential equations. 2020. arXiv:2011.13456.
- Daniel Soudry, Elad Hoffer, Mor Shpigel Nacson, Suriya Gunasekar, and Nathan Srebro. The implicit bias of gradient descent on separable data, 2024. URL <https://arxiv.org/abs/1710.10345>.
- Jann Spiess, Guido Imbens, and Amar Venugopal. Double and single descent in causal inference with an application to high-dimensional synthetic control, 2023. URL <https://arxiv.org/abs/2305.00700>.
- Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. Dropout: A simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research*, 15(56):1929–1958, 2014a. URL <https://jmlr.org/papers/volume15/srivastava14a/srivastava14a.pdf>.

- Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. Dropout: A simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research*, 15:1929–1958, 2014b. URL <https://www.jmlr.org/papers/v15/srivastava14a.html>.
- Tom F. Sterkenburg. Statistical learning theory and occam’s razor: The core argument. *Minds and Machines*, 35(1), November 2024. ISSN 1572-8641. doi: 10.1007/s11023-024-09703-y. URL <http://dx.doi.org/10.1007/s11023-024-09703-y>.
- Leon Sterling and Ehud Shapiro. *The Art of Prolog: Advanced Programming Techniques*. MIT Press, 2nd edition, 1994.
- Lucy A. Suchman. *Plans and Situated Actions: The Problem of Human–Machine Communication*. Cambridge University Press, 1987. ISBN 0521388473.
- Masashi Sugiyama. *Introduction to Statistical Machine Learning*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 2015. ISBN 9780128023501.
- B. Sury. Weierstrass’ theorem — leaving no stone unturned. *Lecture Notes / Expository Review*, 2019. Expository article comparing different proofs including Bernstein.
- Richard S. Sutton. The bitter lesson. Web essay / blog post, 2019. URL <https://www.incompleteideas.net/IncIdeas/BitterLesson.html>.
- Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2nd edition, 2018. Available online: <https://incompleteideas.net/book/the-book-2nd.html>.
- Lars Pålsson Syll. Post-real economics — a severe case of mathiness. Blog post, Heterodox Economic Blogs, 2024.
- Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, Scott Reed, Dragomir Anguelov, Dumitru Erhan, Vincent Vanhoucke, and Andrew Rabinovich. Going deeper with convolutions. *arXiv preprint*, 2014. URL <https://arxiv.org/abs/1409.4842>.
- Csaba Szepesvári. *Algorithms for Reinforcement Learning*. Morgan & Claypool, 2010. Concise, theory-oriented graduate text.
- Roberto Tamassia and Luca Vismara. A case study in algorithm engineering for geometric computing. In *International journal of computational geometry and applications*, 1997. URL <https://api.semanticscholar.org/CorpusID:5770657>.
- James H. Tanis, Chris Giannella, and Adrian V. Mariano. Introduction to Graph Neural Networks: A Starting Point for Machine Learning Engineers, 2024.
- Matus Telgarsky. Benefits of depth in neural networks. In *Proceedings of COLT 2016*, 2016. URL <https://arxiv.org/abs/1602.04485>.
- Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023a. URL <https://arxiv.org/abs/2302.13971>. Meta’s open foundation models ranging from 7B to 65B parameters.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023b. URL <https://arxiv.org/abs/2307.09288>. Updated and improved LLaMA models with better performance.
- Mark K. Transtrum, Gus L. W. Hart, Tyler J. Jarvis, and Jared P. Whitehead. egad! double descent is explained by generalized aliasing decomposition, 2025. URL <https://arxiv.org/abs/2408.08294>.

- Lan V. Truong. On rademacher complexity-based generalization bounds for deep learning, 2025. URL <https://arxiv.org/abs/2208.04284>.
- Paul Tseng. Convergence of a block coordinate descent method for nondifferentiable minimization. *Journal of Optimization Theory and Applications*, 109(3):475–494, 2001. doi: 10.1007/s10957-001-0006-4.
- Onur Tunalı. Empirical Rademacher Complexity and Its Implications to Deep Learning. <https://www.onurtunali.com/ml/2019/02/01/empirical-rademacher-complexity-and-its-implications-to-deep-learning.html>, February 2019.
- L. G. Valiant. A theory of the learnable. *Commun. ACM*, 27(11):1134–1142, November 1984. ISSN 0001-0782. doi: 10.1145/1968.1972. URL <https://doi.org/10.1145/1968.1972>.
- Gido M. van de Ven, Nicholas Soures, and Dhireesha Kudithipudi. Continual learning and catastrophic forgetting, 2024. URL <https://arxiv.org/abs/2403.05175>.
- Vladimir Vapnik. *The Nature of Statistical Learning Theory*. Springer: New York, 1999.
- Vladimir N. Vapnik and Alexey Y. Chervonenkis. On the uniform convergence of relative frequencies of events to their probabilities. *Theory of Probability & Its Applications*, 16(2):264–280, 1971.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30, pages 5998–6008. Neural Information Processing Systems, 2017a. URL <https://arxiv.org/abs/1706.03762>. NeurIPS 2017. The foundational paper introducing the Transformer architecture.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017b. URL <https://arxiv.org/abs/1706.03762>.
- Petar Veličković. Everything is connected: Graph neural networks. *Current Opinion in Structural Biology*, 79:102538, April 2023. ISSN 0959-440X. doi: 10.1016/j.sbi.2023.102538.
- Martin J. Wainwright and Michael I. Jordan. Graphical models, exponential families, and variational inference. *Foundations and Trends in Machine Learning*, 1(1-2):1–305, 2008.
- Huibing Wang, Jinbo Xiong, Zhiqiang Yao, Mingwei Lin, and Jun Ren. Research survey on support vector machine. In *Proceedings of the 10th EAI International Conference on Mobile Multimedia Communications (MOBIMEDIA)*. EAI, December 2017. doi: 10.4108/eai.13-7-2017.2270596.
- Jingyuan Wang, Kai Feng, and Junjie Wu. Svm-based deep stacking networks. In *Proceedings of the Twenty-Eighth International Joint Conference on Artificial Intelligence (IJCAI-19)*, pages 5262–5268, 2019. doi: 10.24963/ijcai.2019/731. URL <https://www.ijcai.org/proceedings/2019/731>.
- Christopher J. C. H. Watkins and Peter Dayan. Q-learning. *Machine Learning*, 8:279–292, 1992.
- Tobias Wegel, Geelon So, Junhyung Park, and Fanny Yang. On the sample complexity of semi-supervised multi-objective learning, 2025.
- Ingo Wegener. *The Complexity of Boolean Functions*. John Wiley & Sons, Chichester, UK, 1987.
- Colin Wei, Zeyuan Allen-Zhu Shen, and Tengyu Ma. Theoretical understanding of deep learning. *Communications of the ACM*, 65(11):98–106, 2022a.

- Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, Ed H. Chi, Tatsunori Hashimoto, Oriol Vinyals, Percy Liang, Jeff Dean, and William Fedus. Emergent abilities of large language models, 2022b. URL <https://arxiv.org/abs/2206.07682>.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35:24824–24837, 2022c. URL <https://arxiv.org/abs/2201.11903>. NeurIPS 2022. Demonstrated how step-by-step reasoning improves LLM performance.
- Gail Weiss, Yoav Goldberg, and Eran Yahav. On the practical computational power of finite precision rnns for language recognition. In *Proceedings of ACL (short papers) / arXiv*, 2018. URL <https://arxiv.org/abs/1805.04908>. arXiv:1805.04908.
- Paul J. Werbos. *Beyond Regression: New Tools for Prediction and Analysis in the Behavioral Sciences*. PhD thesis, Harvard University, 1974.
- Bernard Widrow and Marcian E. Hoff. Adaptive switching circuits. In *1960 IRE WESCON Convention Record, Part 4*, pages 96–104, New York, 1960a. Institute of Radio Engineers.
- Bernard Widrow and Ted Hoff. Adaline: An adaptive ‘adaline’ neuron and adaptive switching circuits. Technical report, Stanford University, 1960b.
- Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. In *Machine Learning*, volume 8, pages 229–256, 1992. Introduces REINFORCE / policy gradient ideas.
- Stephen J Wright. Coordinate descent algorithms. *Mathematical Programming*, 151(1):3–34, 2015. doi: 10.1007/s10107-015-0892-3.
- Lin Xiao and Tong Zhang. A proximal stochastic gradient method with progressive variance reduction. *SIAM Journal on Optimization*, 24(4):2057–2075, 2014. doi: 10.1137/140961791.
- An Xu and Maxim Raginsky. Information-theoretic analysis of generalization capability of learning algorithms. In *NeurIPS*, 2017a. Mutual information bounds.
- Ankit Xu and Maxim Raginsky. Information-theoretic analysis of generalization capability of learning algorithms. *Advances in Neural Information Processing Systems*, 30:2524–2533, 2017b.
- Greg Yang. Scaling limits of wide neural networks with weight sharing: Gaussian process behavior, gradient independence, and neural tangent kernel derivation. *arXiv preprint arXiv:1902.04760*, 2019.
- H. H. Yang and Shun-ichi Amari. Complexity issues in natural gradient descent method for training multilayer perceptrons. *Neural Computation*, 10(8):2137–2157, 1998. doi: 10.1162/089976698300017007.
- Tian-Le Yang and Joe Suzuki. Dropout drops double descent. *Japanese Journal of Statistics and Data Science*, 7(2): 615–632, March 2024. ISSN 2520-8764. doi: 10.1007/s42081-024-00242-5. URL <http://dx.doi.org/10.1007/s42081-024-00242-5>.
- Zitong Yang, Yaodong Yu, Chong You, Jacob Steinhardt, and Yi Ma. Rethinking Bias-Variance Trade-off for Generalization of Neural Networks, December 2020. URL <http://arxiv.org/abs/2002.11328>. arXiv:2002.11328 [cs, stat].
- Chiyuan Zhang, Samy Bengio, Moritz Hardt, Benjamin Recht, and Oriol Vinyals. Understanding deep learning requires rethinking generalization, 2017.

- Jiawei Zhang. Gradient Descent based Optimization Algorithms for Deep Learning Models Training, March 2019. URL <http://arxiv.org/abs/1903.03614>. arXiv:1903.03614 [cs].
- Rui Zhang and Hong Wang. The equivalence relationship between kernel functions based on svm and four-layer functional networks. In De-Shuang Huang, Vito Bevilacqua, Prashan Premaratne, and Praveen Gupta, editors, *Intelligent Computing Methodologies*, volume 8589 of *Lecture Notes in Computer Science*, pages 81–89. Springer, 2014. doi: 10.1007/978-3-319-09339-0_10. URL https://doi.org/10.1007/978-3-319-09339-0_10.
- Tong Zhang. *Mathematical Analysis of Machine Learning Algorithms*. Cambridge University Press, 2023. doi: 10.1017/9781009093057. URL <https://www.cambridge.org/core/books/mathematical-analysis-of-machine-learning-algorithms/EB9BABB05A5C312F19C38E5A01A5ECFC>.
- Ya Shi Zhang. Manipulating sparse double descent, 2024. URL <https://arxiv.org/abs/2401.10686>.
- Wayne Xin Zhao, Kun Zhou, Junyi Li, Tianyi Tang, Xiaolei Wang, Yupeng Hou, Yingqian Min, Beichen Zhang, Junjie Zhang, Zican Dong, et al. A survey of large language models. *arXiv preprint arXiv:2303.18223*, March 2023. URL <https://arxiv.org/abs/2303.18223>. Comprehensive survey covering pre-training, adaptation, utilization, and evaluation of LLMs. Updated through March 2025.
- Wayne Xin Zhao et al. Large language models: A survey. *arXiv preprint arXiv:2402.06196*, February 2024. URL <https://arxiv.org/abs/2402.06196>. Reviews prominent LLM families (GPT, LLaMA, PaLM) and discusses techniques for building and augmenting LLMs.
- Martin Zinkevich. Online convex programming and generalized infinitesimal gradient ascent. In *Proceedings of the 20th International Conference on Machine Learning (ICML)*, pages 928–936, 2003. URL <https://www.cs.cmu.edu/~zinkevich/publications/onlineconvexprog.pdf>.
- Liu Ziyin and Masahito Ueda. Exact phase transitions in deep learning, 2022. URL <https://arxiv.org/abs/2205.12510>.
- Liu Ziyin and Masahito Ueda. Zeroth, first, and second-order phase transitions in deep neural networks. *Phys. Rev. Res.*, 5:043243, Dec 2023. doi: 10.1103/PhysRevResearch.5.043243. URL <https://link.aps.org/doi/10.1103/PhysRevResearch.5.043243>.
- Difan Zou, Yuan Cao, Dongruo Zhou, and Quanquan Gu. Stochastic gradient descent optimizes over-parameterized deep relu networks. In *Machine Learning Research (ICML)*, pages 6962–6972, 2019.